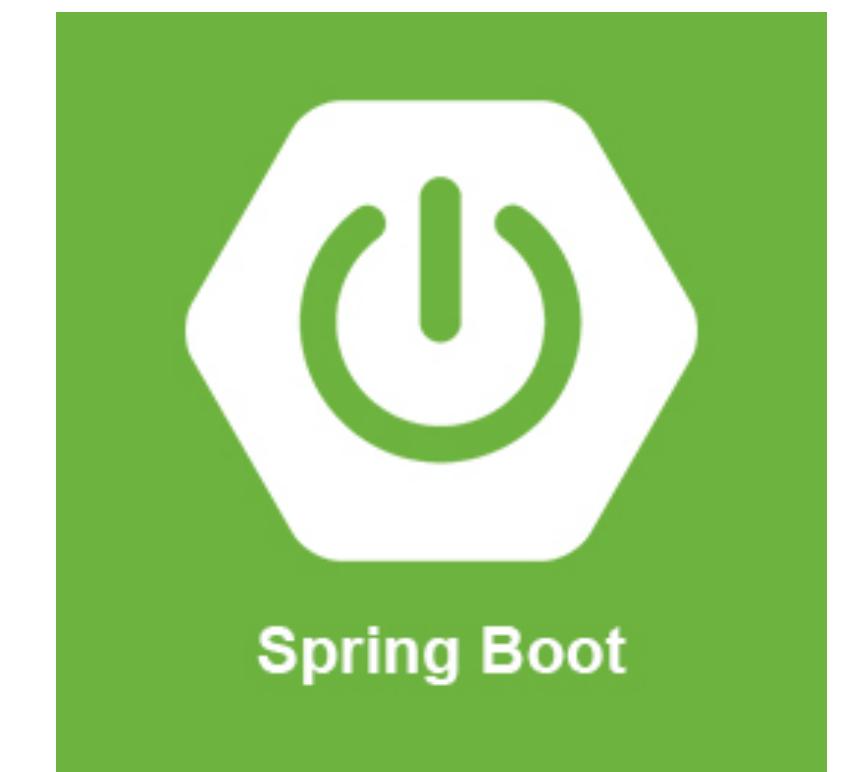


# Kafka For Developers Using Spring Boot



**Dilip Sundarraj**

# About Me

- Dilip
- Building Software's since 2008
- Teaching in **UDEMY** Since 2016

# Whats Covered?

- Introduction to Kafka and internals of Kafka
- Building Enterprise standard Kafka Clients using **Spring-Kafka/  
SpringBoot**
- Resilient Kafka Client applications using **Error-Handling/Retry/Recovery**
- Writing Unit/Integration tests using **JUnit**

# Targeted Audience

- Focused for developers
- Interested in learning the internals of Kafka
- Interested in building Kafka Clients using Spring Boot
- Interested in building Enterprise standard Kafka client applications using Spring boot

# Source Code

**Thank You !**

# **Introduction to Apache Kafka**

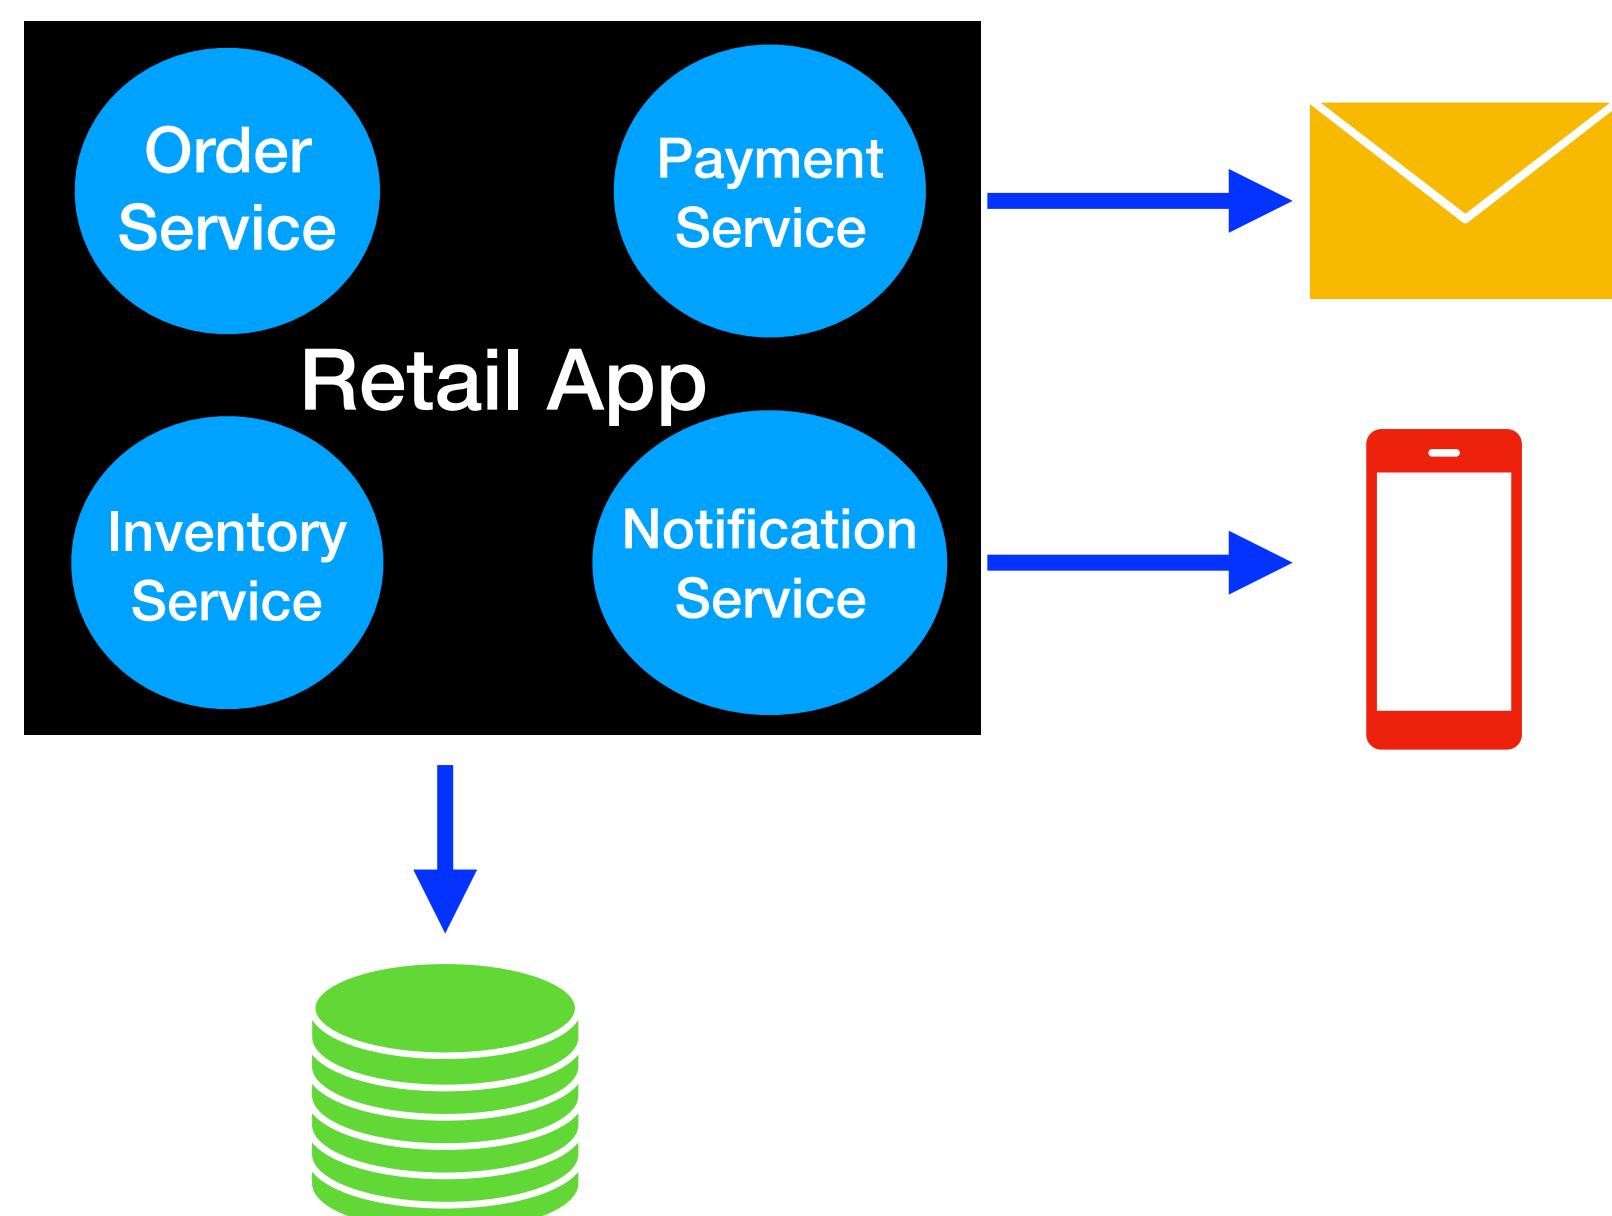
# **Prerequisites**

# Course Prerequisites

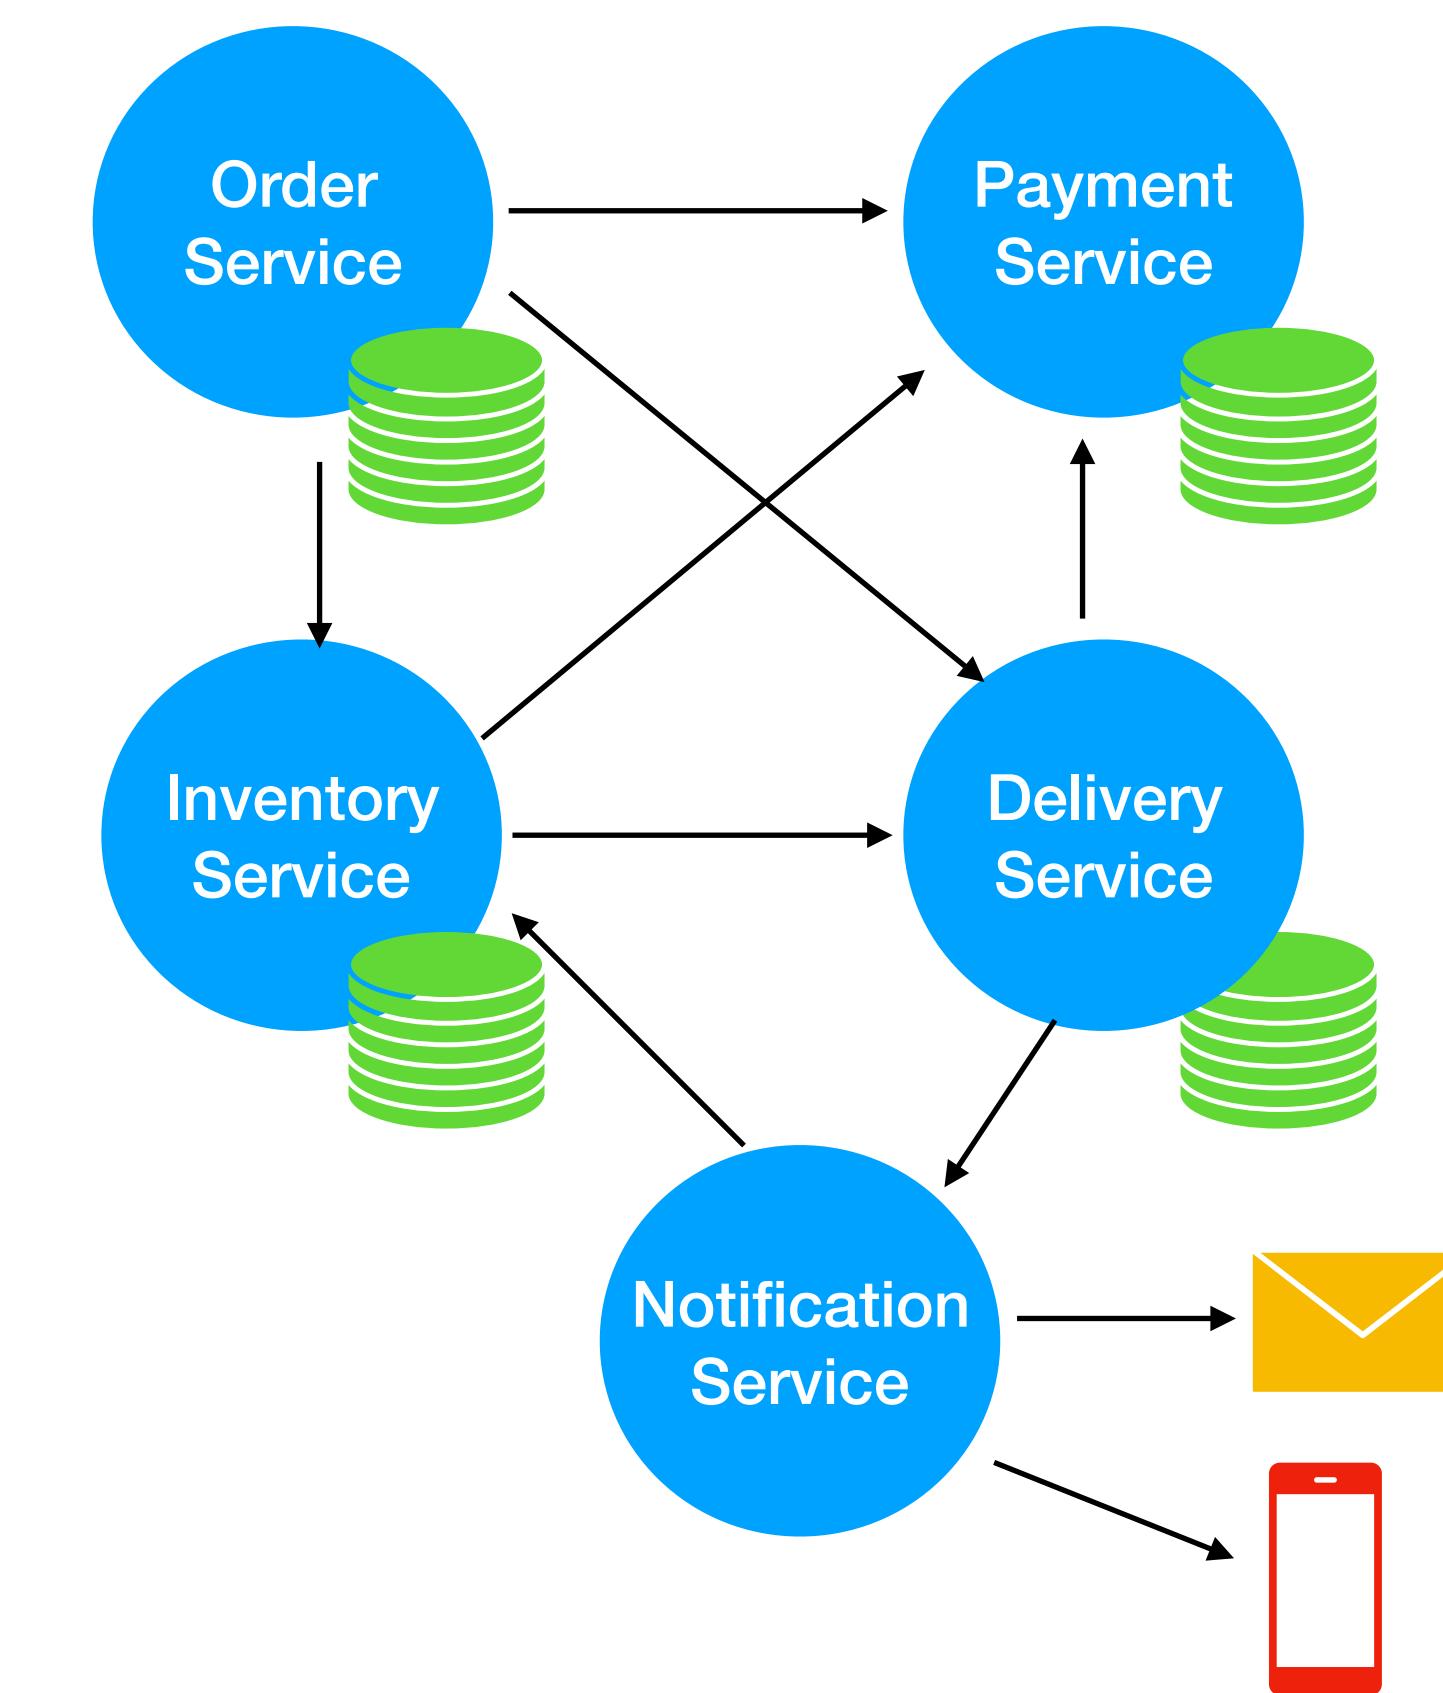
- Prior Knowledge or Working Experience with **Spring Boot/Framework**
- Knowledge about building **Kafka Clients** using Producer and Consumer API
- Knowledge about building **RESTFUL APIs** using Spring Boot
- Experience working with **Spring Data JPA**
- Automated tests using **JUnit**
- Experience Working with **Mockito**
- **Java 11 or Higher** is needed
- **IntelliJ , Eclipse** or any other IDE is needed

# Software Development

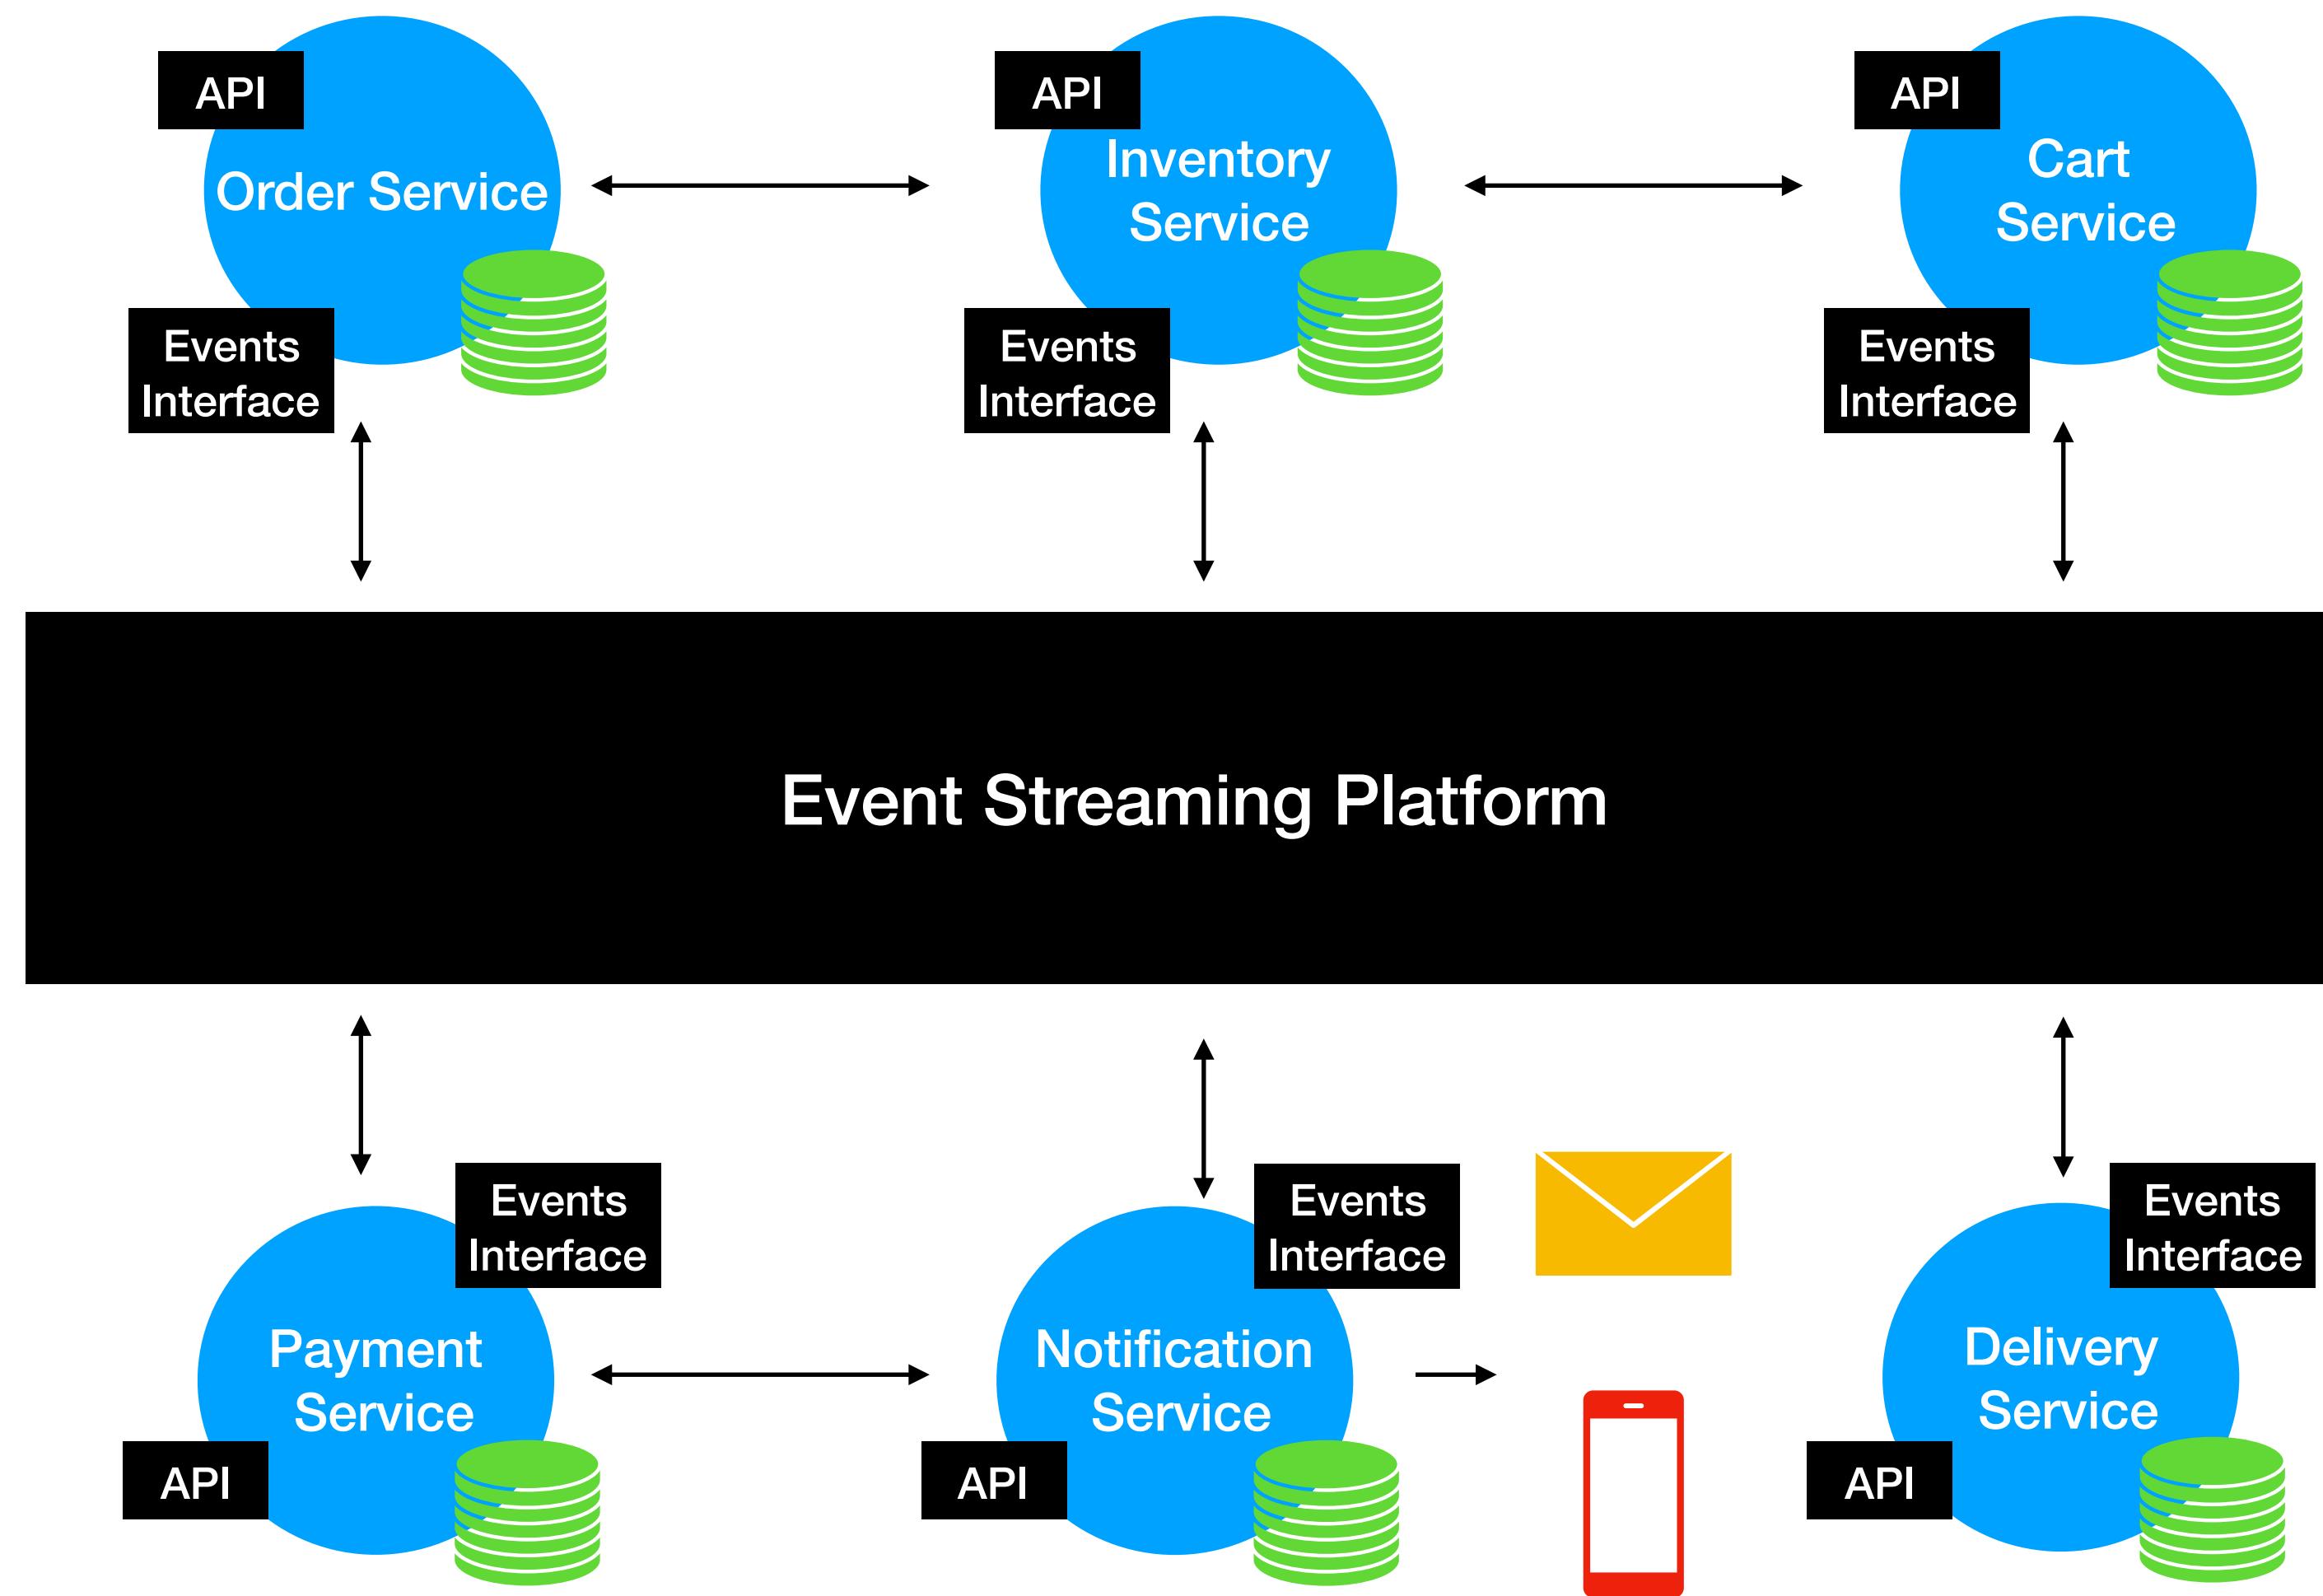
Past



Current

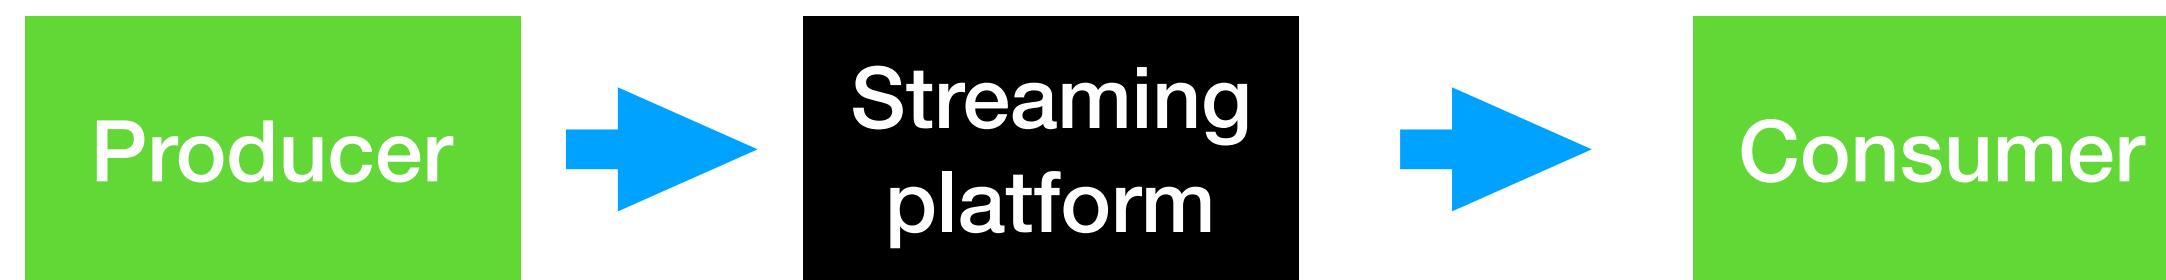


# MicroServices Architecture



# What is an Event Streaming Platform?

- Producers and Consumers subscribe to a stream of records

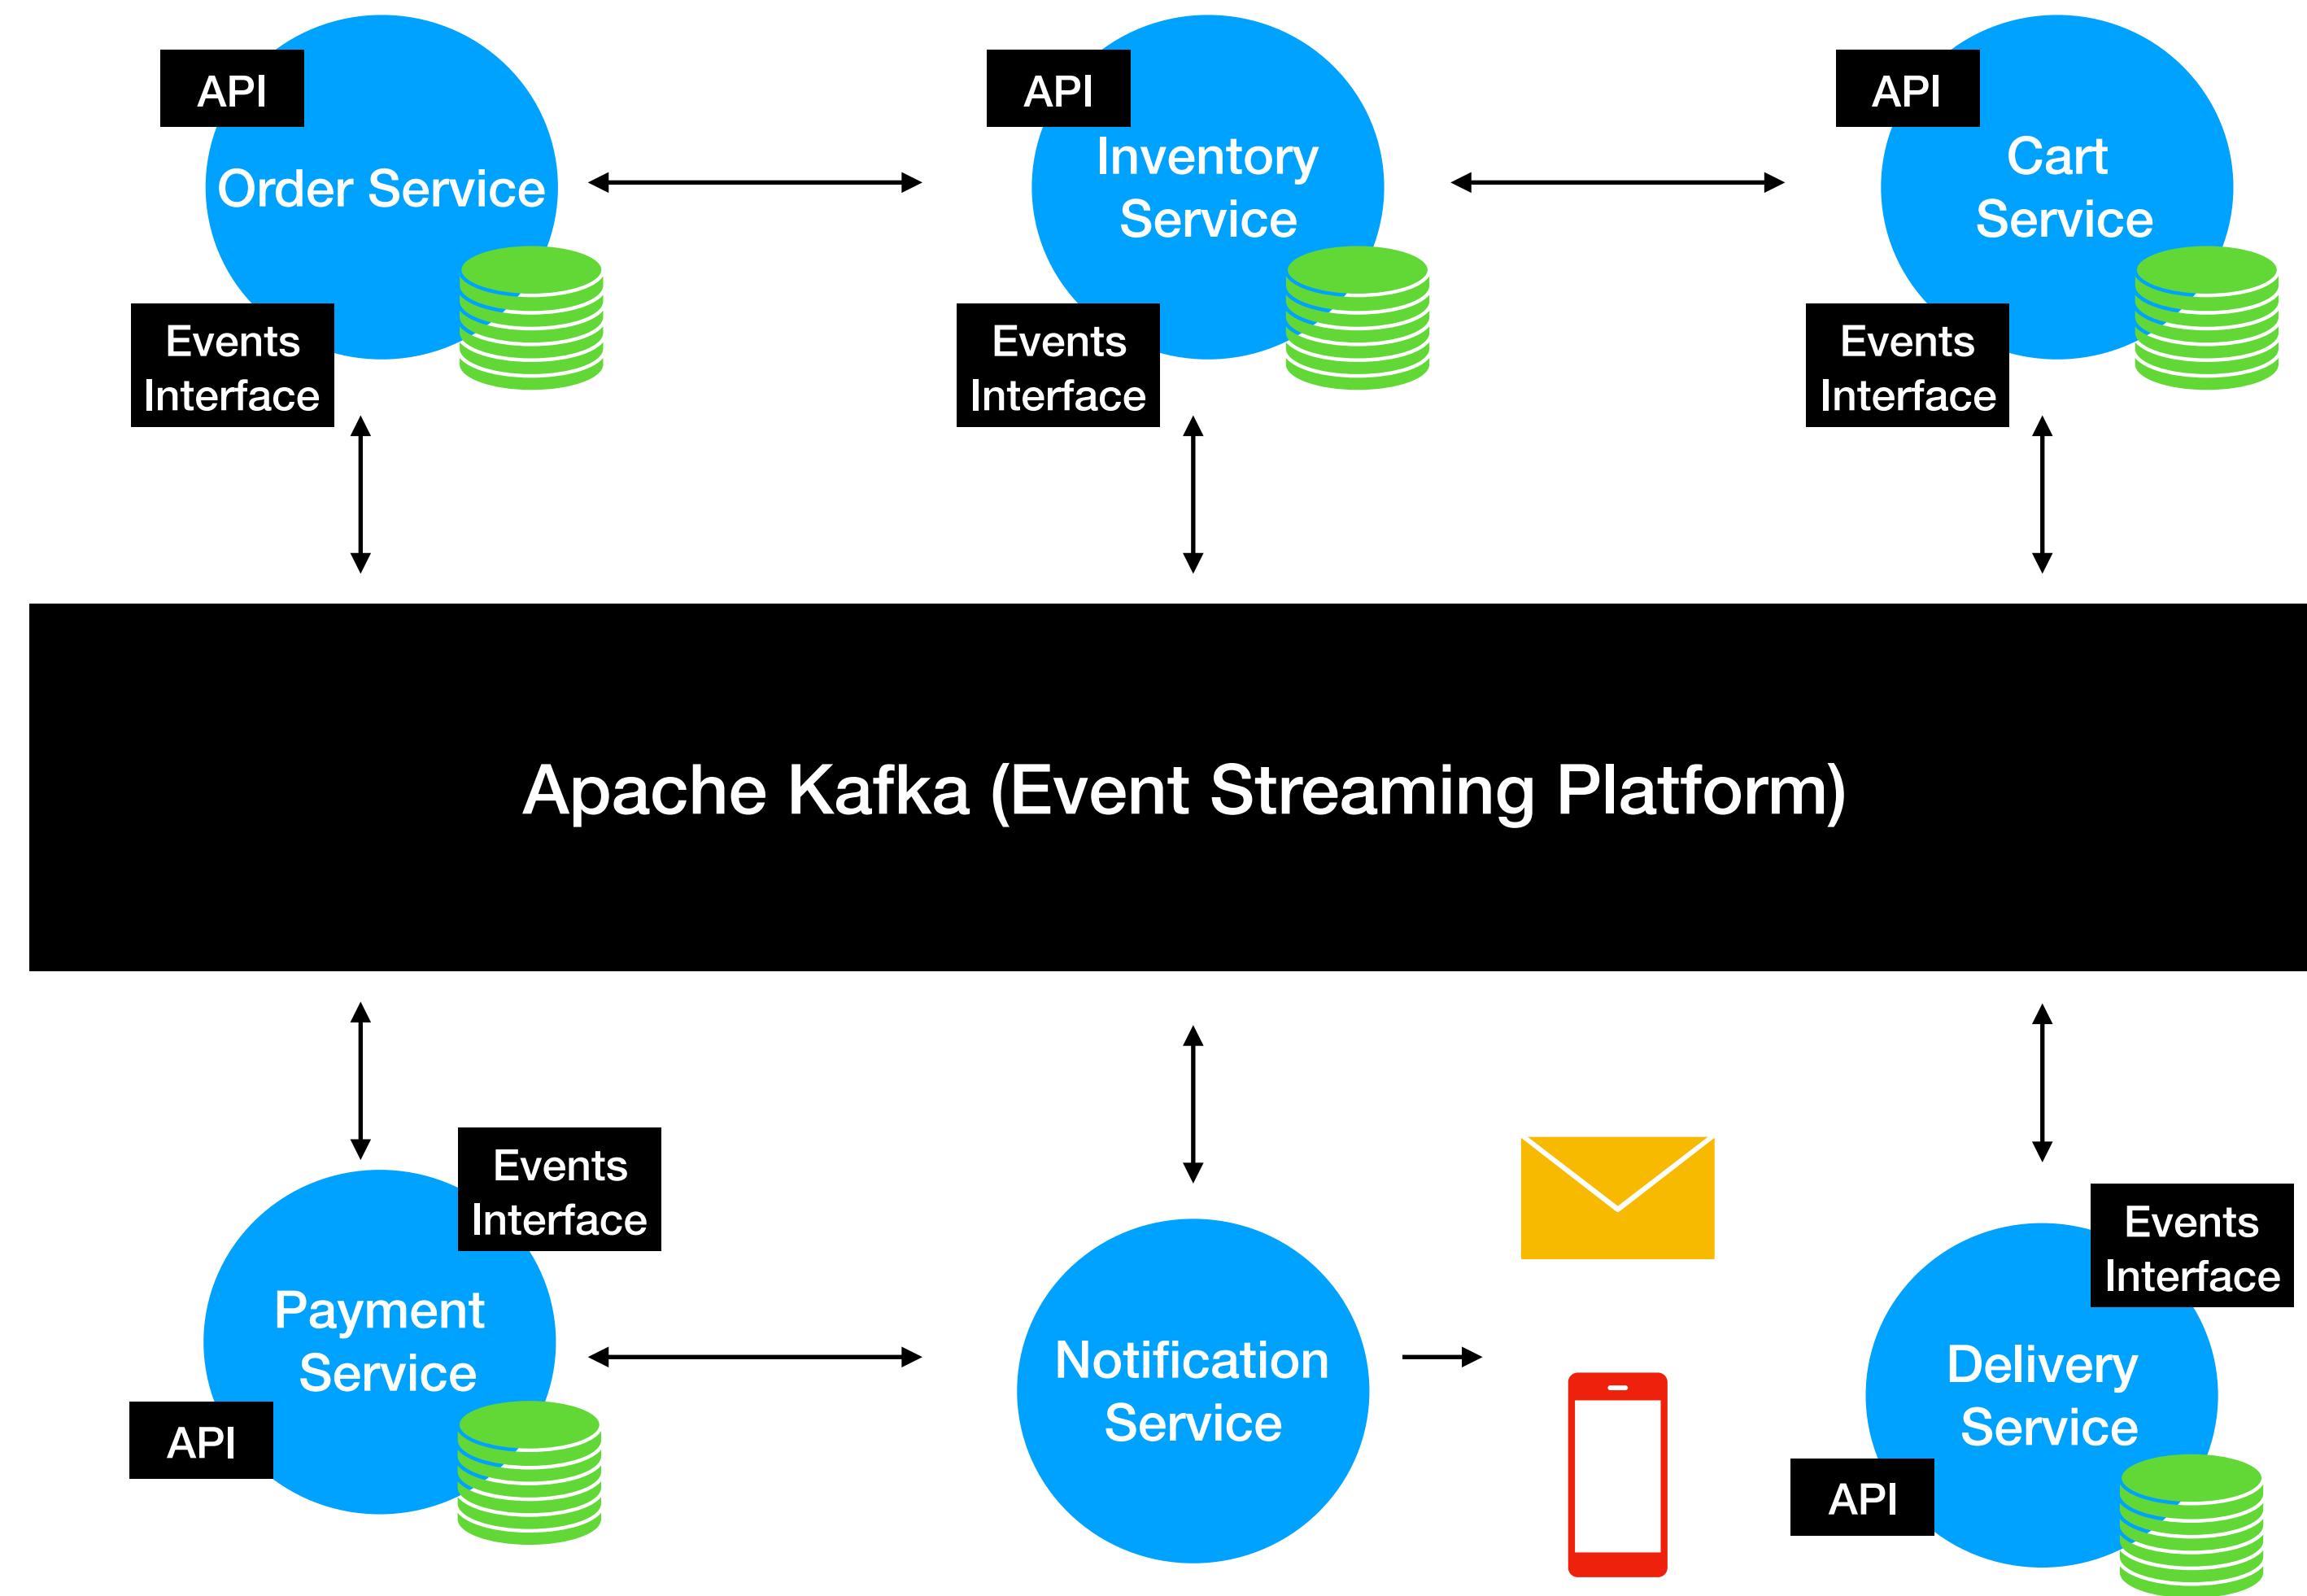


- Store stream of Events



- Analyze and Process Events as they occur

# Apache Kafka (Event Streaming Platform)



# Traditional Messaging System

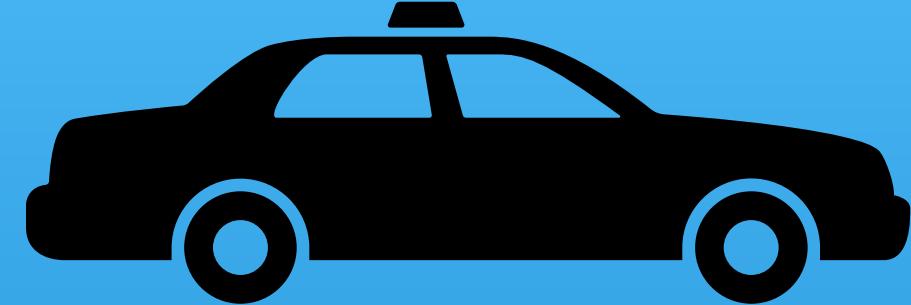
- Transient Message Persistence
- Brokers responsibility to keep track of consumed messages
- Target a specific Consumer
- Not a distributed system

# Kafka Streaming Platform

- Stores events based on a retention time. Events are Immutable
- Consumers Responsibility to keep track of consumed messages
- Any Consumer can access a message from the broker
- It's a distributed streaming system

# Kafka Use Cases

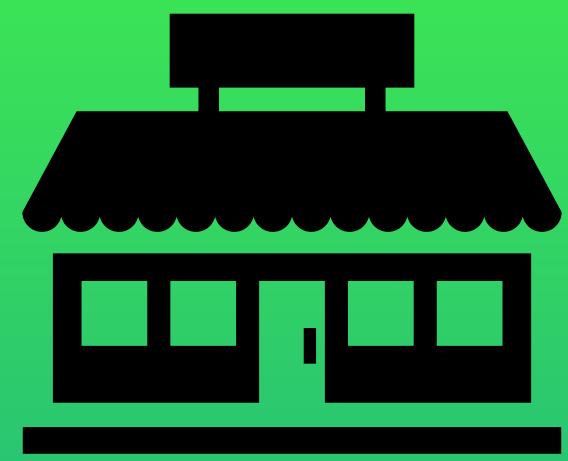
## Transportation



Driver-Rider Notifications

Food Delivery Notifications

## Retail



Sale Notifications

RealTime Purchase  
recommendations

Tracking Online Order  
Deliveries

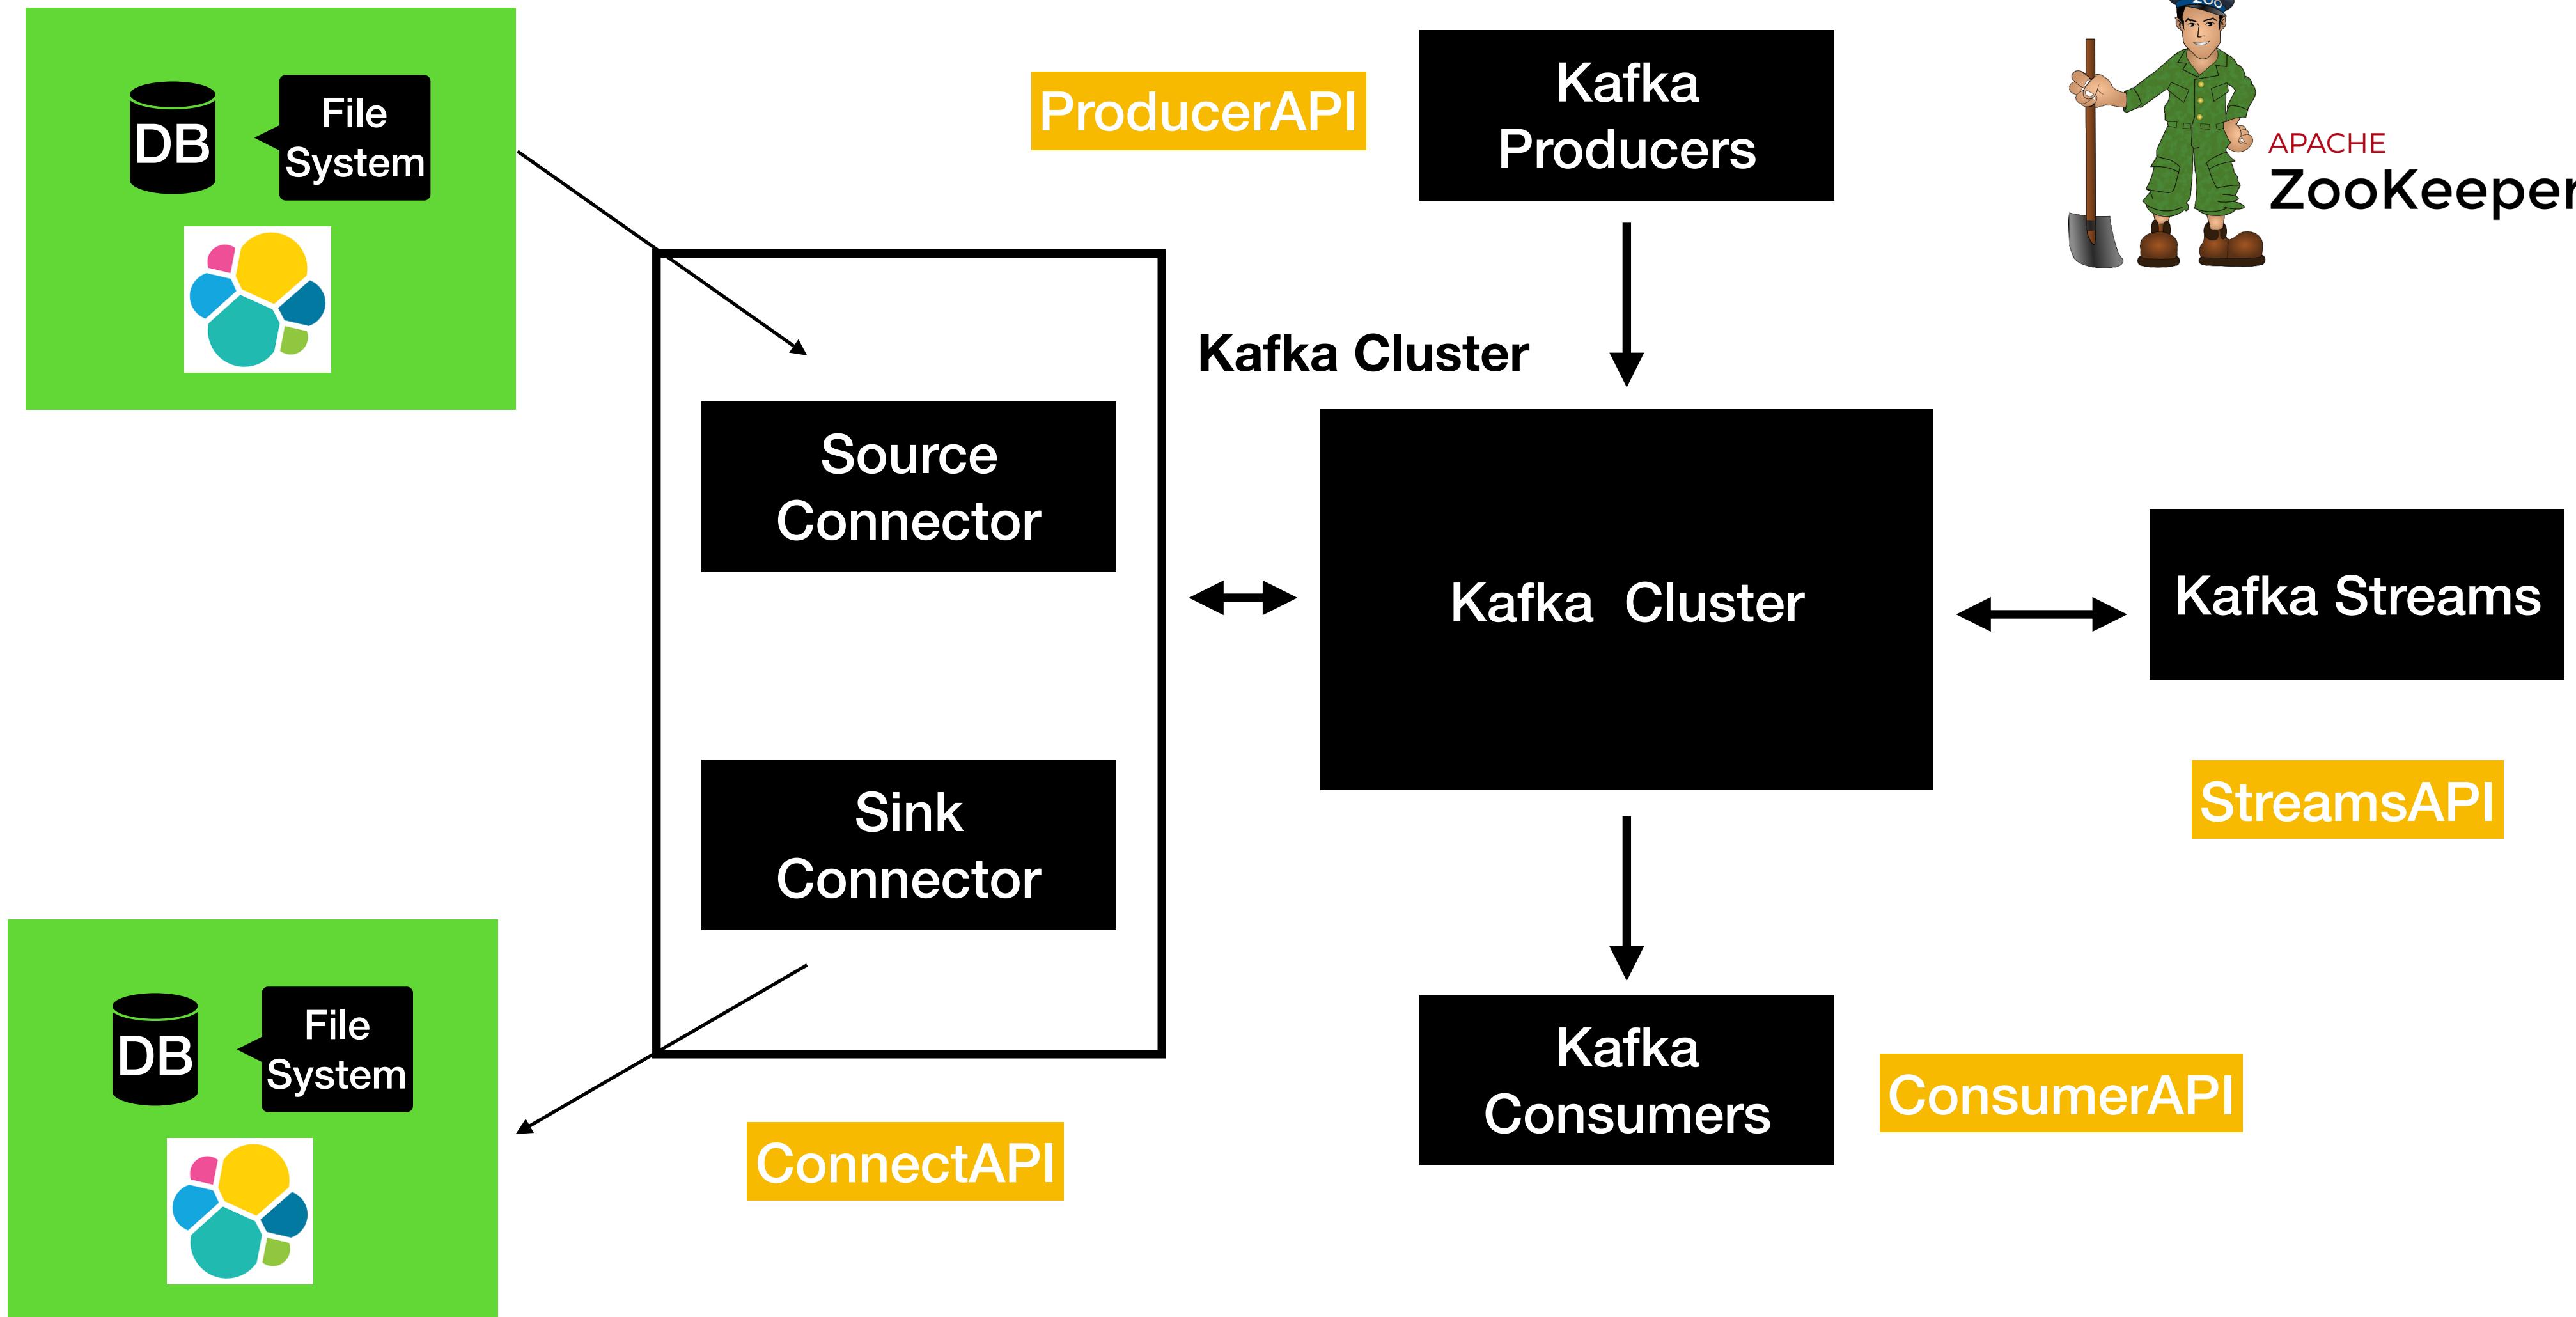
## Banking



Fraud Transactions

New Feature/Product  
notifications

# Kafka Terminology & Client APIs



# **Download Kafka**

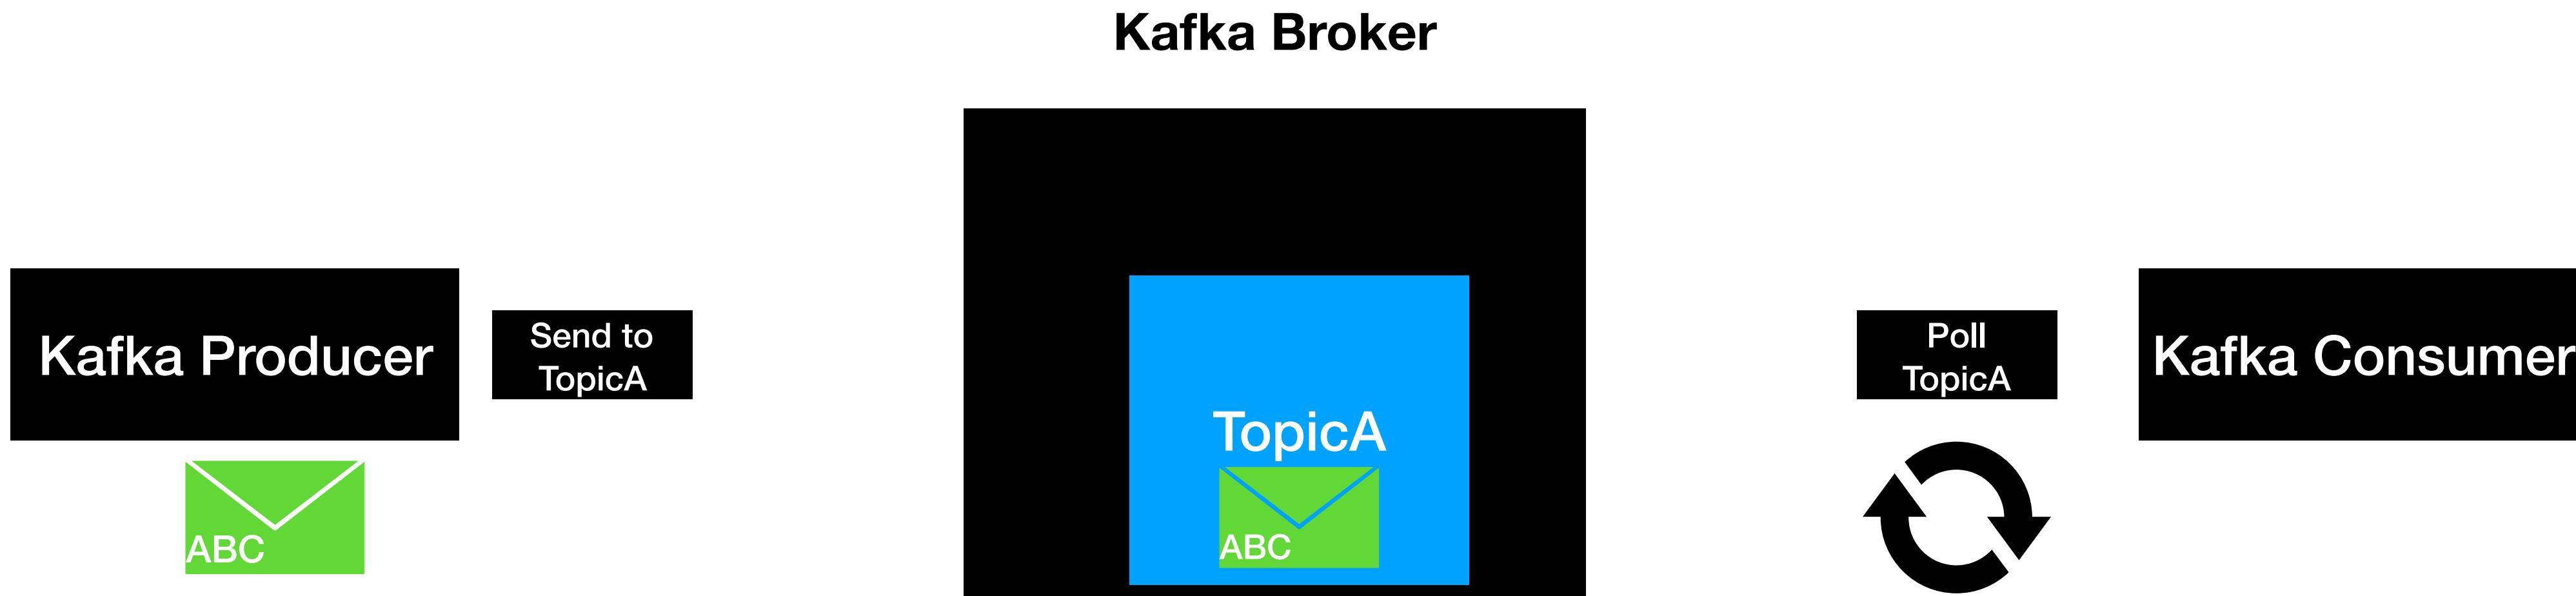
# Kafka Topics & Partitions

# Kafka Topics

- Topic is an **Entity** in Kafka with a name

# Kafka Topics

- Topic is an **Entity** in Kafka with a name

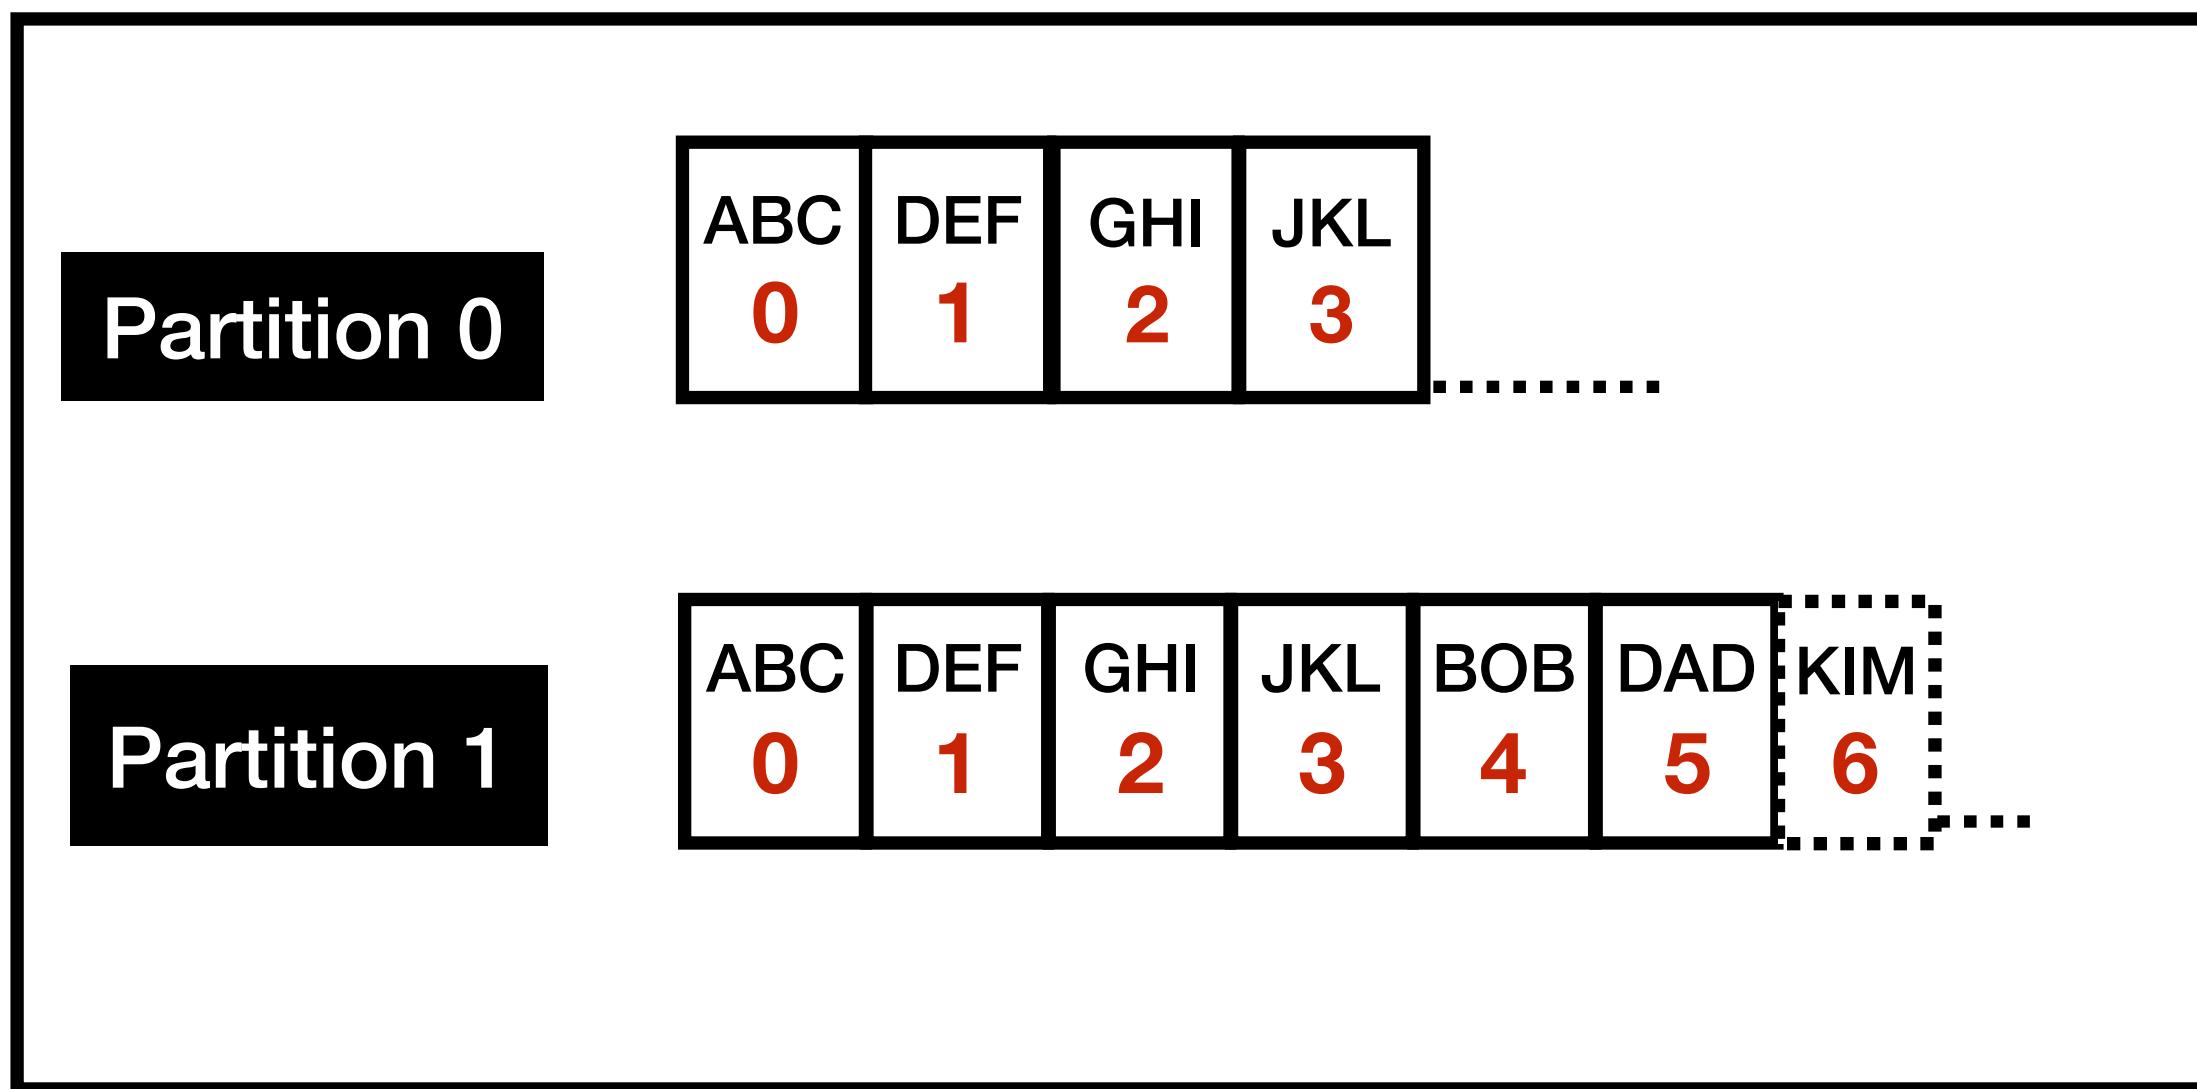


# Topic and Partitions

- Partition is where the message lives inside the topic
- Each Topic will be create with one or more partitions

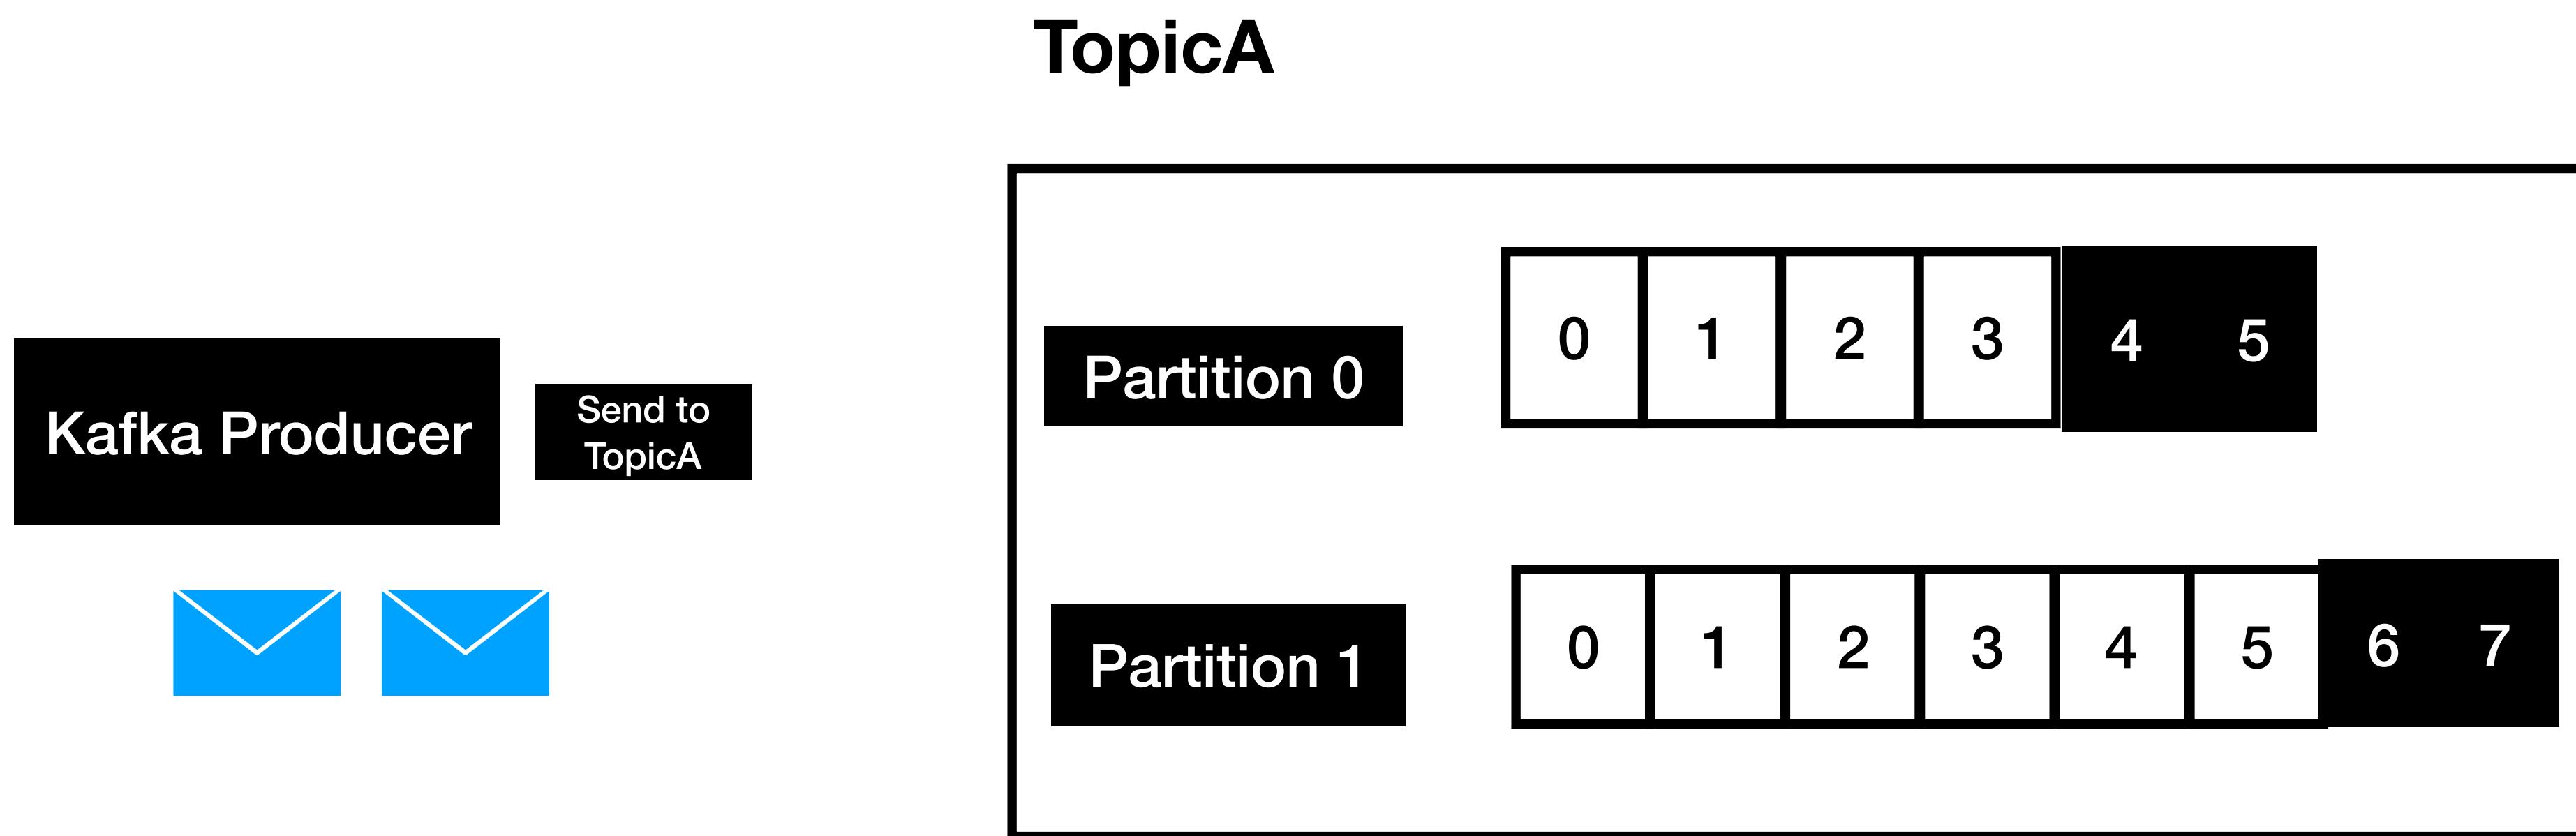
# Topic and Partitions

TopicA



- Each Partition is an ordered , immutable sequence of records
- Each record is assigned a sequential number called **offset**
- Each partition is independent of each other
- Ordering is guaranteed only at the partition level
- Partition continuously grows as new records are produced
- All the records are persisted in a commit log in the file system where Kafka is installed

# Topics and Partitions

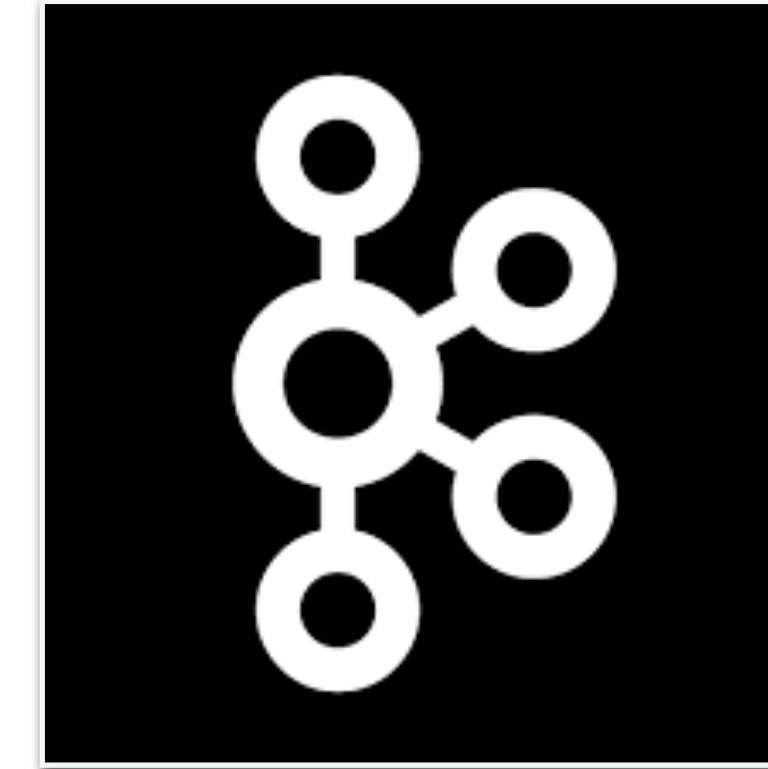
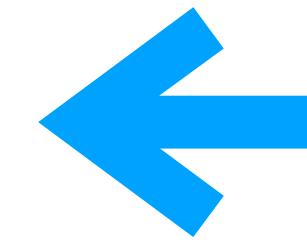


**Setting up  
Zookeeper  
&  
Kafka Broker**

# Setting up Kafka in Local



APACHE  
**ZooKeeper**™



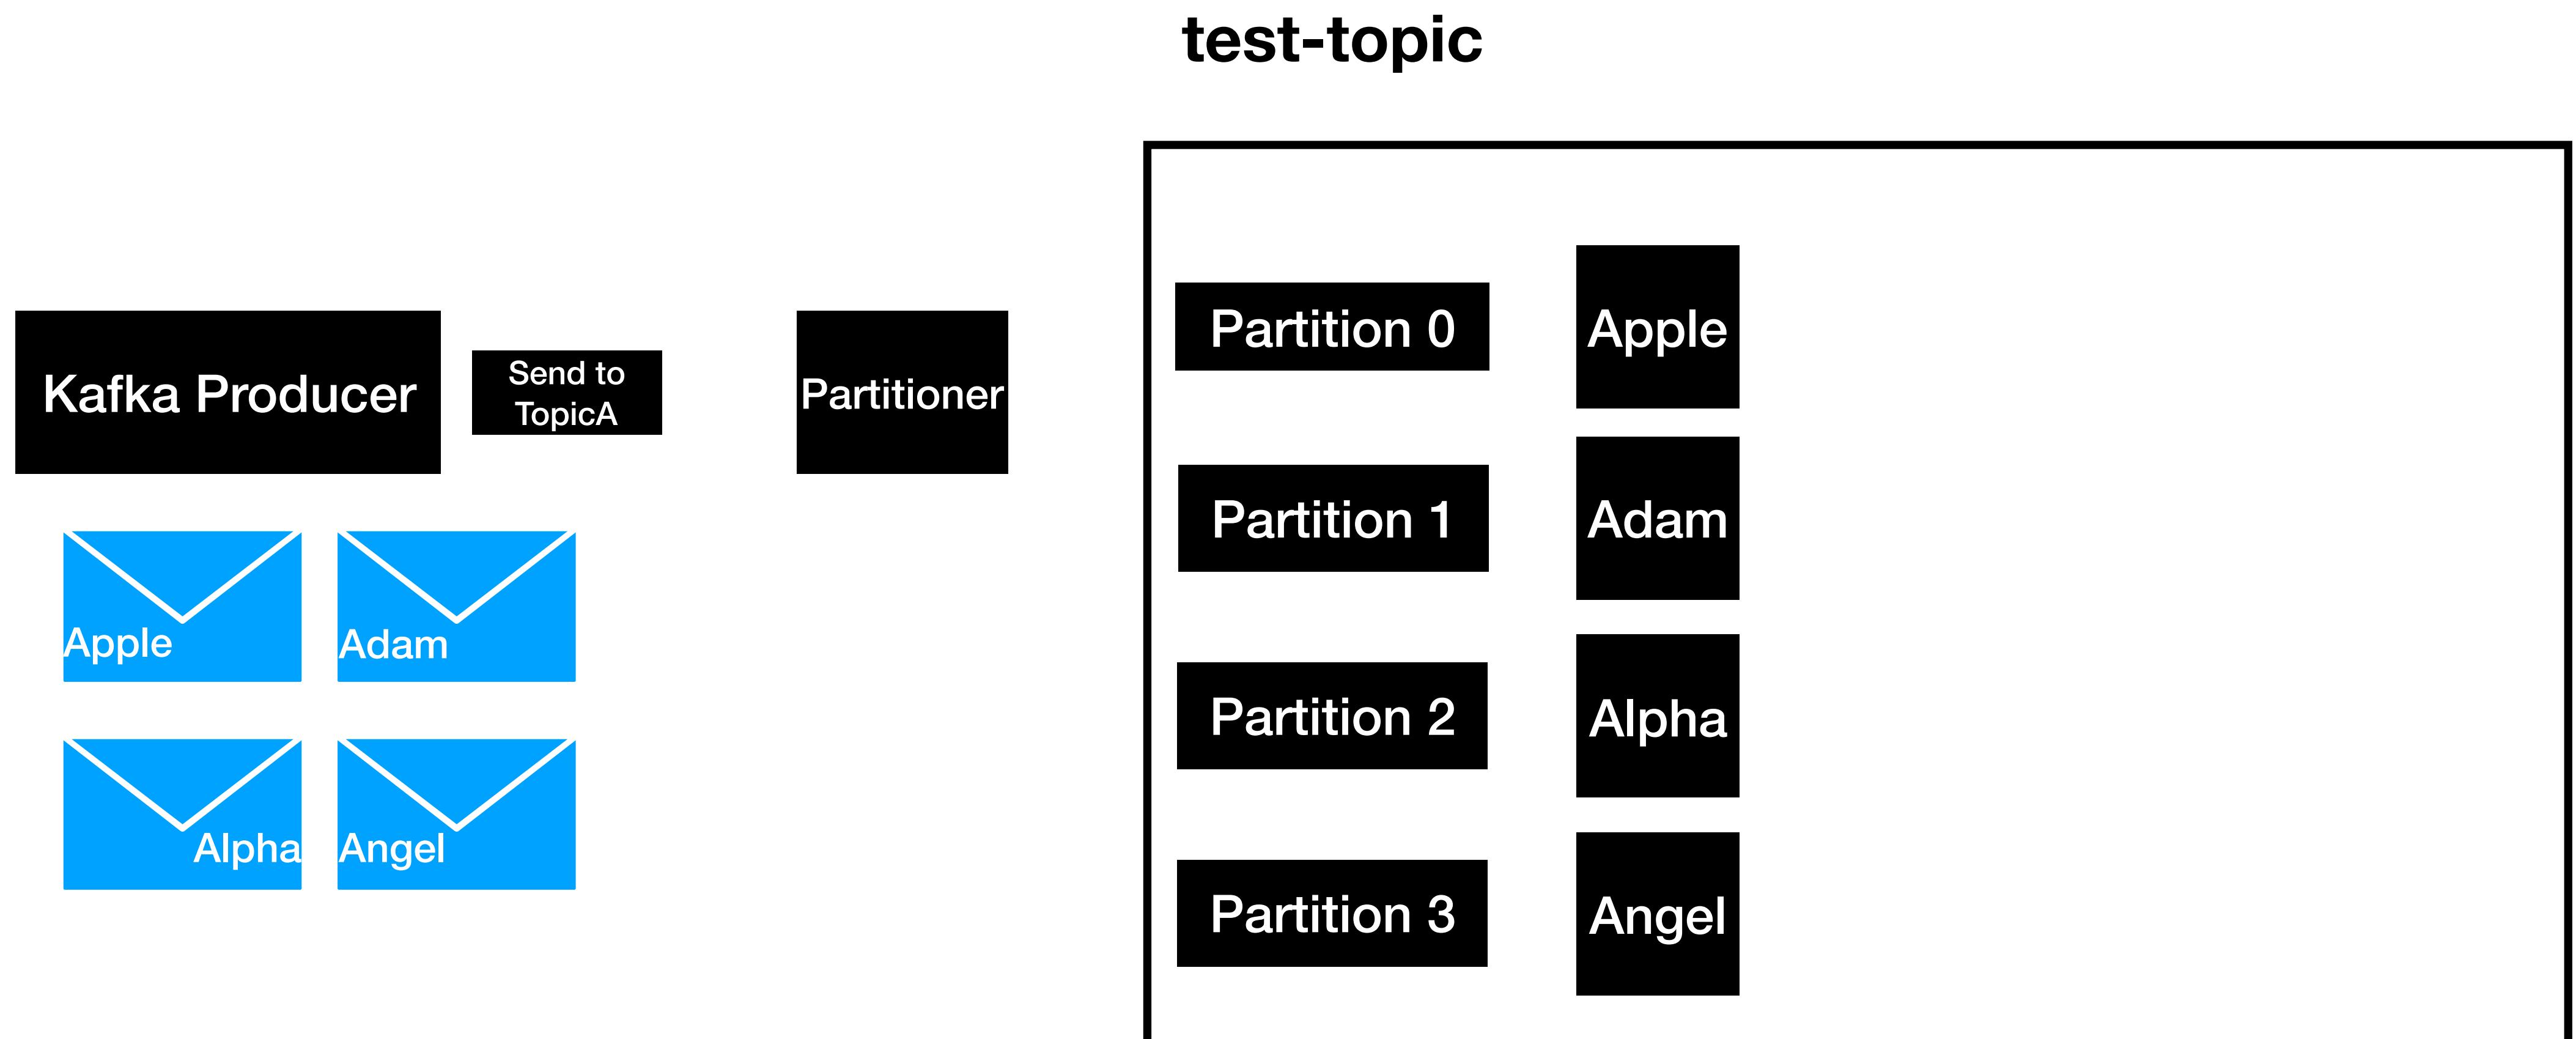
Broker registered  
with zookeeper

Sending  
Kafka Messages  
With  
Key and Value

# Kafka Message

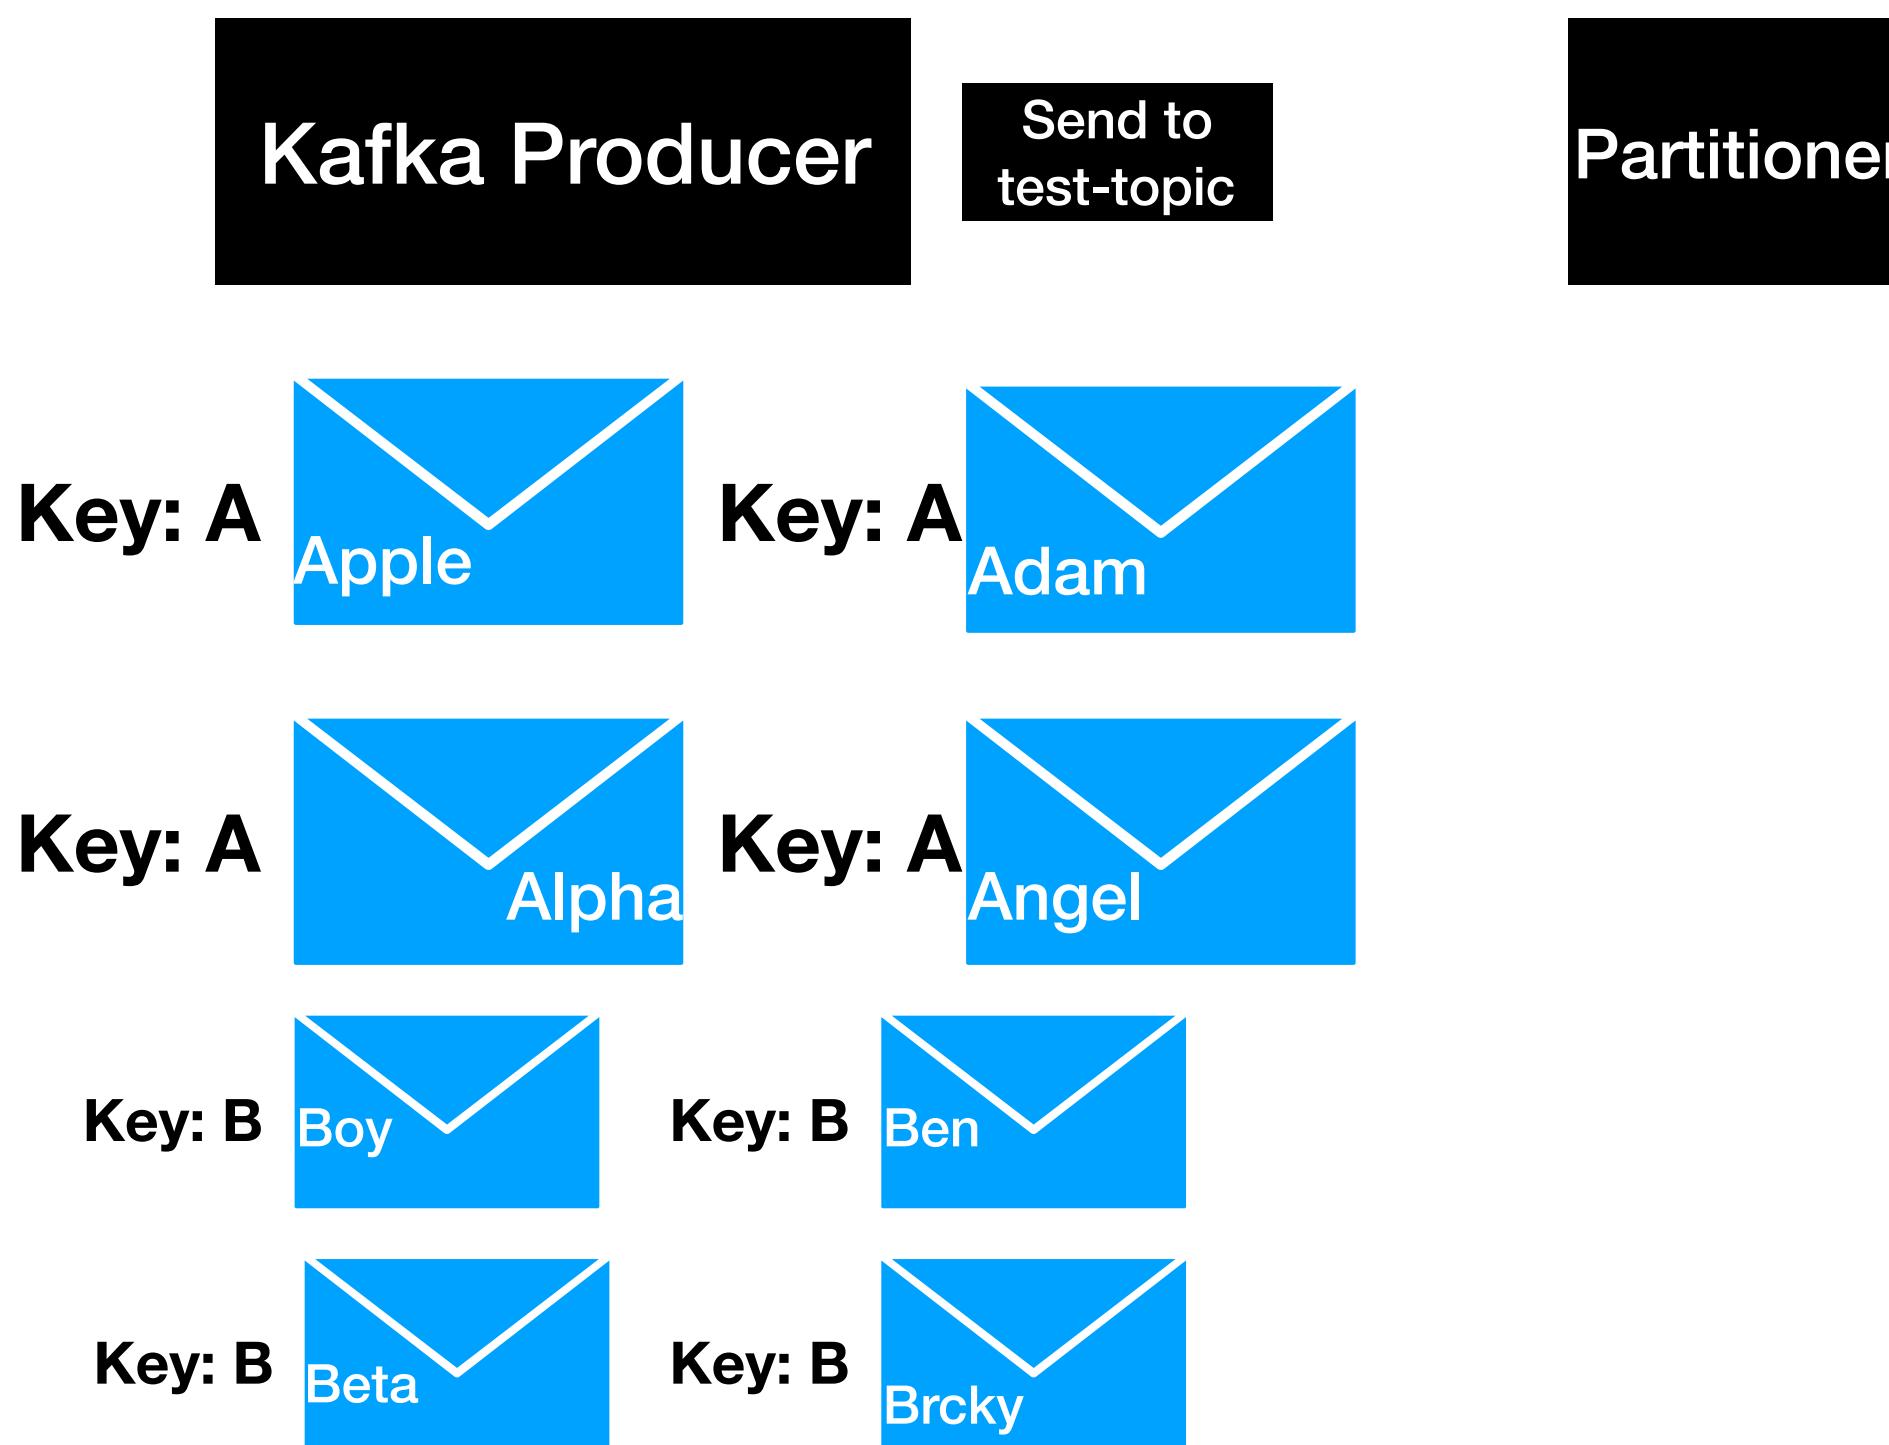
- Kafka Message these sent from producer has two properties
  - Key (optional)
  - Value

# Sending Message Without Key



# Sending Message With Key

**test-topic**



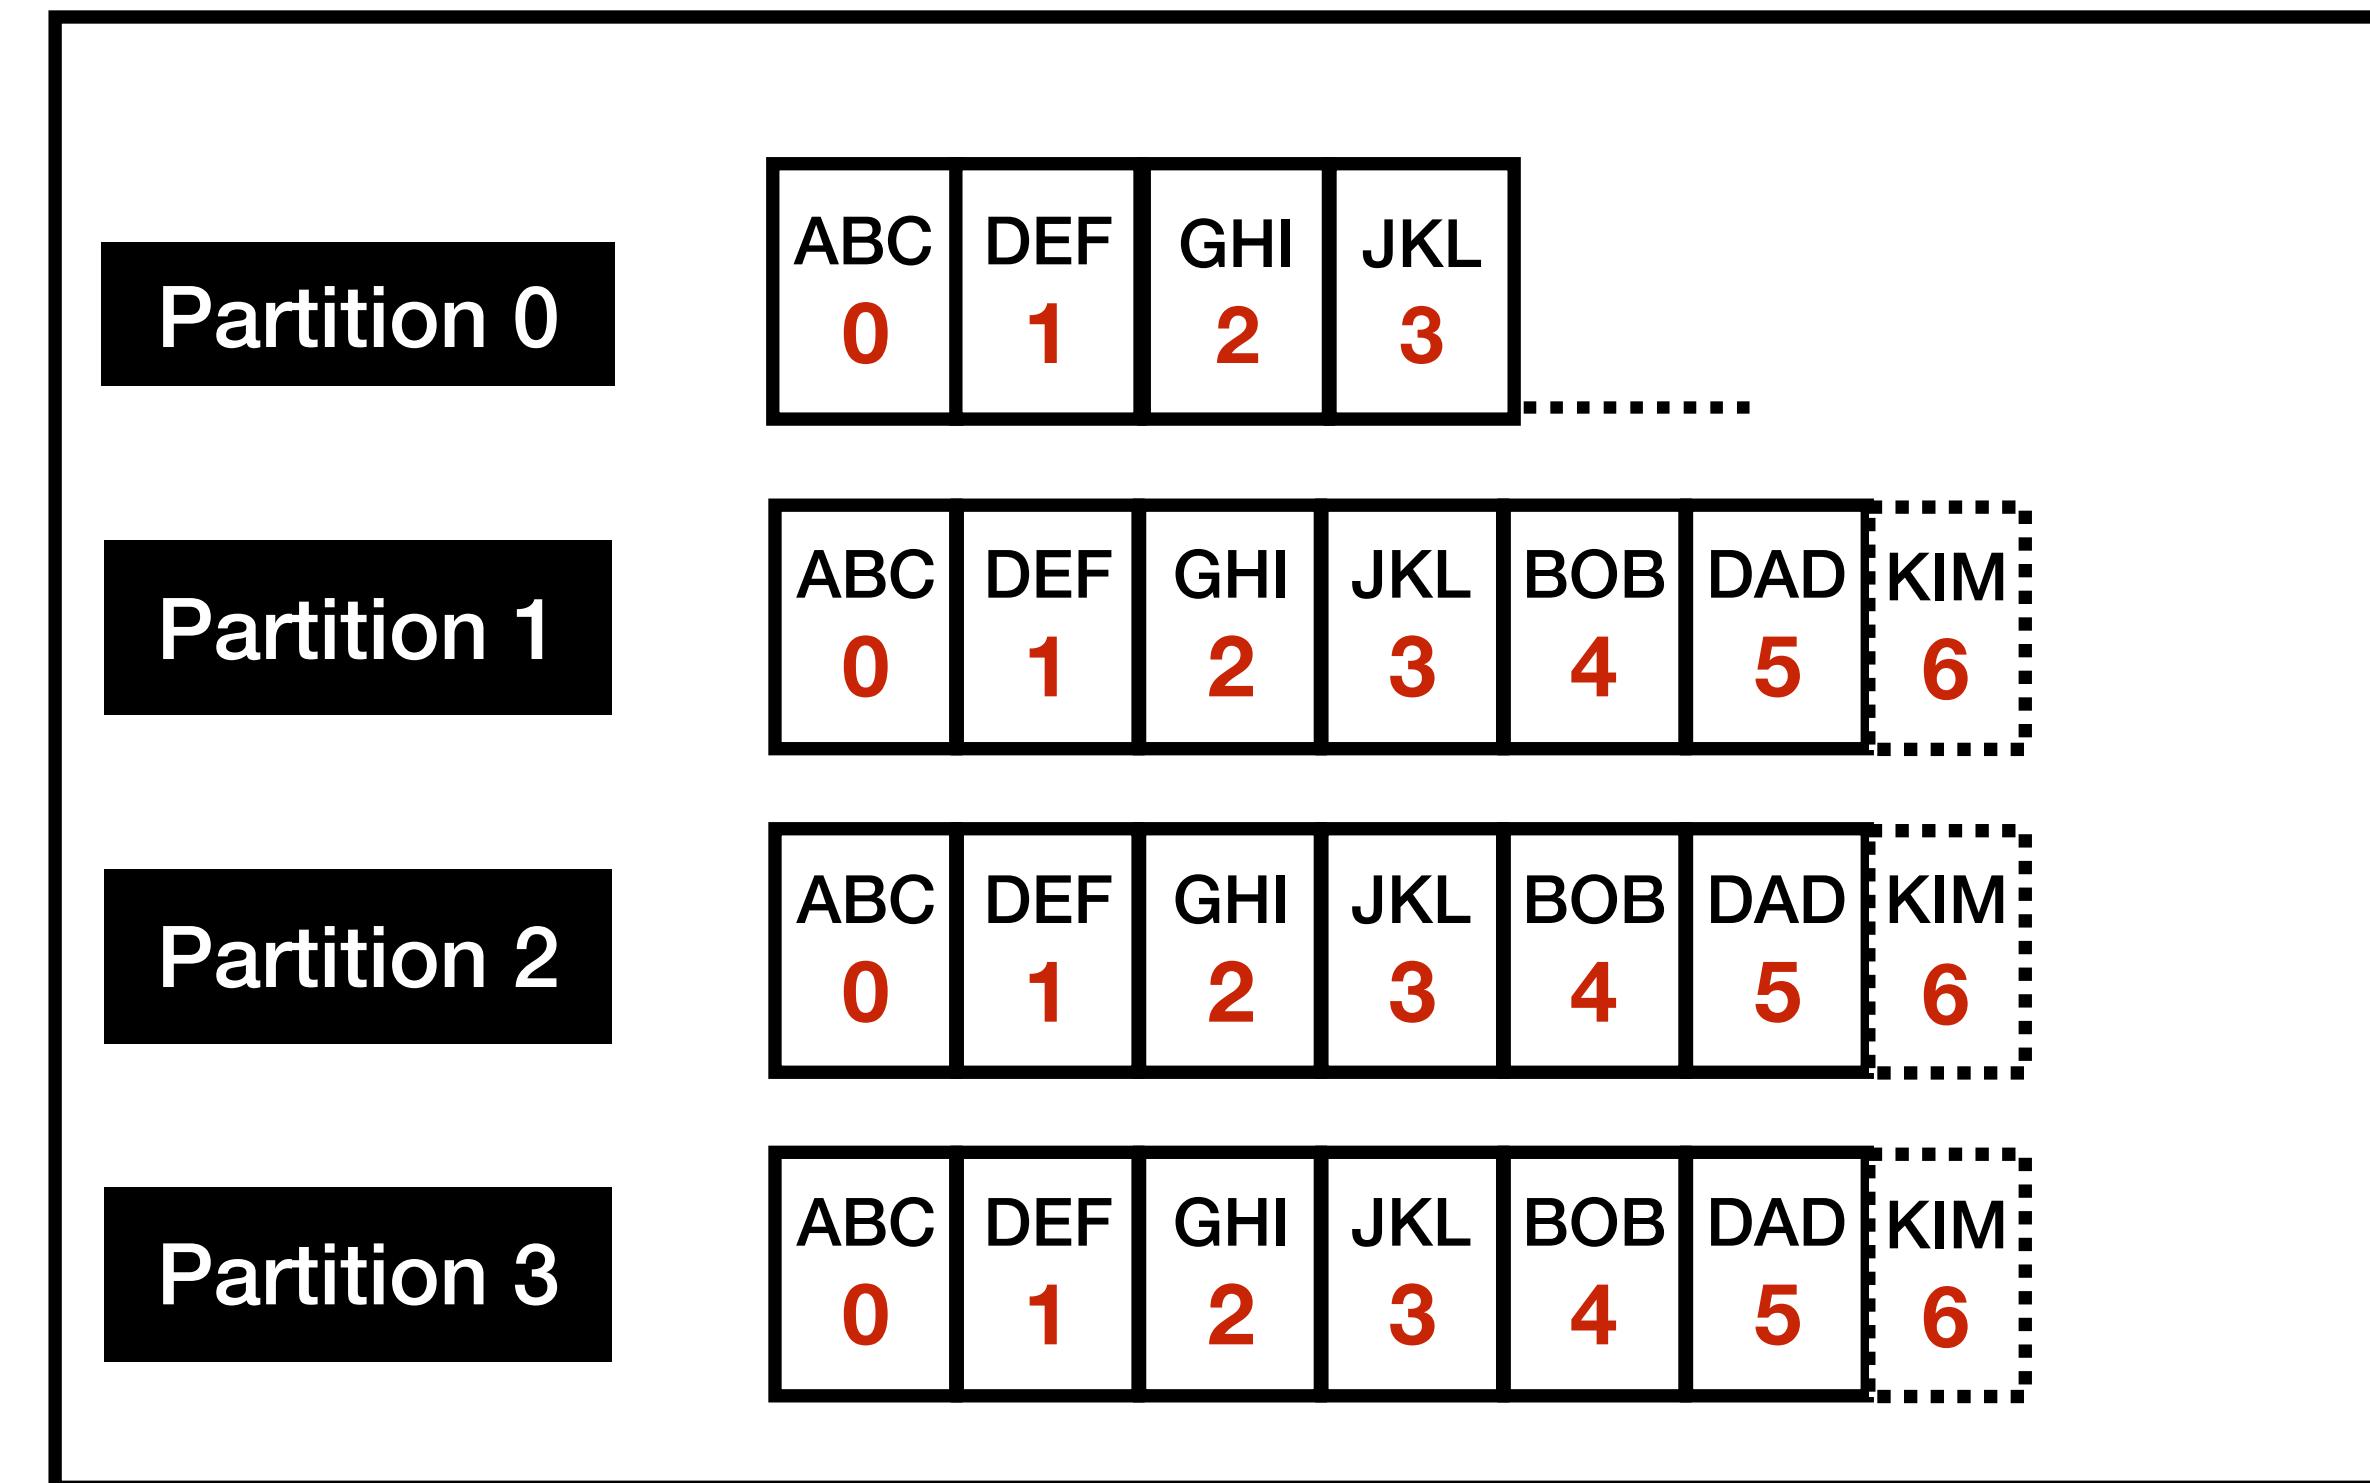
# Consumer Offsets

# Consumer Offsets

- Consumer have three options to read
  - from-beginning
  - latest
  - specific offset

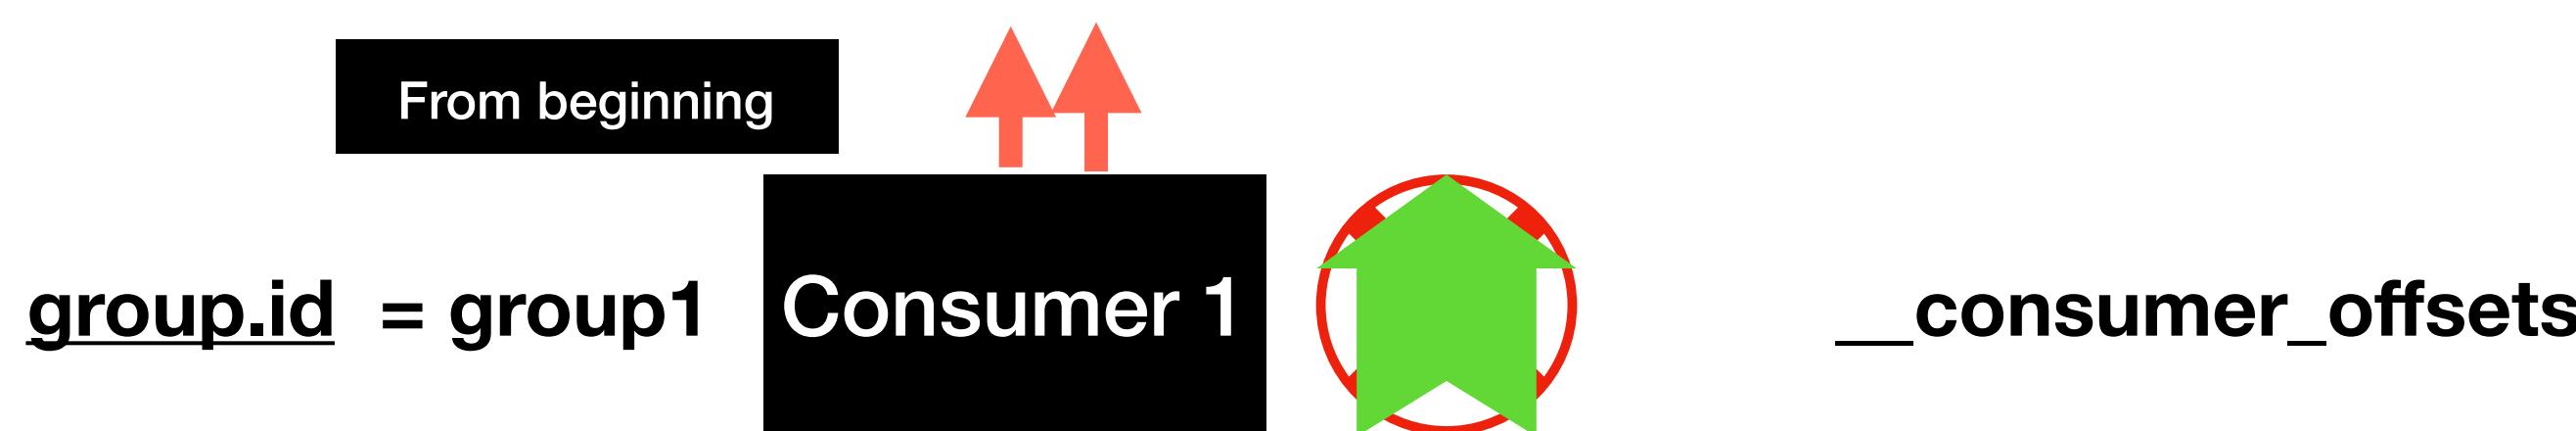
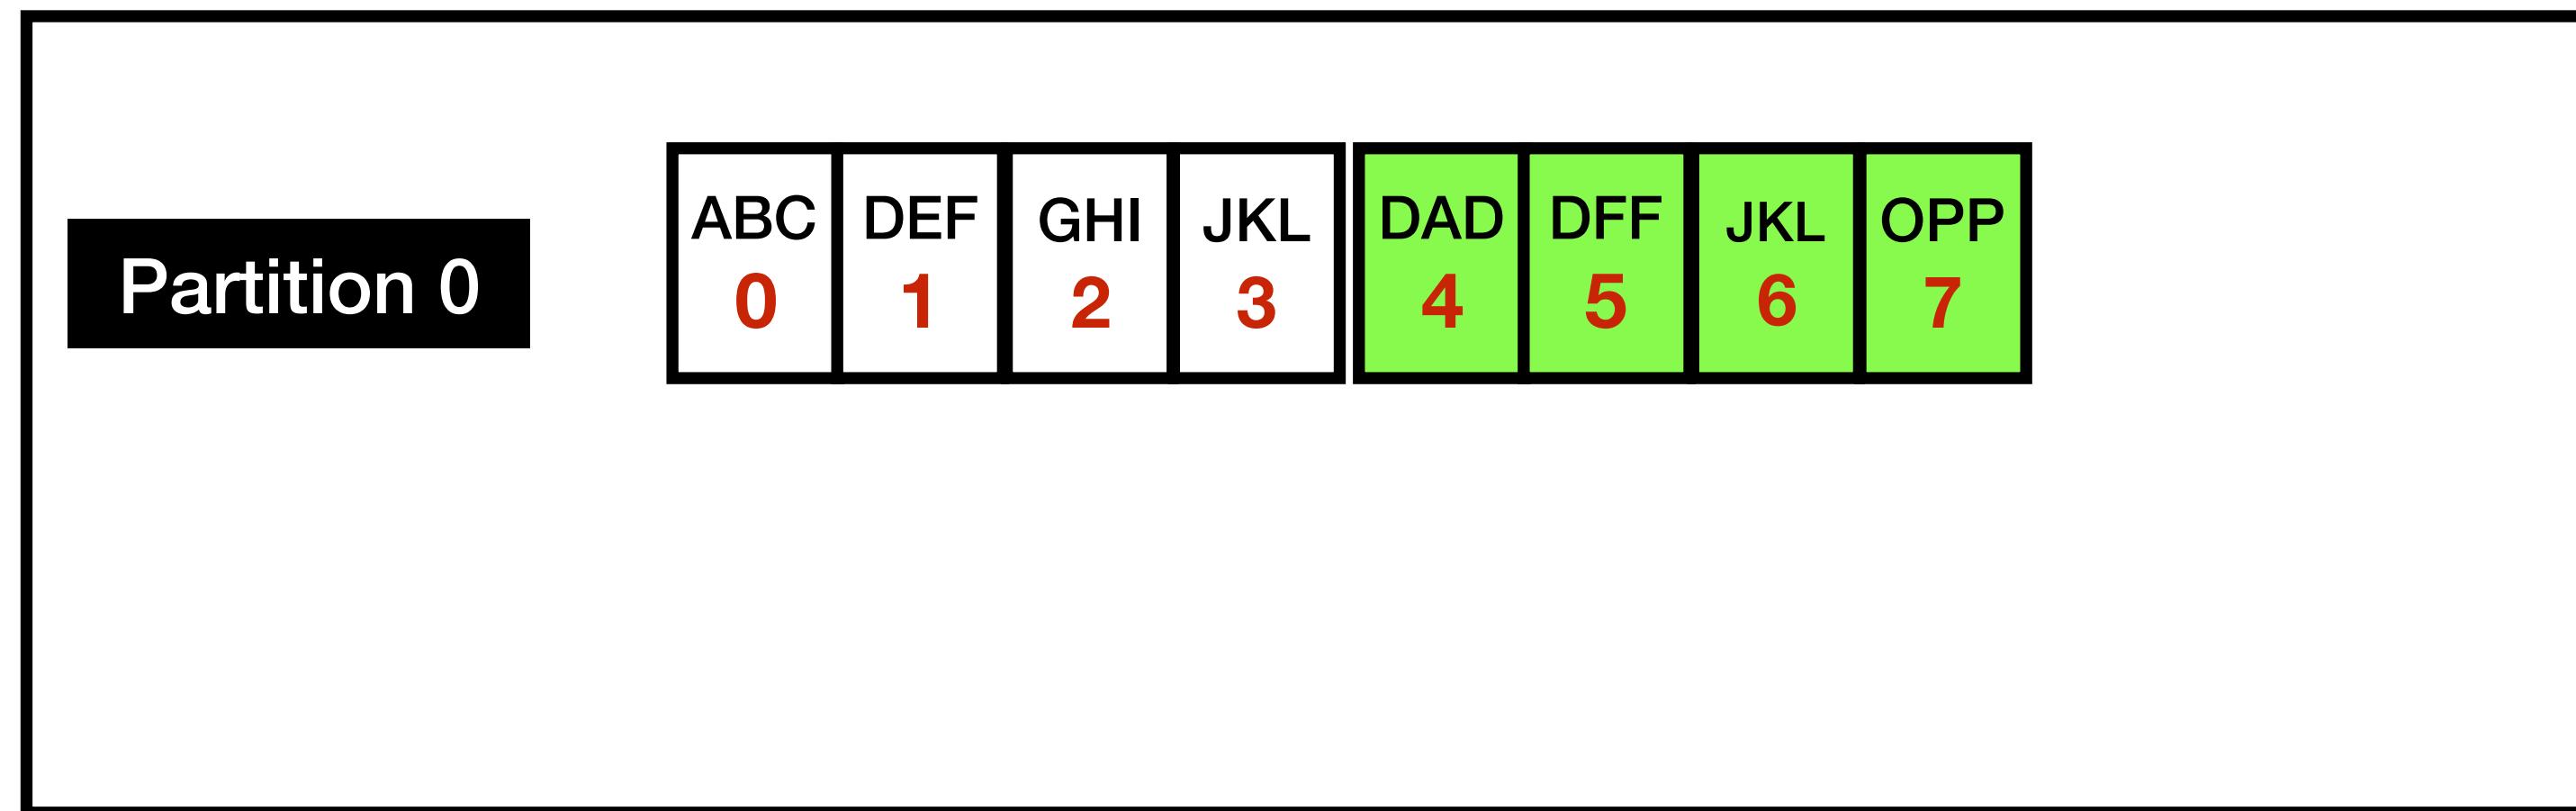
This option will be done through only programmatic way

**test-topic**



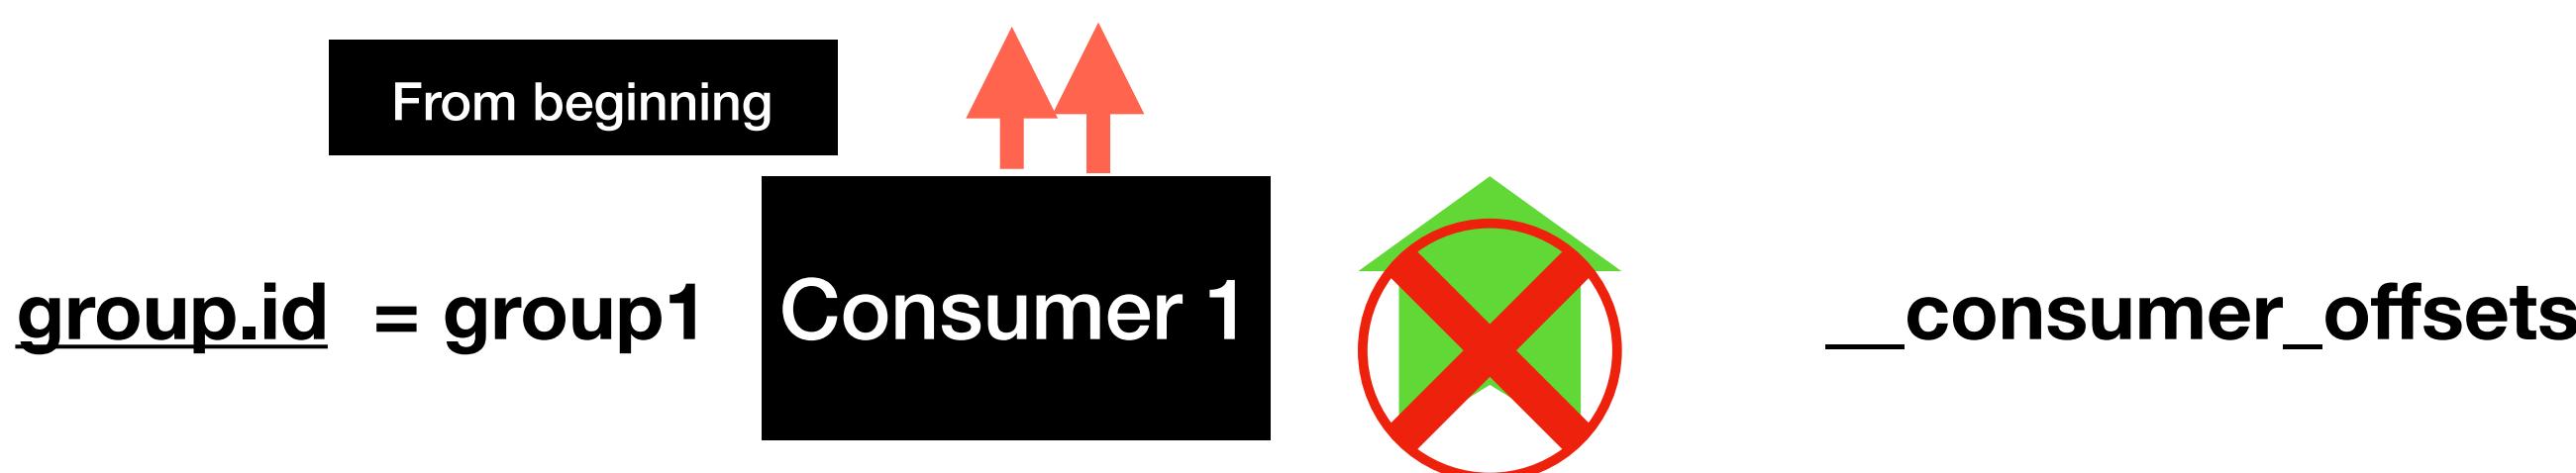
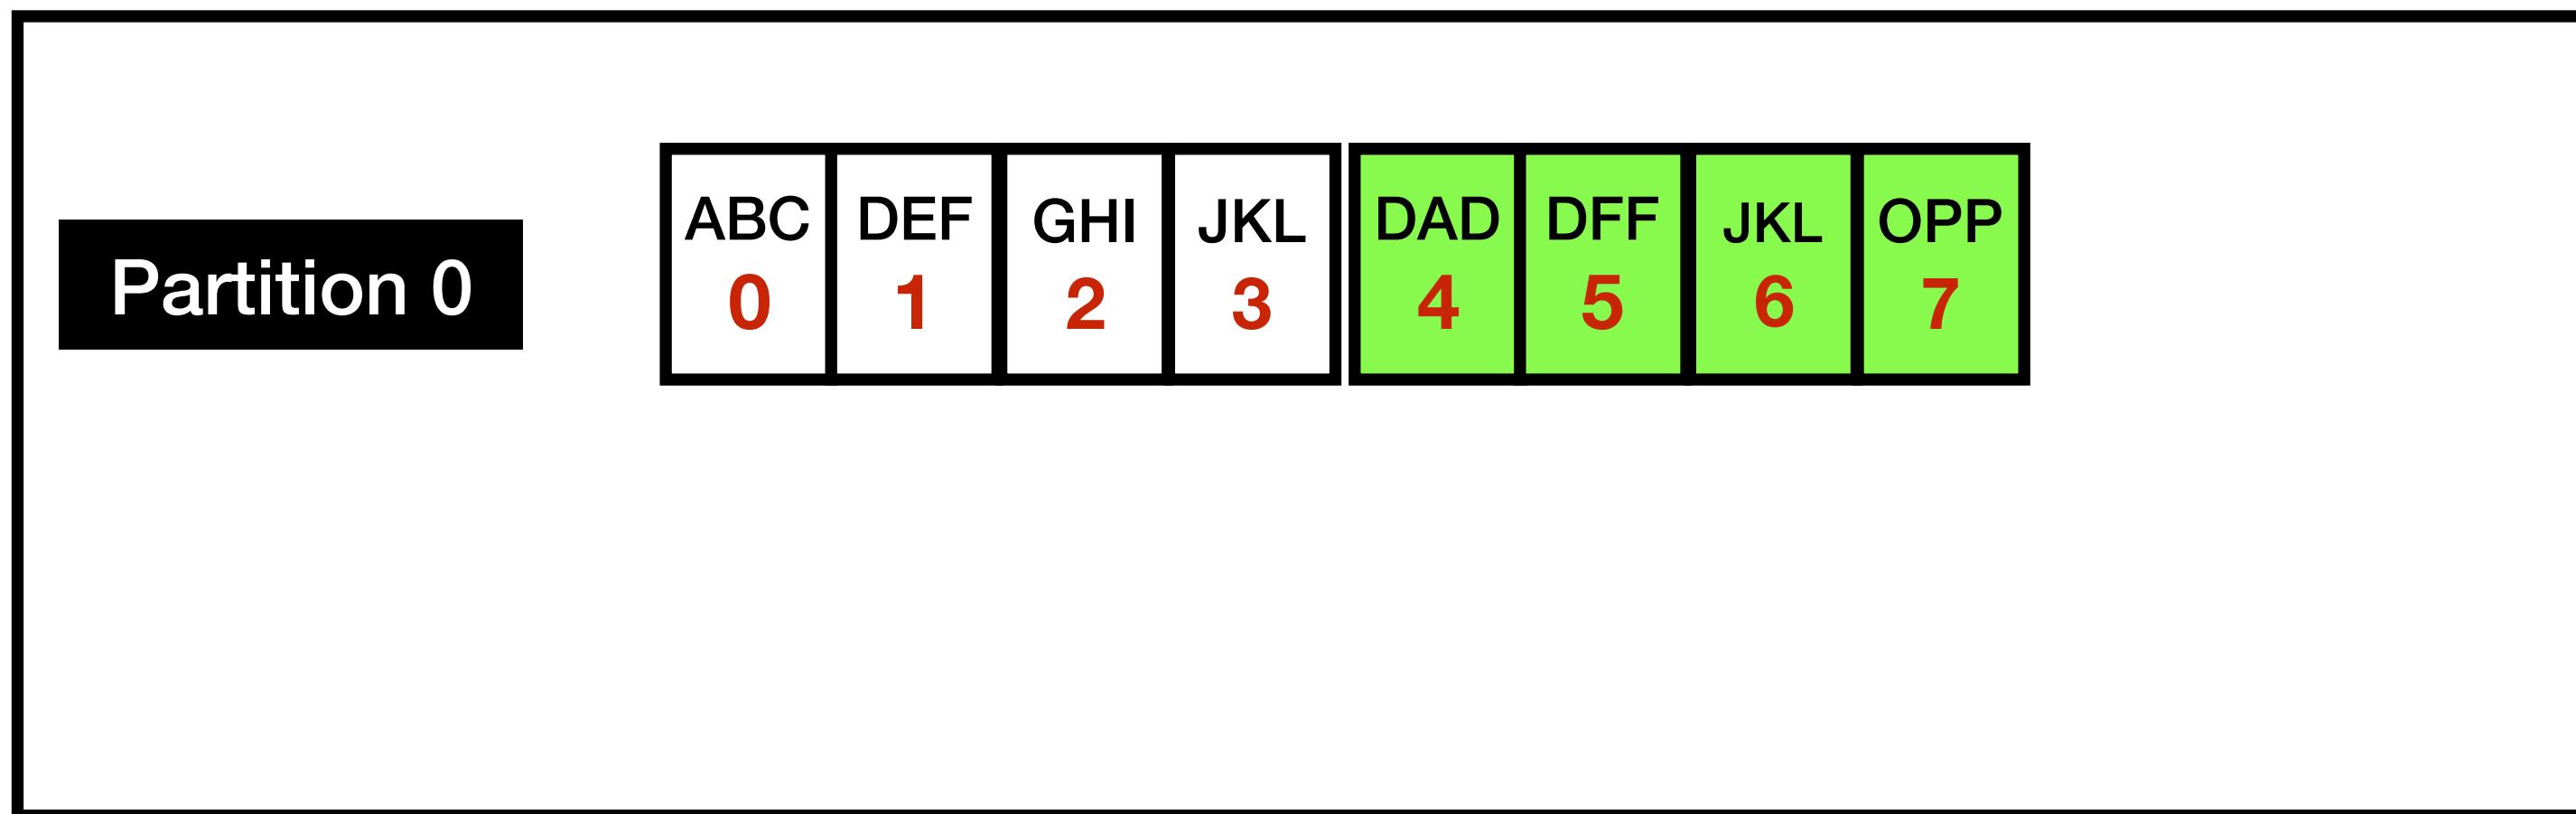
# Consumer Offsets

test-topic



# Consumer Offsets

test-topic



# Consumer Offset

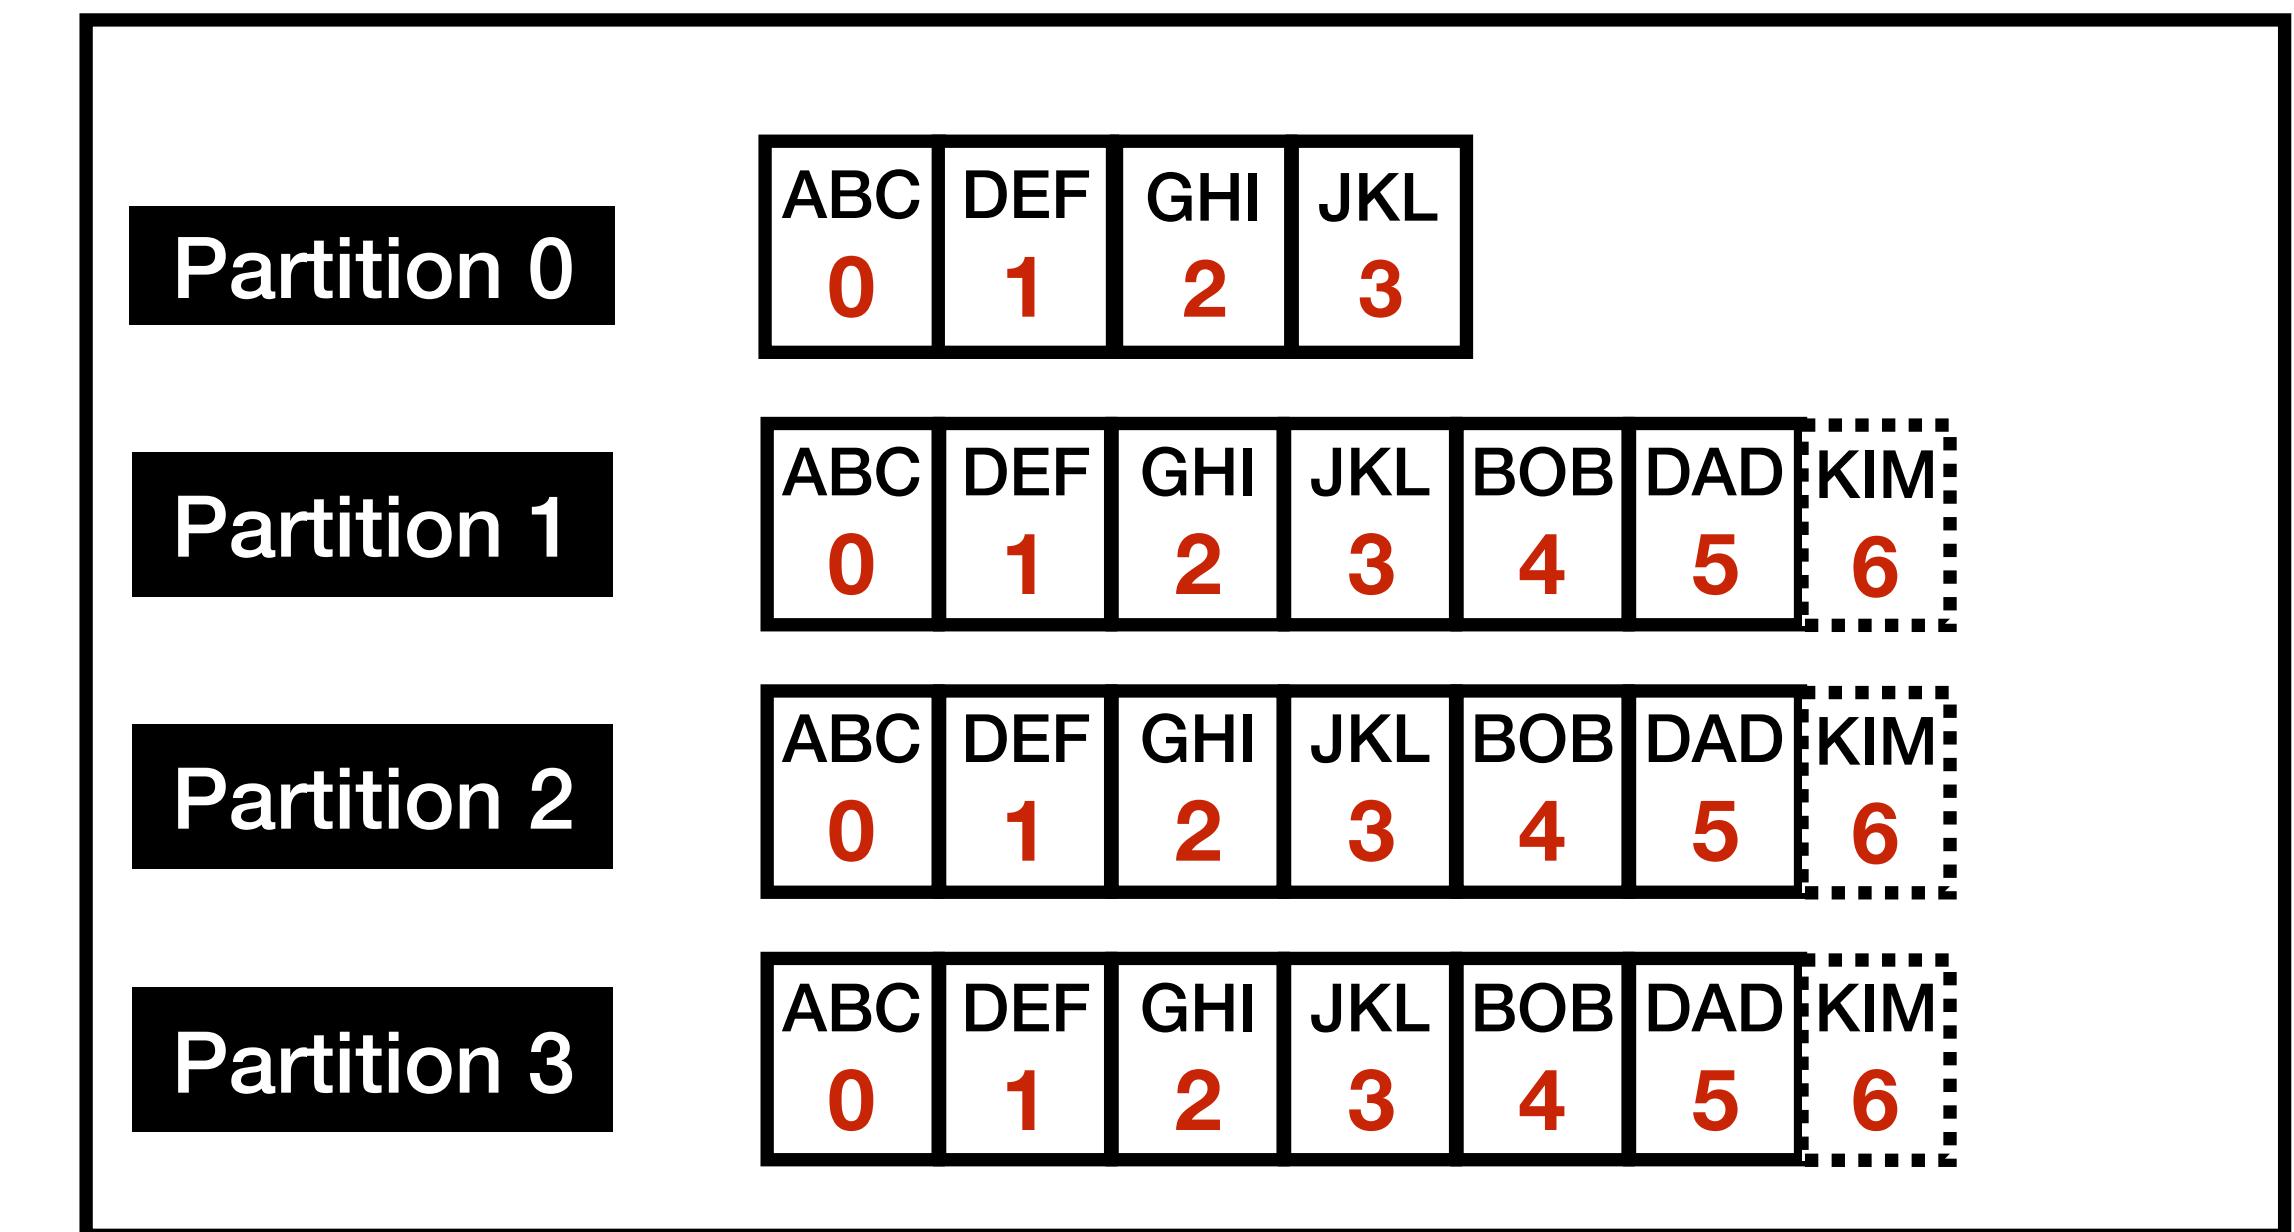
- Consumer offsets behaves like a bookmark for the consumer to start reading the messages from the point it left off.

# **Consumer Groups**

# Consumer Groups

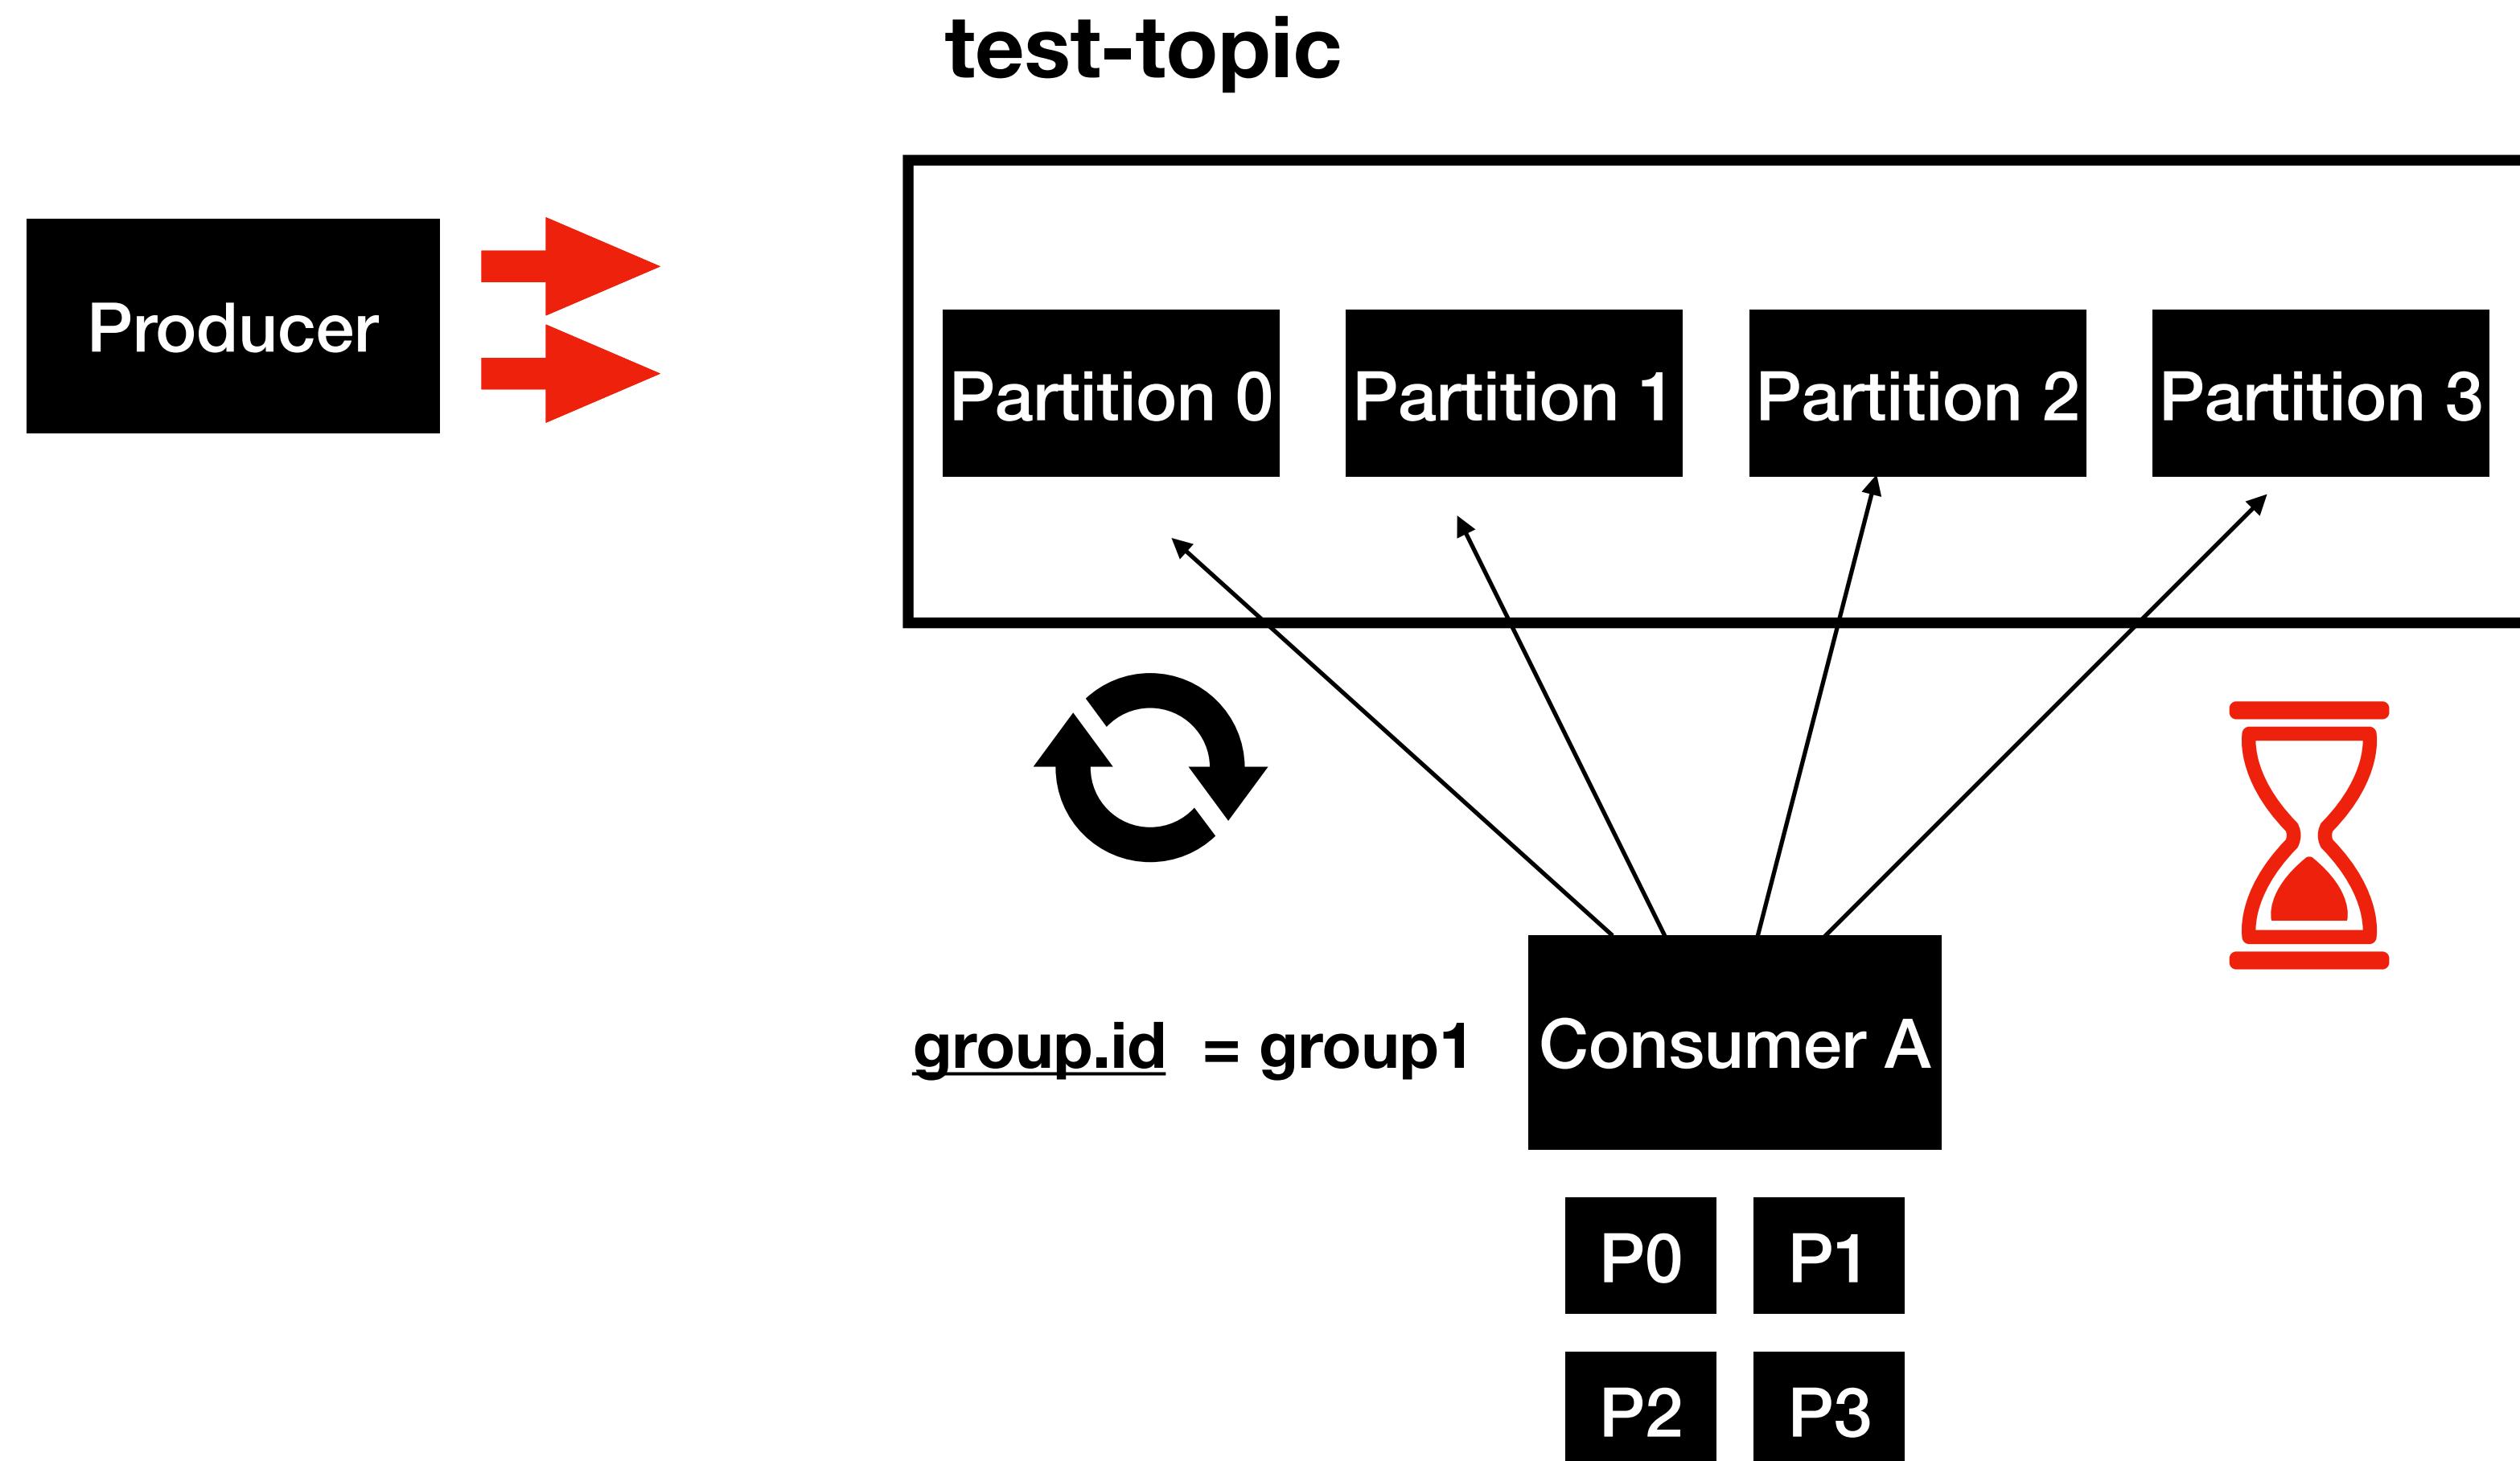
- **group.id** is mandatory
- **group.id** plays a major role when it comes to scalable message consumption.

**test-topic**

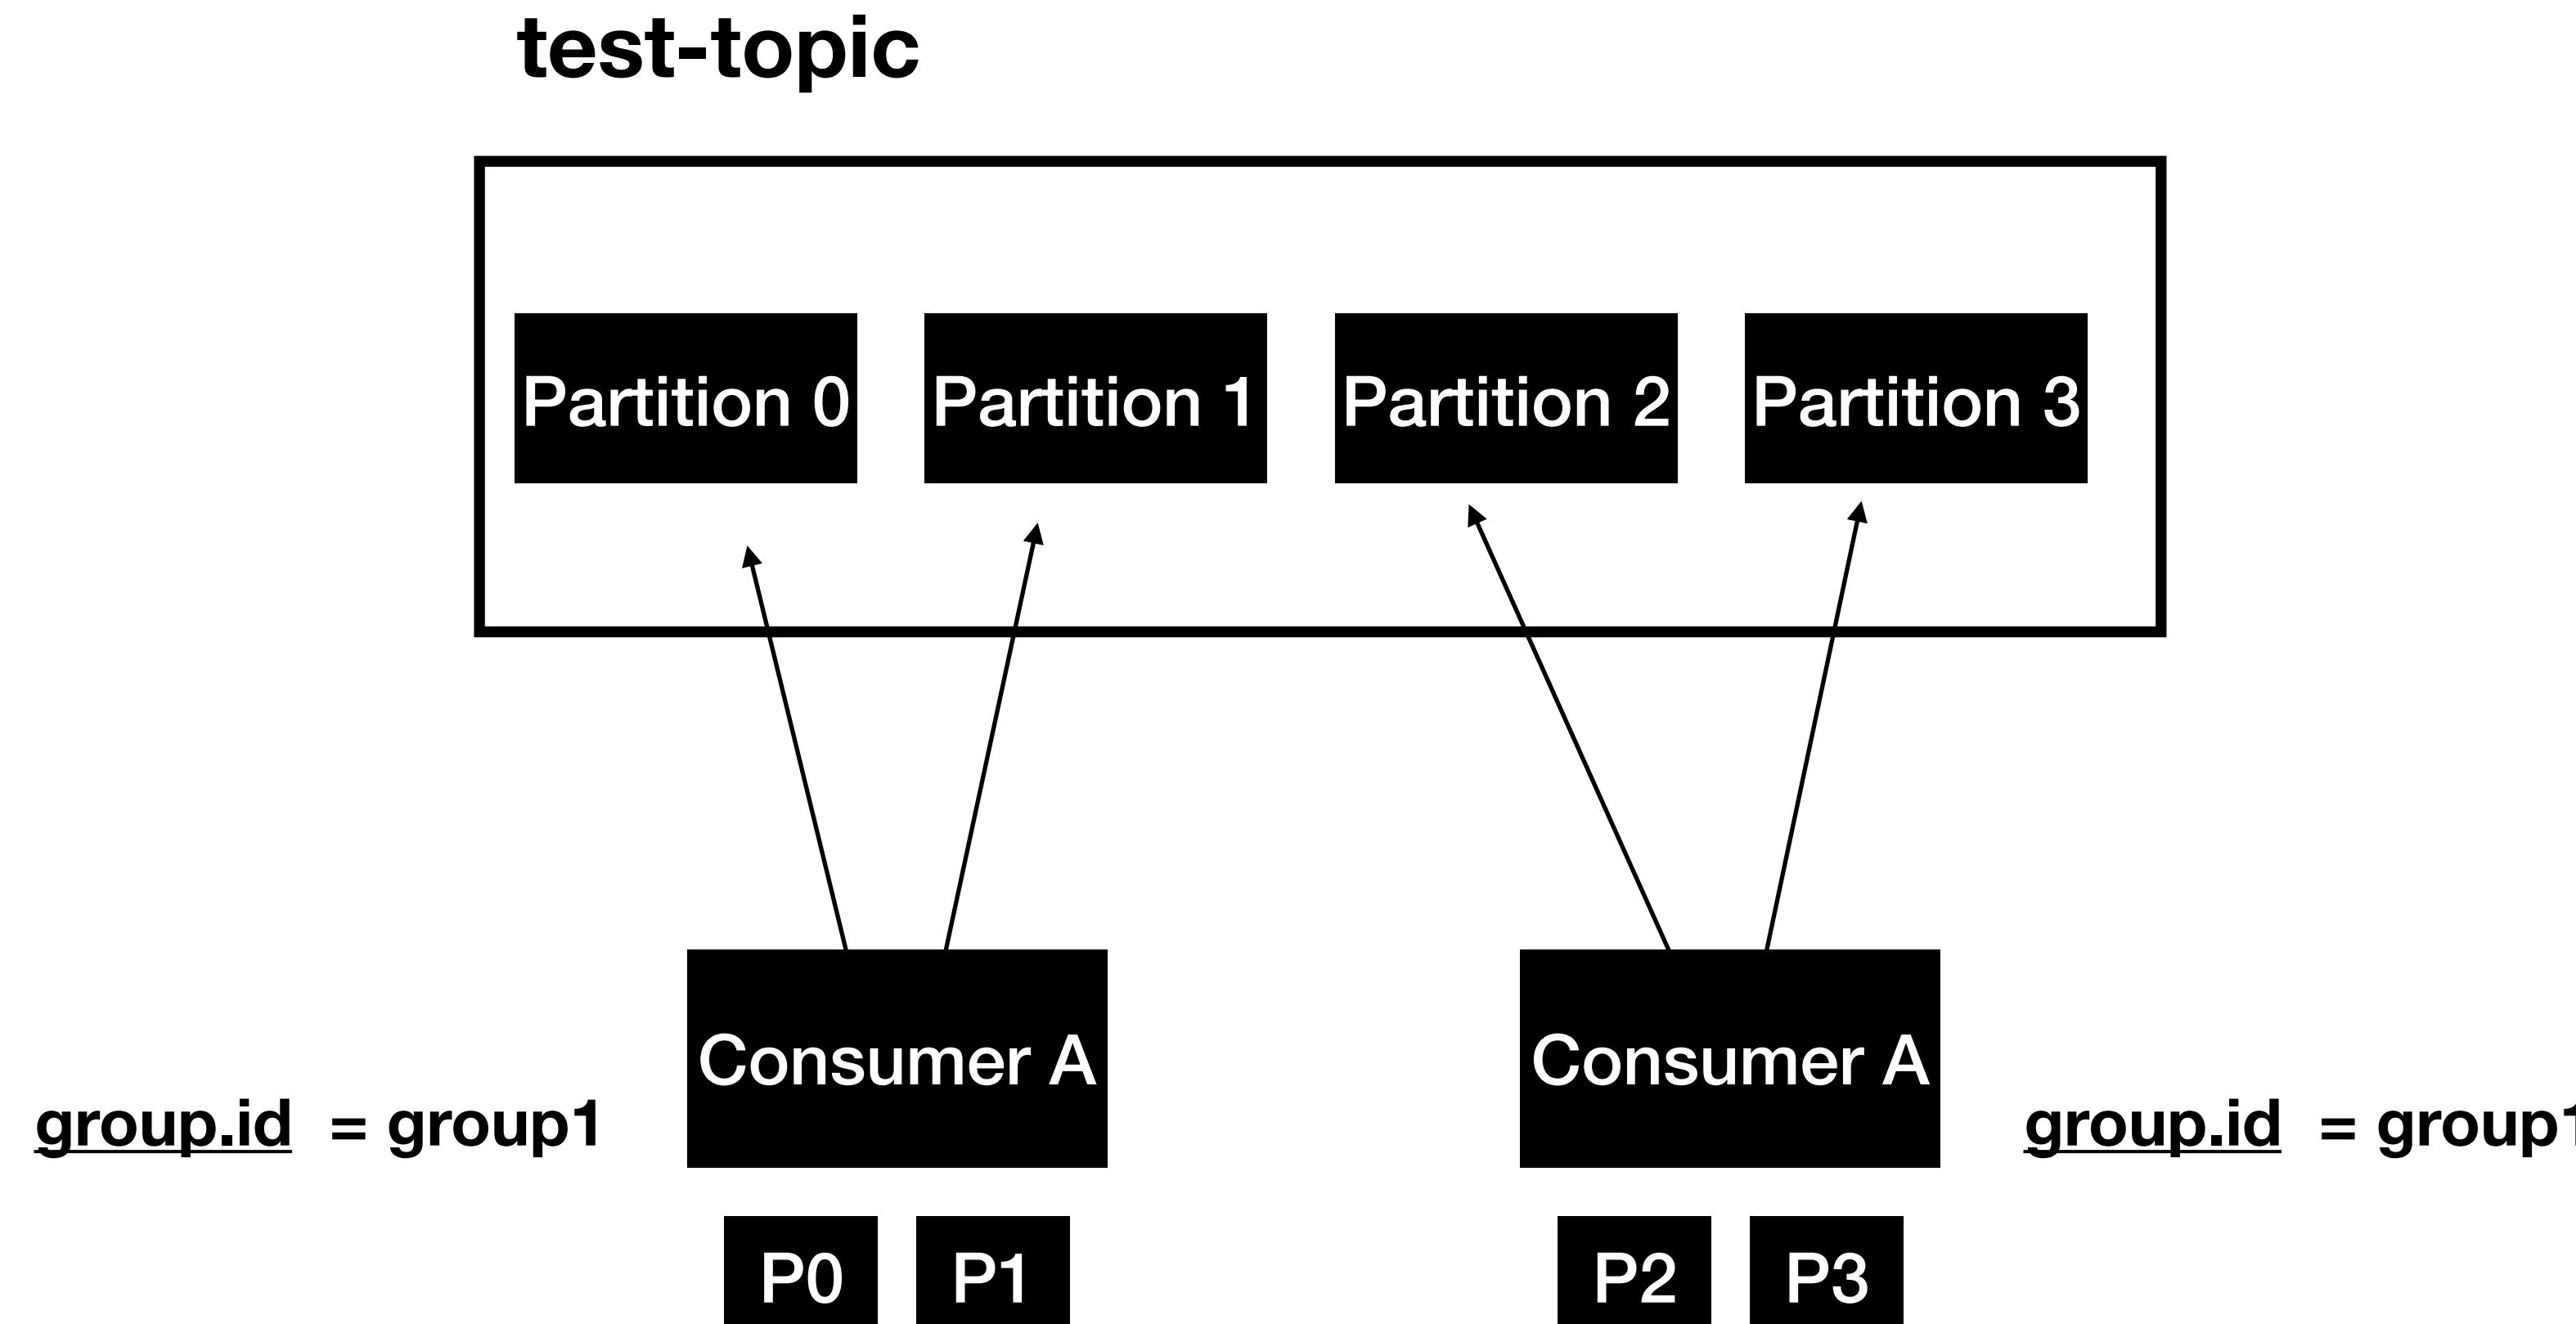


**group.id = group1**      **Consumer 1**

# Consumer Groups

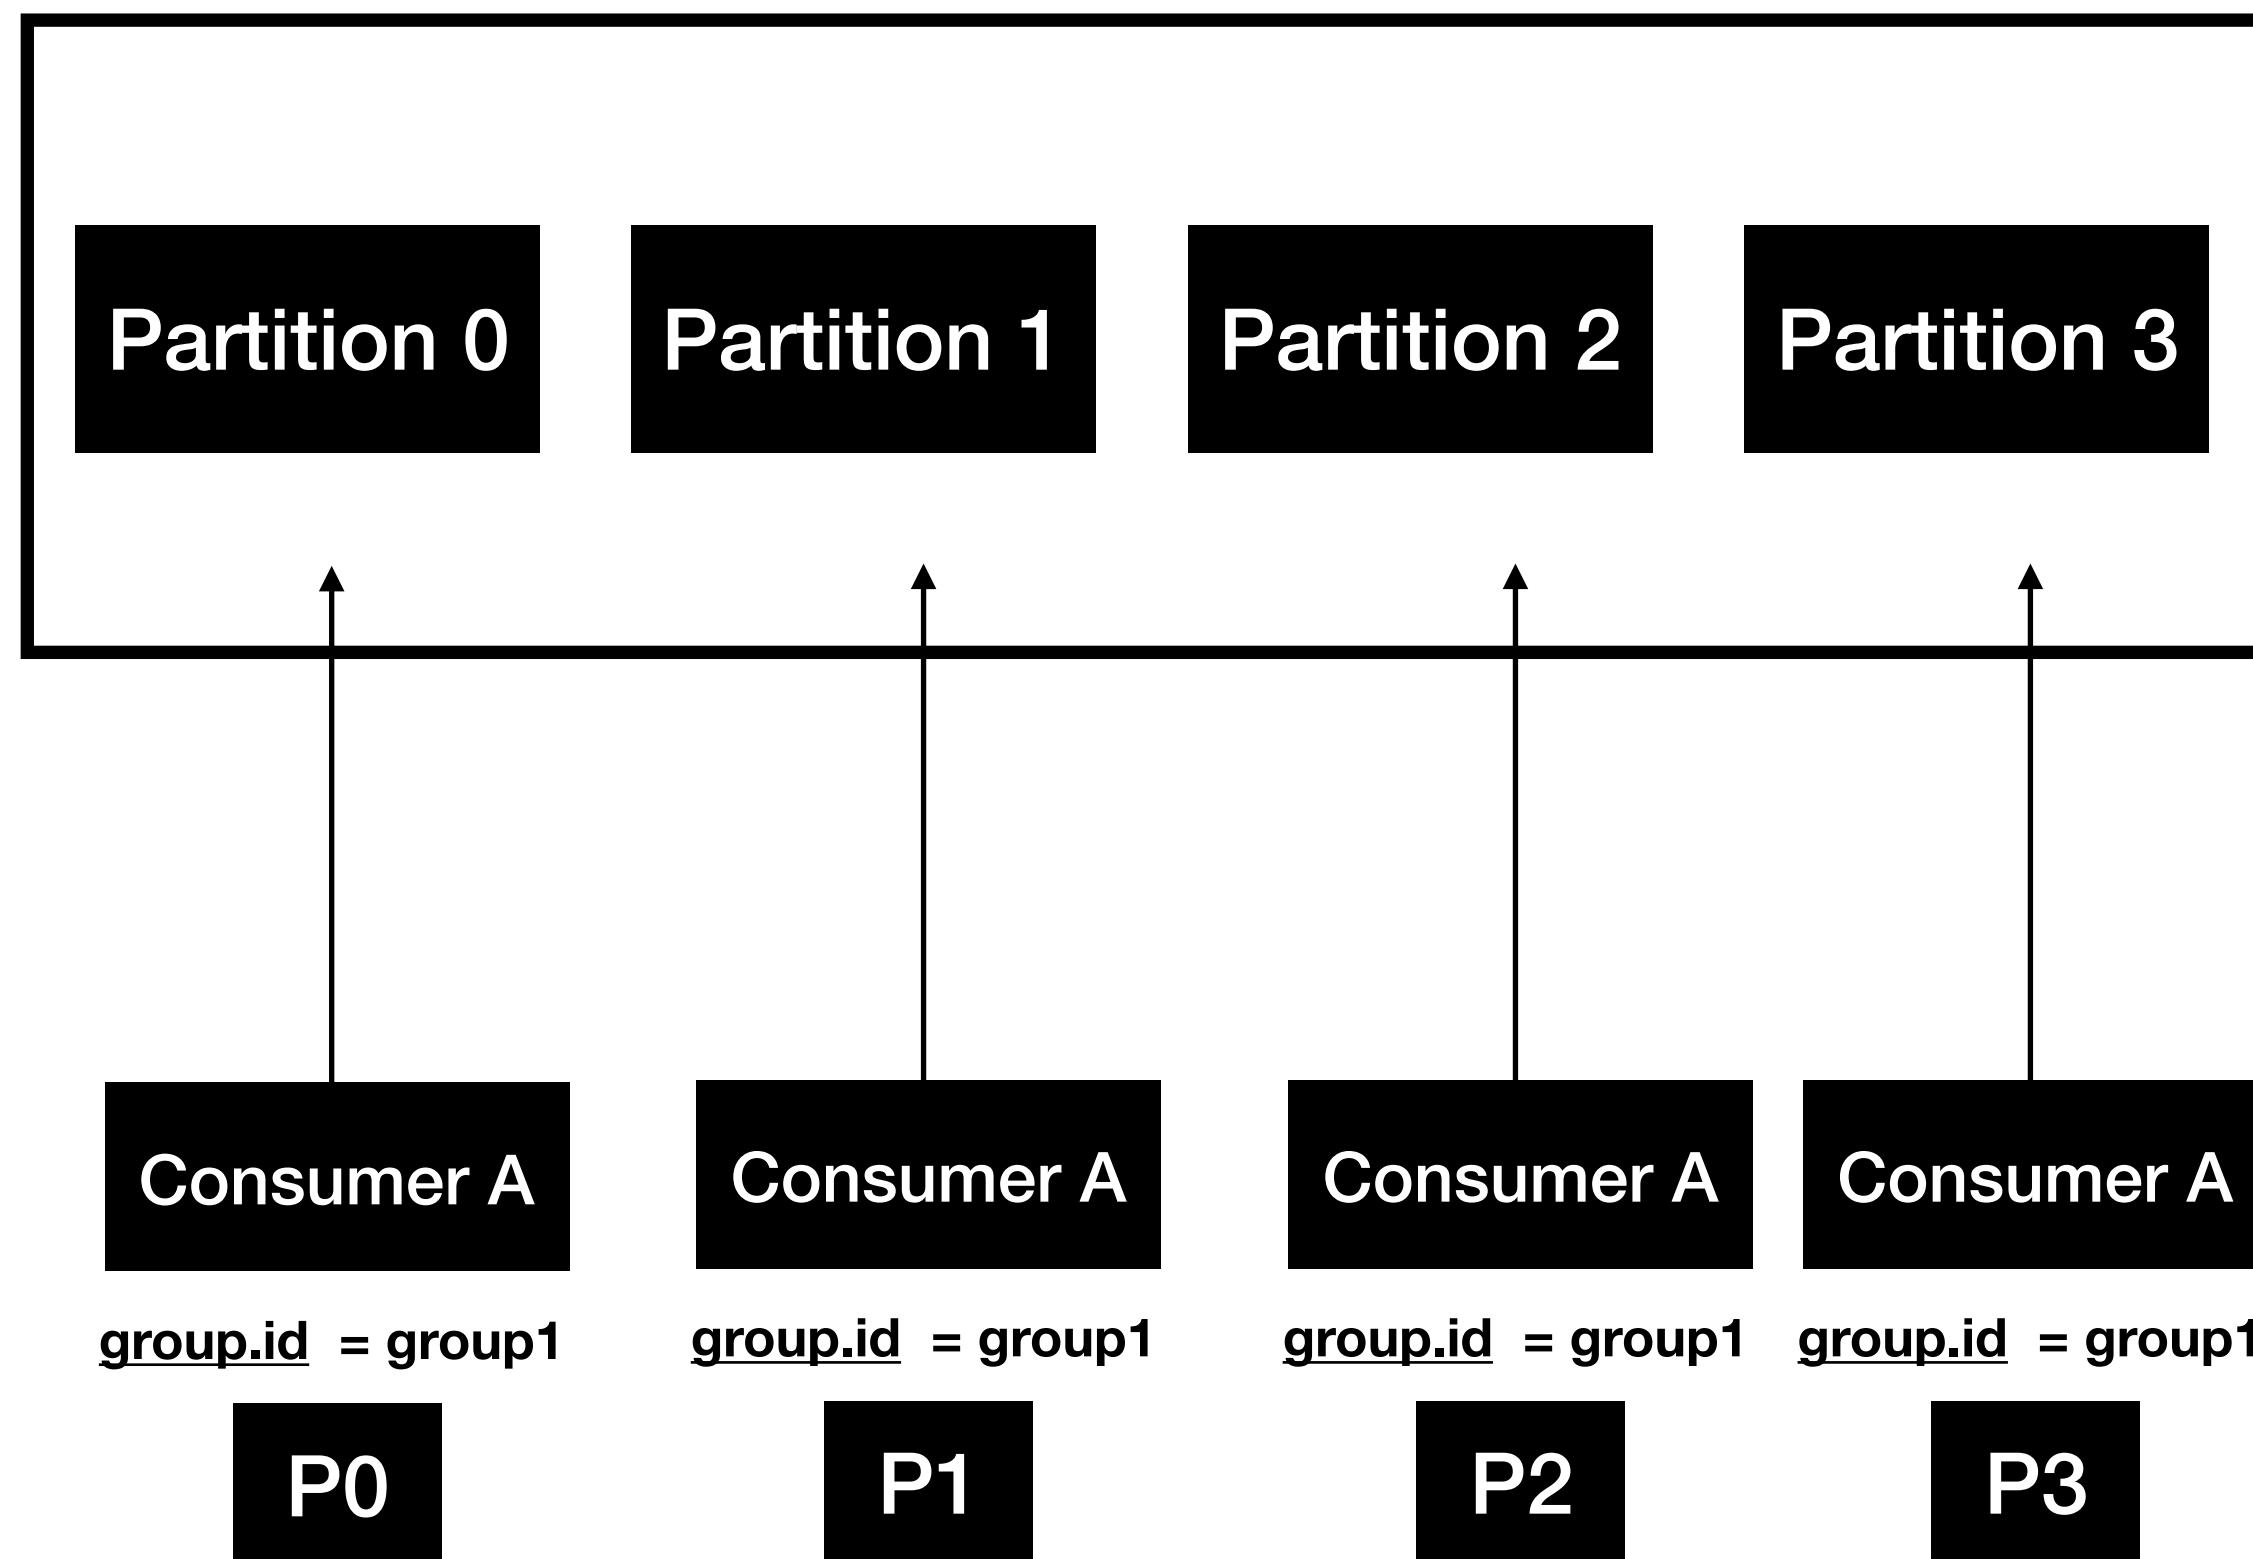


# Consumer Groups



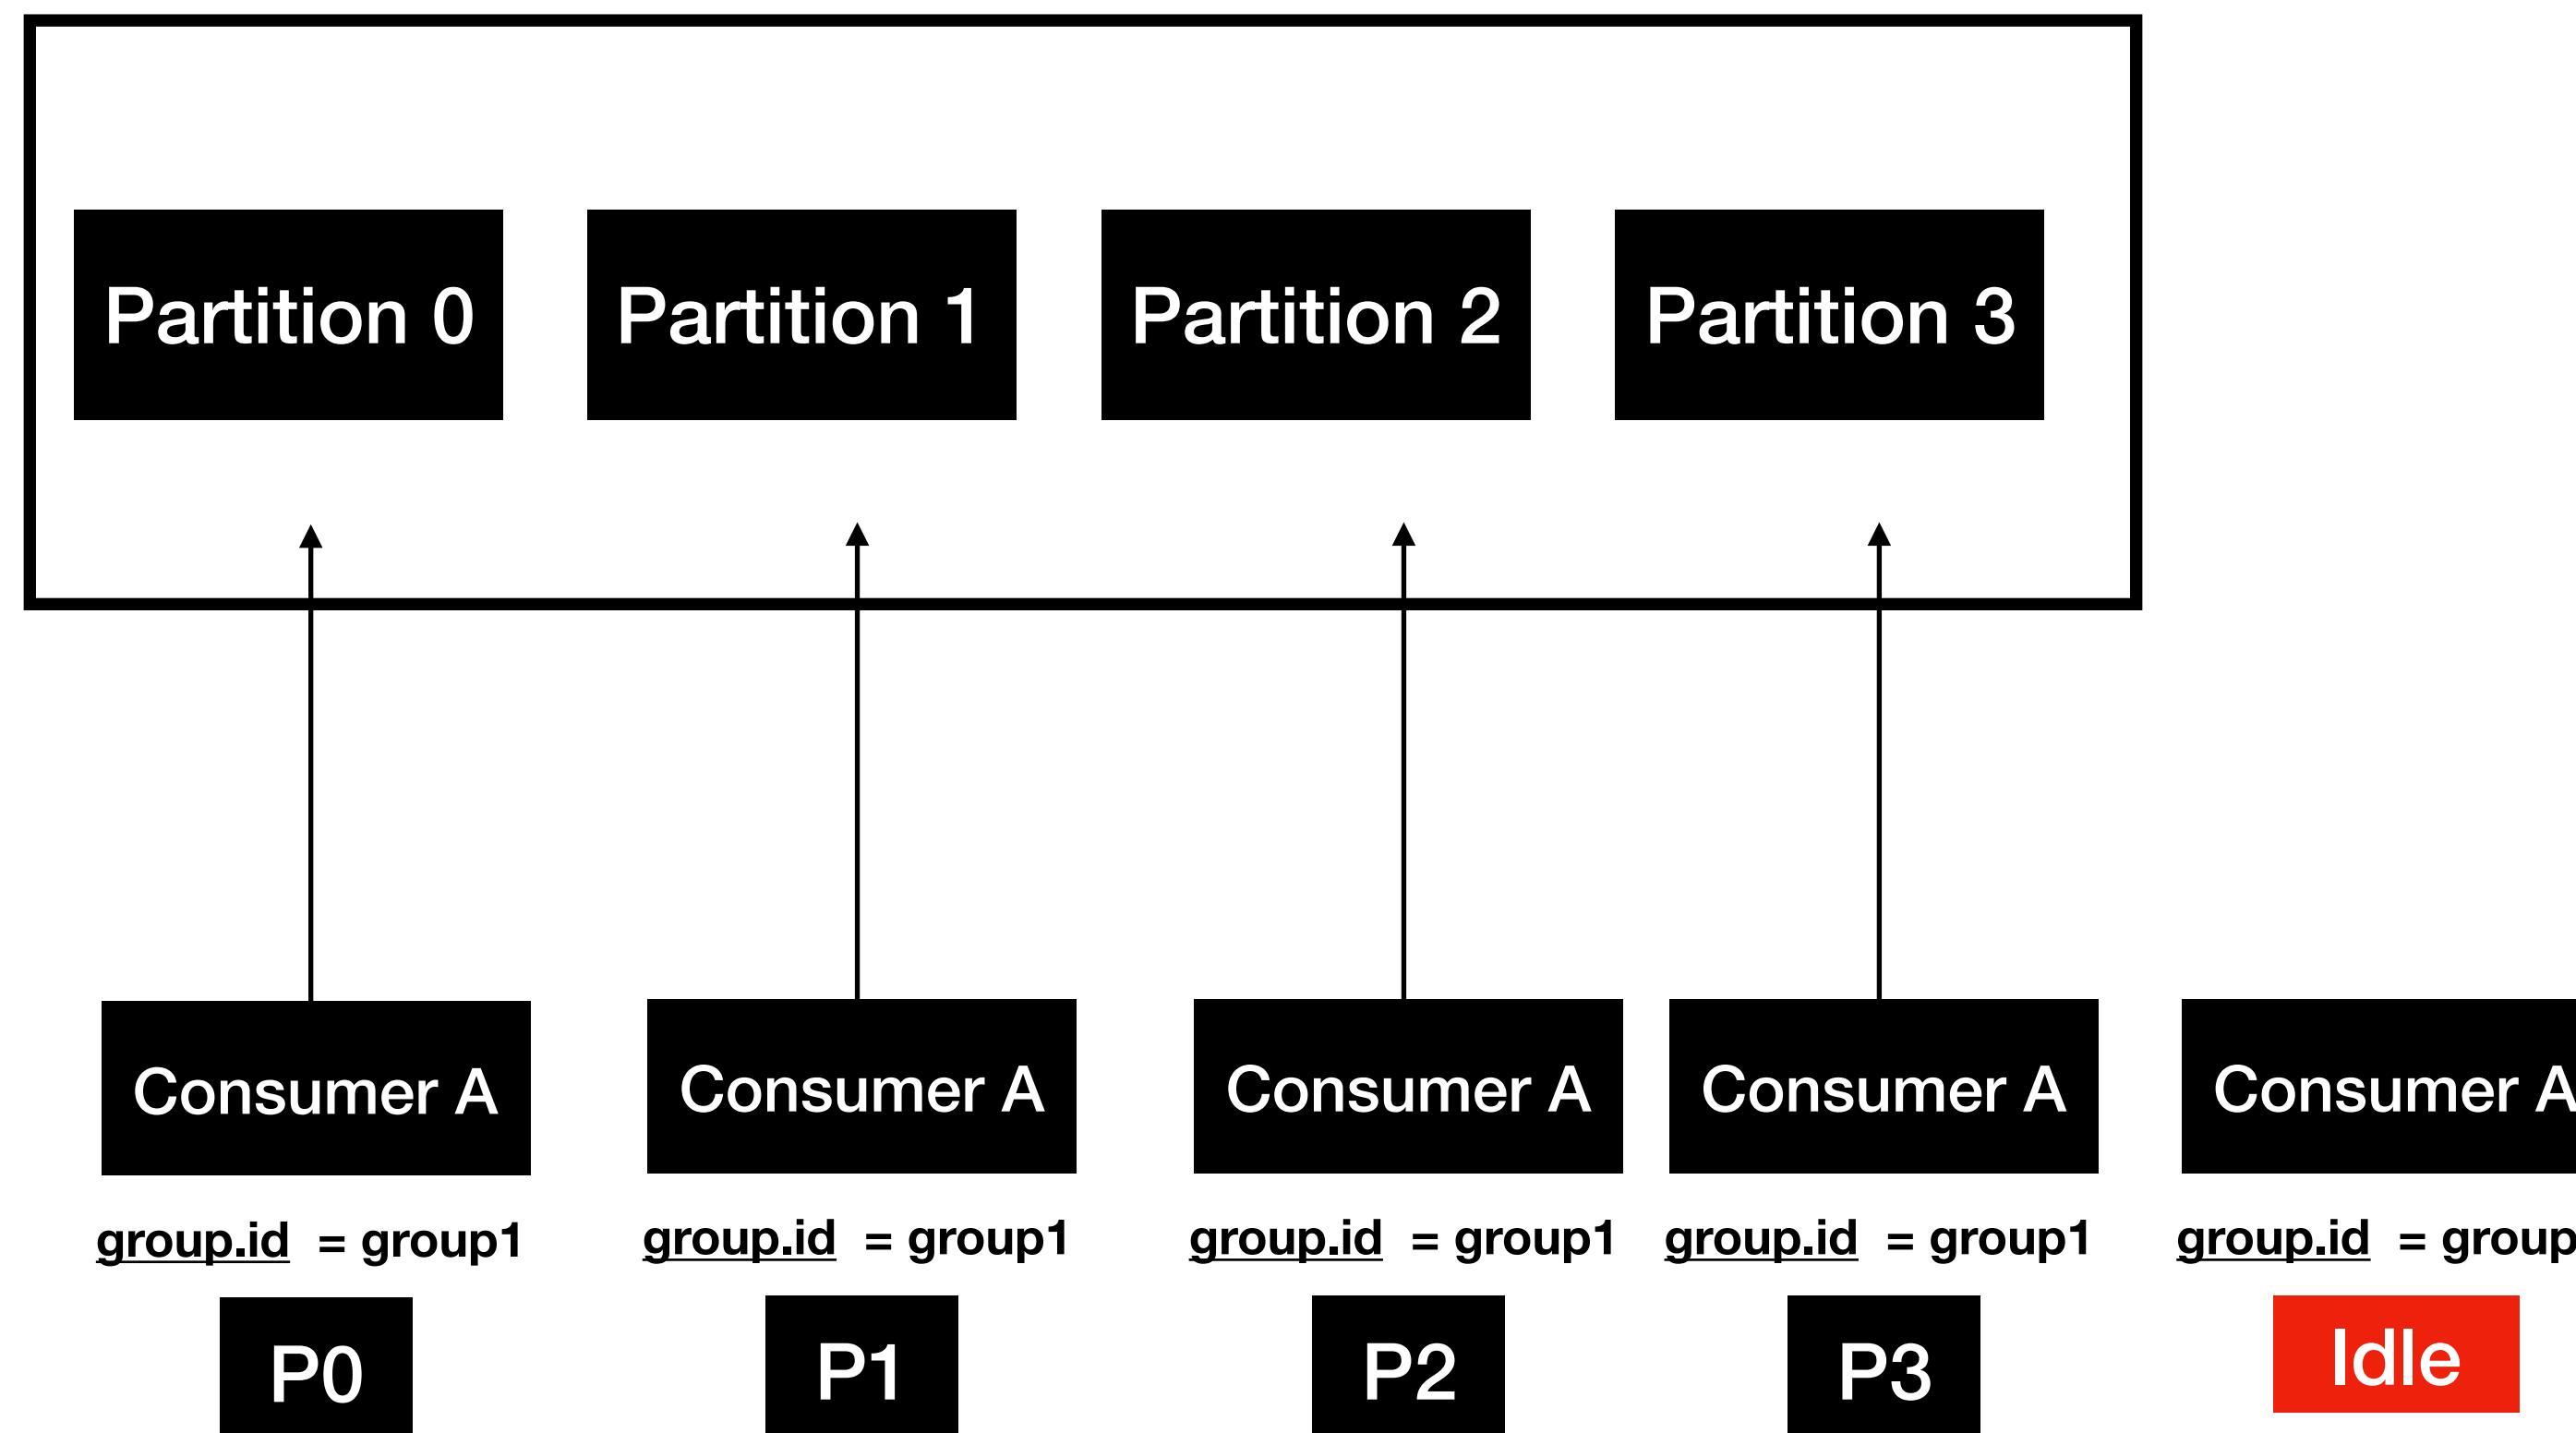
# Consumer Groups

**test-topic**



# Consumer Groups

test-topic

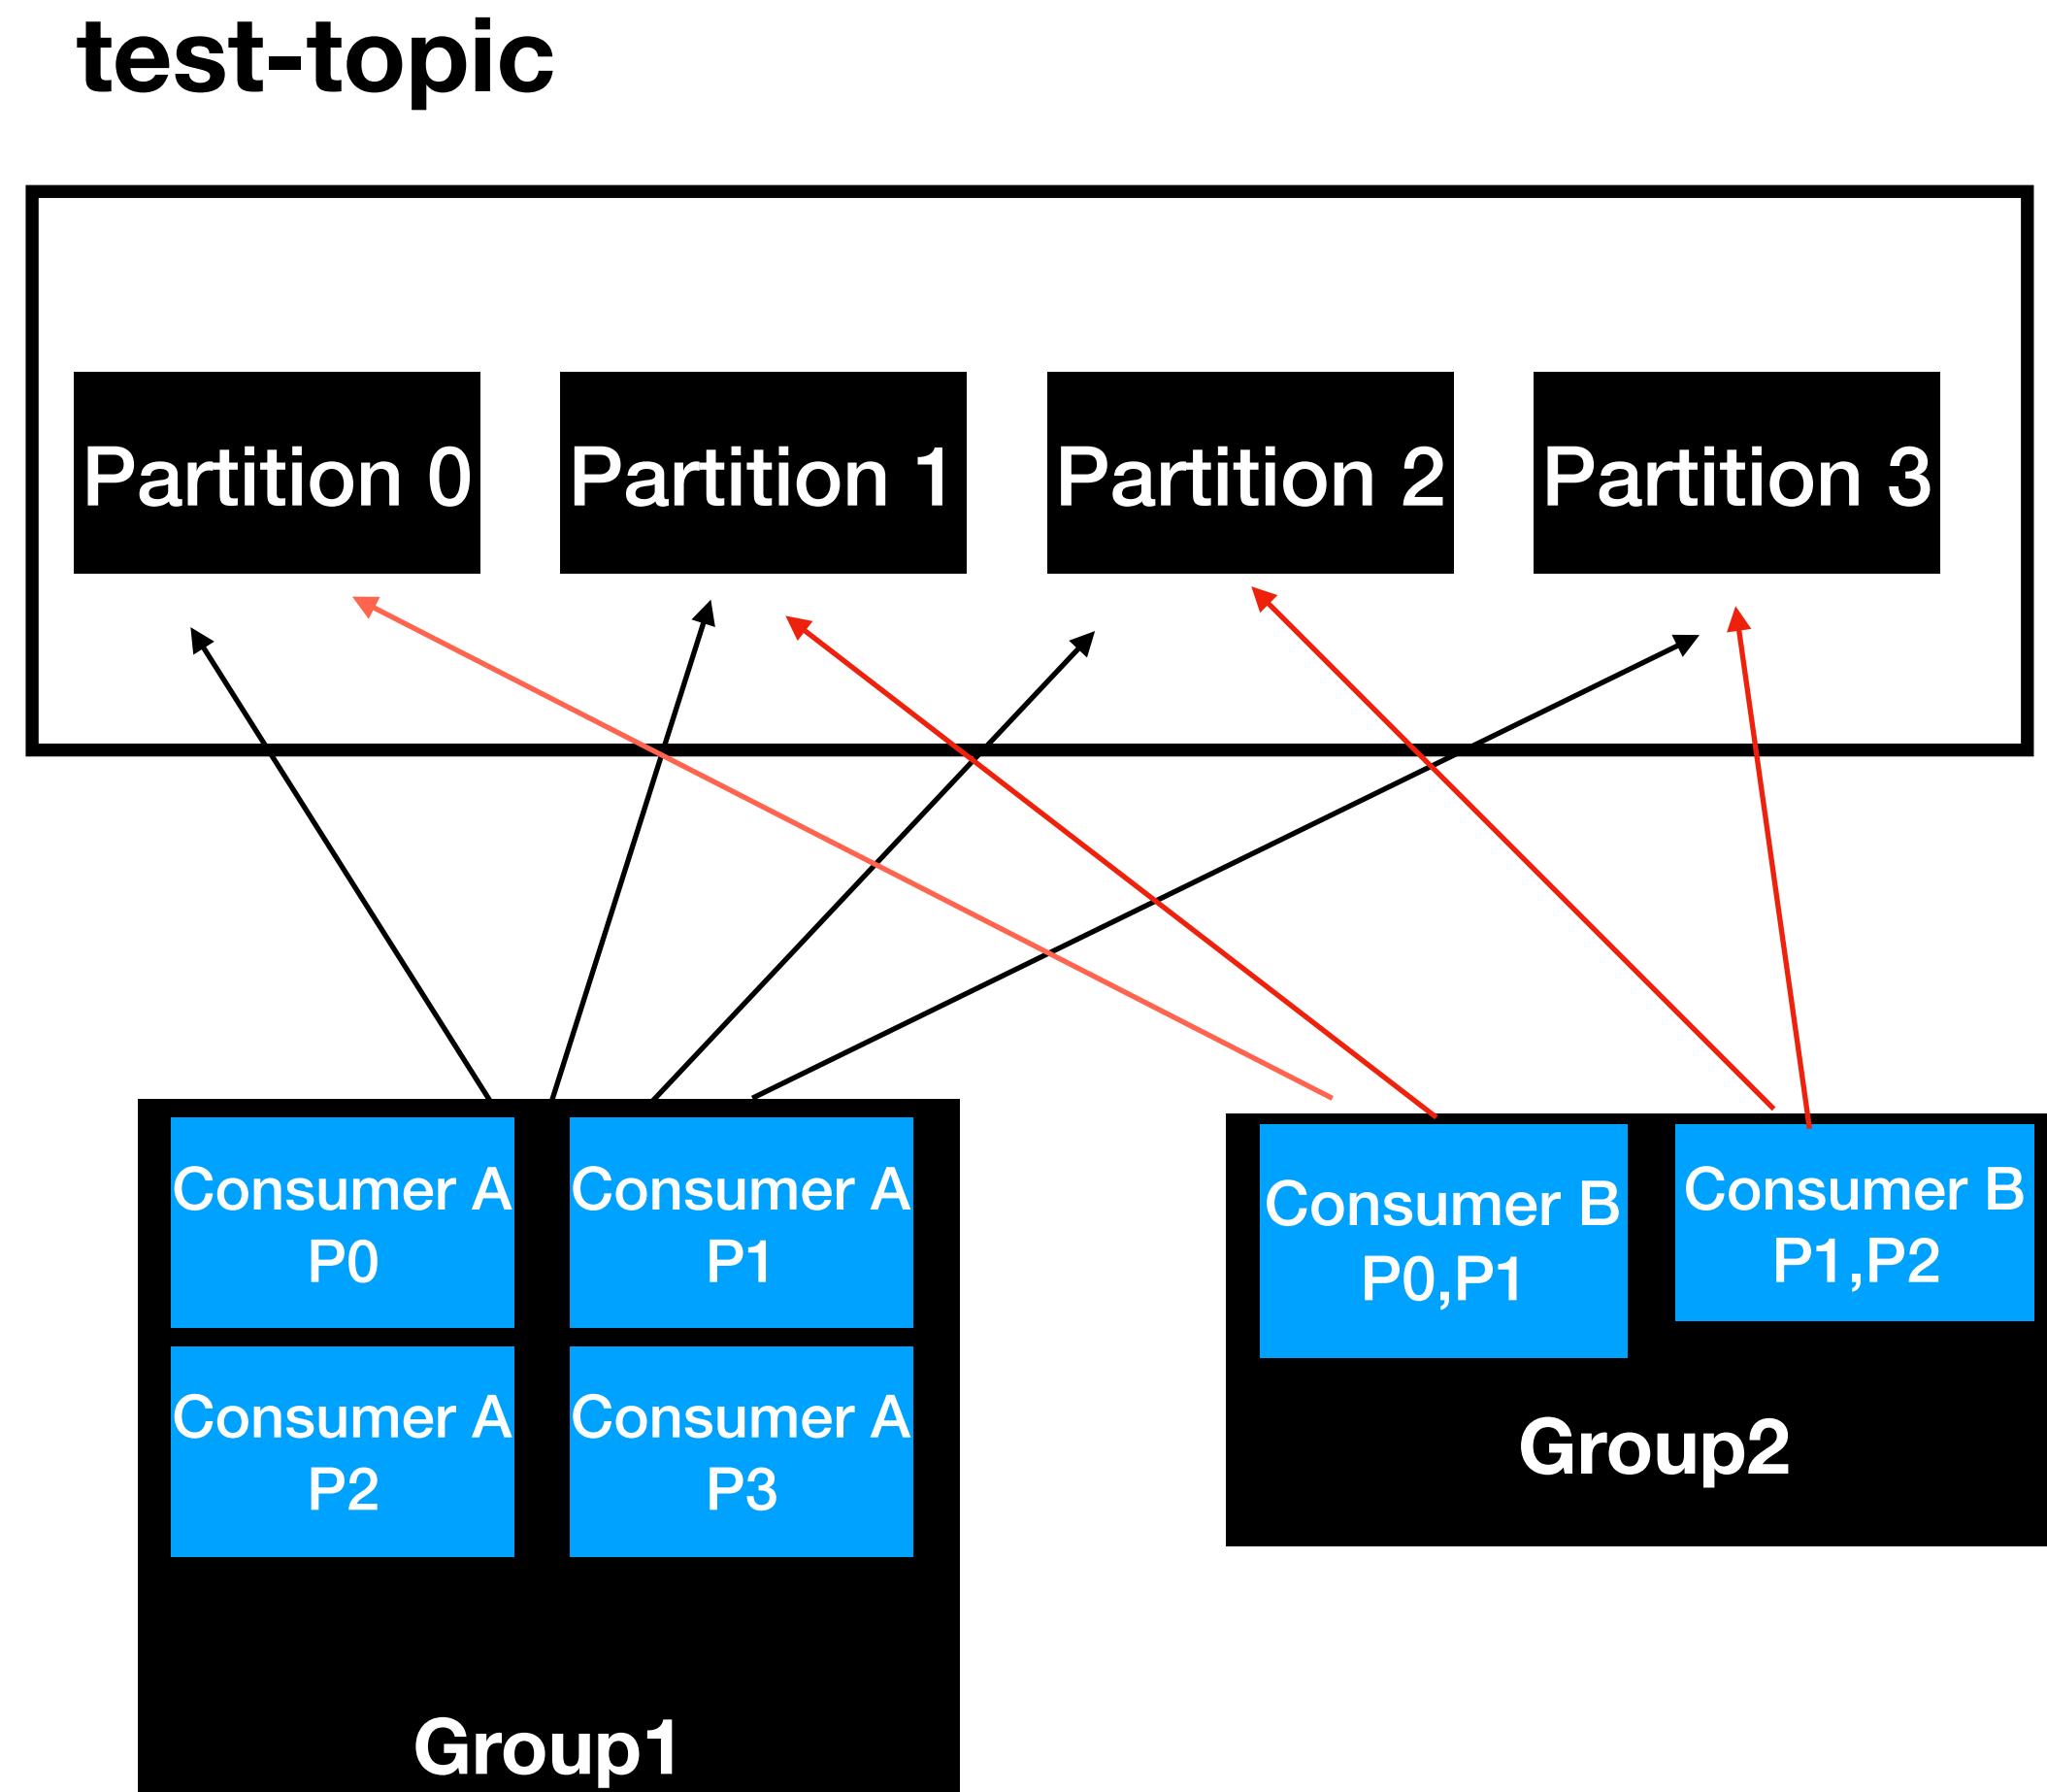


# Consumer Groups

In order to test this scenario, open two consumer group window pointing to same consumer group.

Start the console producer to produce message. In this case, your message will go to respective consumer group and based on the number of partitions.

since, we have created two console consumer pointing to same consumer group, each consumer group will have two partitions.



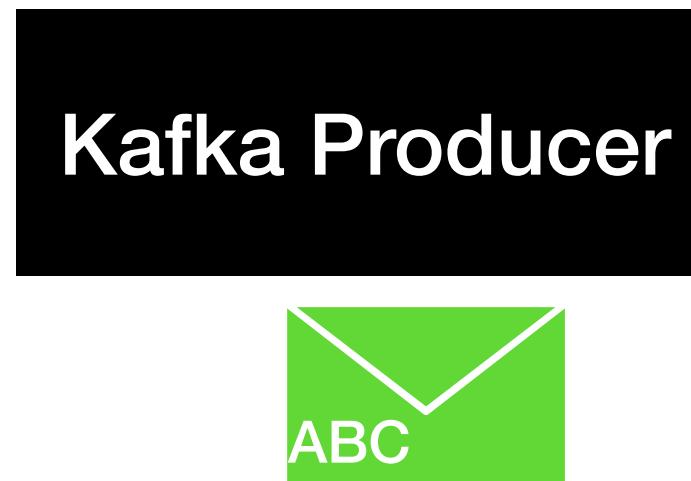
# Consumer Groups : Summary

- Consumer Groups are used for scalable message consumption
- Each different application will have a unique consumer group
- Who manages the consumer group?
  - Kafka Broker manages the consumer-groups
  - Kafka Broker acts as a Group Co-ordinator

# **Commit Log & Retention Policy**

## Producer:

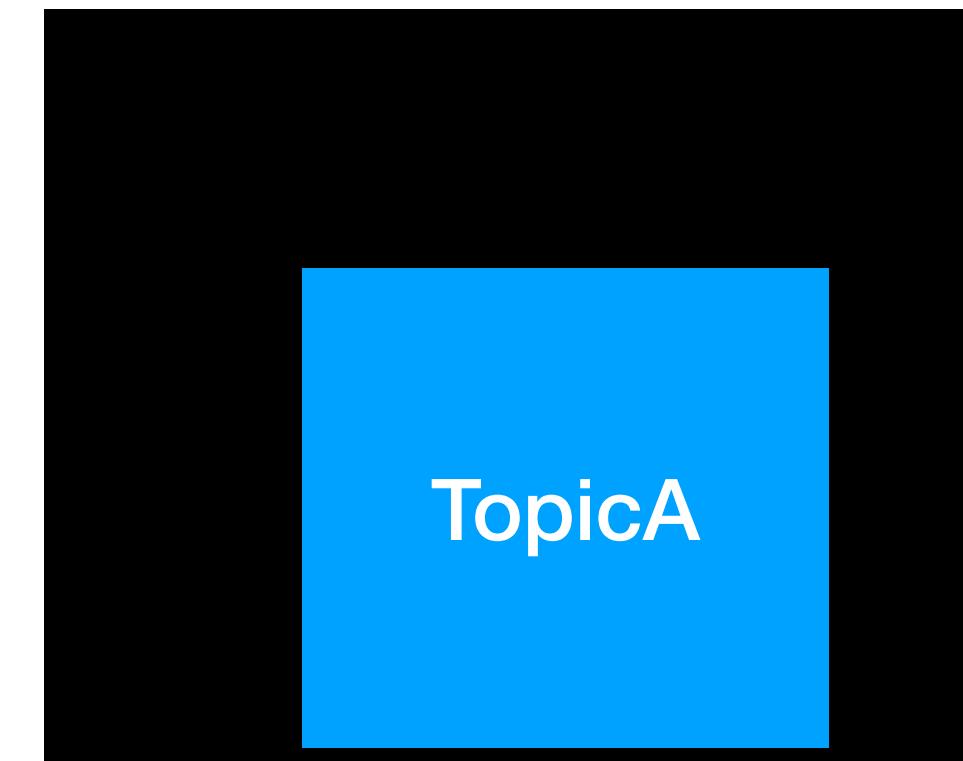
- \* Kafka Producer will produce a message.
- \* Each message will send to a Topic in the Kafka Broker.
- \* Once message reached to respective topic in Broker, message will be stored in the file system in a binary format (bytes)
- \* It is called as Commit Log
- \* If you have more than one partition, then log file will be created per partition level.
- \* These log files are also known as partition commit log
- \* Each message - will be recorded in the log file



Send to  
TopicA

# Commit Log

## Kafka Broker



Poll  
TopicA



## Consumer:

- \* When the consumer who is polling for new records can only see the records that are committed to the file system.
- \* As the new records are published on the topic then the records get appended to the log file.

Navigate to /bin directory and execute this command: To see more detailed log information.

`./kafka-run-class.sh kafka.tools.DumpLogSegments --deep-iteration --files /tmp/kafka-logs/test-topic-0/00000000000000000000.log`

# Retention Policy

- Determines how long the message is retained ?
- Configured using the property **log.retention.hours** in **server.properties** file
- Default retention period is **168 hours** (7 days)

# **Kafka**

**as a**

# **Distributed Streaming**

# **System**

# **Apache Kafka® is a *distributed streaming platform***

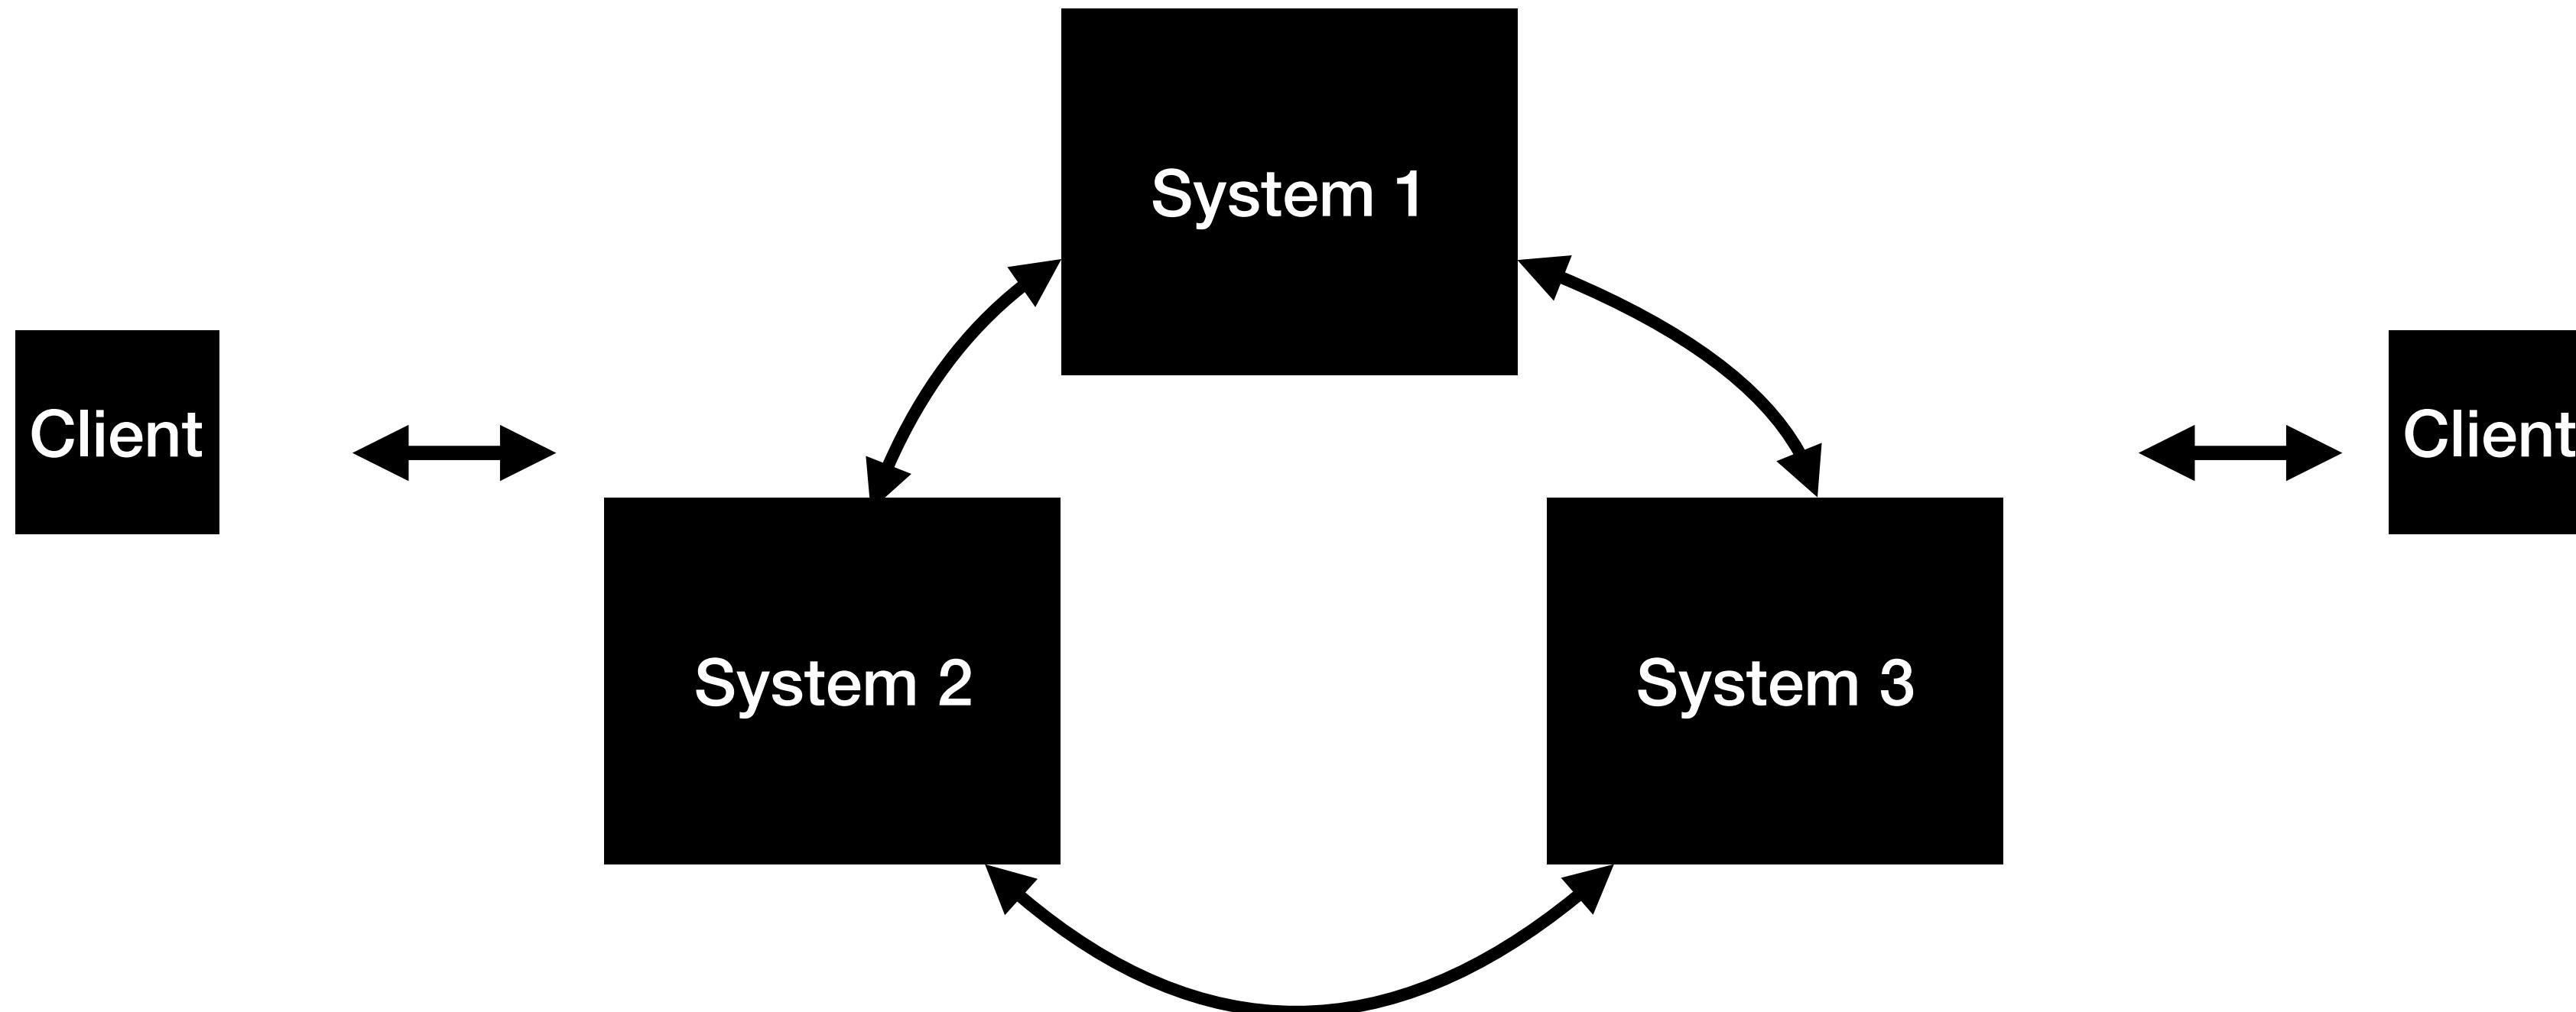
Notes:

What is a distributed system and its characteristics?

Distributed system in general are a collection of systems work and interact together in order to deliver the functionality.

# What is a Distributed System?

- Distributed systems are a collection of systems working together to deliver a value



# Characteristics of Distributed System

- Availability and Fault Tolerance

Let's say one of the system is down.

Still this won't impact the overall availability of the systems.

Even in this case the client request will be handled gracefully.

- Reliable Work Distribution

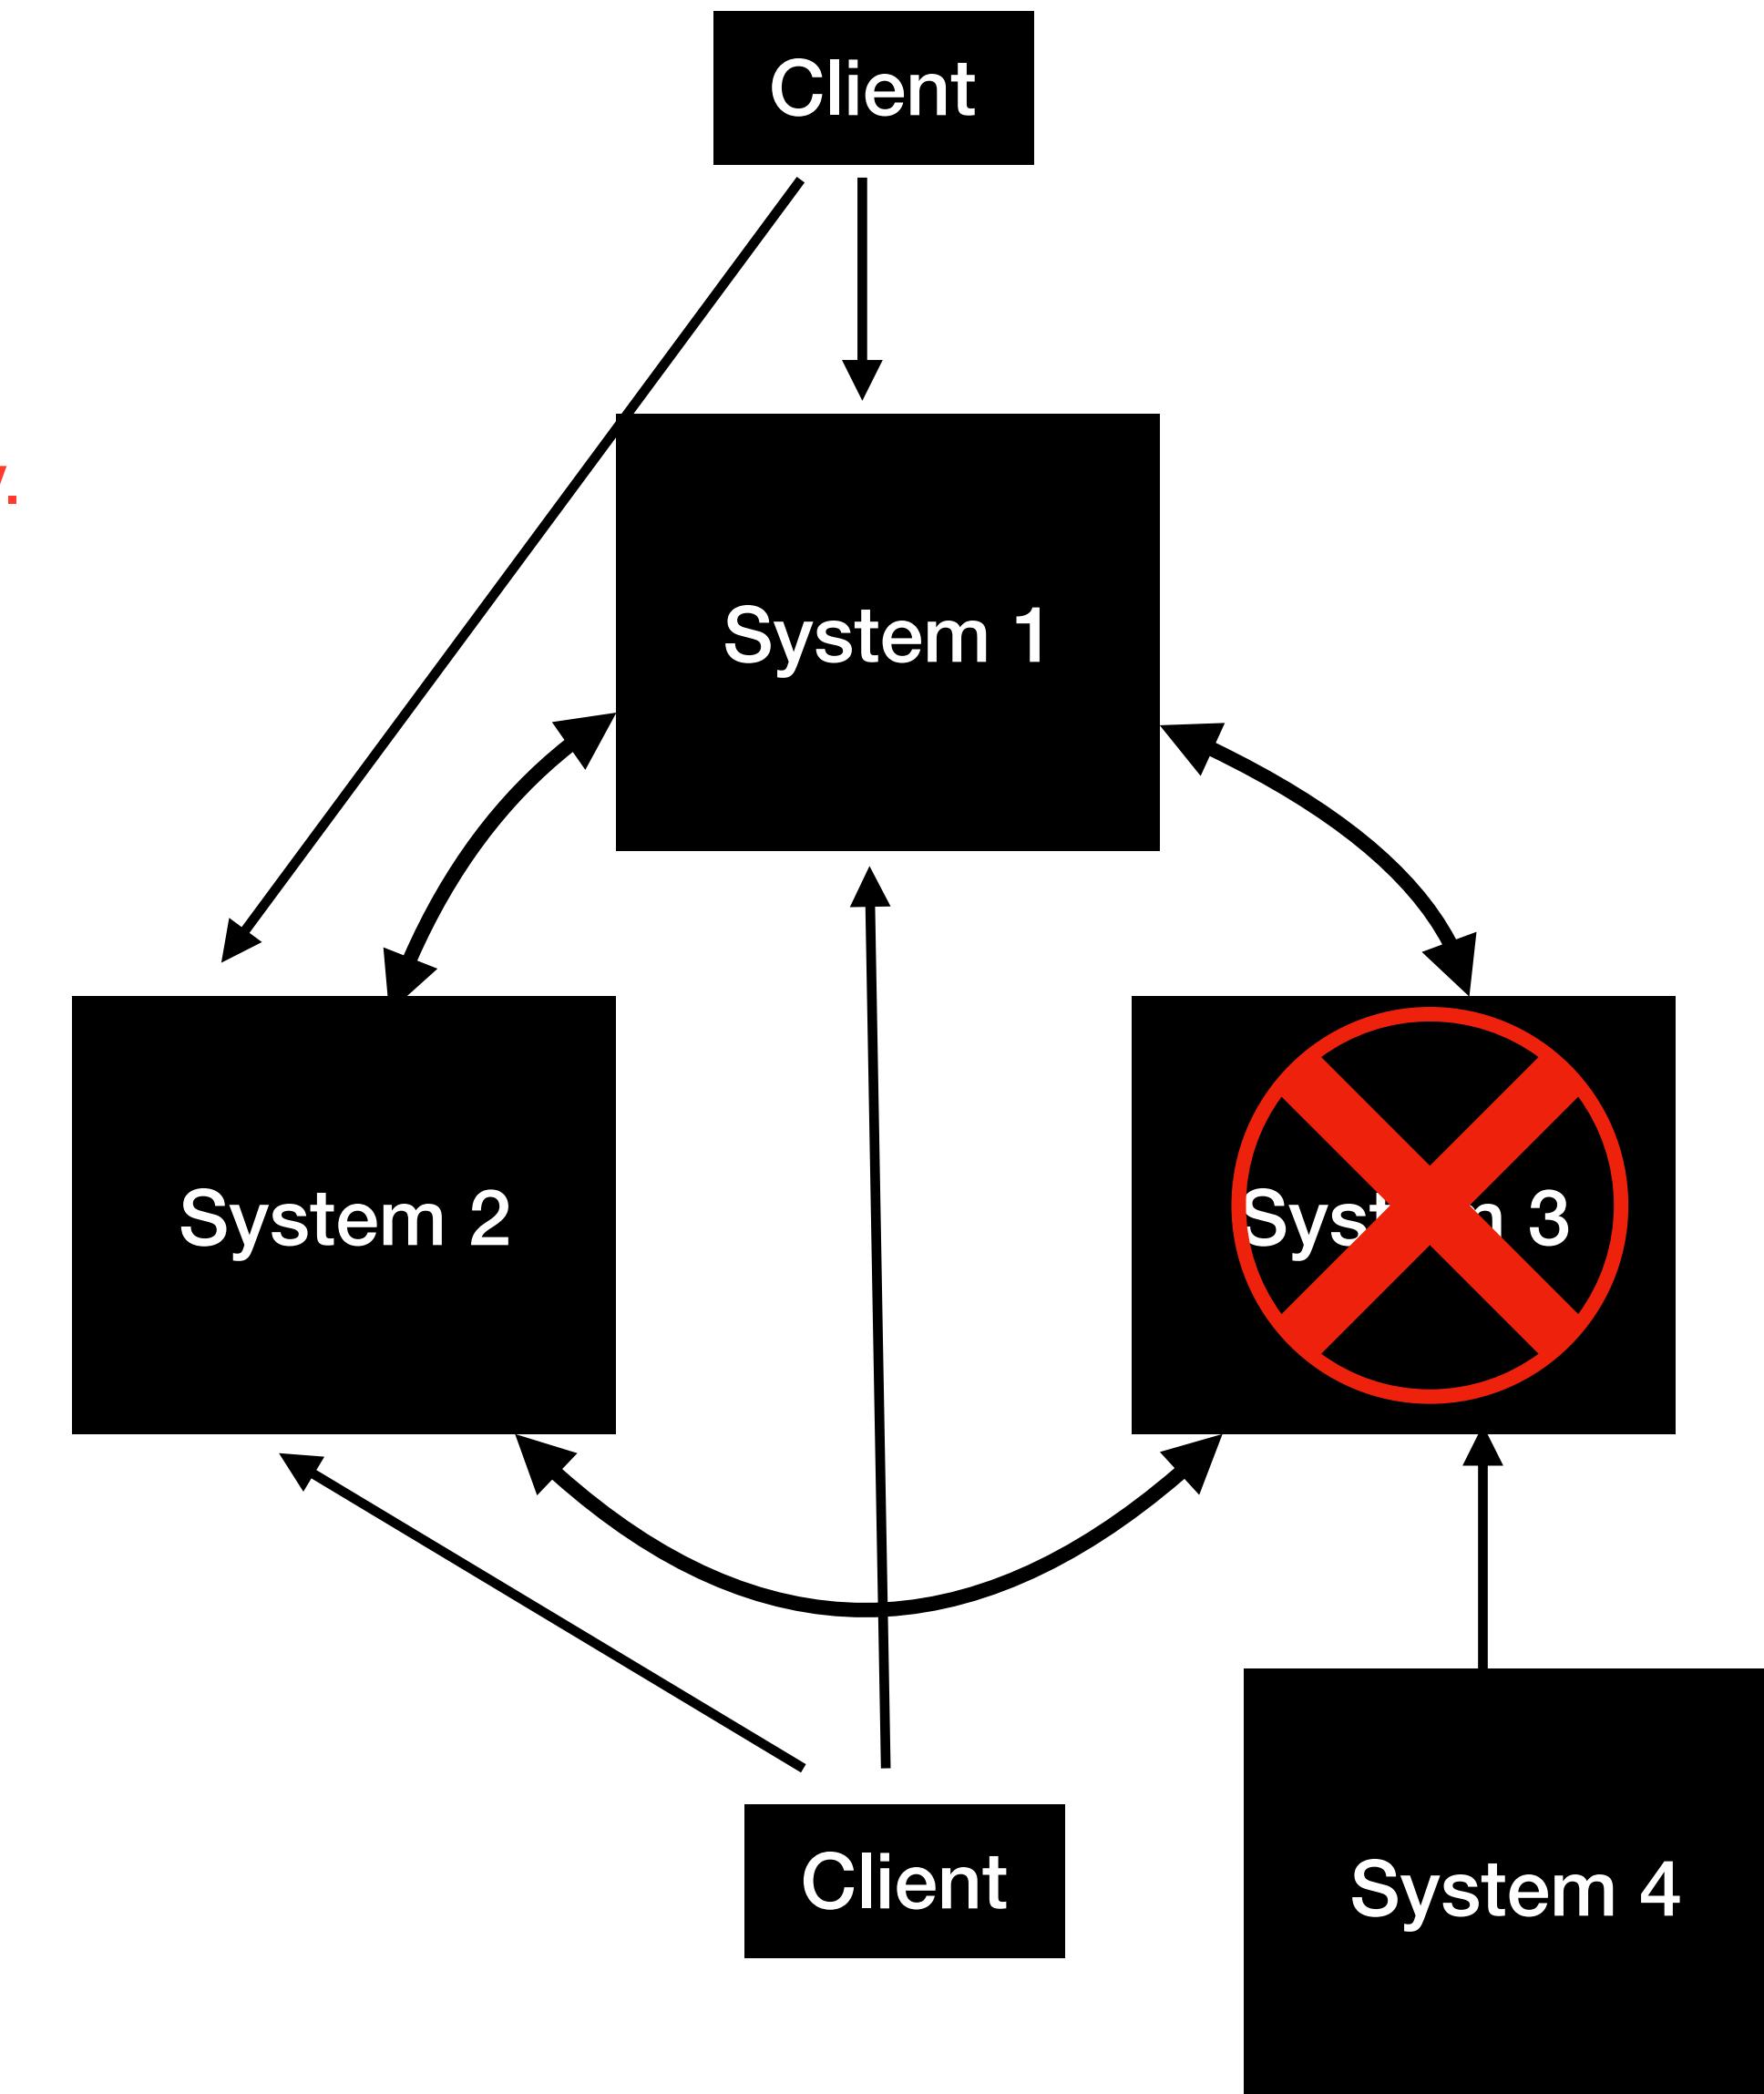
The requests that received from the clients in general are equally distributed between the available systems.

- Easily Scalable

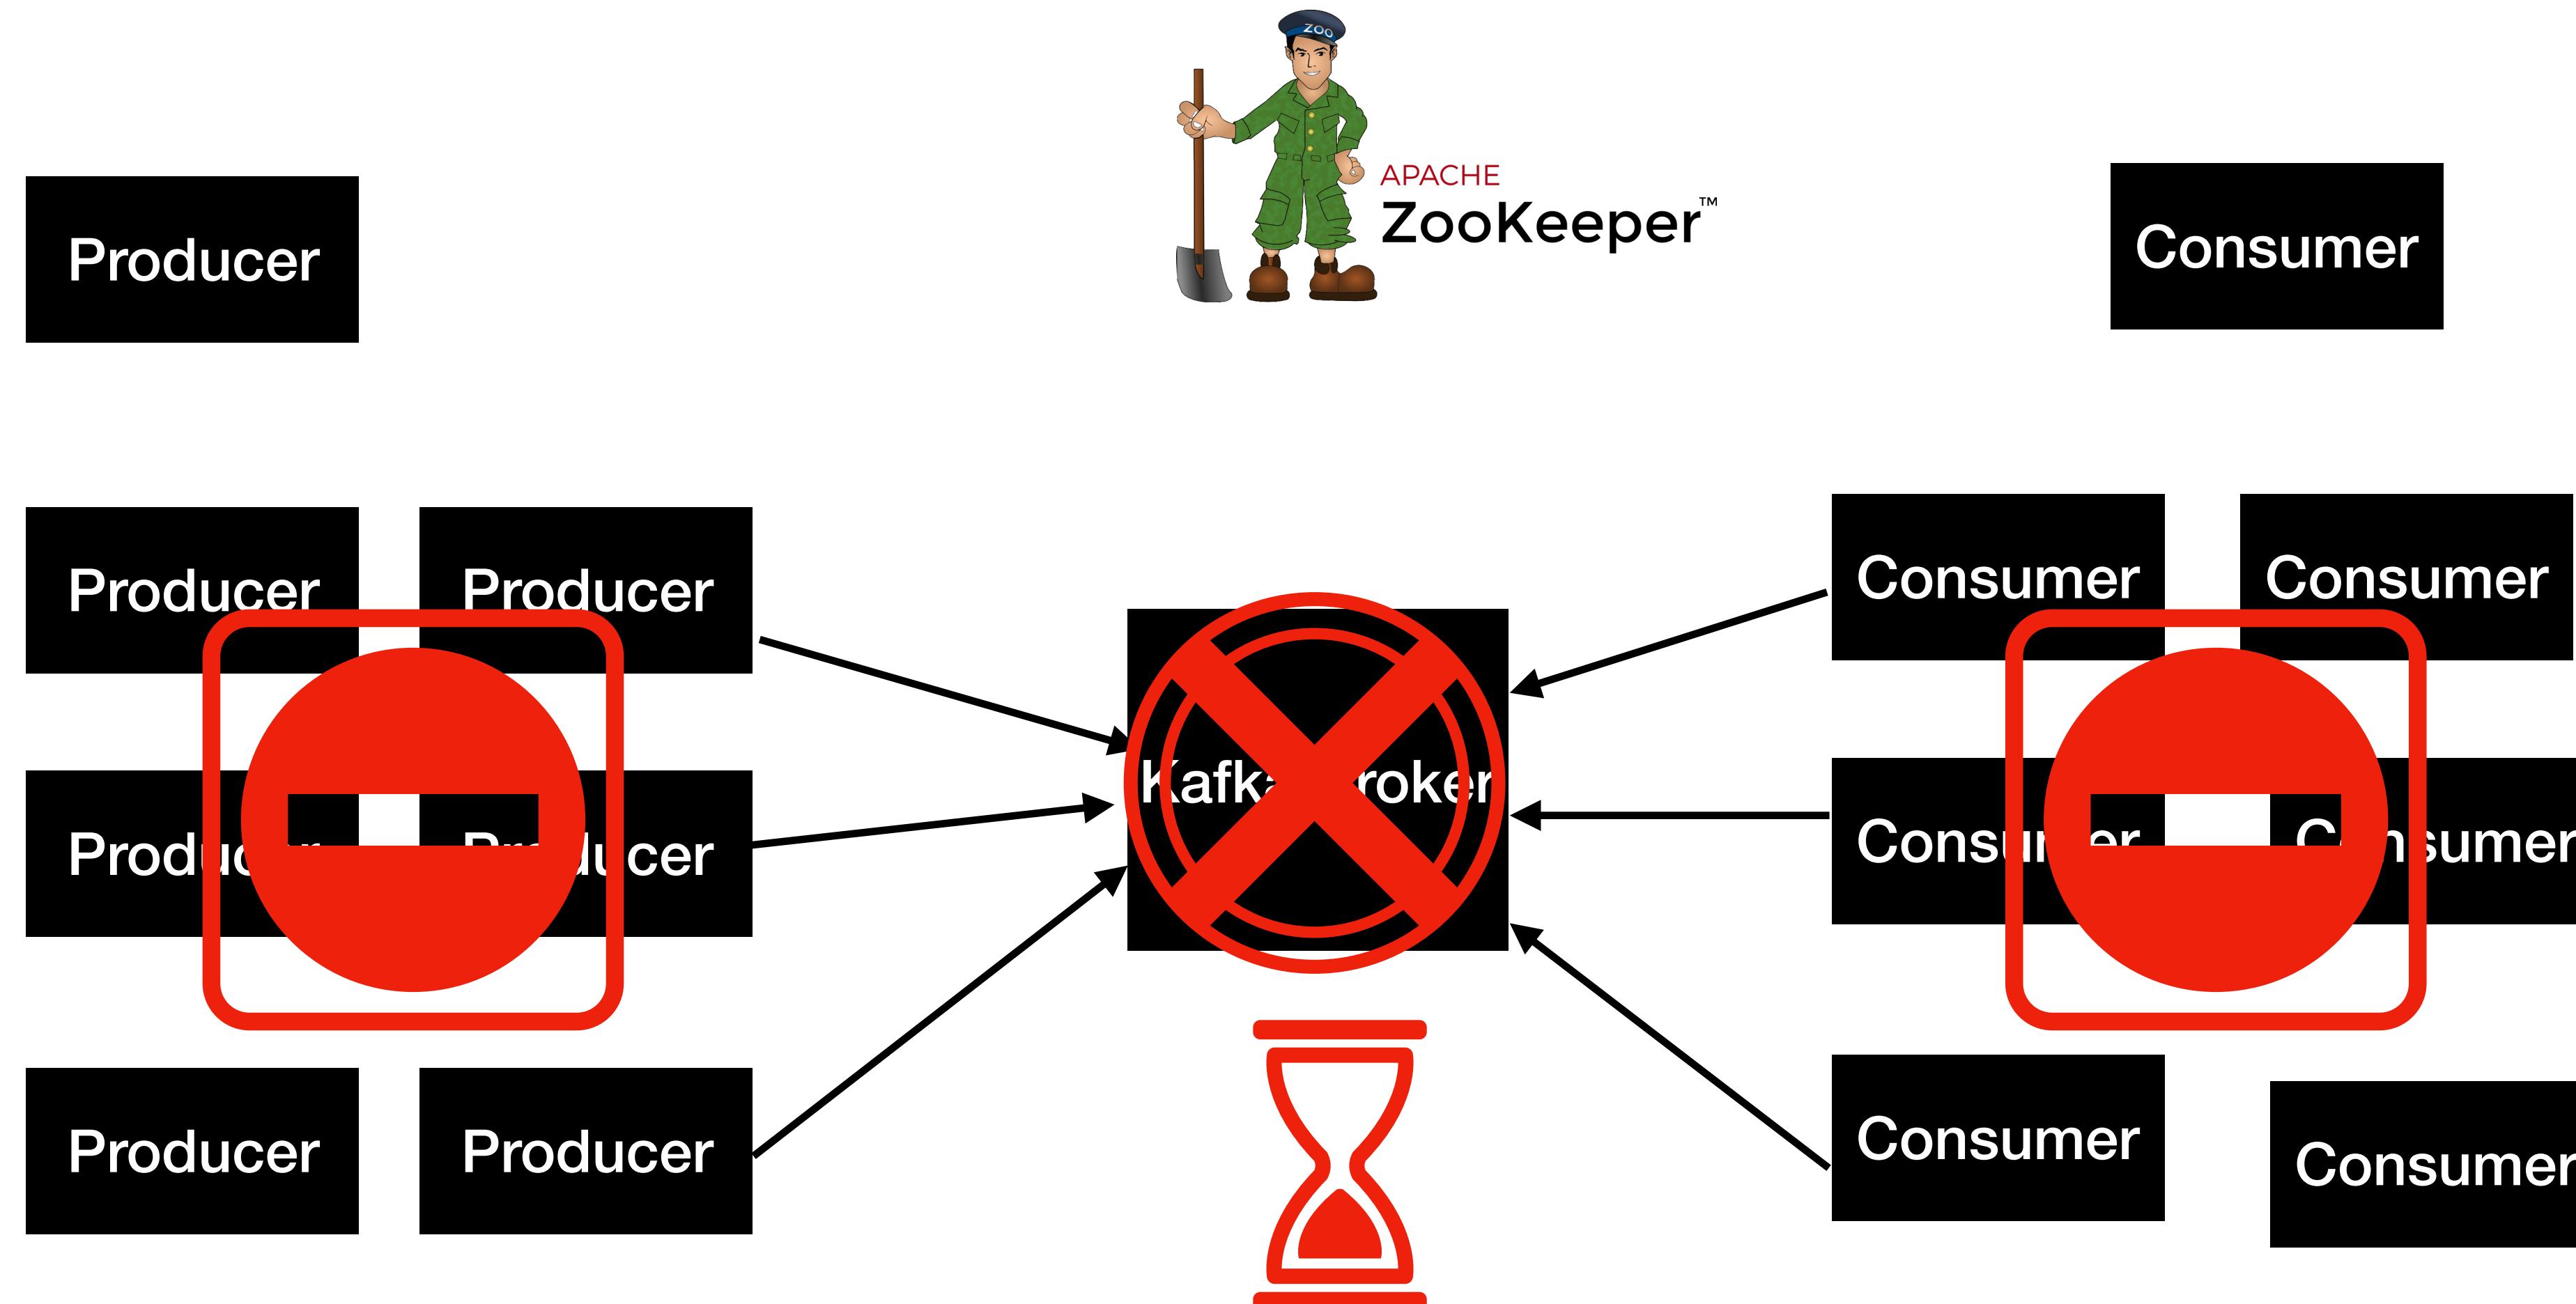
Adding new systems to the existing setup is really easy.

- Handling Concurrency is fairly easy

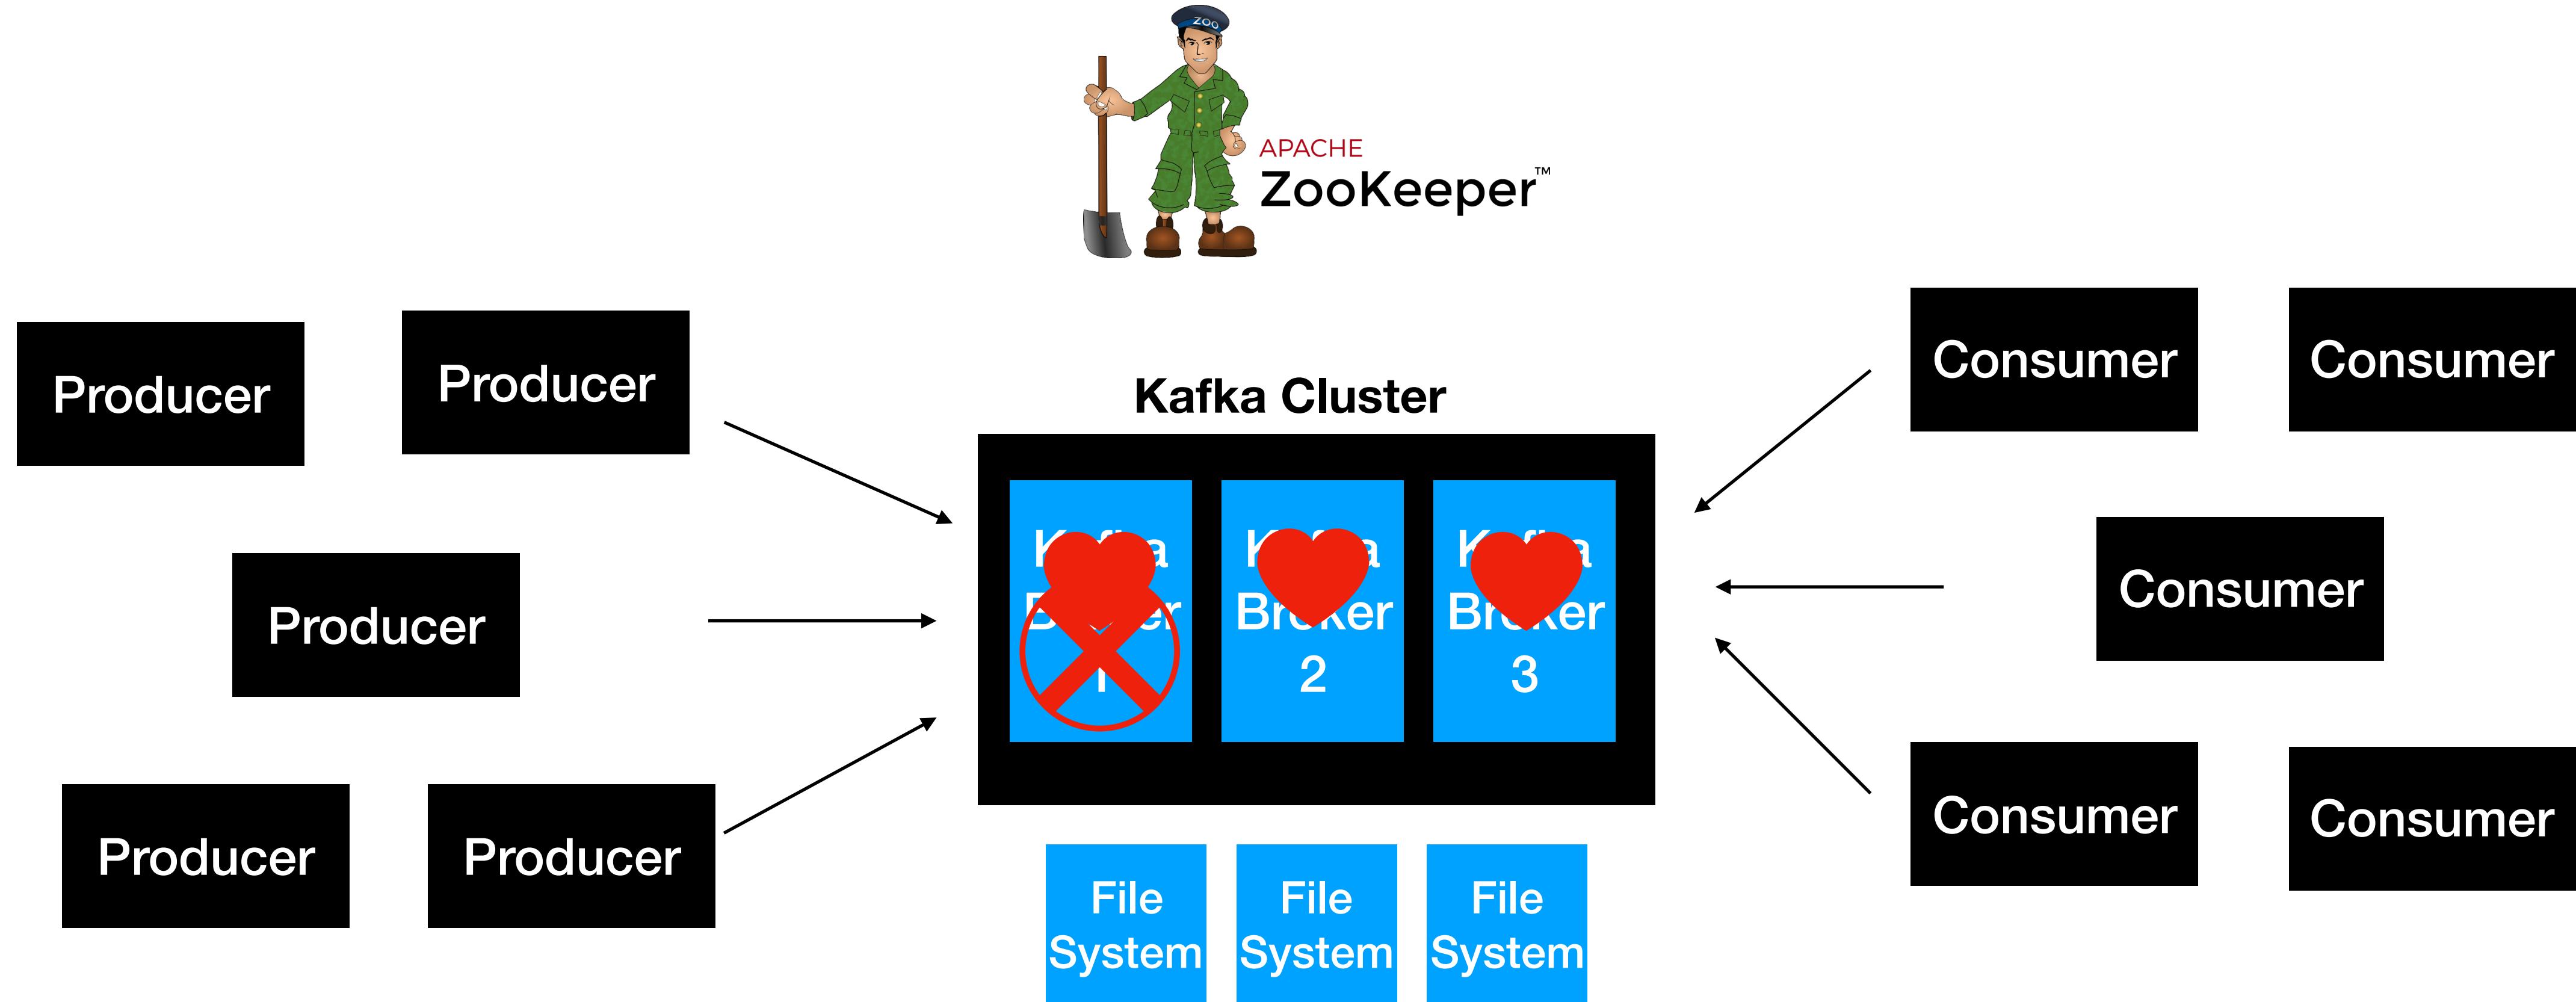
Which can be handled fairly easy but distributed systems.



# Kafka as a Distributed System



# Kafka as a Distributed System



- Client requests are distributed between brokers
- Easy to scale by adding more brokers based on the need
- Handles data loss using Replication

# **SetUp Kafka Cluster Using Three Brokers**

# Start Kafka Broker

```
./kafka-server-start.sh ../config/server.properties
```

# Setting up Kafka Cluster

- New **server.properties** files with the new broker details.

```
broker.id=<unique-broker-id>
listeners=PLAINTEXT://localhost:<unique-port>
log.dirs=/tmp/<unique-kafka-folder>
auto.create.topics.enable=false(optional)
```

## Example: **server-1.properties**

```
broker.id=1
listeners=PLAINTEXT://localhost:9093
log.dirs=/tmp/kafka-logs-1
auto.create.topics.enable=false(optional)
```

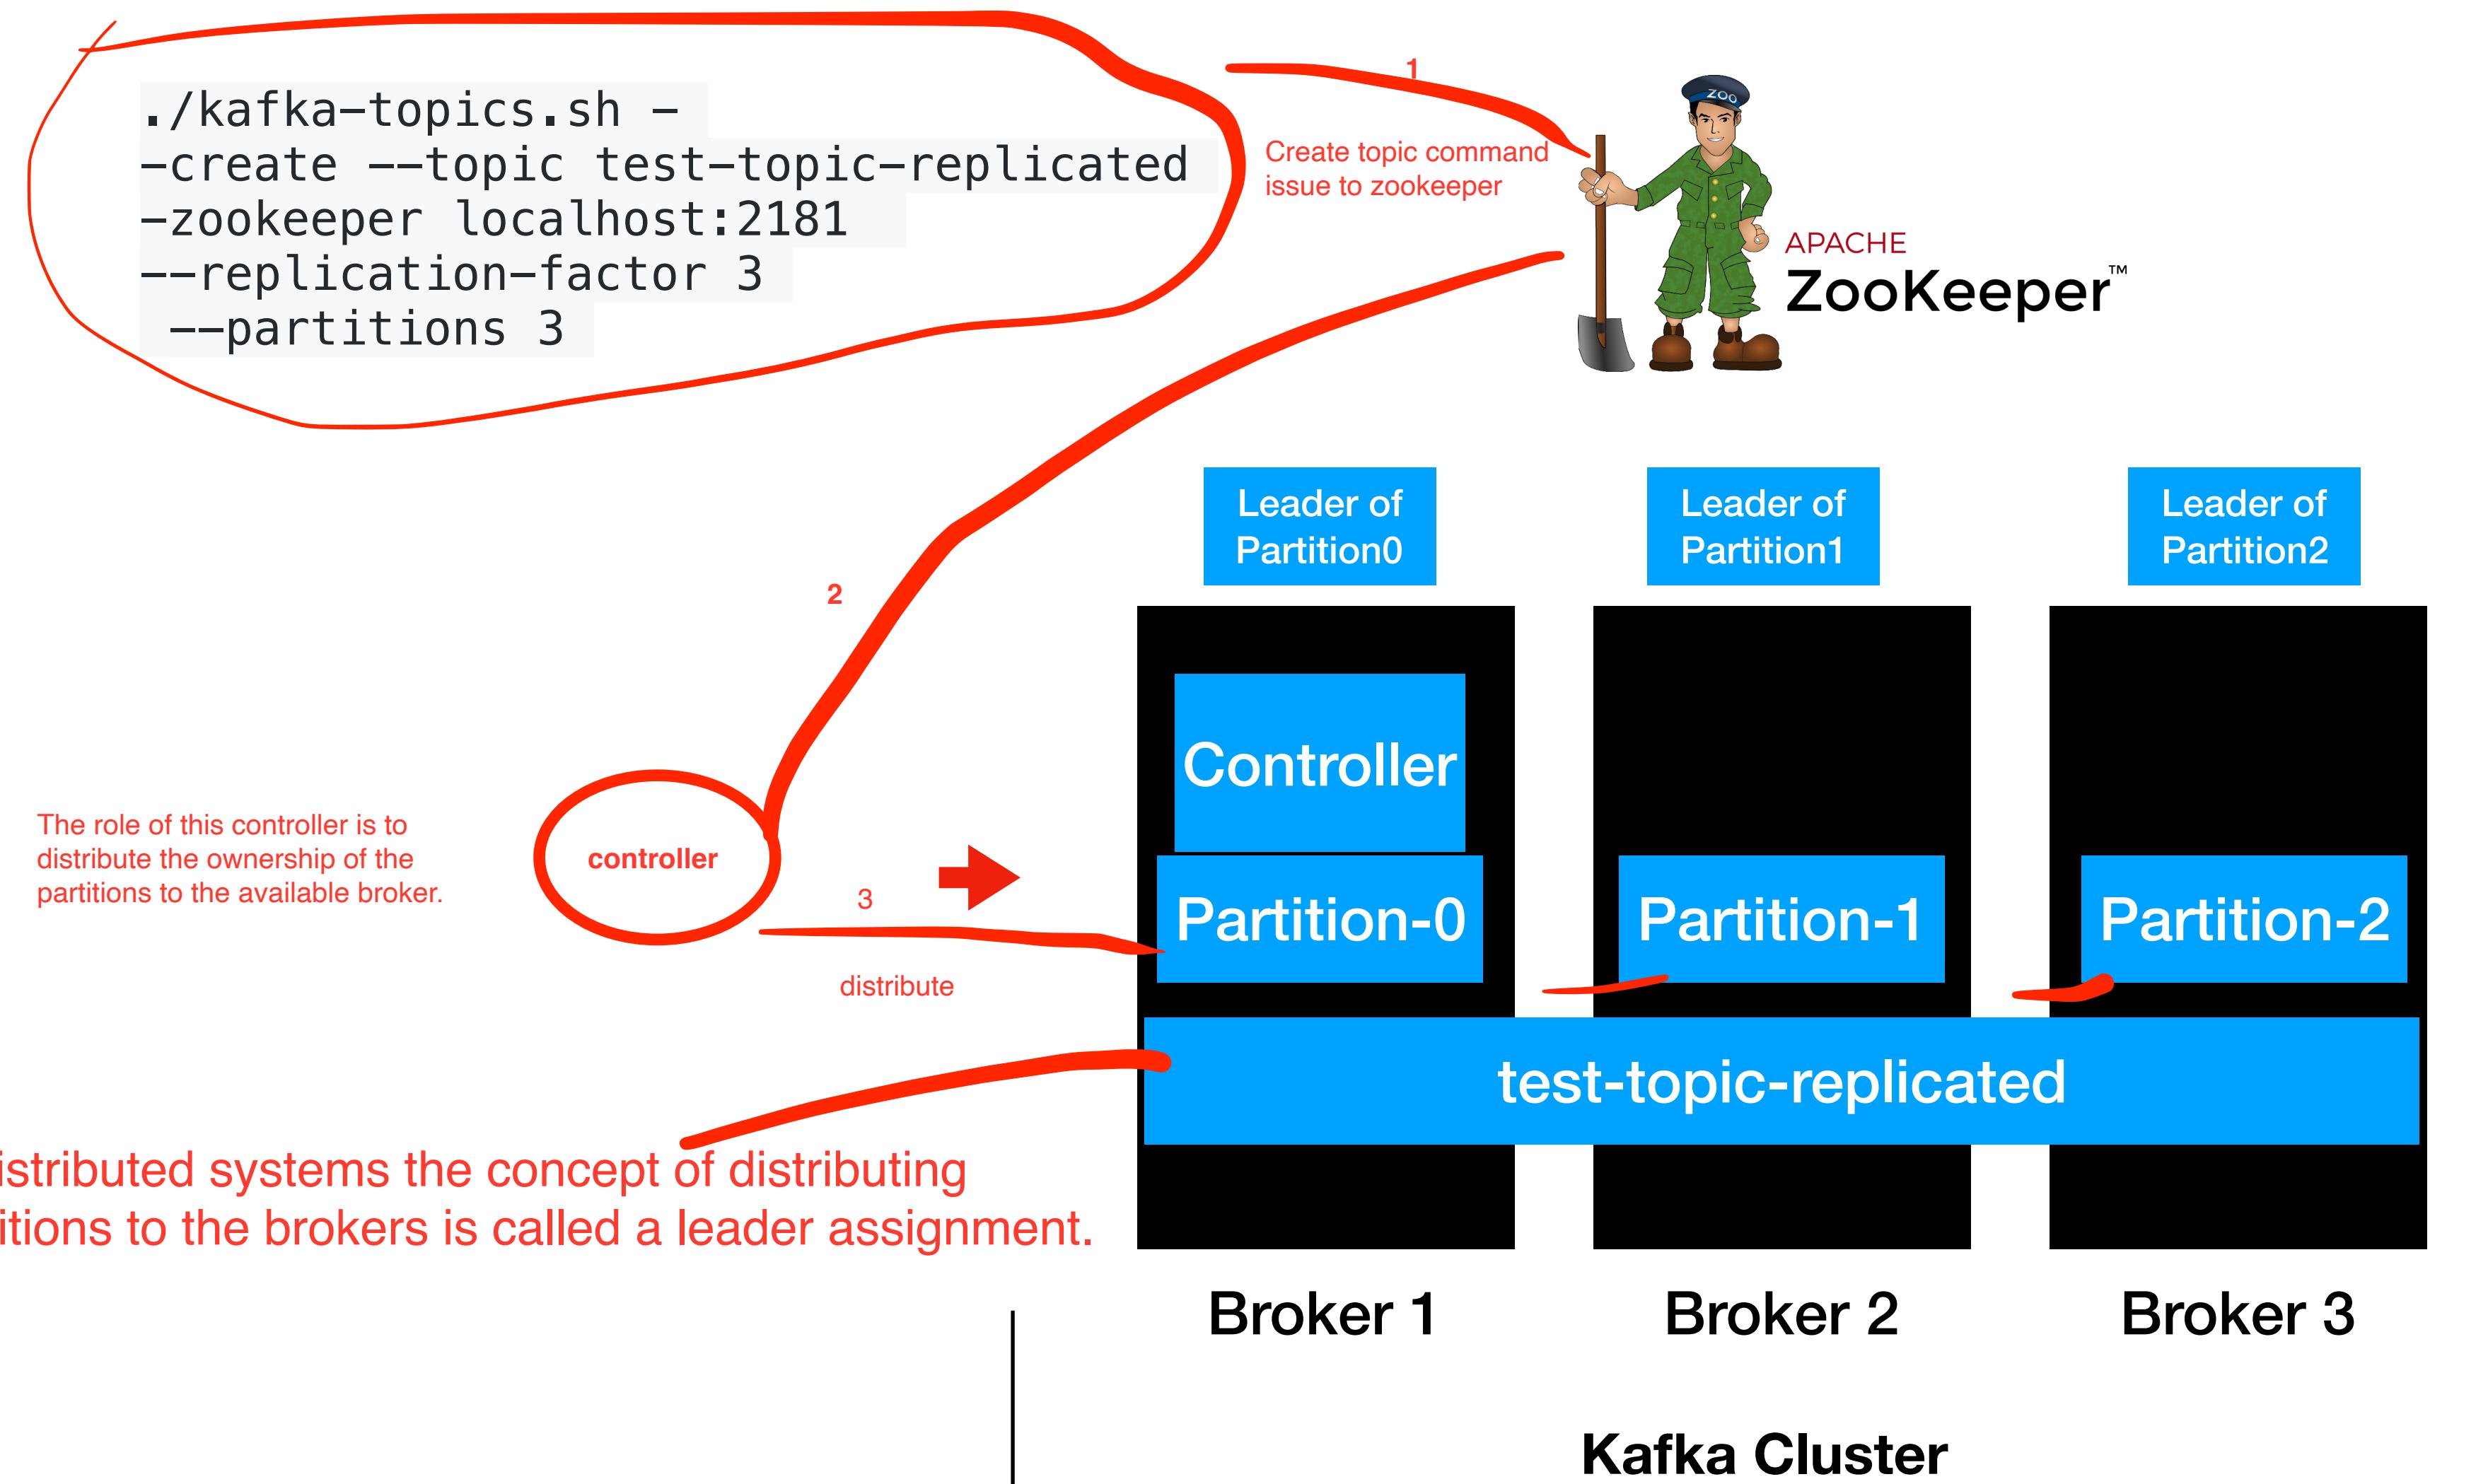
Similarly create “server-2” properties which contains broker id=2 and port 9094 and different log directory

Start all 3 Kafka brokers  
Once all the brokers are started, start create topic

```
$bin> ./kafka-topics.sh --create --topic test-topic-replicated -zookeeper localhost:2181 --replication-factor 3 --partitions 3  
here, “—replication-factor 3” is the number which is <= number of clusters. In ours case we have 3 borkers.
```

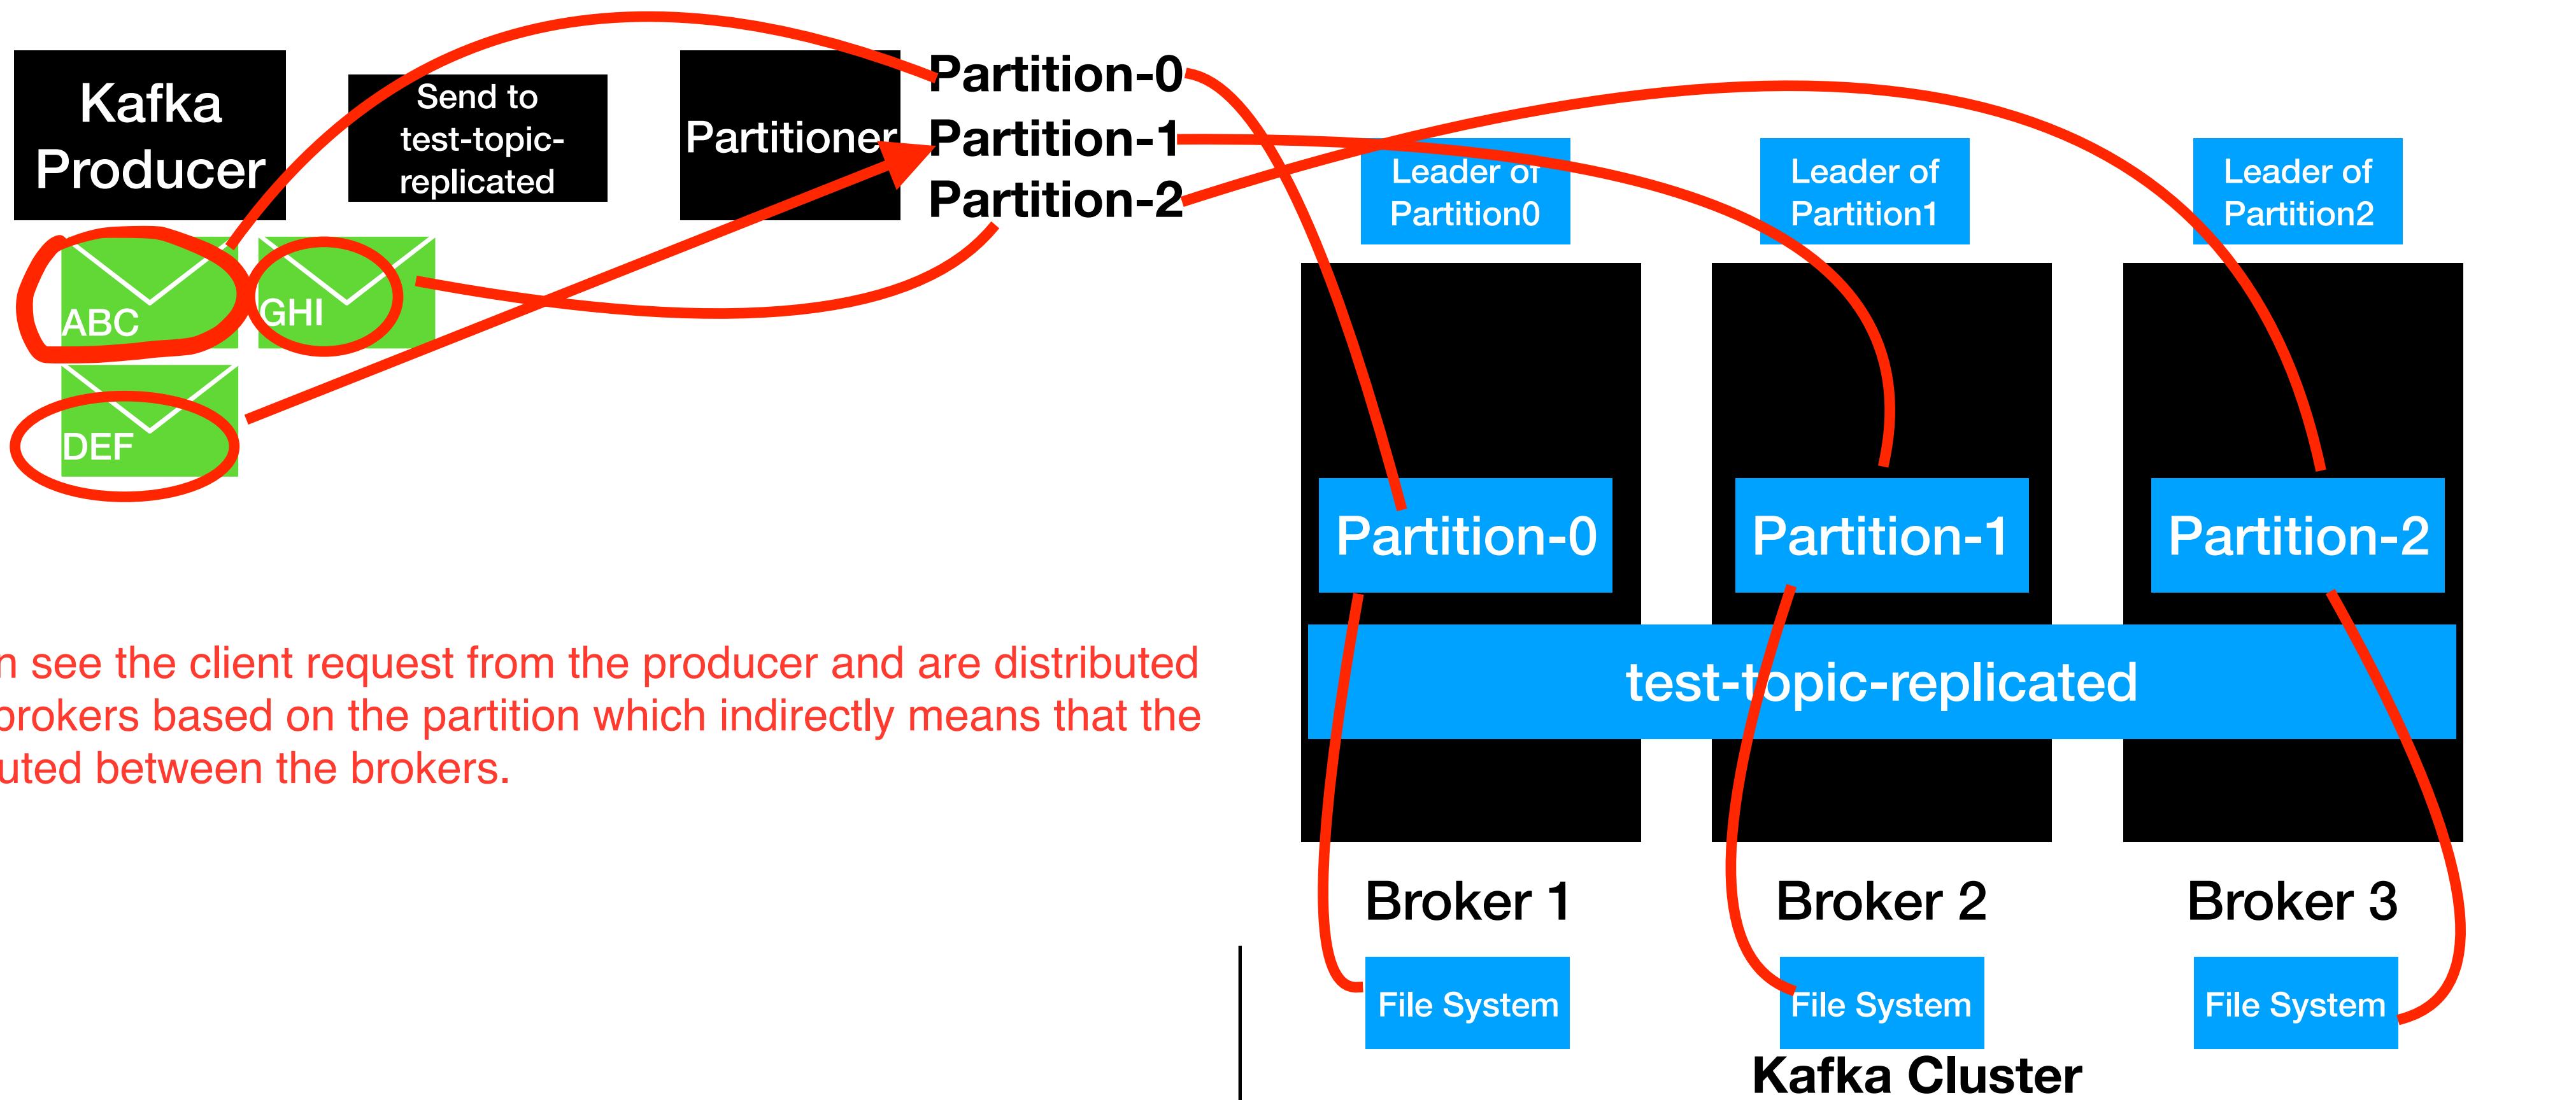
# **How Kafka Distributes the Client Requests?**

# How Topics are distributed?



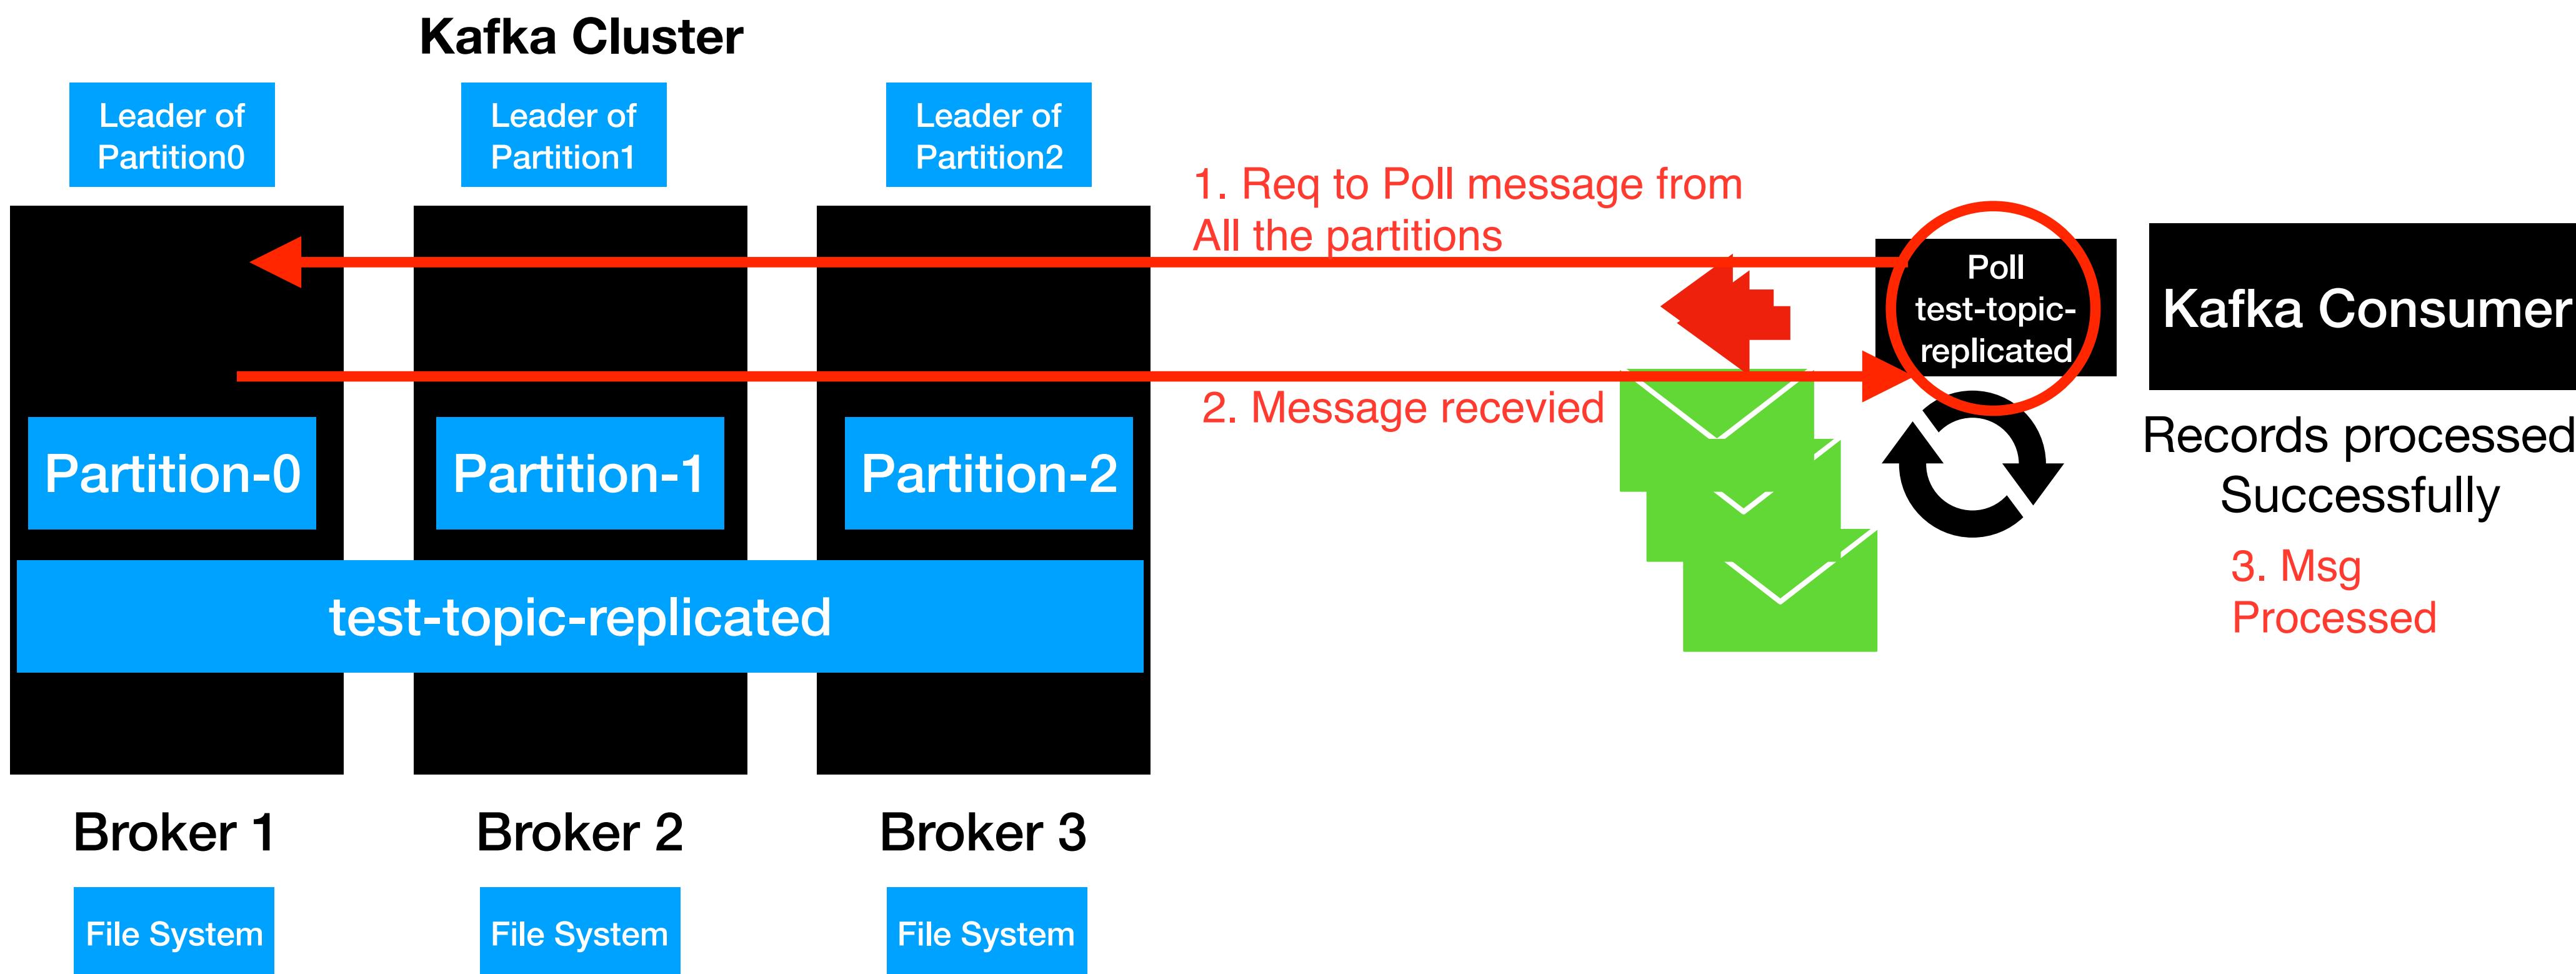
# How Kafka Distributes Client Requests?

## Kafka Producer



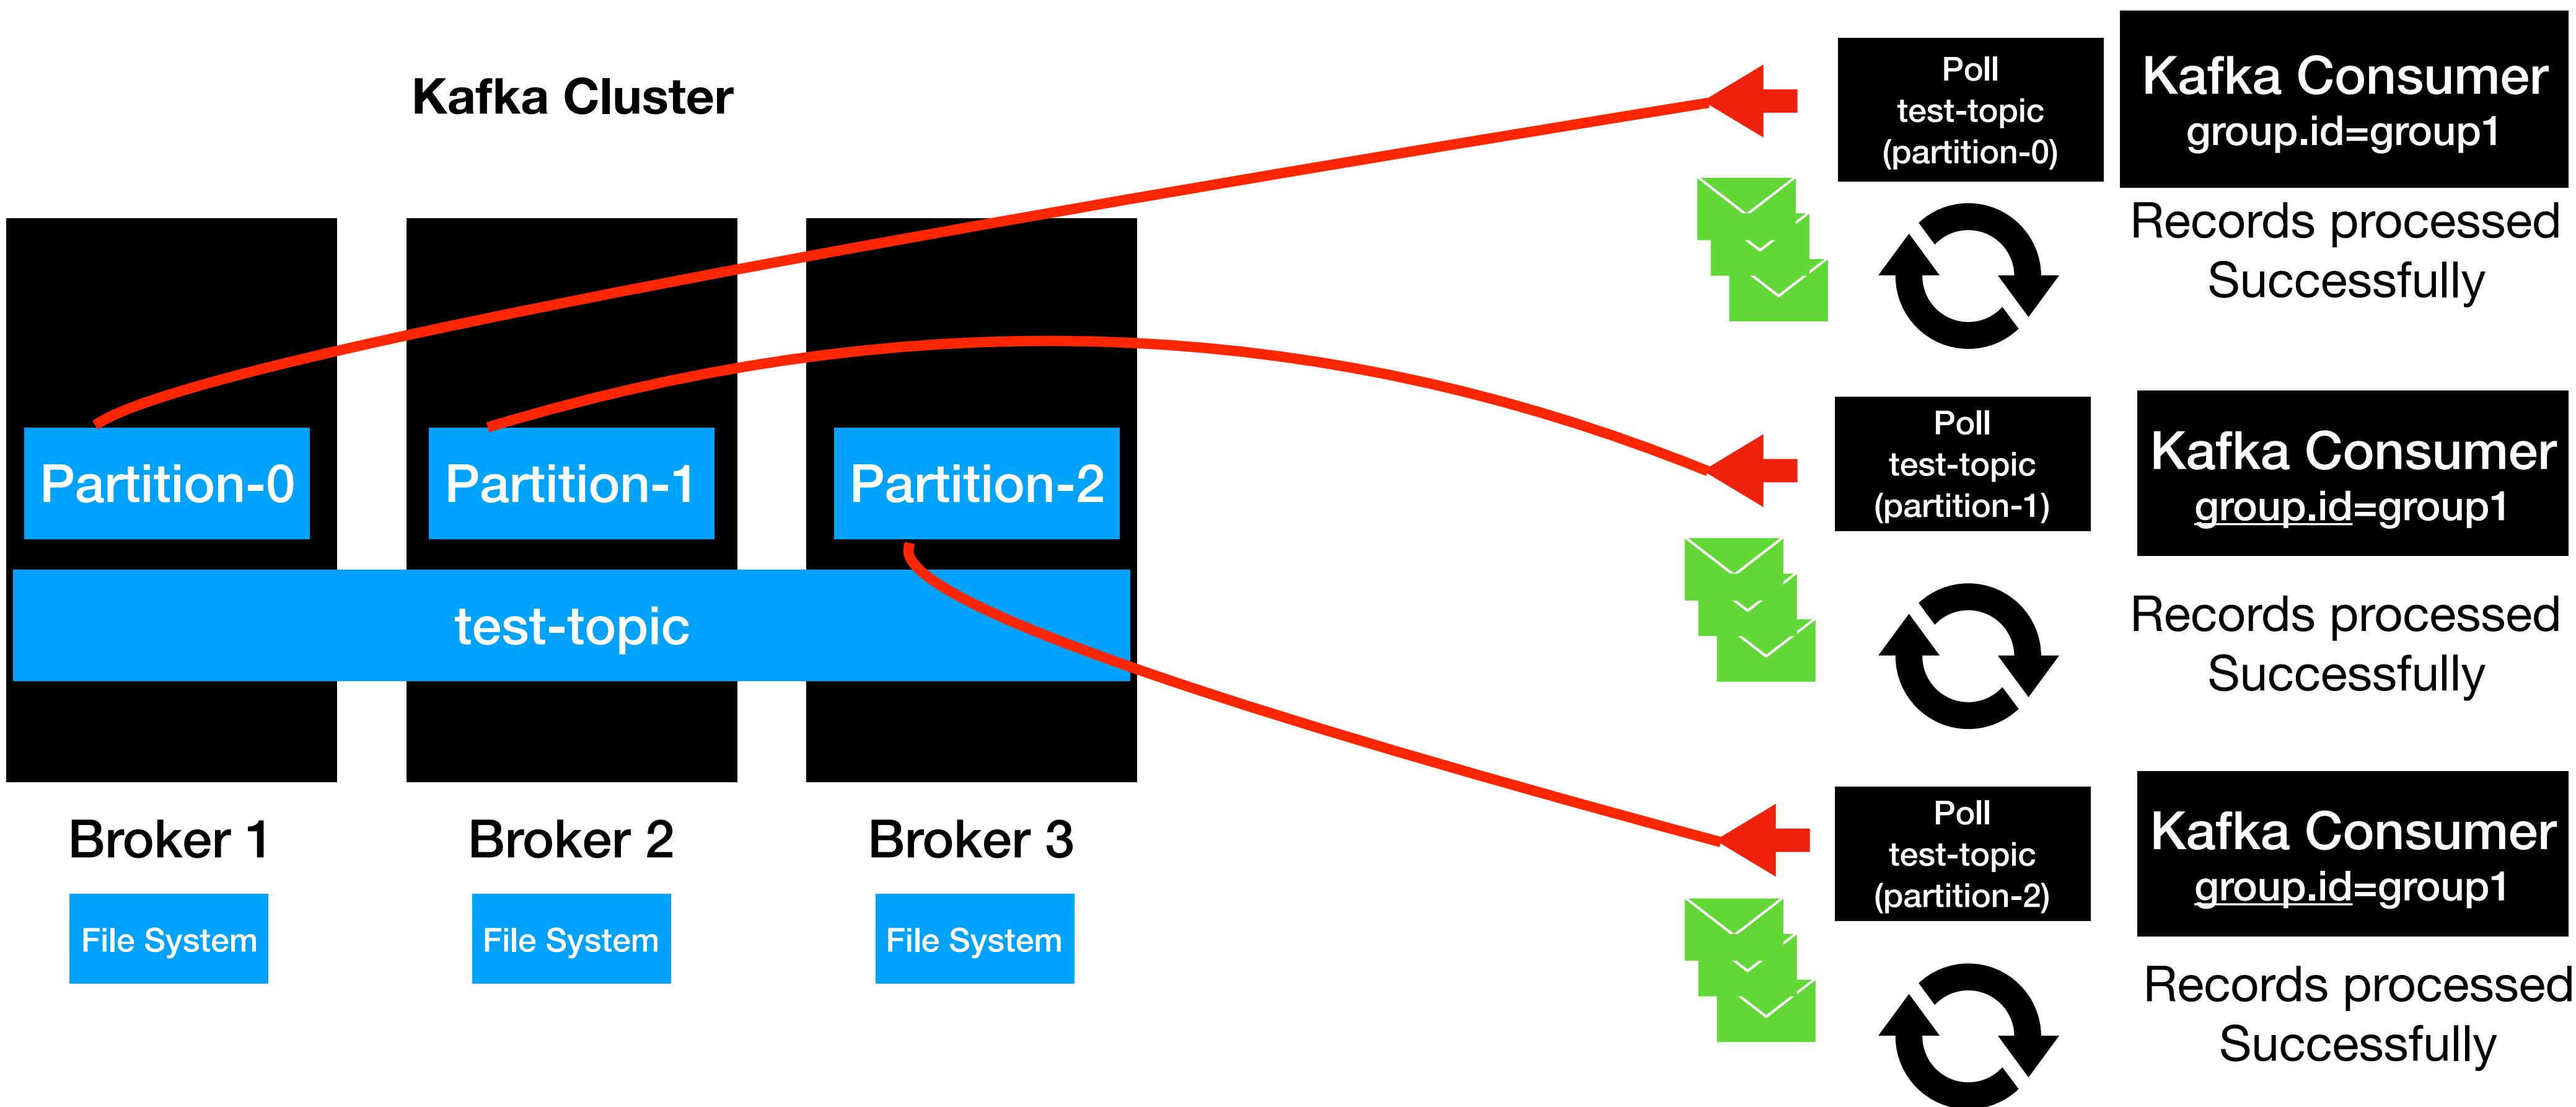
# How Kafka Distributes Client Requests?

## Kafka Consumer



# How Kafka Distributes Client Requests?

## Kafka Consumer Groups

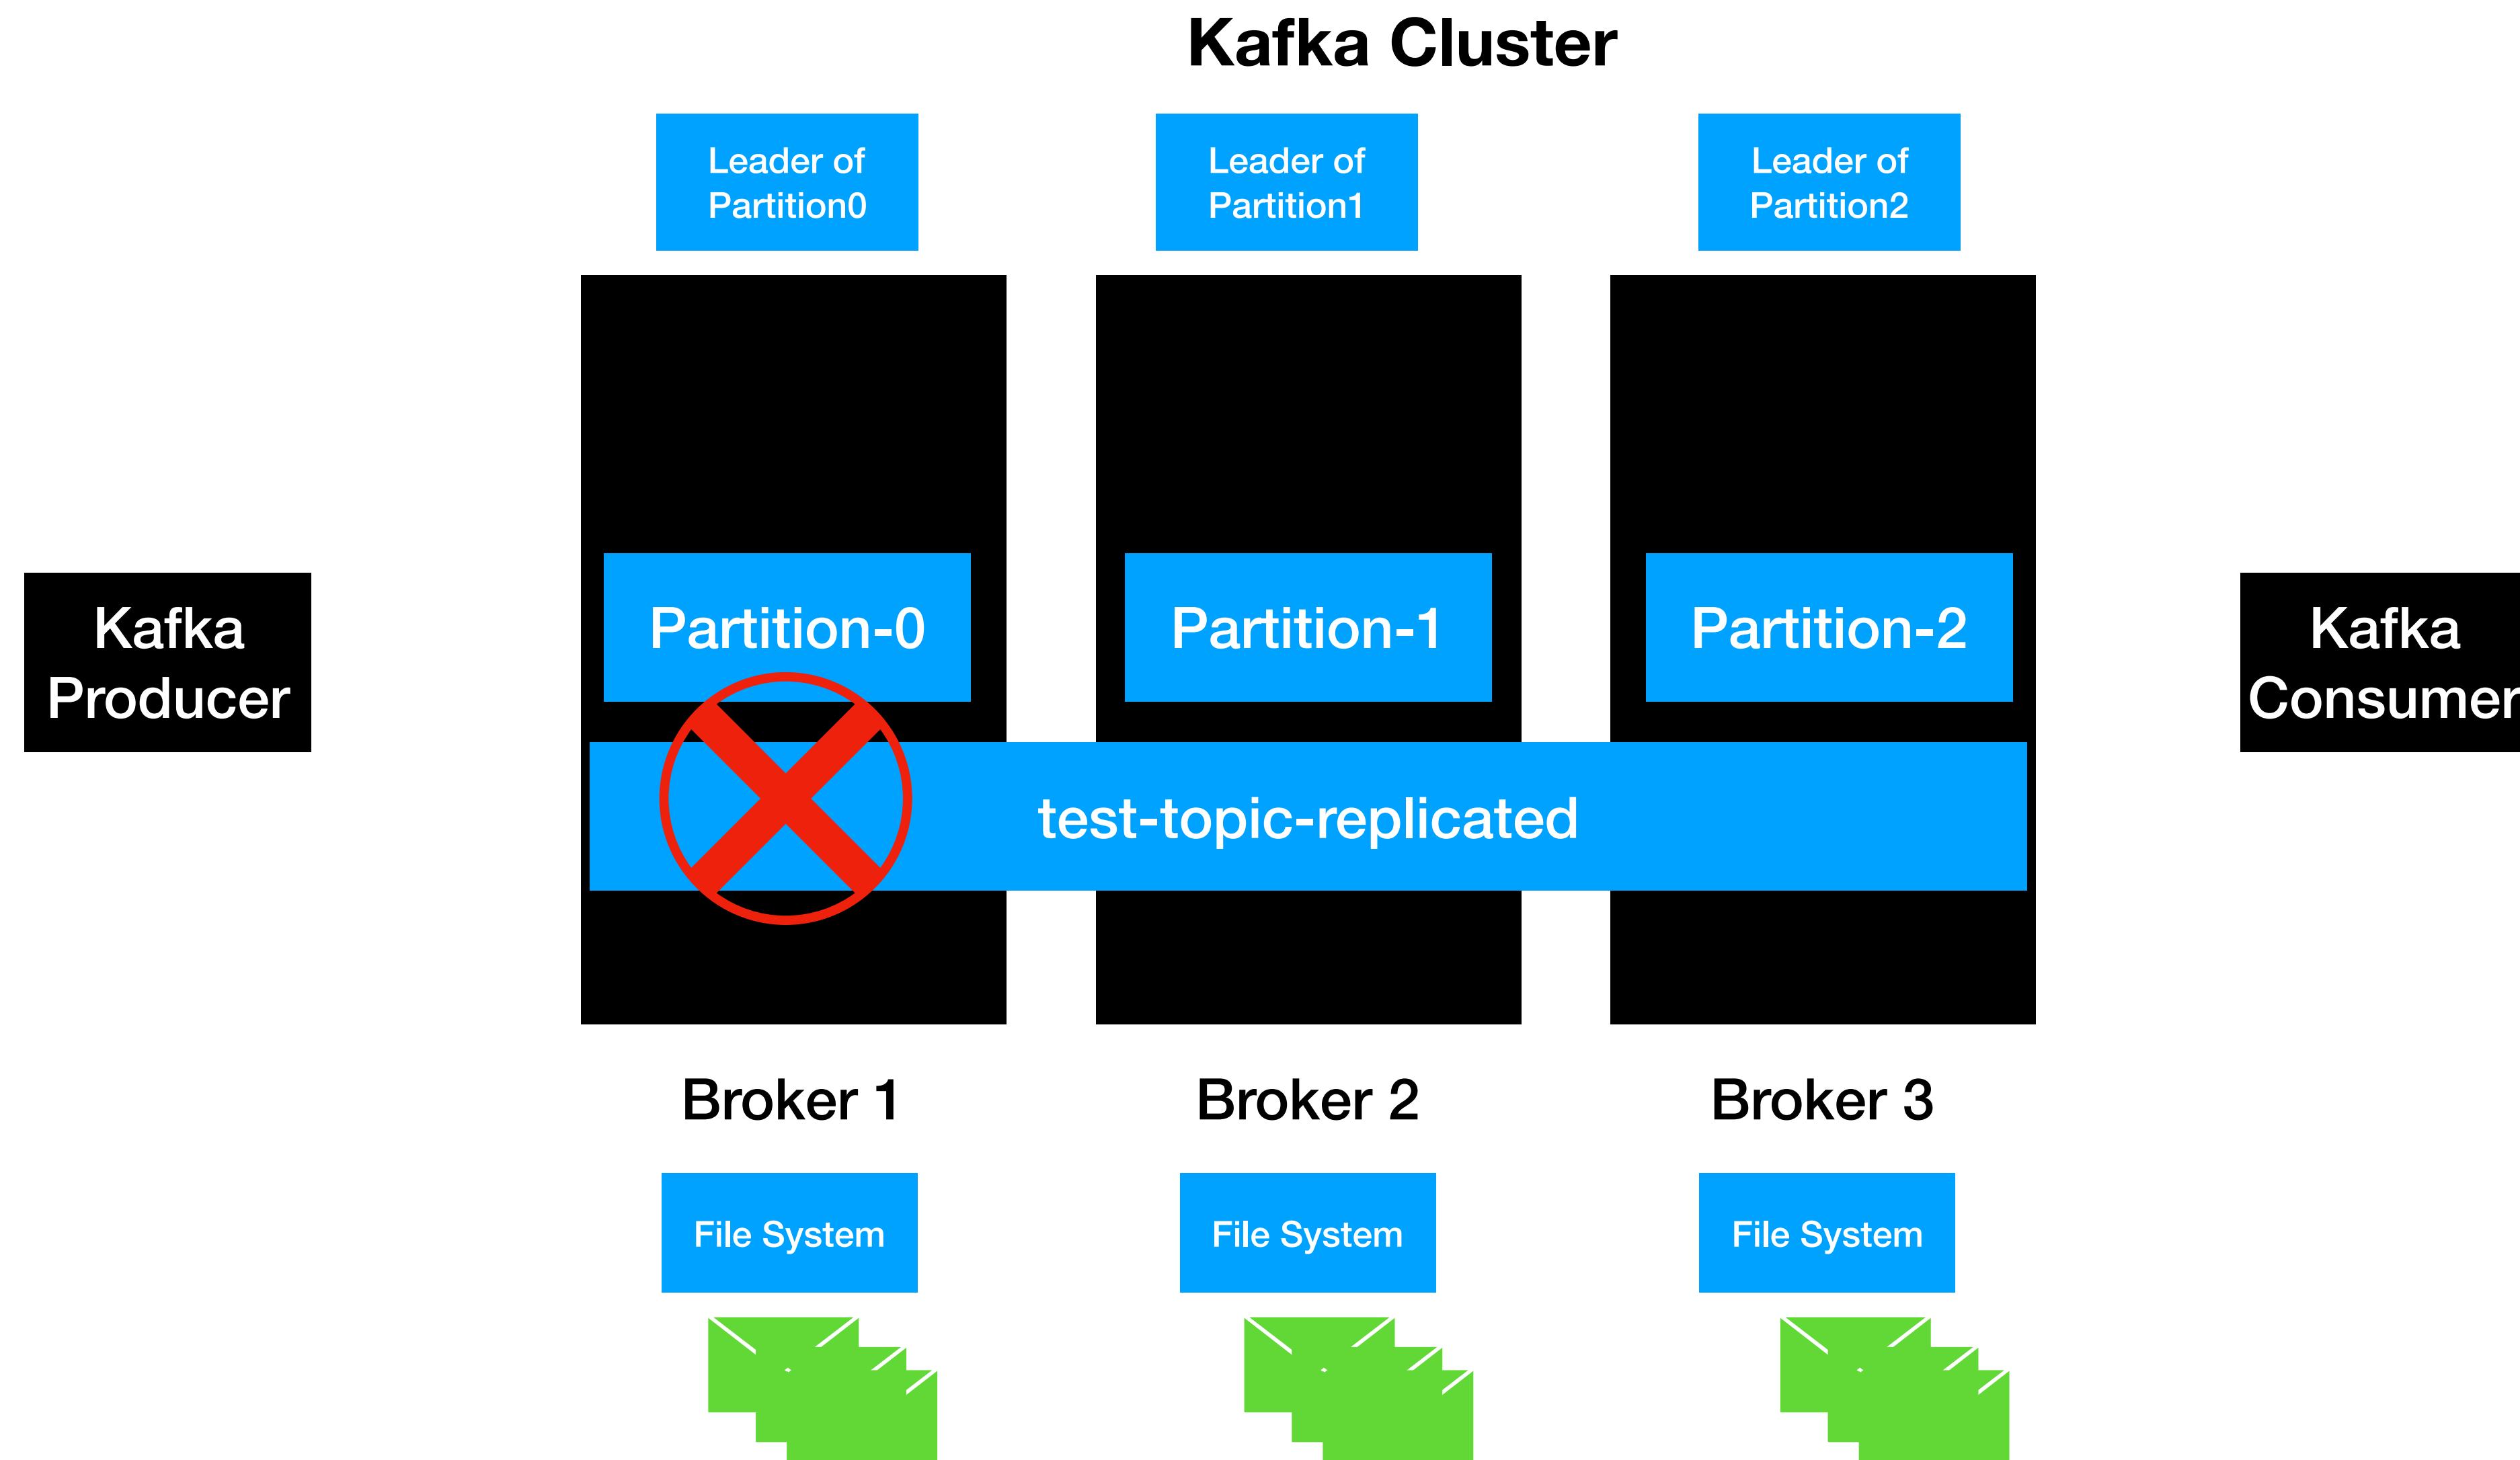


# **Summary : How Kafka Distributes the Client Requests?**

- Partition leaders are assigned during topic Creation
- Clients will only invoke leader of the partition to produce and consume data
  - Load is evenly distributed between the brokers

# **How Kafka handles Data Loss ?**

# How Kafka handles Data loss?

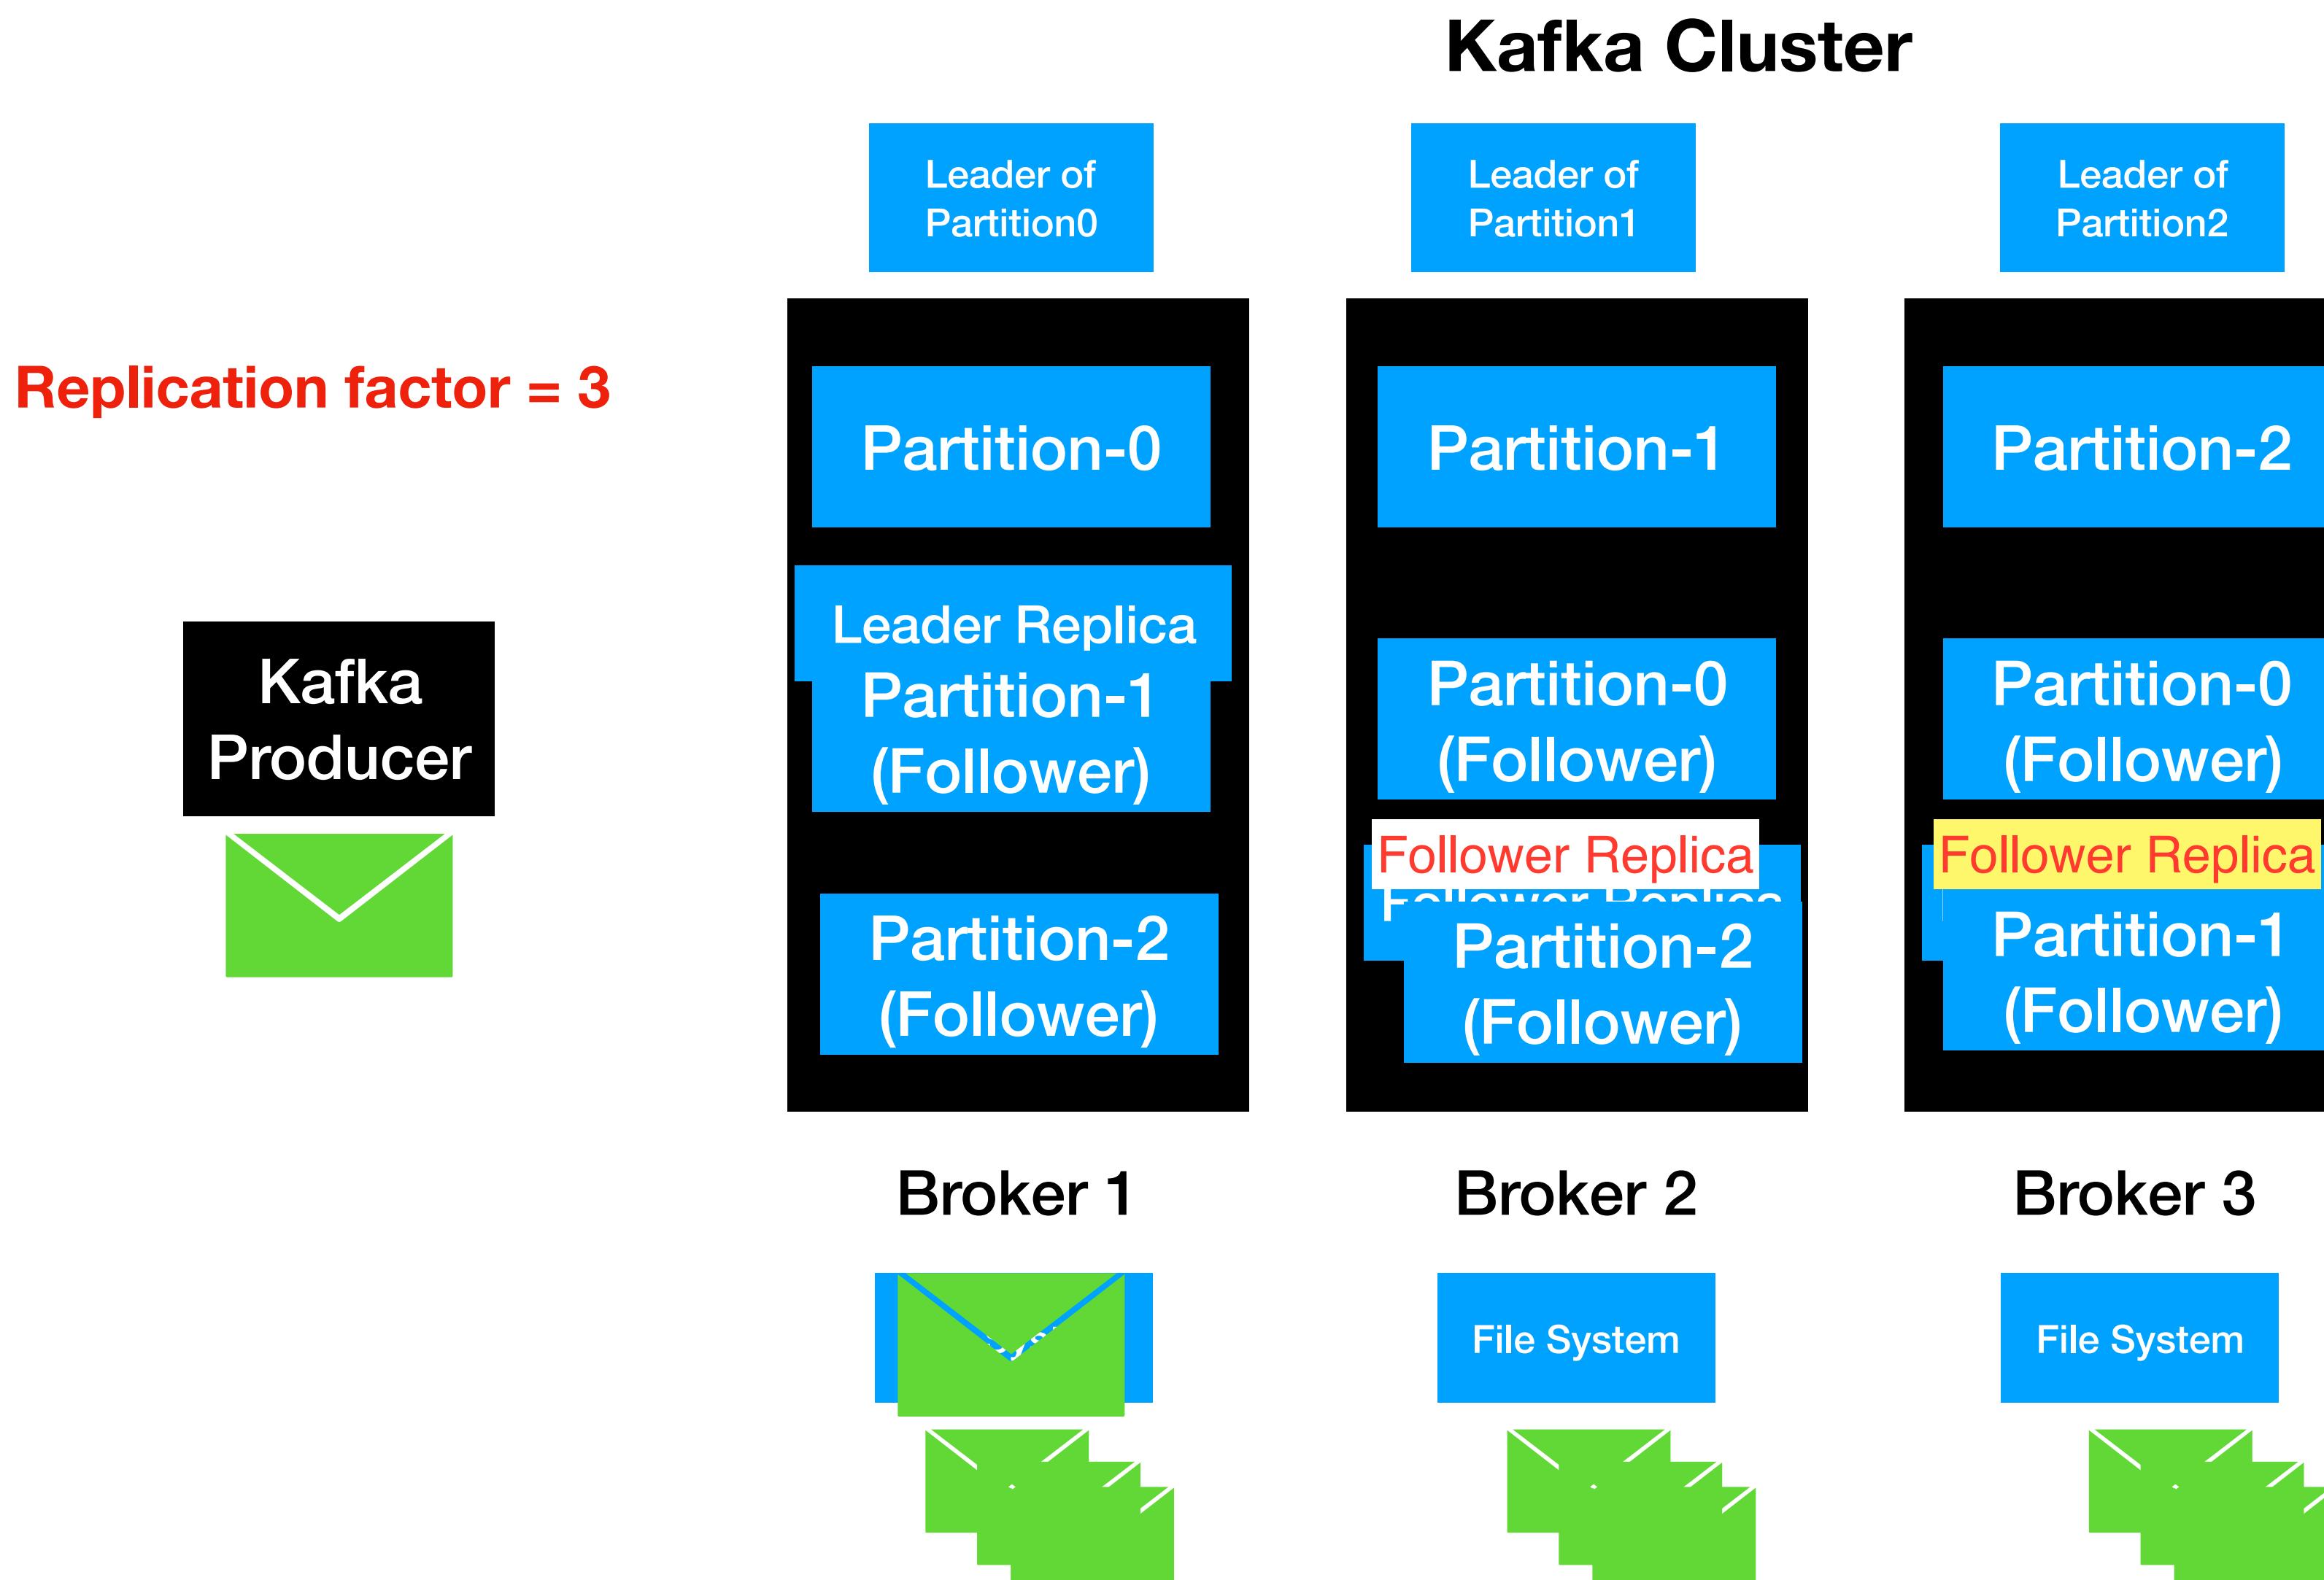


# Replication

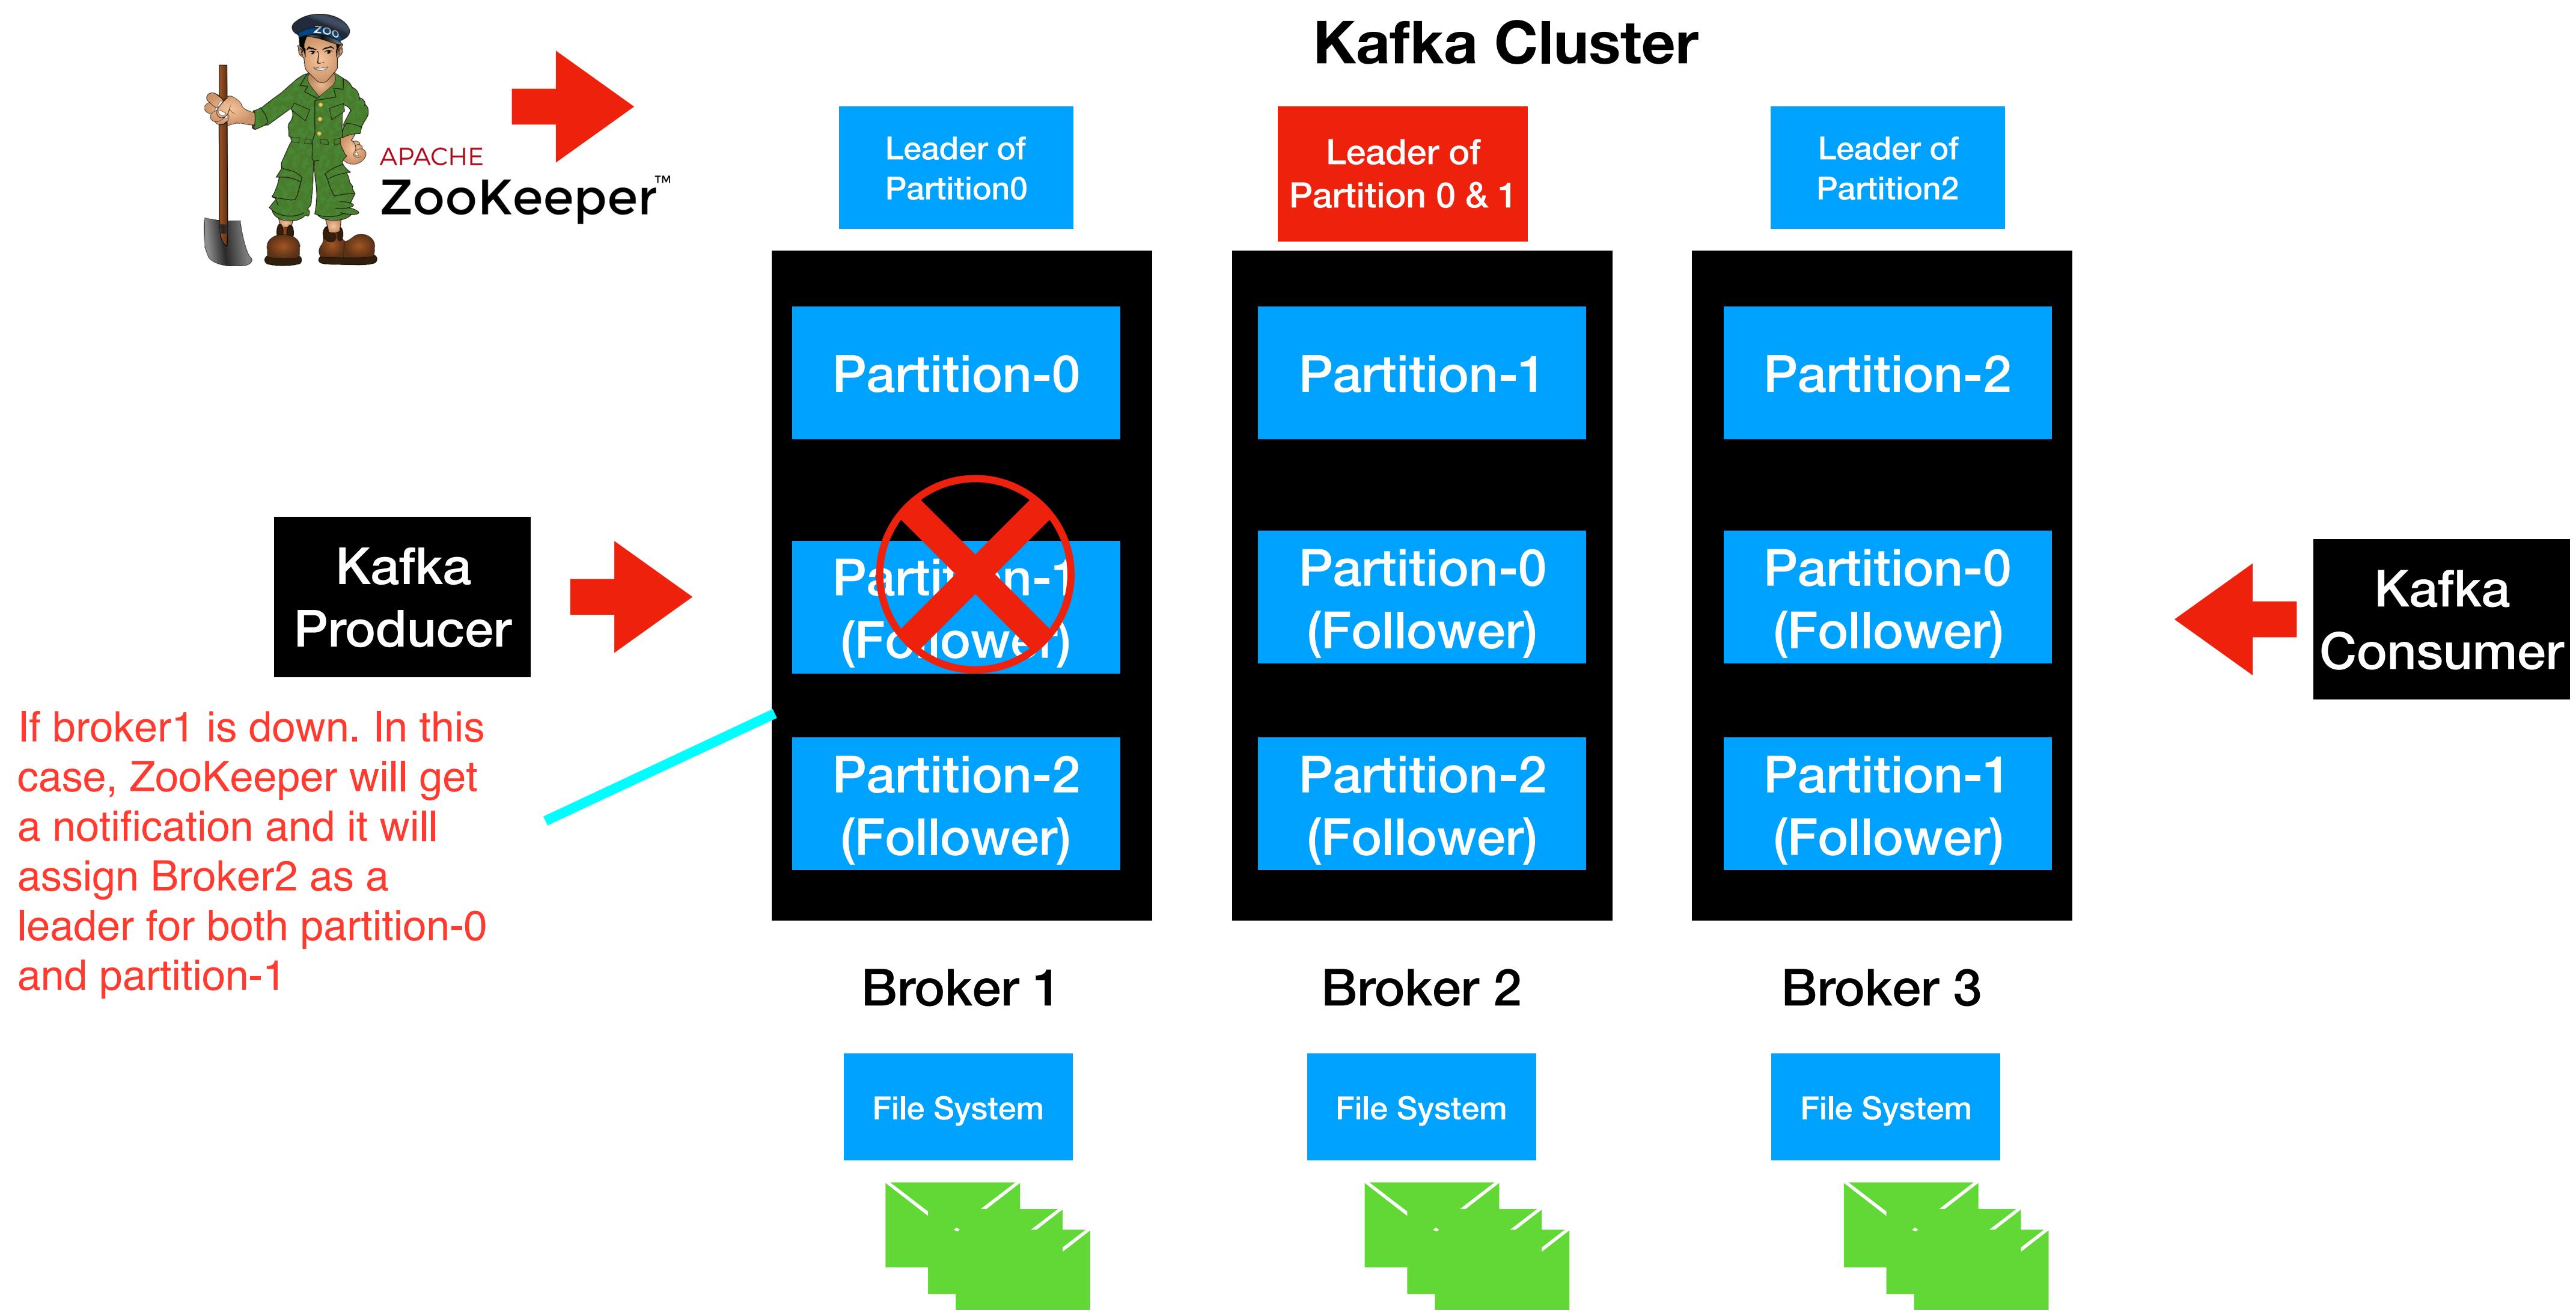
```
./kafka-topics.sh -  
-create --topic test-topic-replicated  
-zookeeper localhost:2181  
--replication-factor 3  
--partitions 3
```

So replication factors equal to number of copies  
of the same message.

# Replication



# Replication



# In-Sync Replica(ISR)

- Represents the number of replica in sync with each other in the cluster
  - Includes both **leader** and **follower** replica
- Recommended value is always greater than 1
- Ideal value is **ISR == Replication Factor**
- This can be controlled by **min.insync.replicas** property
  - It can be set at the **broker** or **topic** level

ISR: will be used to sync the data across all the available broker. Suppose you have 3 brokers, namely 0,1,2. If 2nd broker is down, the In-Sync Replica will point to other brokers automatically based on the signal provided by ZooKeeper.

```
./kafka-topics.sh --zookeeper localhost:2181 --describe --topic test-topic-replicated
```

# Fault Tolerance & Robustness

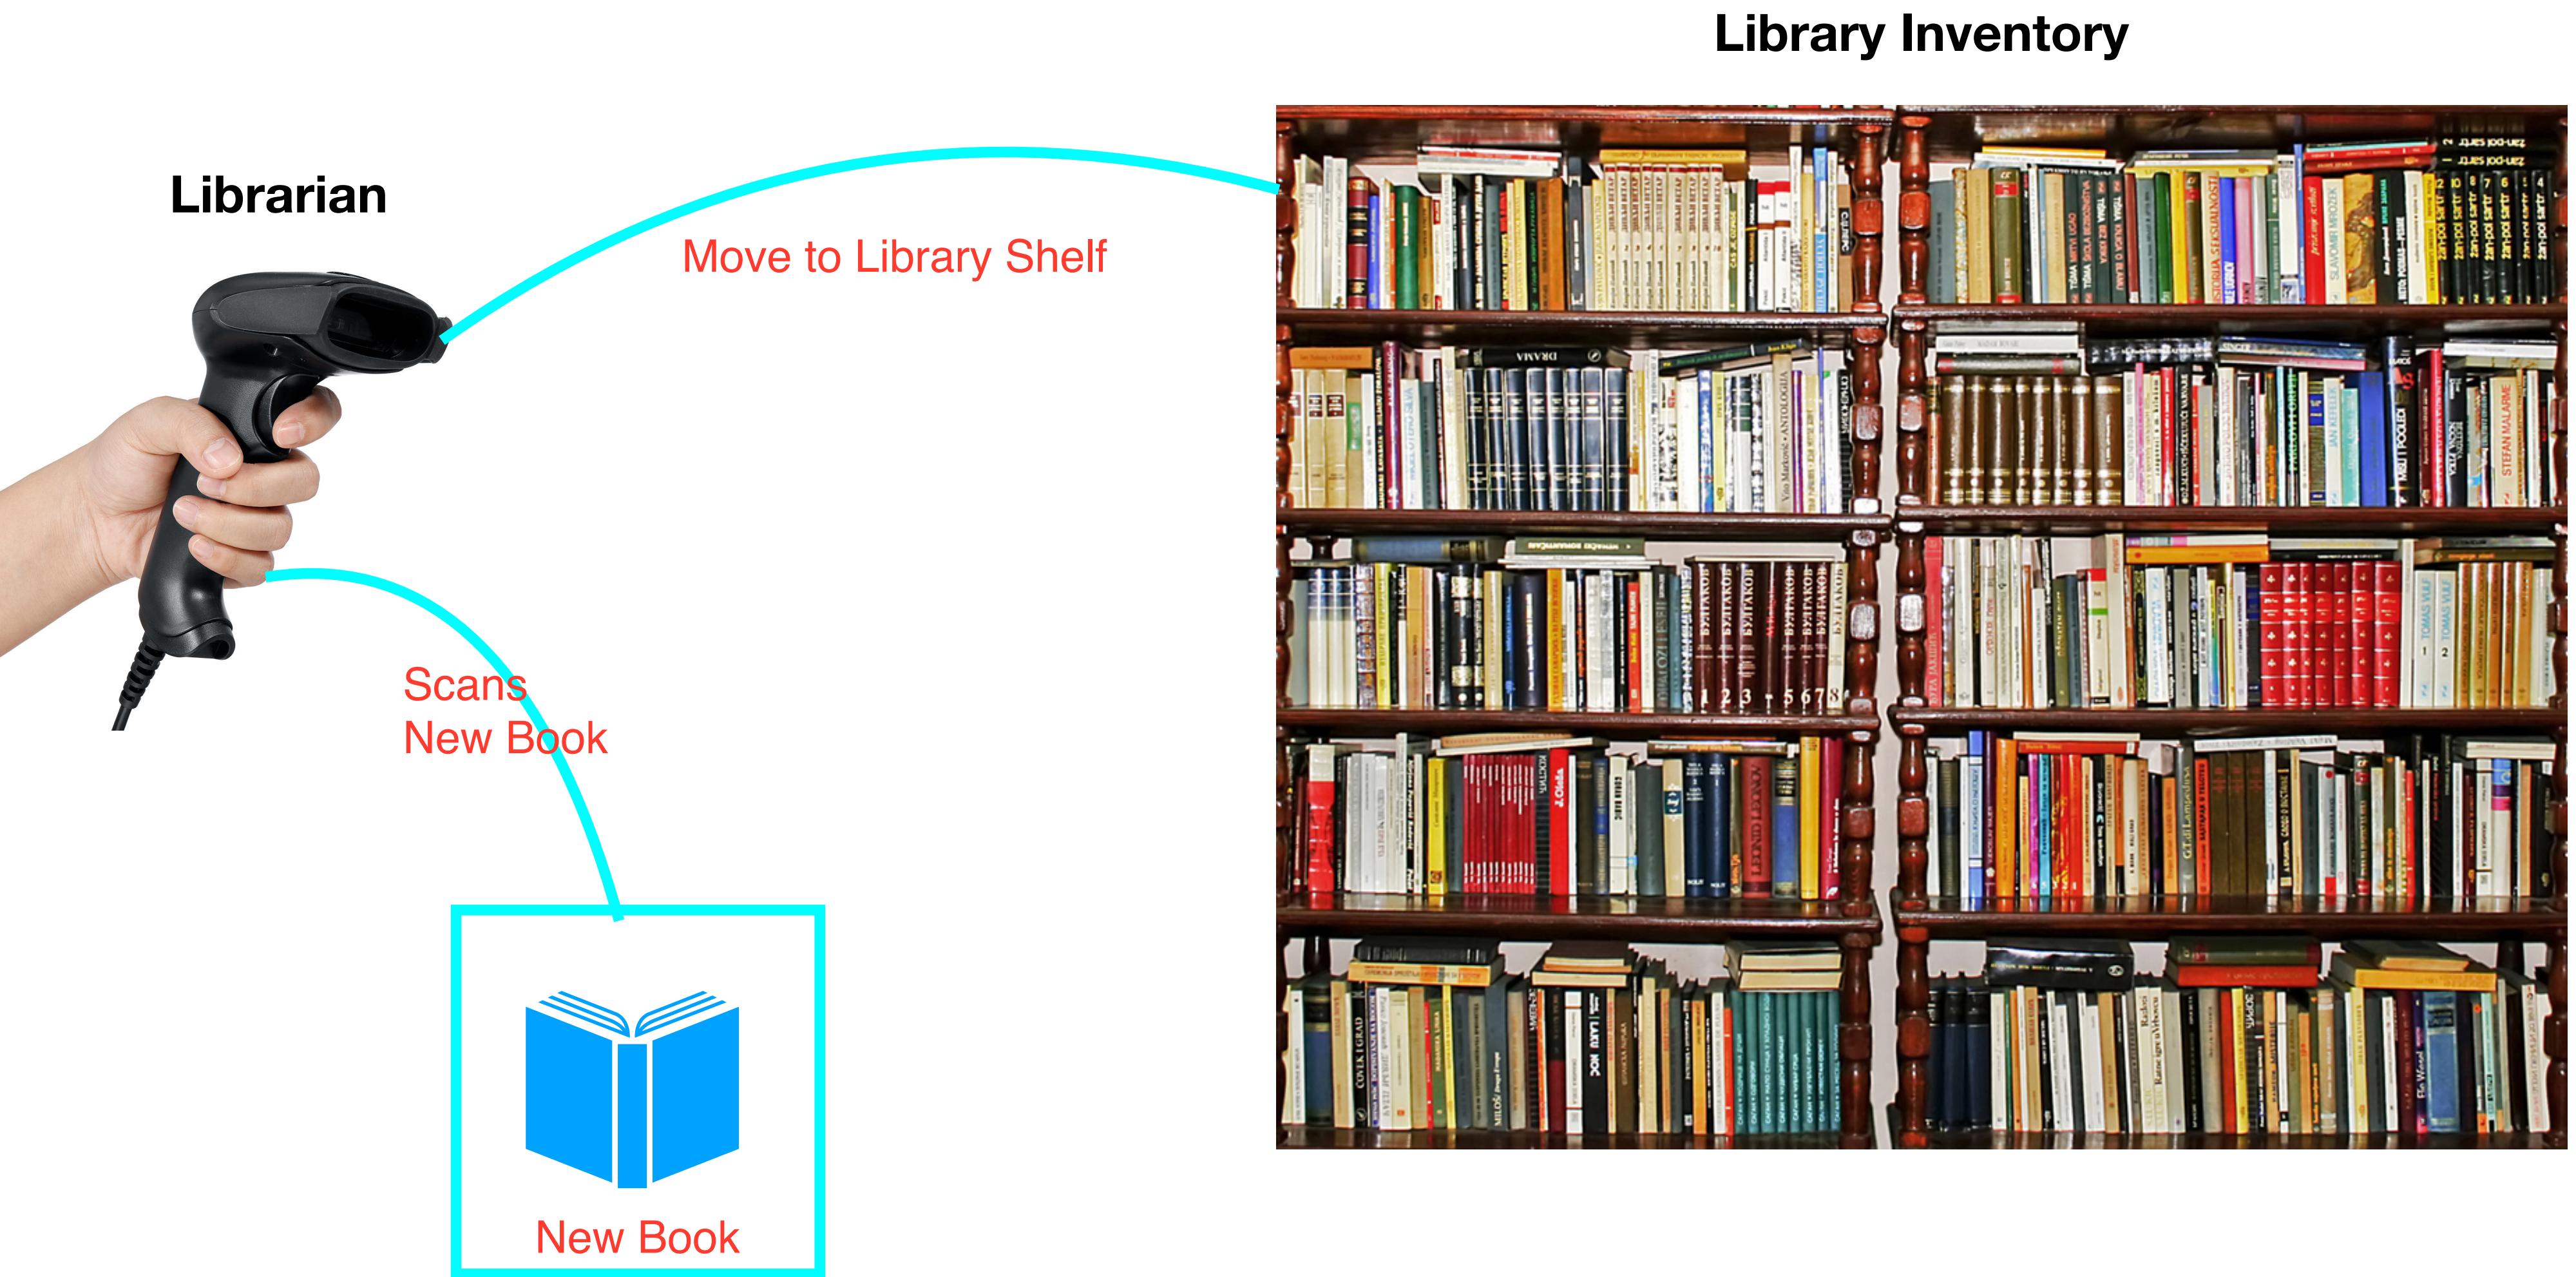
When you have more than one Kafka broker and start producing messages, if consumers connected to your Kafka broker it will start consume messages.  
Suppose, one of the borker is down, it will not impact the consumers. Consumers still can receive the messages. (Robustness)  
Kafka takes care the error handling. From the client perspective you no need to do anything.

# **Application Overview**

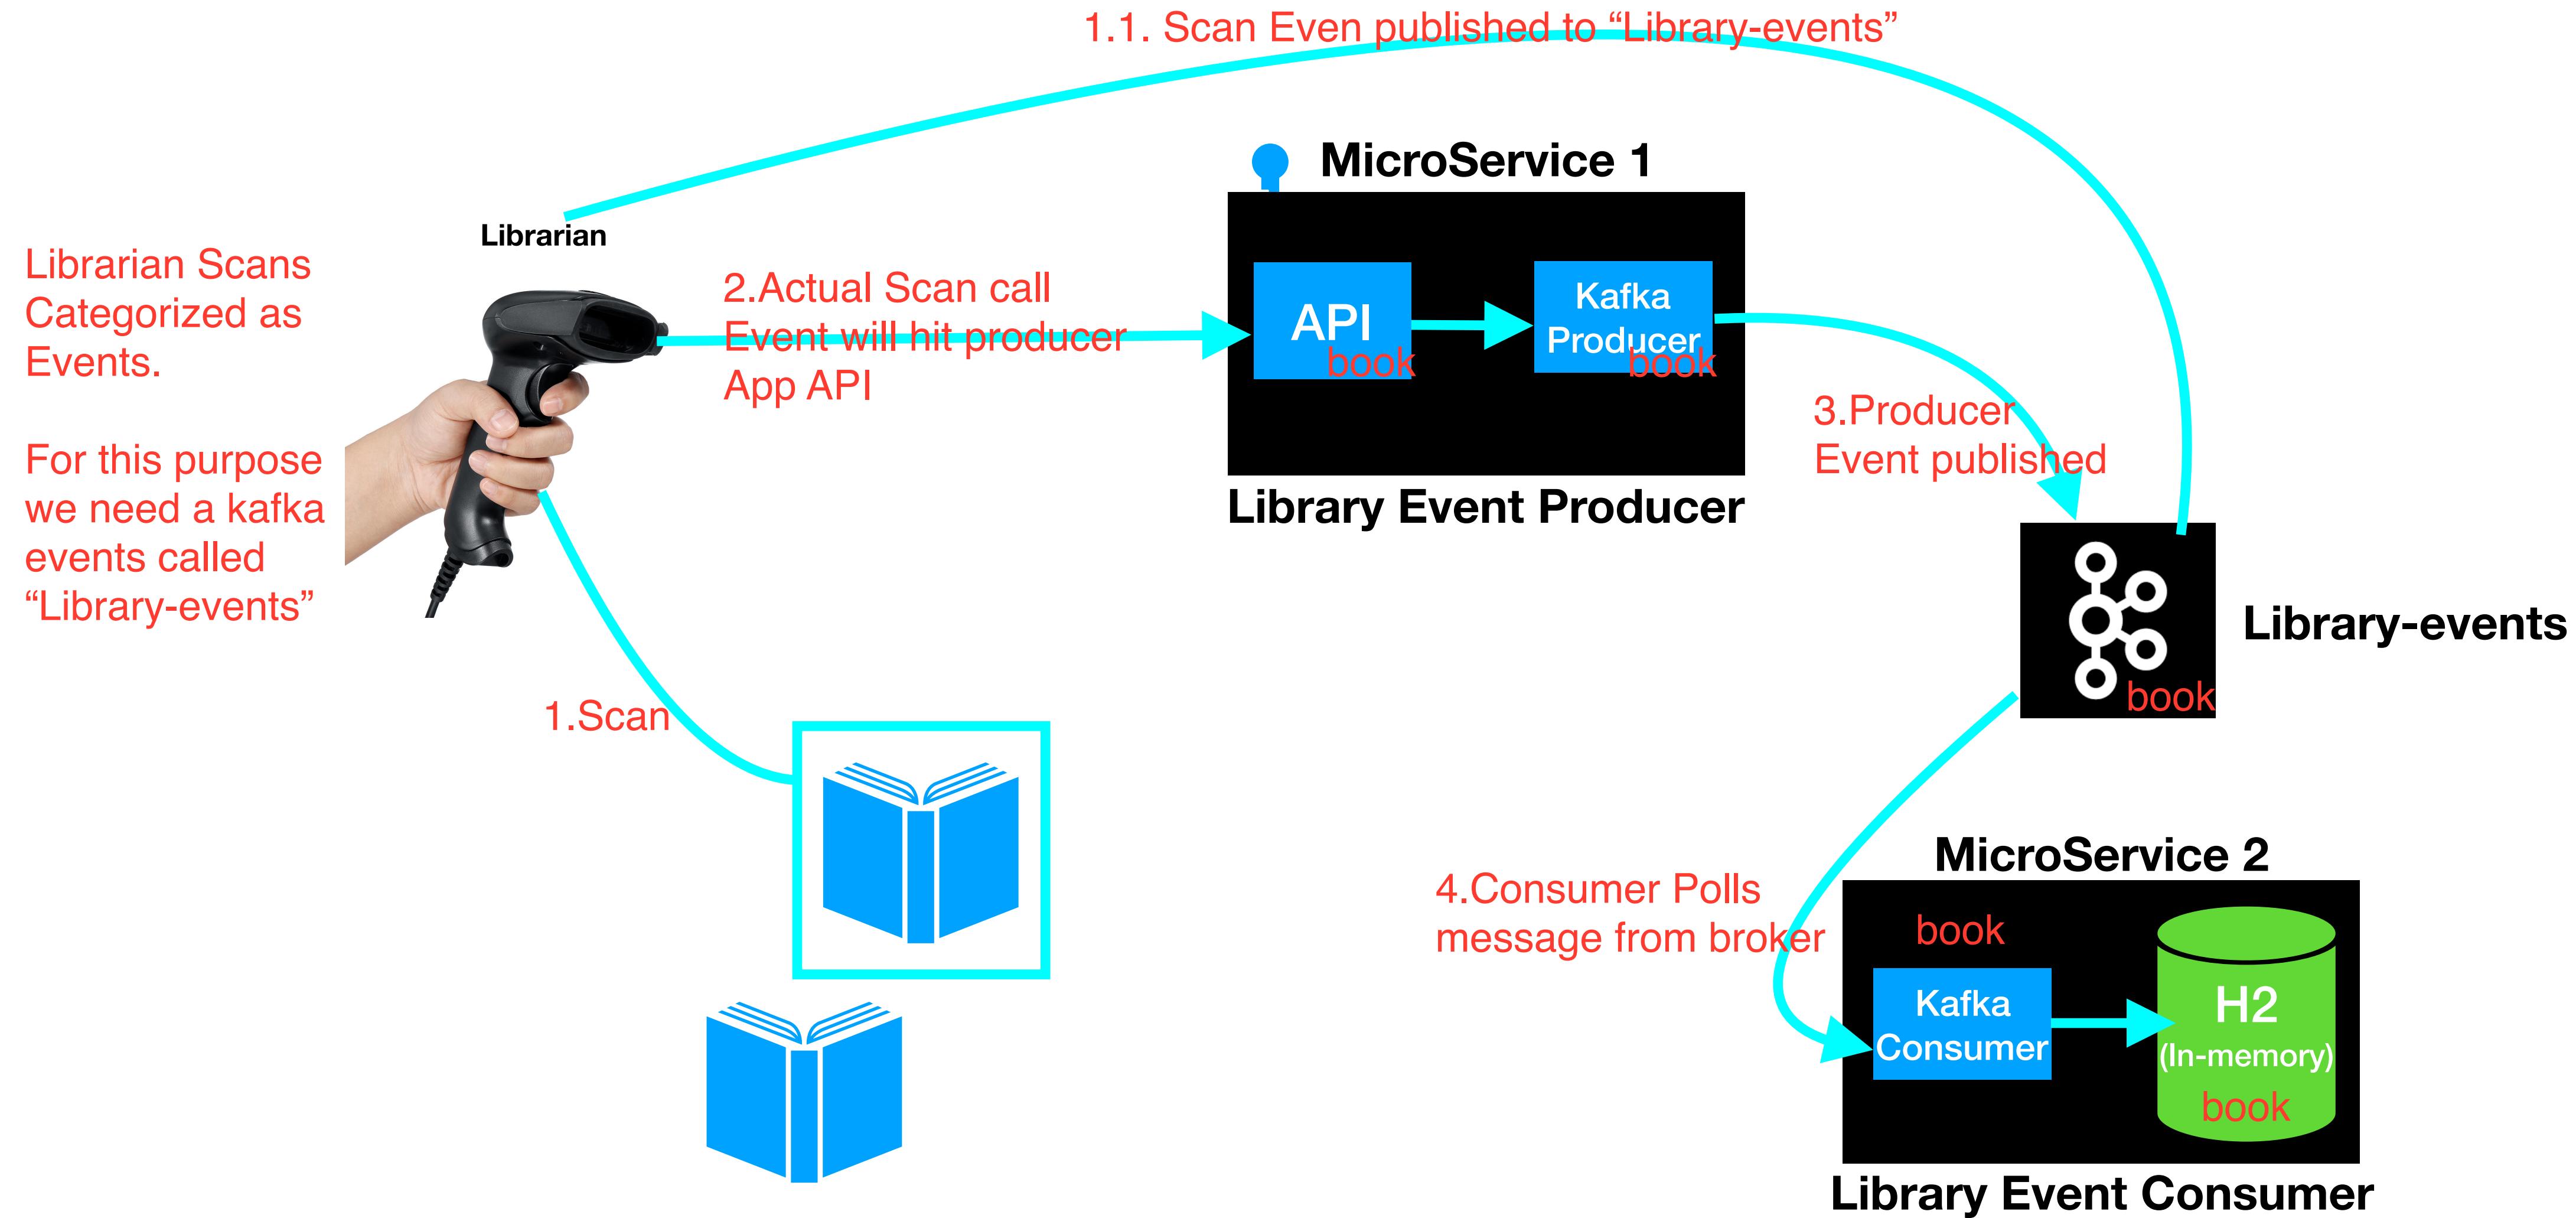
# Library Inventory



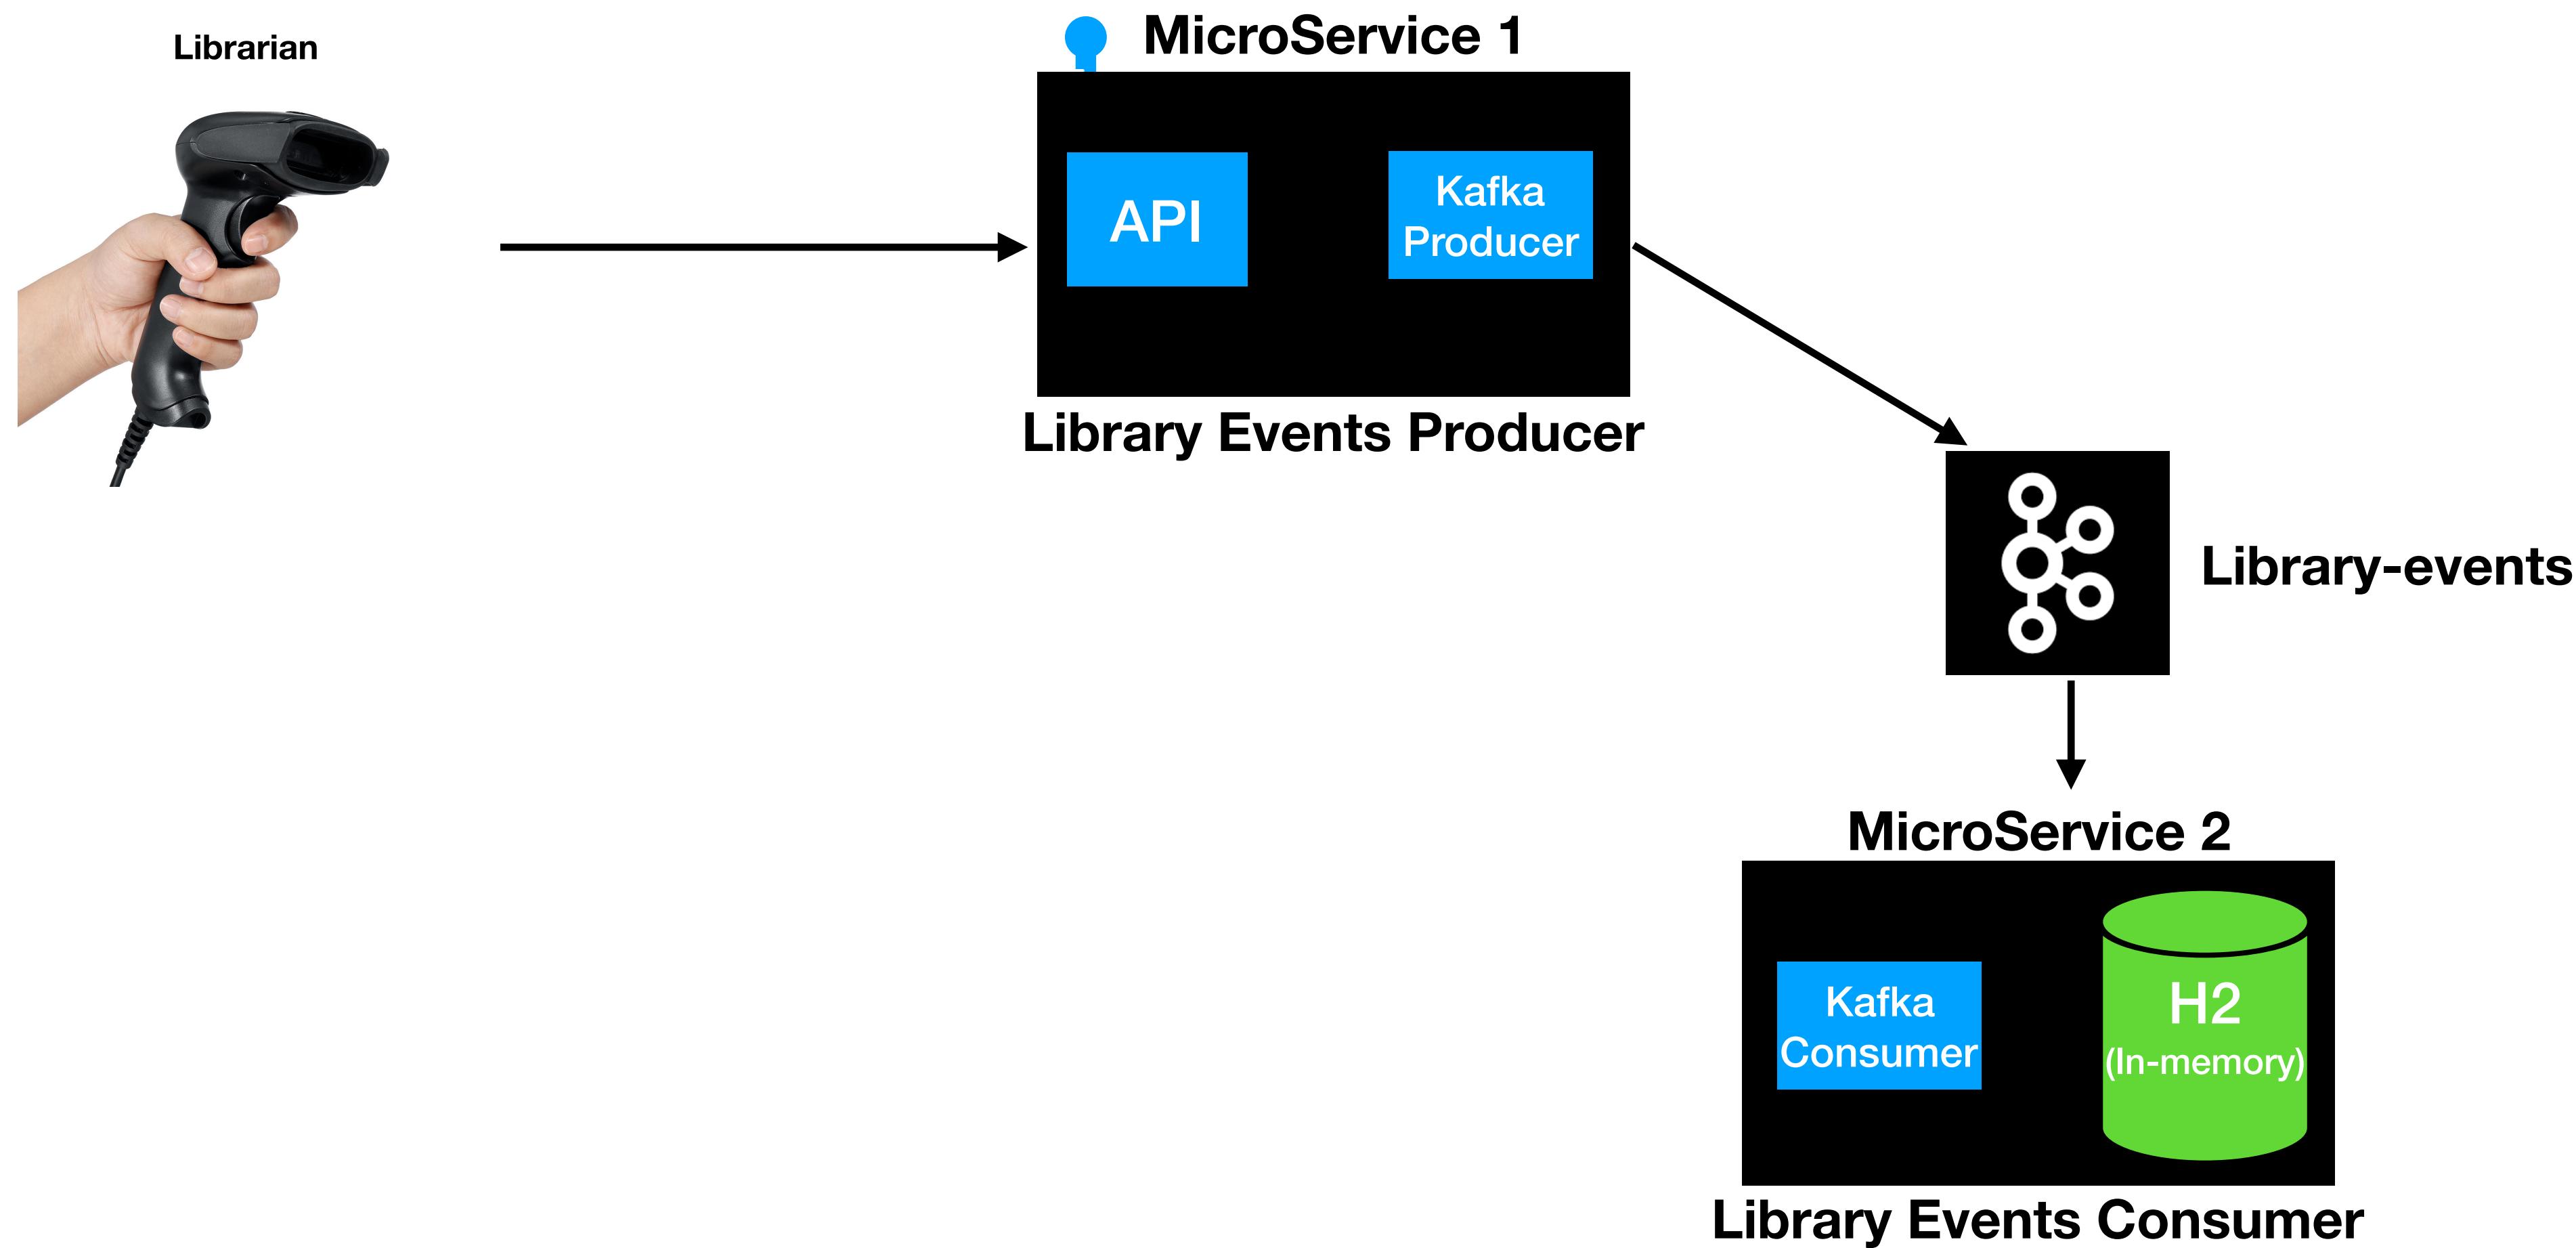
# Library Inventory Flow



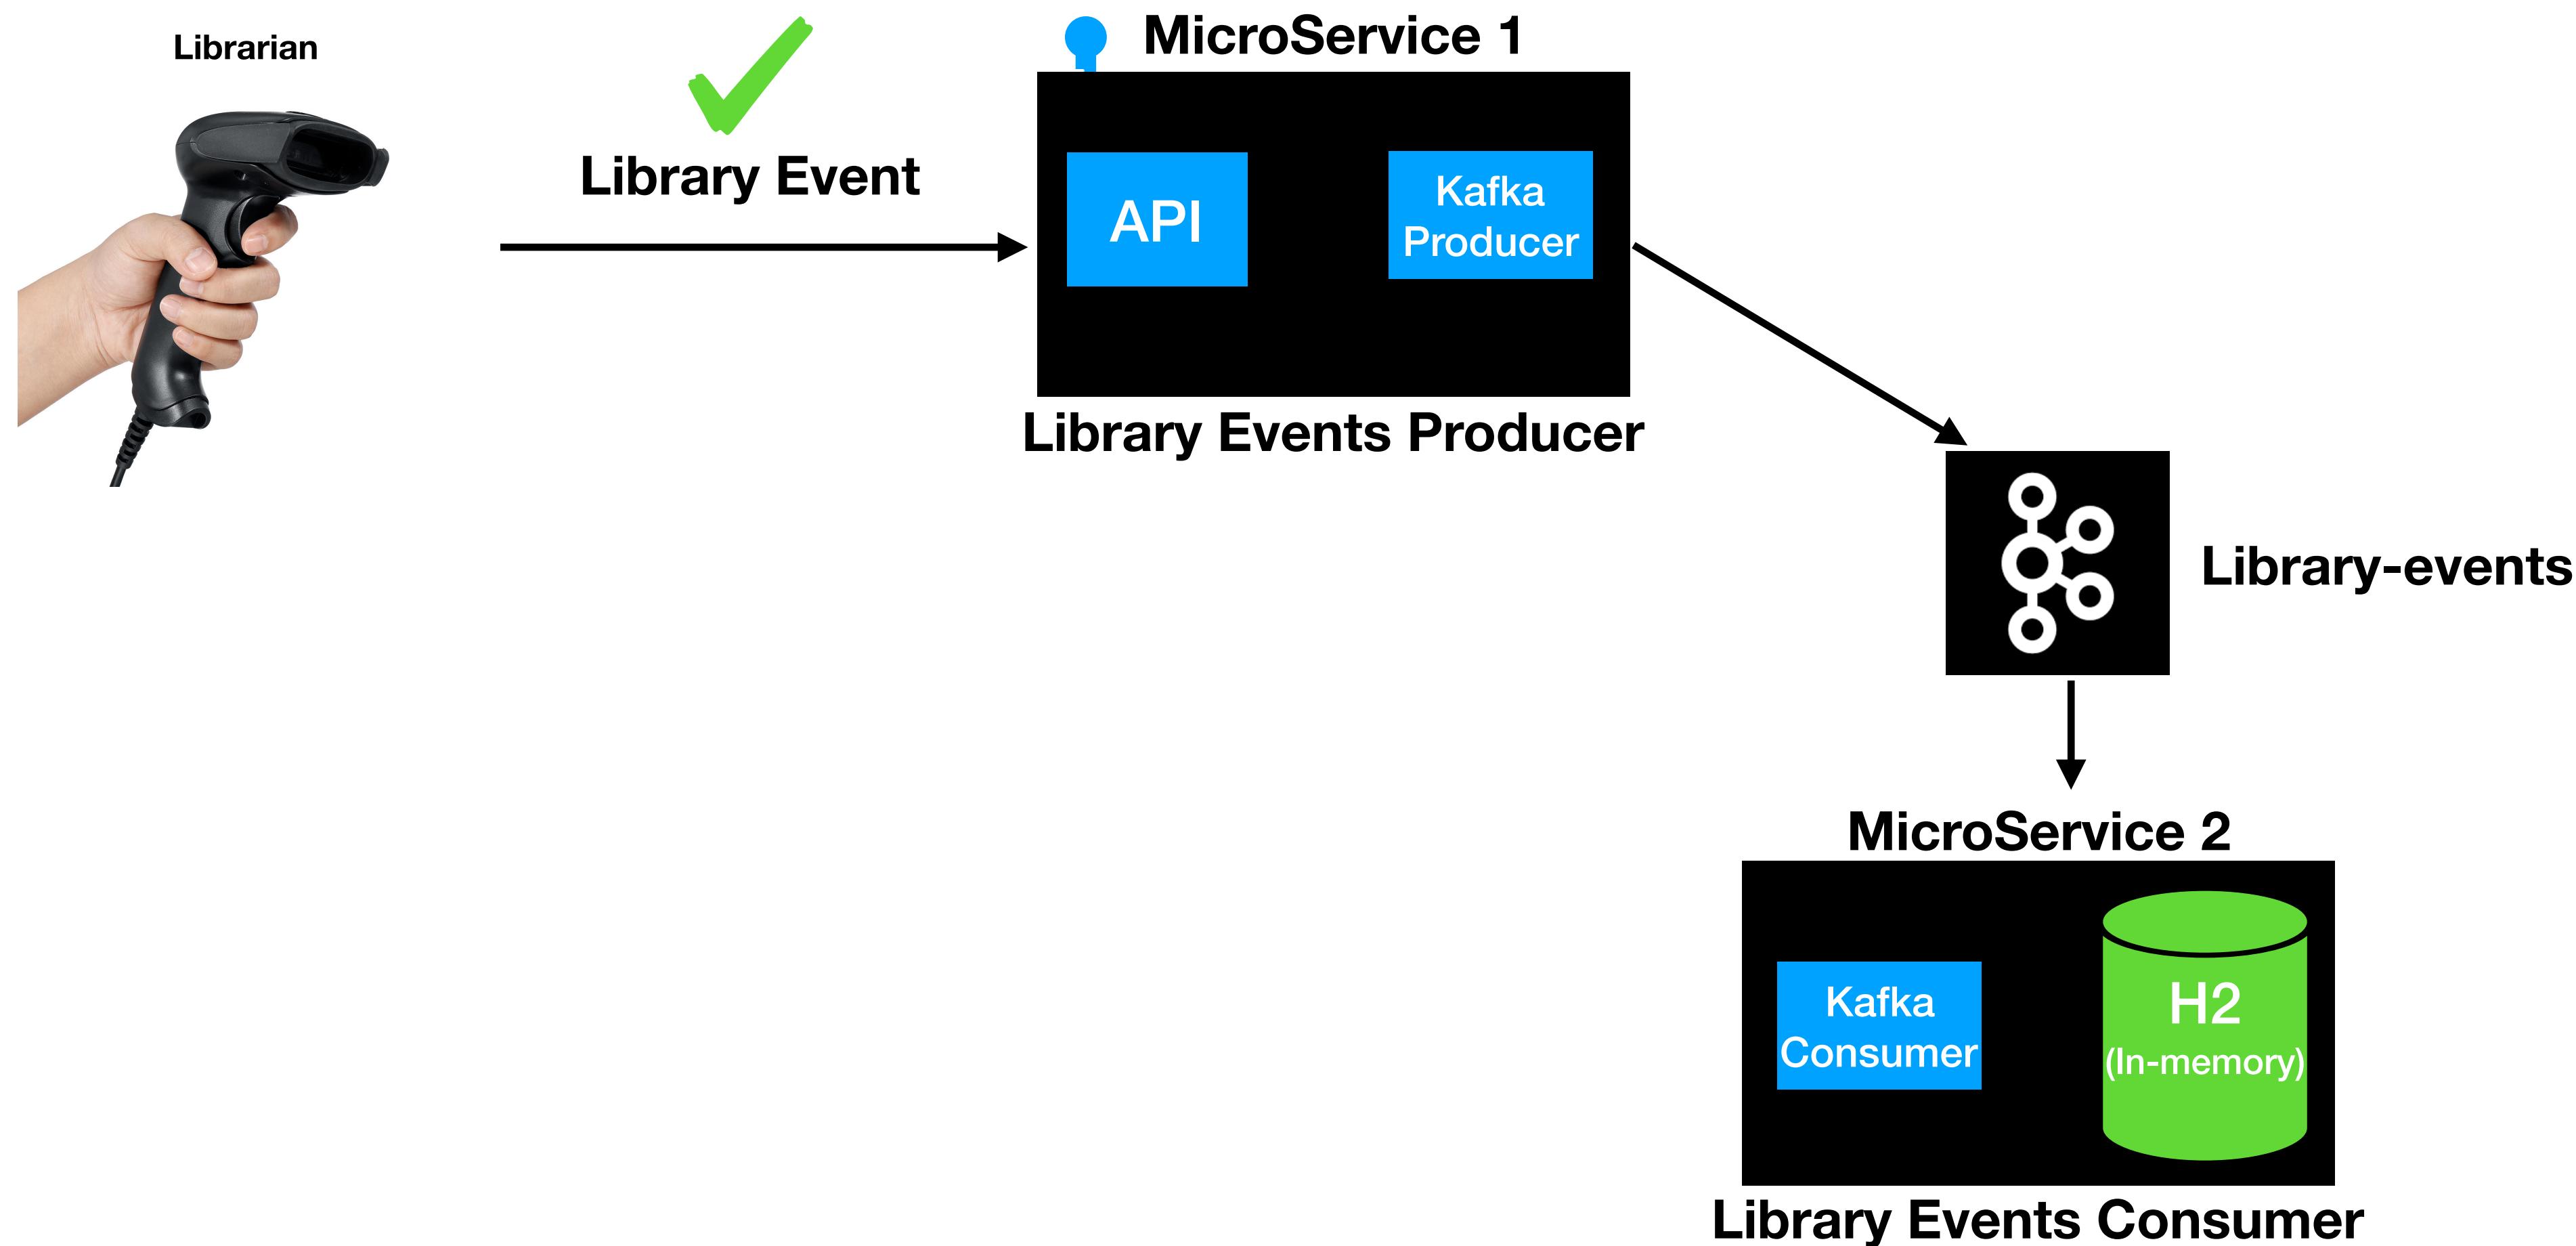
# Library Inventory Architecture



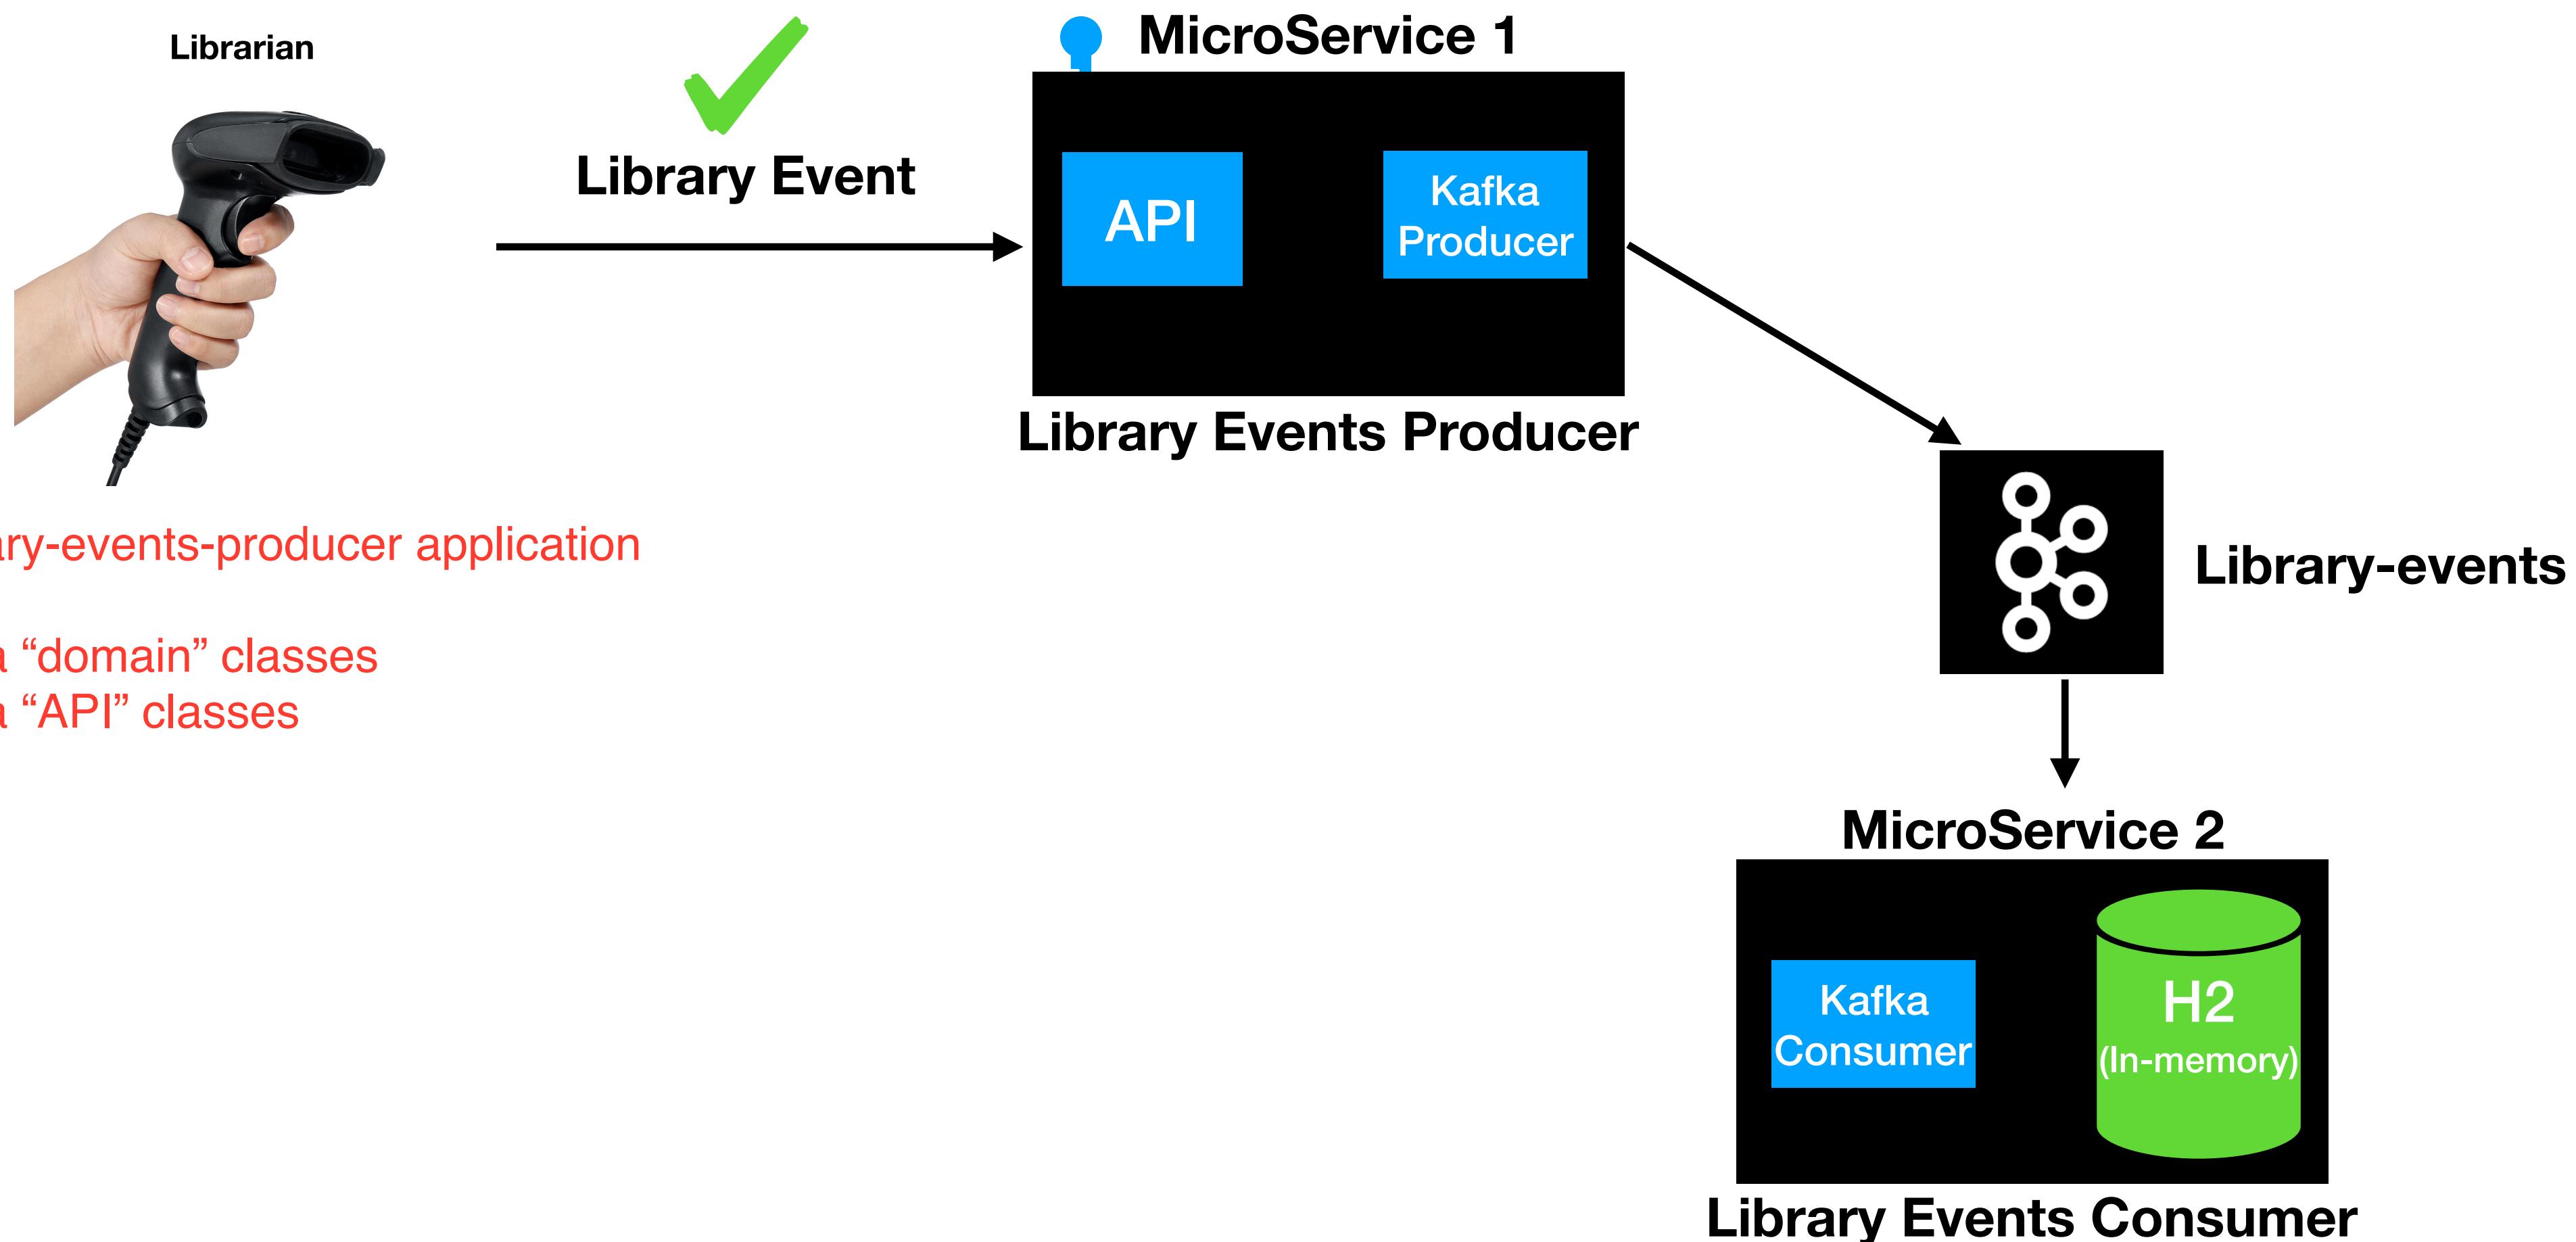
# Library Inventory Architecture



# Library Event Domain



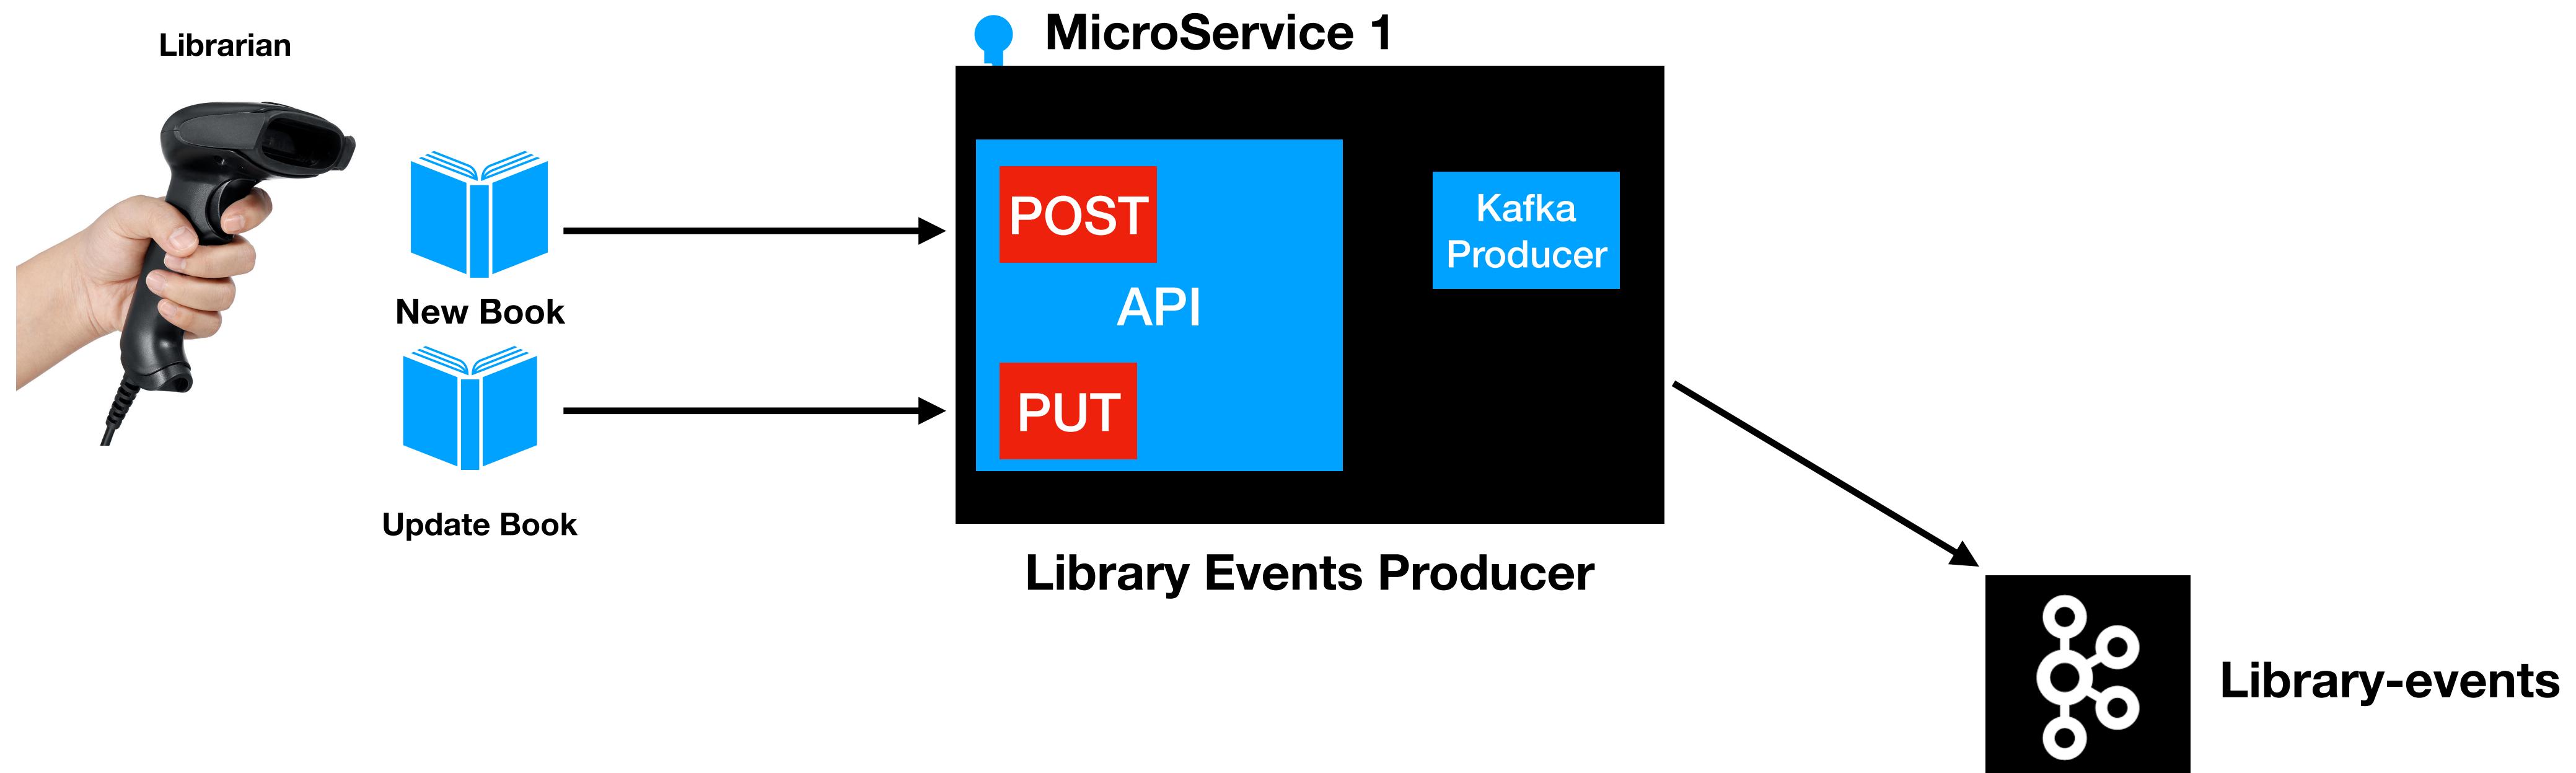
# Library Event Domain



# Library Events Producer API

POST WITH-OUT STATUS

```
curl -i \  
-d '{"libraryEventId":null,"book":{"bookId":456,"bookName":"Kafka Using Spring Boot","bookAuthor":"Dilip"}}' \  
-H "Content-Type: application/json" \  
-X POST http://localhost:8080/v1/library-event
```



# KafkaTemplate

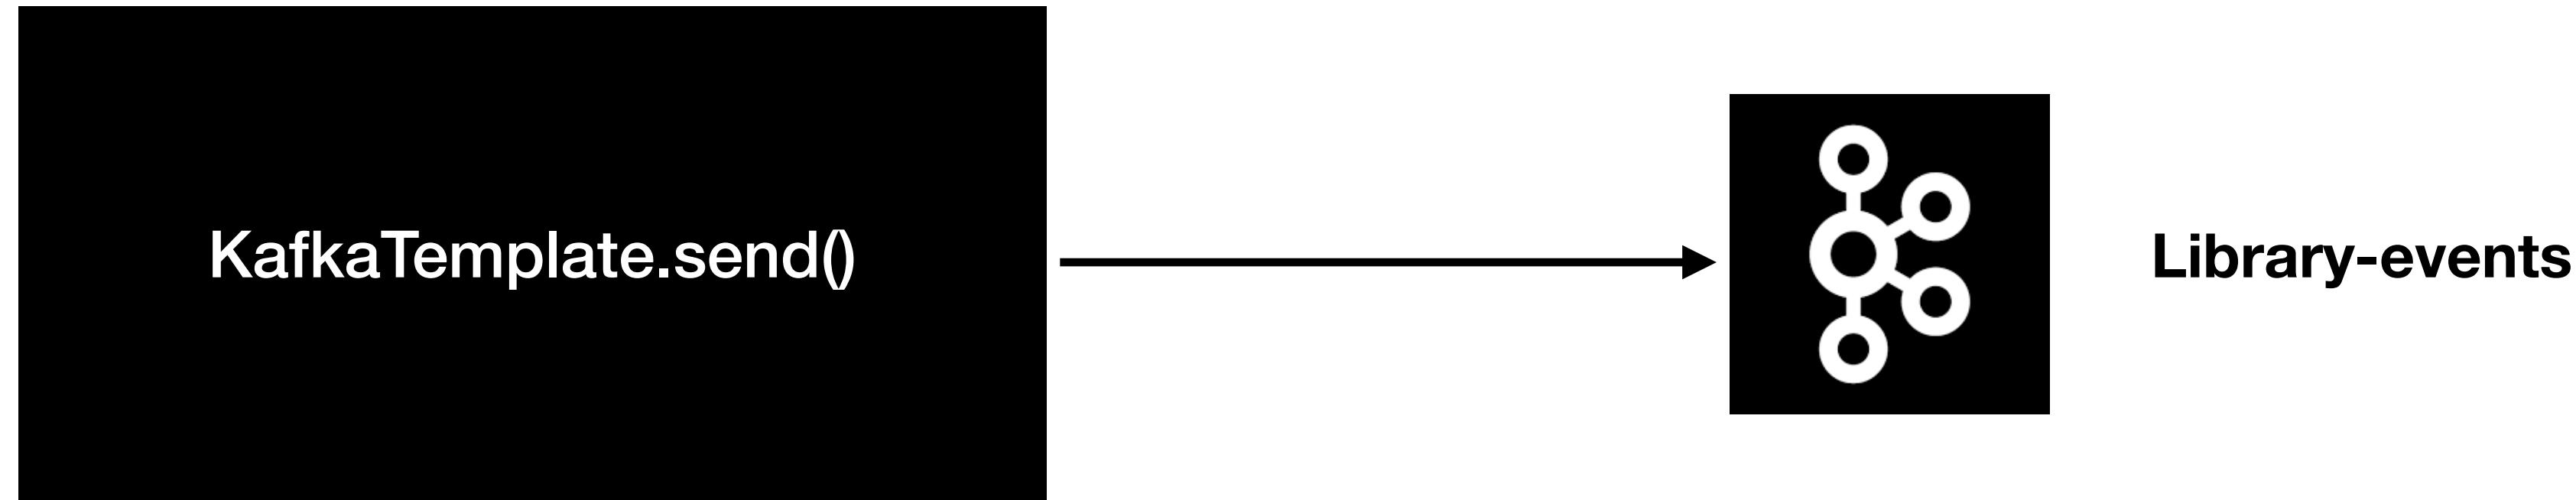
Kafka Producer in Spring

# KafkaTemplate

- Produce records in to Kafka Topic
  - Similar to JDBCTemplate for DB

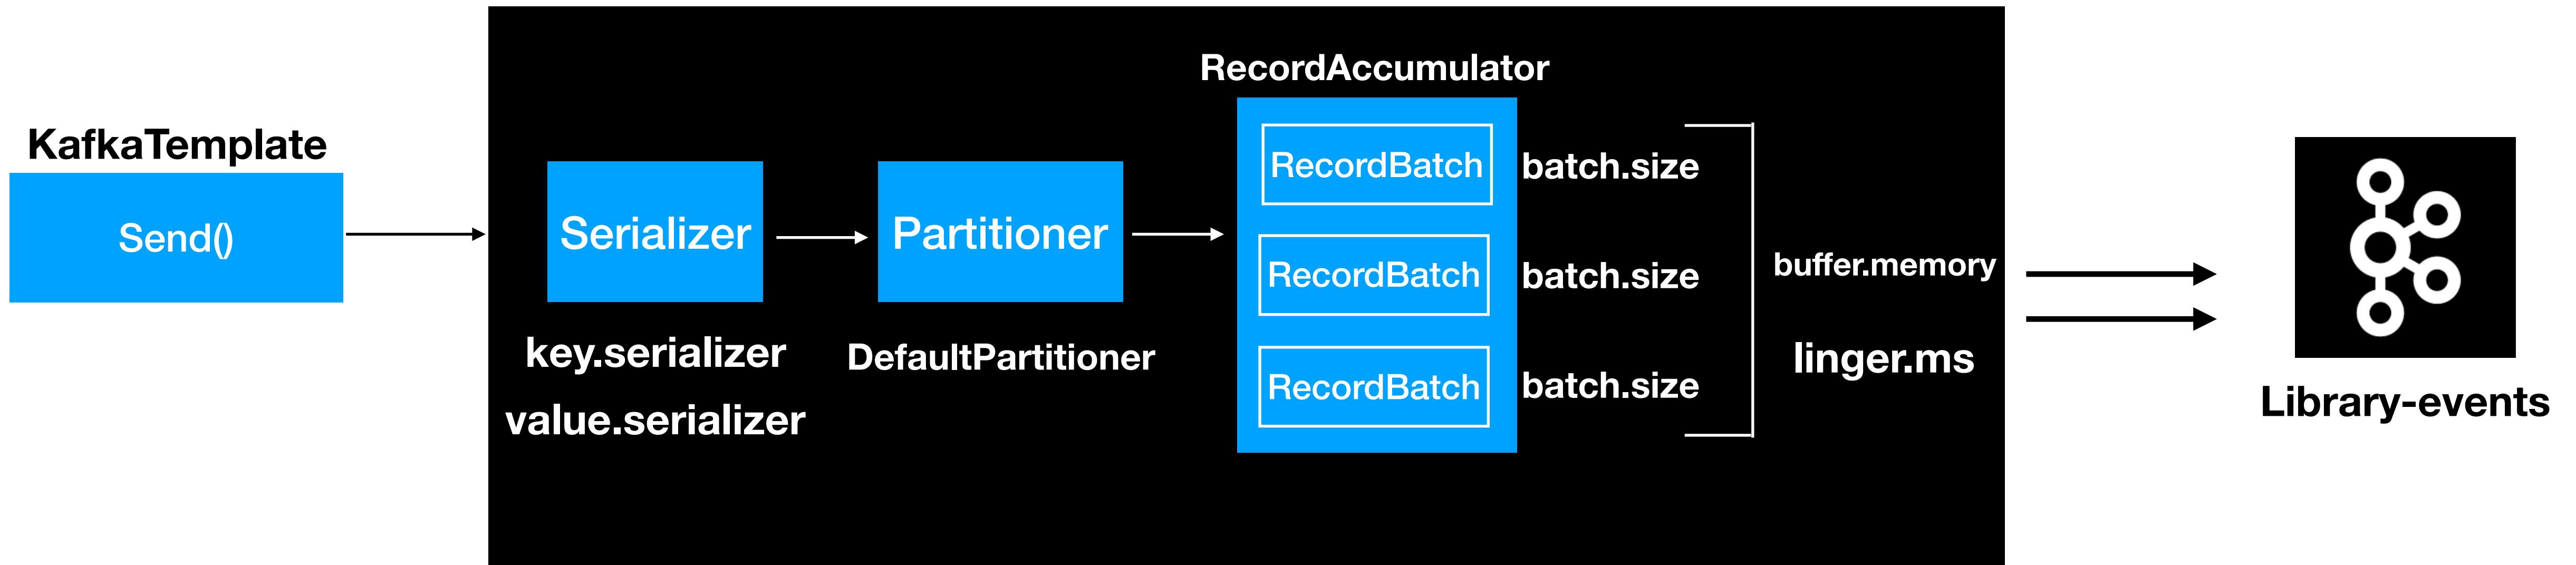
Reference: <https://docs.spring.io/spring-kafka/reference/html/#kafka-template>

# How KafkaTemplate Works ?



# KafkaTemplate.send()

## Behind the Scenes



# Configuring KafkaTemplate

## Mandatory Values:

**bootstrap-servers**: localhost:9092,localhost:9093,localhost:9094

**key-serializer**: org.apache.kafka.common.serialization.IntegerSerializer

**value-serializer**: org.apache.kafka.common.serialization.StringSerializer

Why IntegerSerializer for key is because we are trying to send message as LibraryEvent which contains “libraryEventId” as Integer type.

The actual payload value type is String type. Hence we used StringSerializer

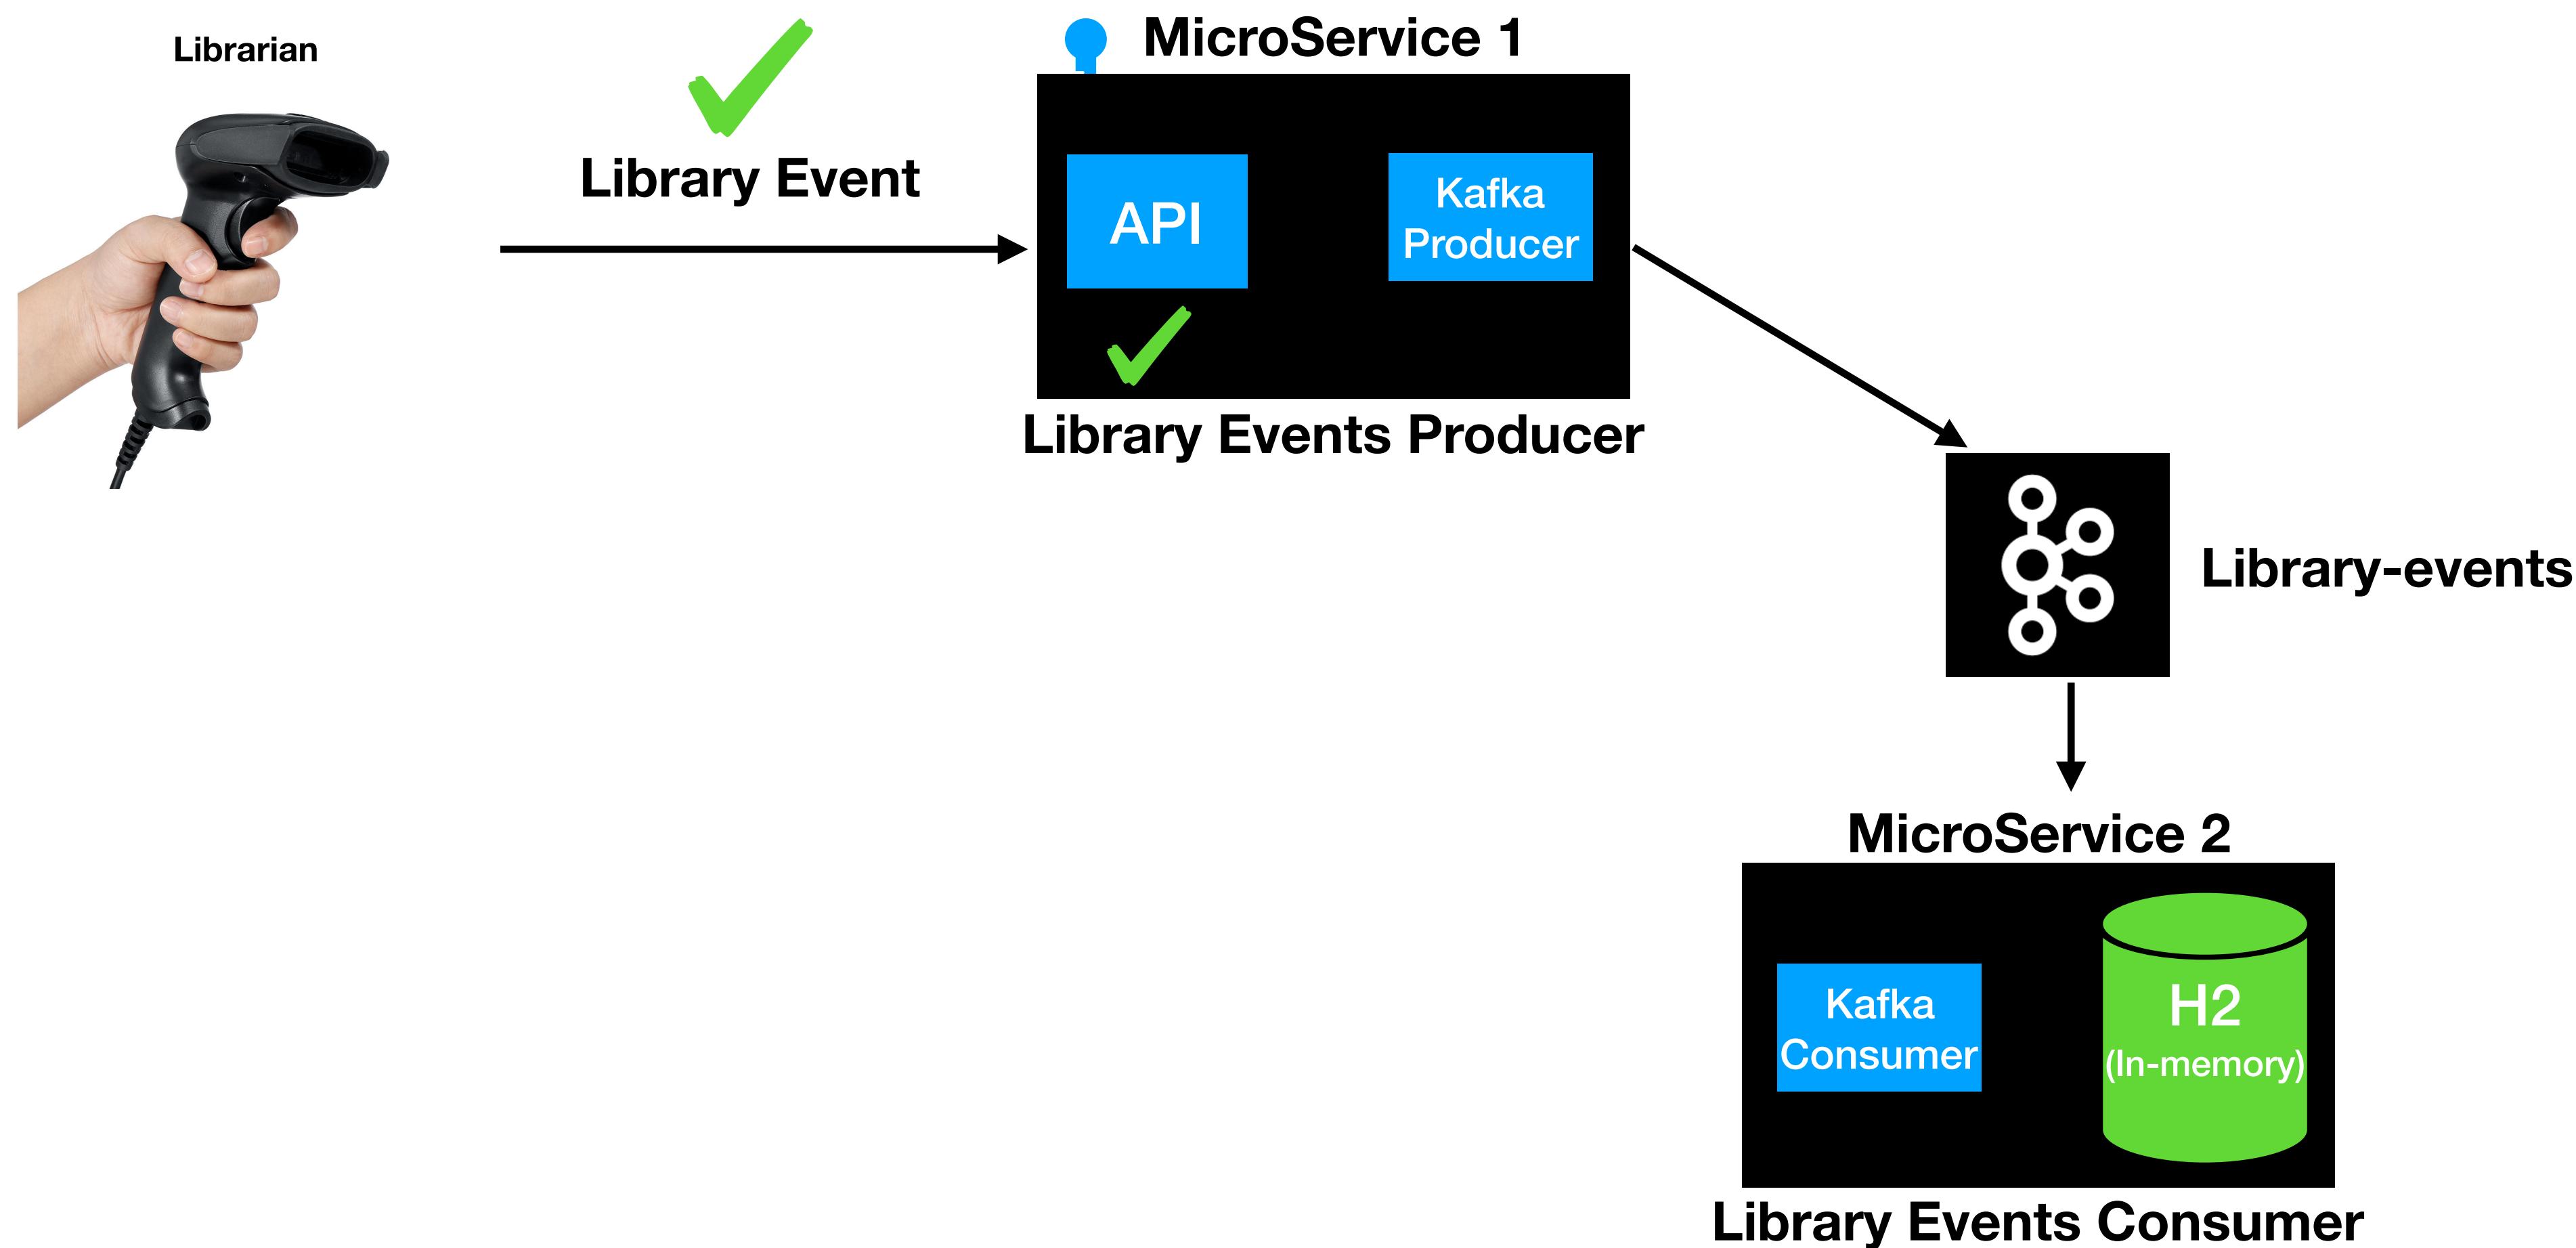
How SpringBoot AutoConfigure the Kafka producer? Refer: spring.factories, KafkaAutoConfiguration.java, KafkaProperties.java

# KafkaTemplate AutoConfiguration

## application.yml

```
spring:  
  profiles: local  
  kafka:  
    producer:  
      bootstrap-servers: localhost:9092,localhost:9093,localhost:9094  
      key-serializer: org.apache.kafka.common.serialization.IntegerSerializer  
      value-serializer: org.apache.kafka.common.serialization.StringSerializer
```

# Library Inventory Architecture



# KafkaAdmin

- Create topics Programmatically
  - Part of the **SpringKafka**
  - How to Create a topic from Code?
    - Create a Bean of type **KafkaAdmin** in SpringConfiguration
    - Create a Bean of type **NewTopic** in SpringConfiguration

Refer: “AutoCreateConfig.java” to create a topic programmatically. After this step, run your Spring Boot Application.

Note: Before start your application, start your ZooKeeper, Broker0,1,2 (server, server-1, server-2)

Once you run, Kafka ZooKeeper would have created a topic “library-events”. To verify this, execute this command

```
$bin>./kafka-topics.sh --zookeeper localhost:2181 --list
```

# **Introduction To Automated Tests**

# Why Automated Tests ?

- Manual testing is time consuming
- Manual testing slows down the development
- Adding new changes are error prone

# What are Automated Tests?

- Automated Tests run against your code base
- Automated Tests run as part of the build
- This is a requirement for todays software development
- Easy to capture bugs
- Types of Automated Tests:
  - UnitTest
  - Integration Tests
  - End to End Tests

# Tools for Automated

- JUnit
- Spock

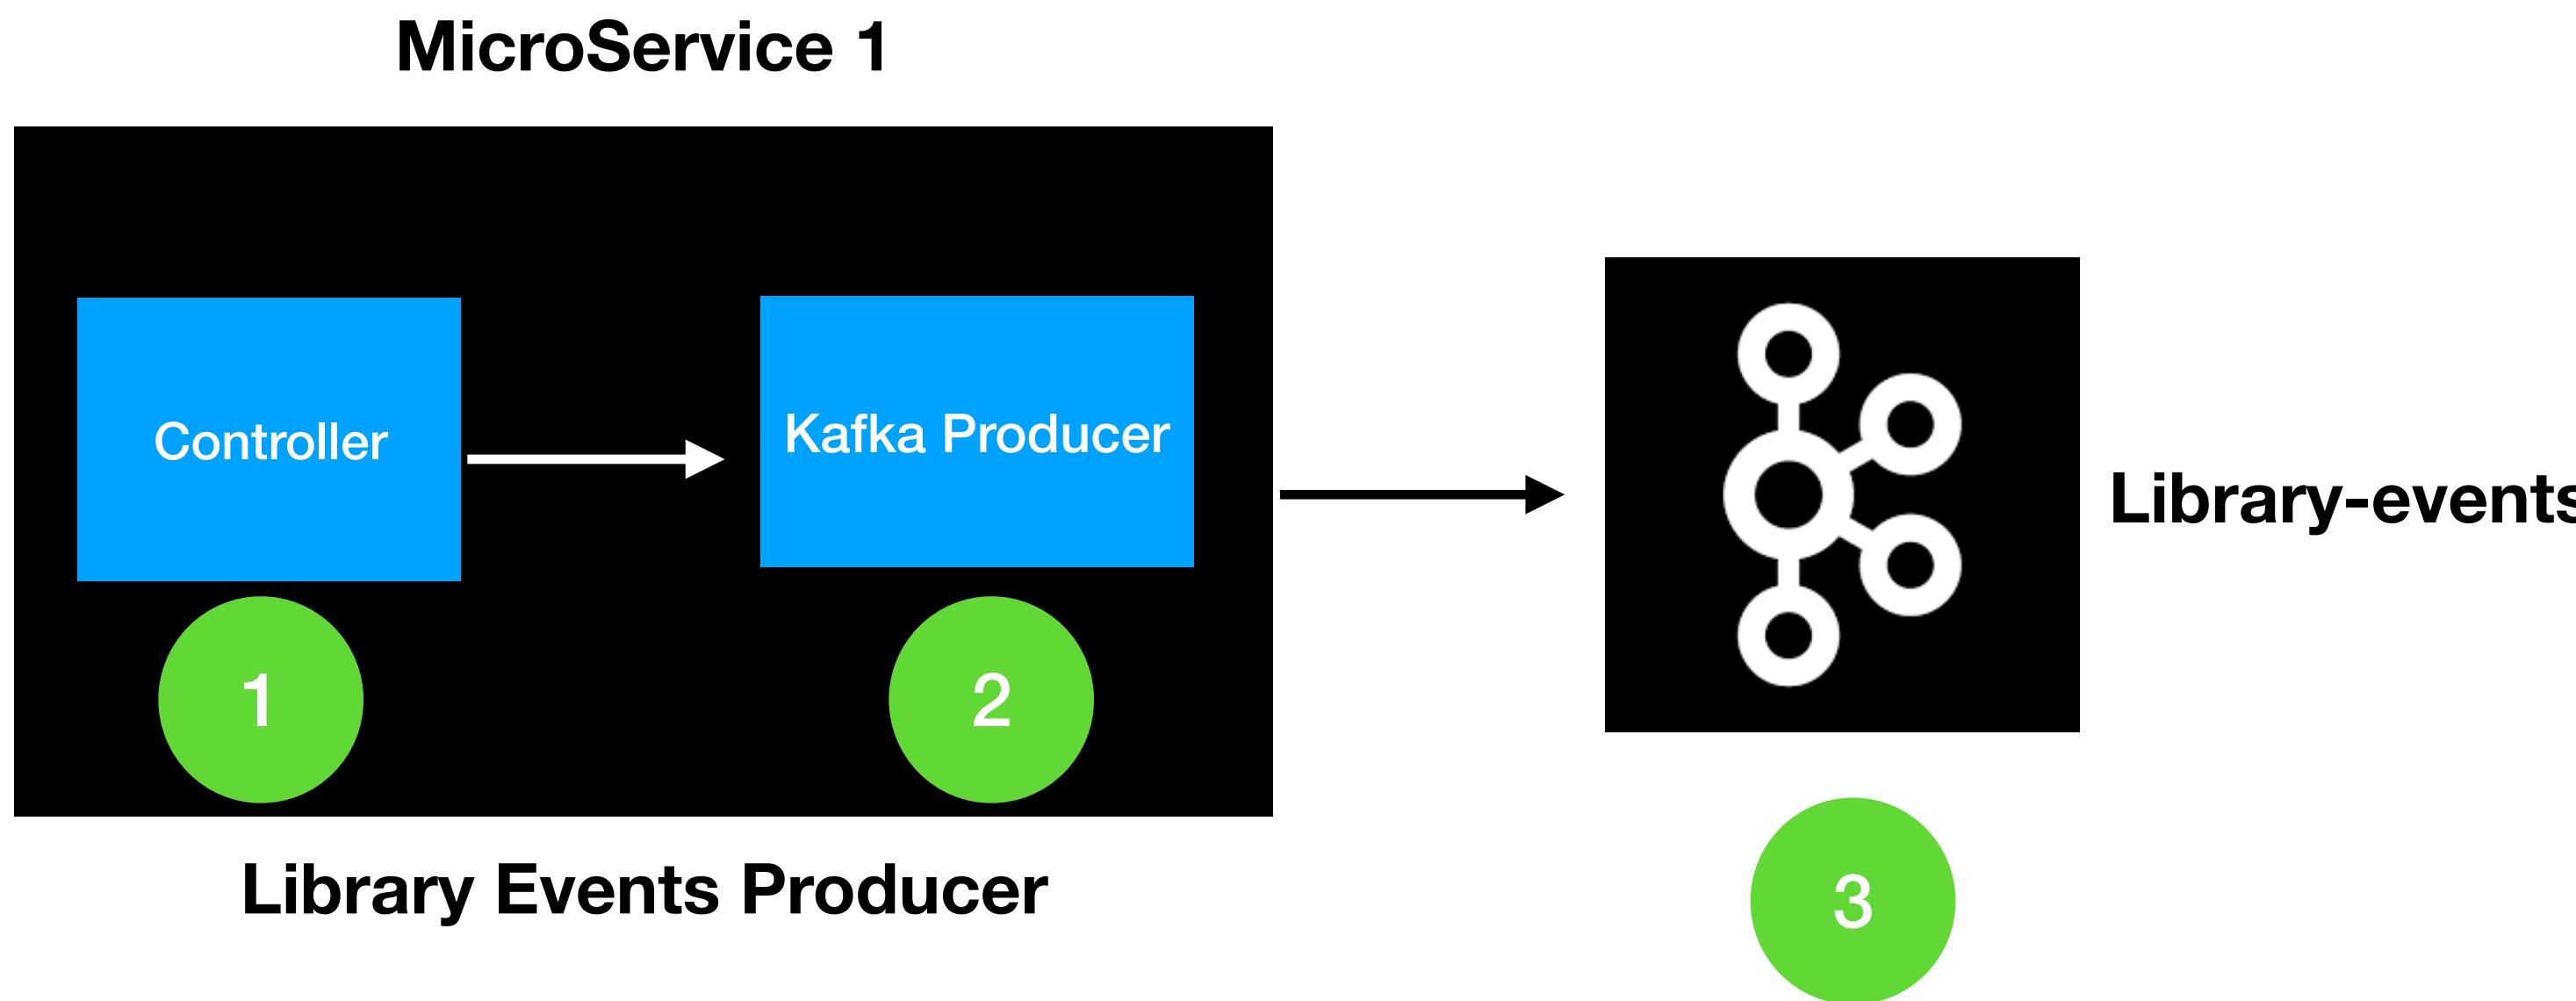
# **Integration Tests**

## **Using**

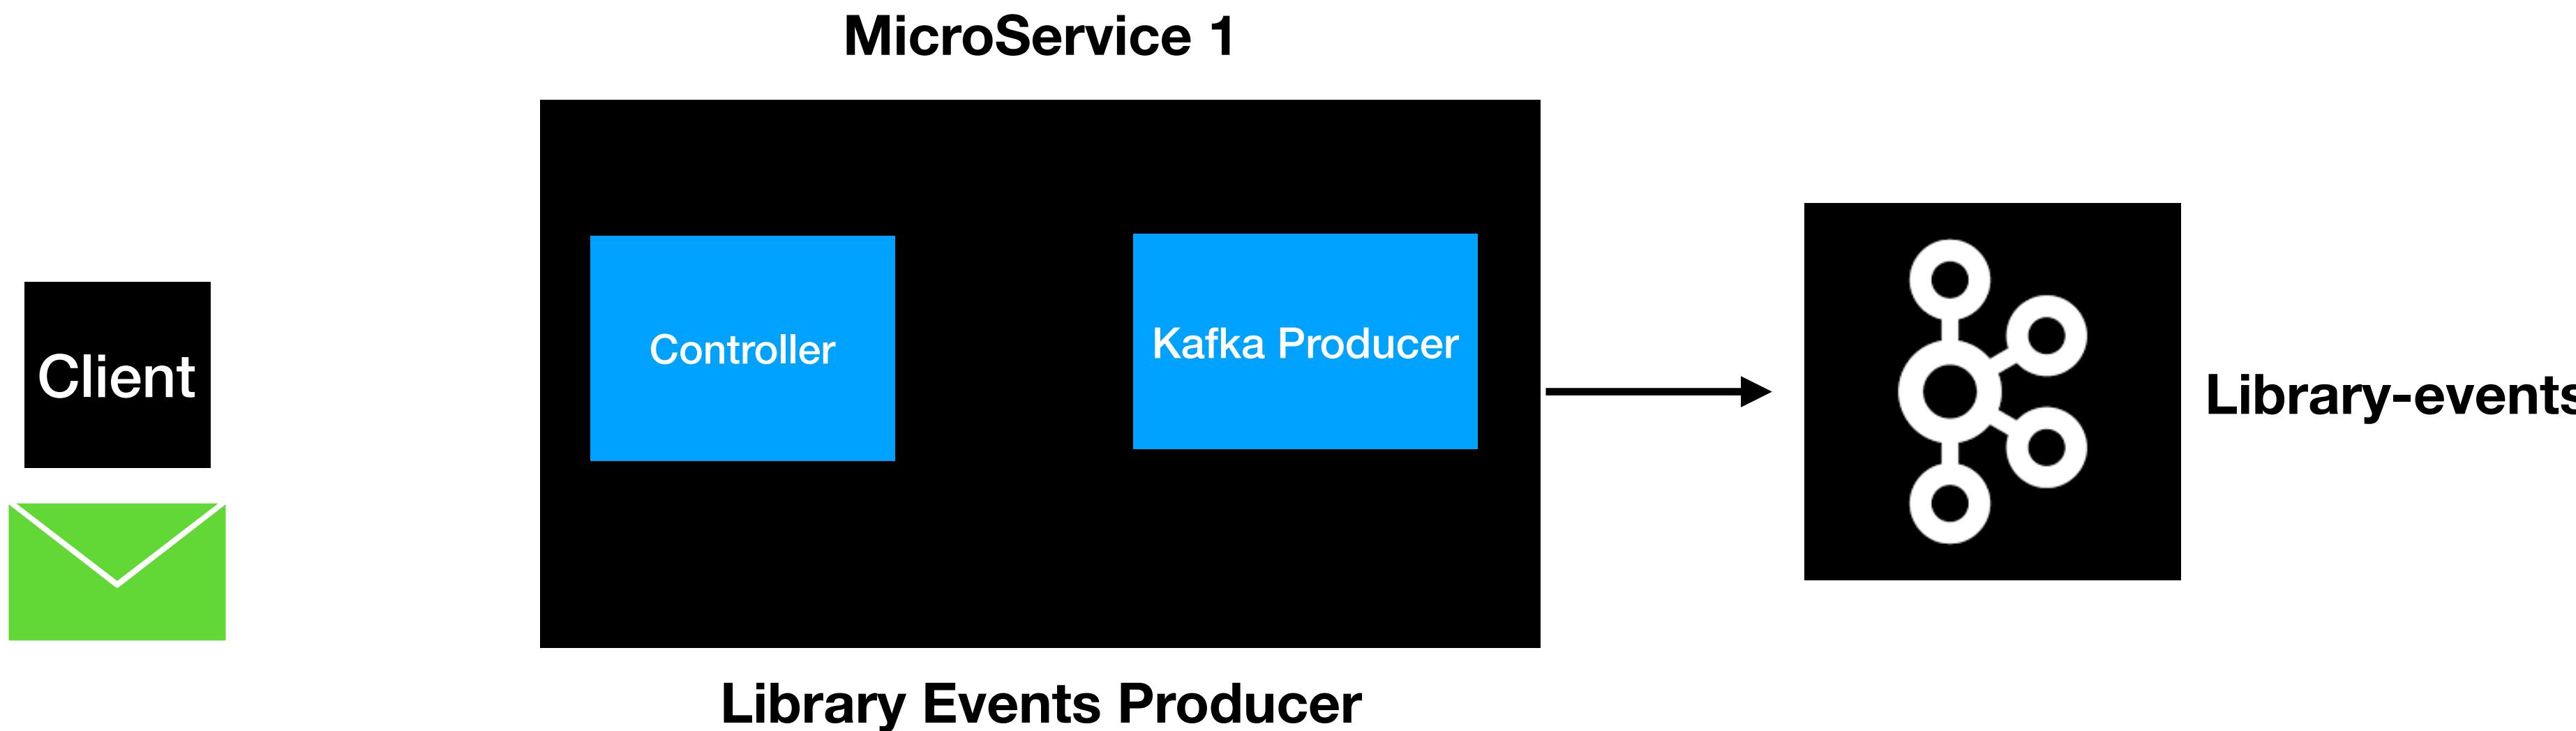
## **JUnit5**

# What is Integration Test?

- Test combines the different layers of the code and verify the behavior is working as expected.



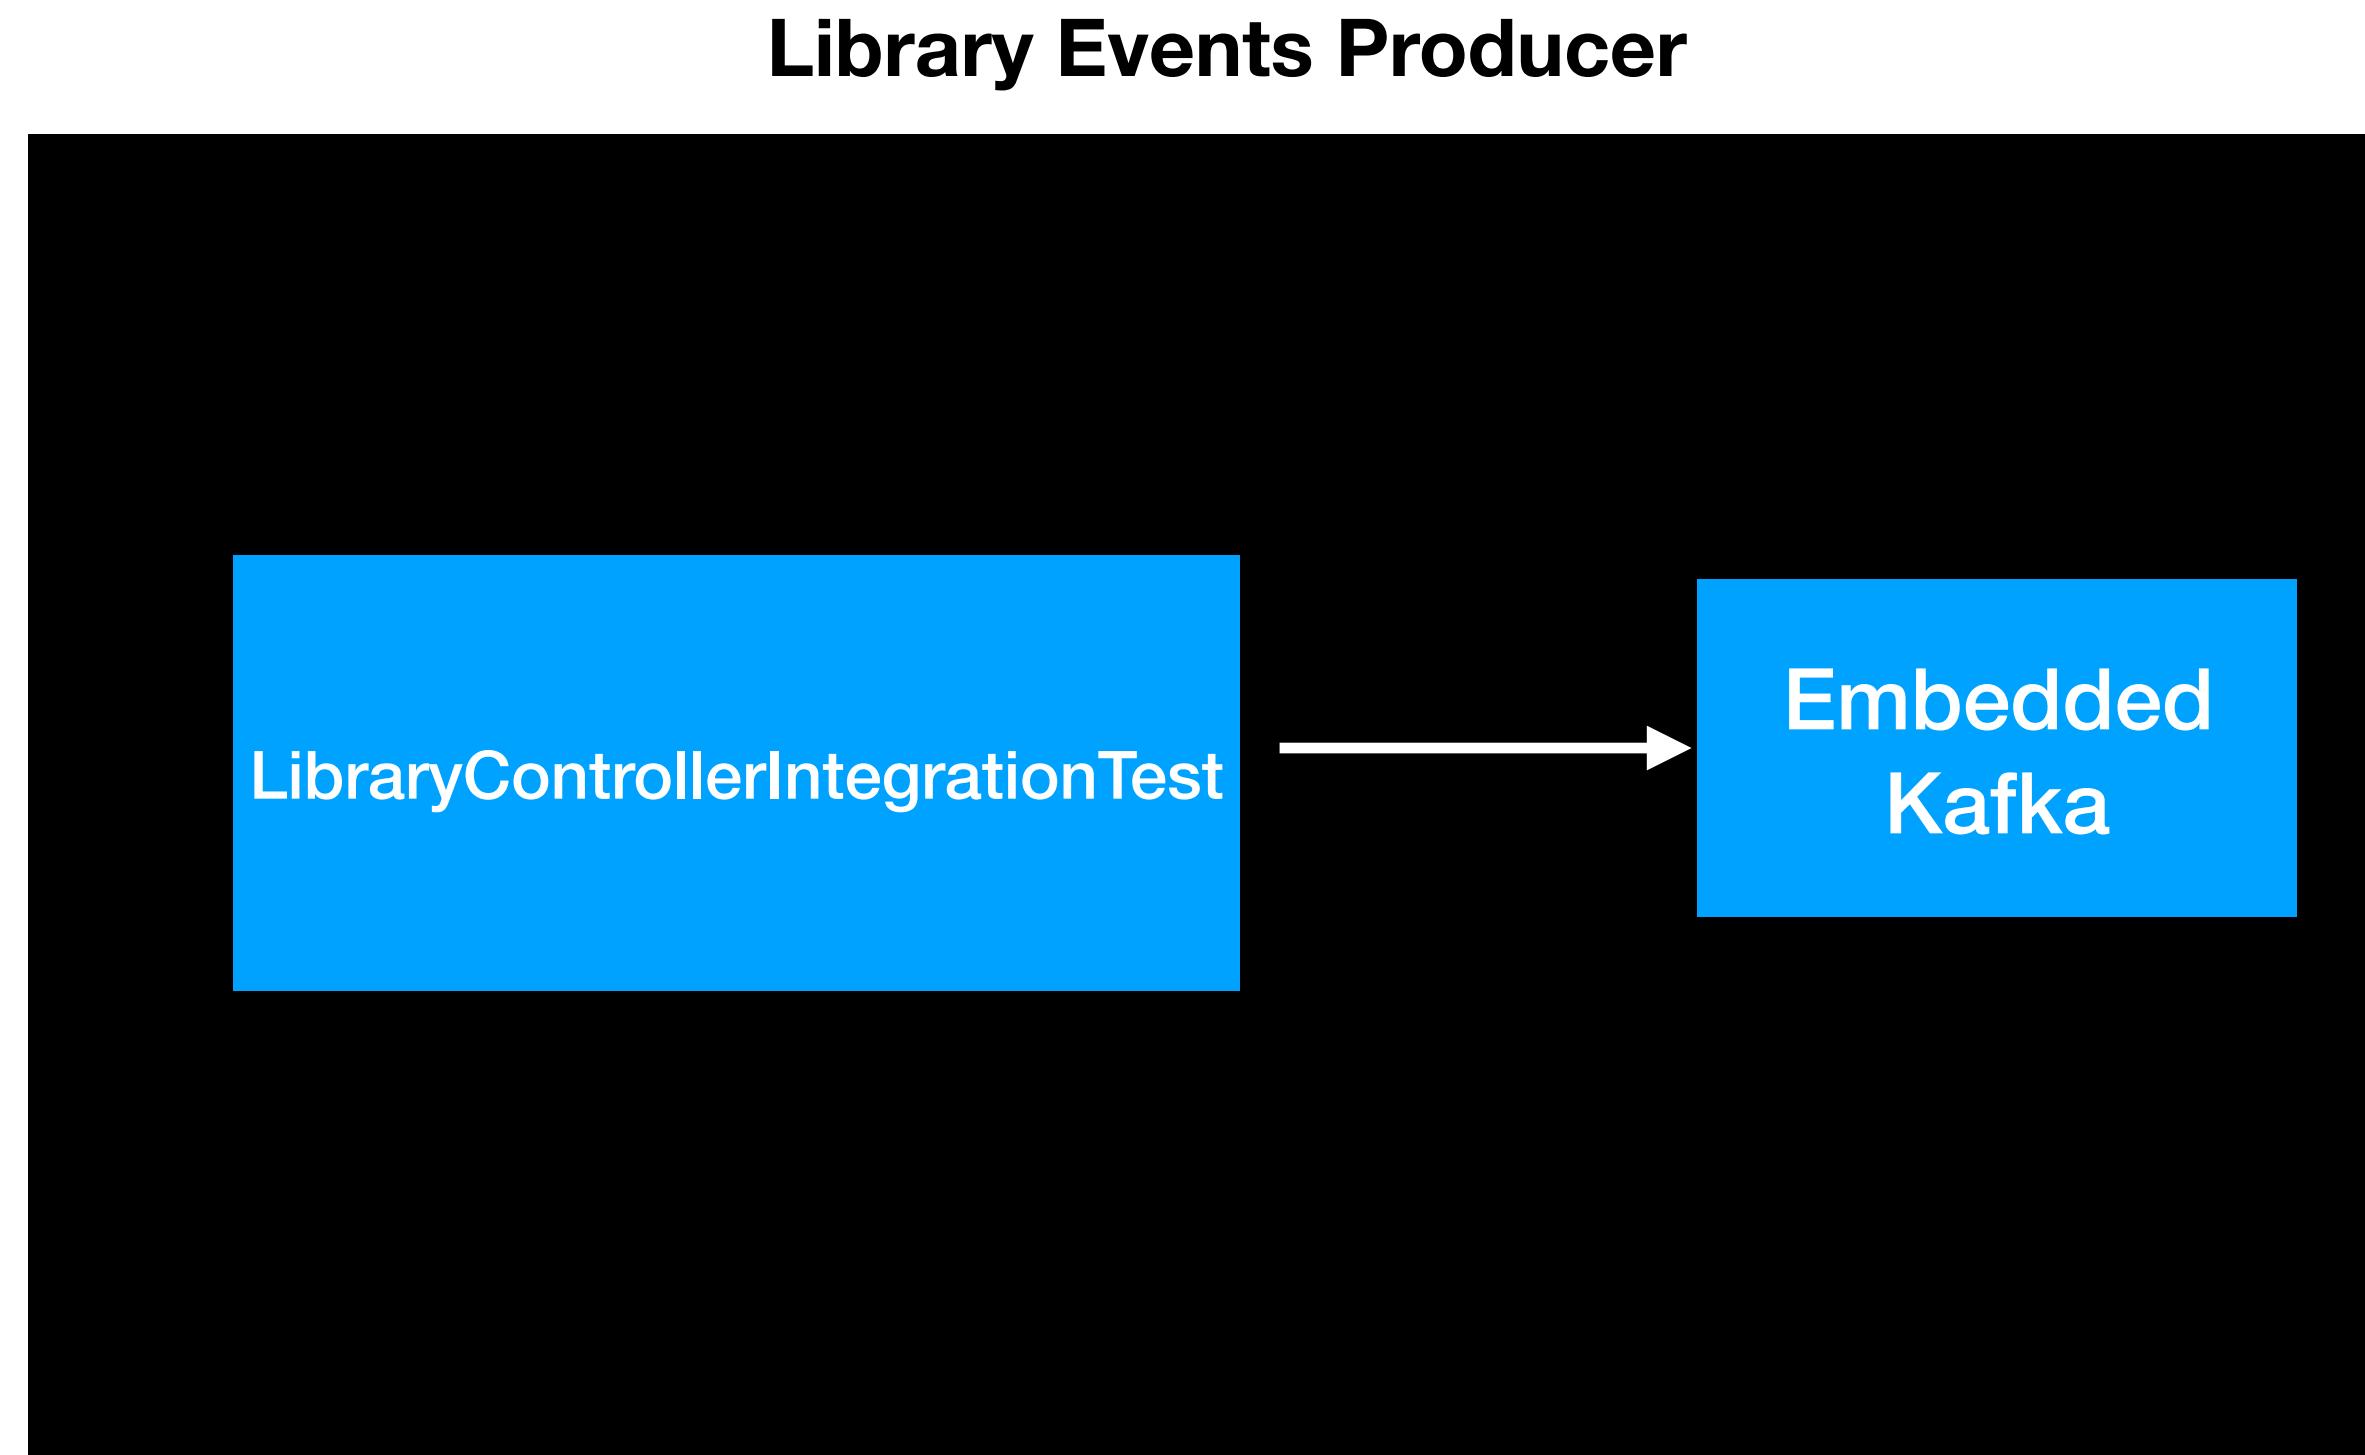
# Integration Test



# Embedded Kafka

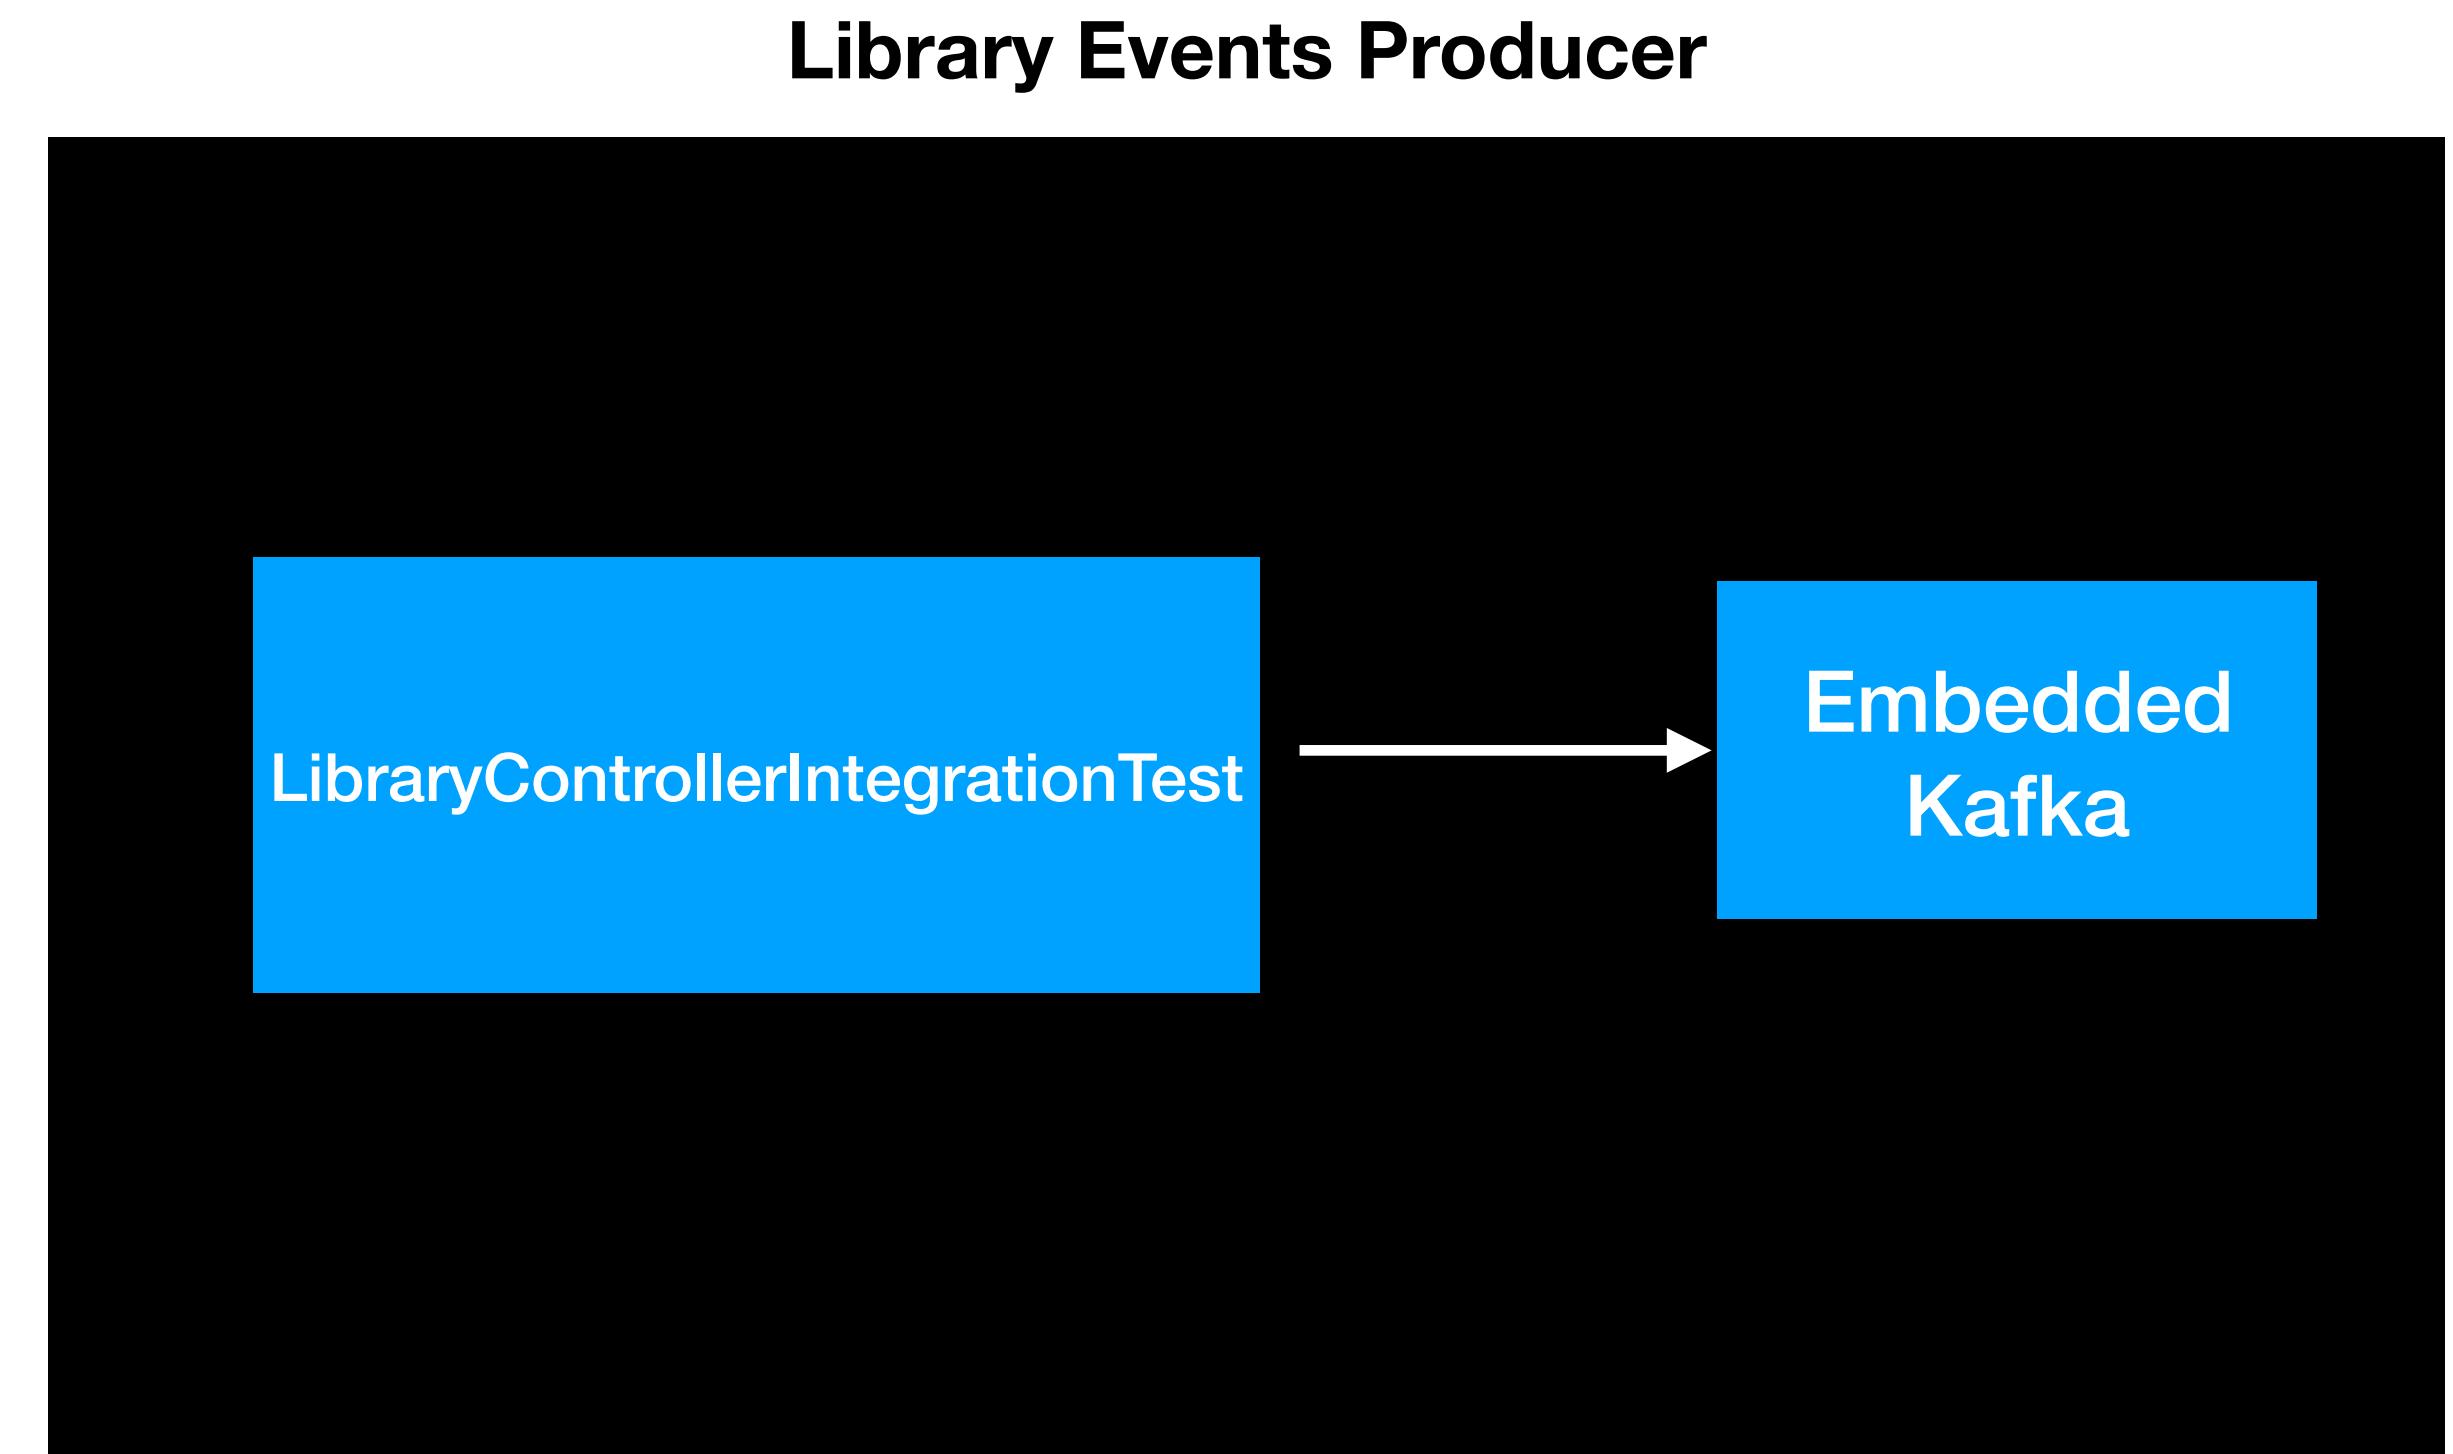
# What is EmbeddedKafka?

- In-Memory Kafka
- Integration Tests can interact with EmbeddedKafka



# Why Embedded Kafka ?

- Easy to write Integration Tests
- Test all the code as like you interact with Kafka



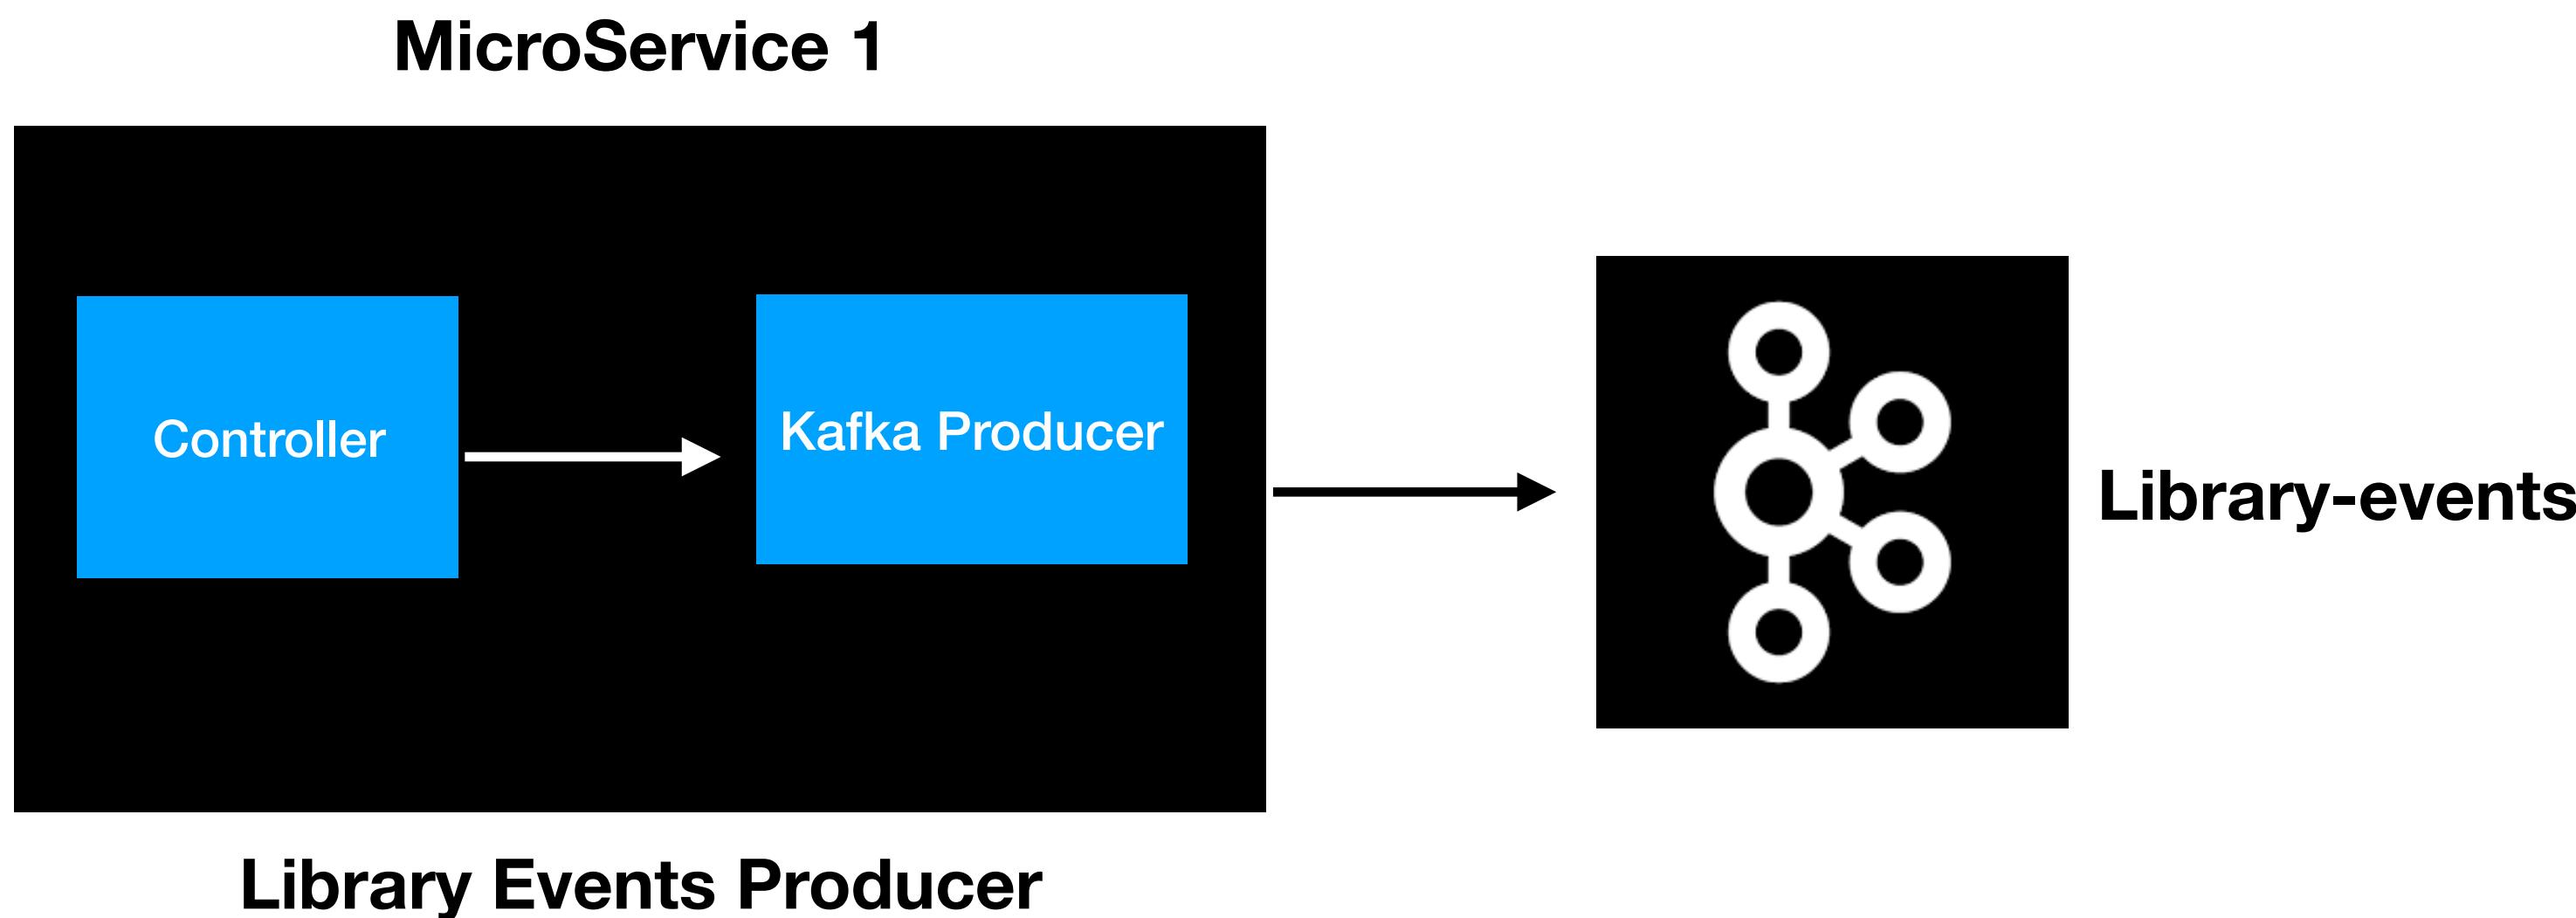
# **Unit Tests**

## **Using**

# **JUnit5**

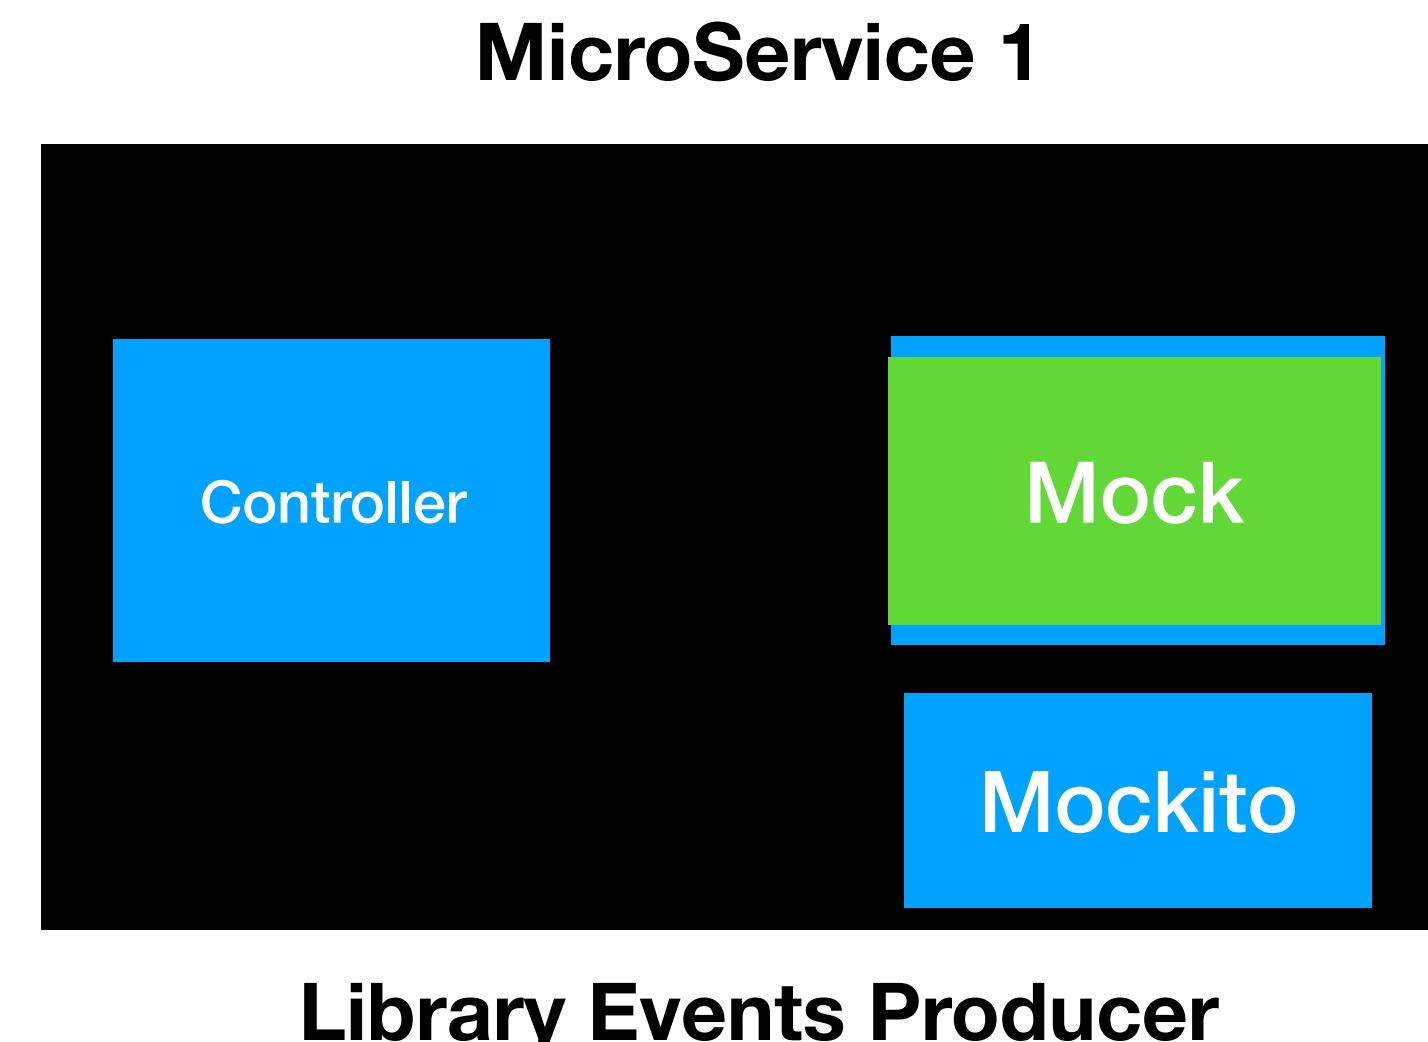
# What is Unit Test?

- Test the just focuses on a single unit (method)
- Mocks the external dependecies



# What is Unit Test?

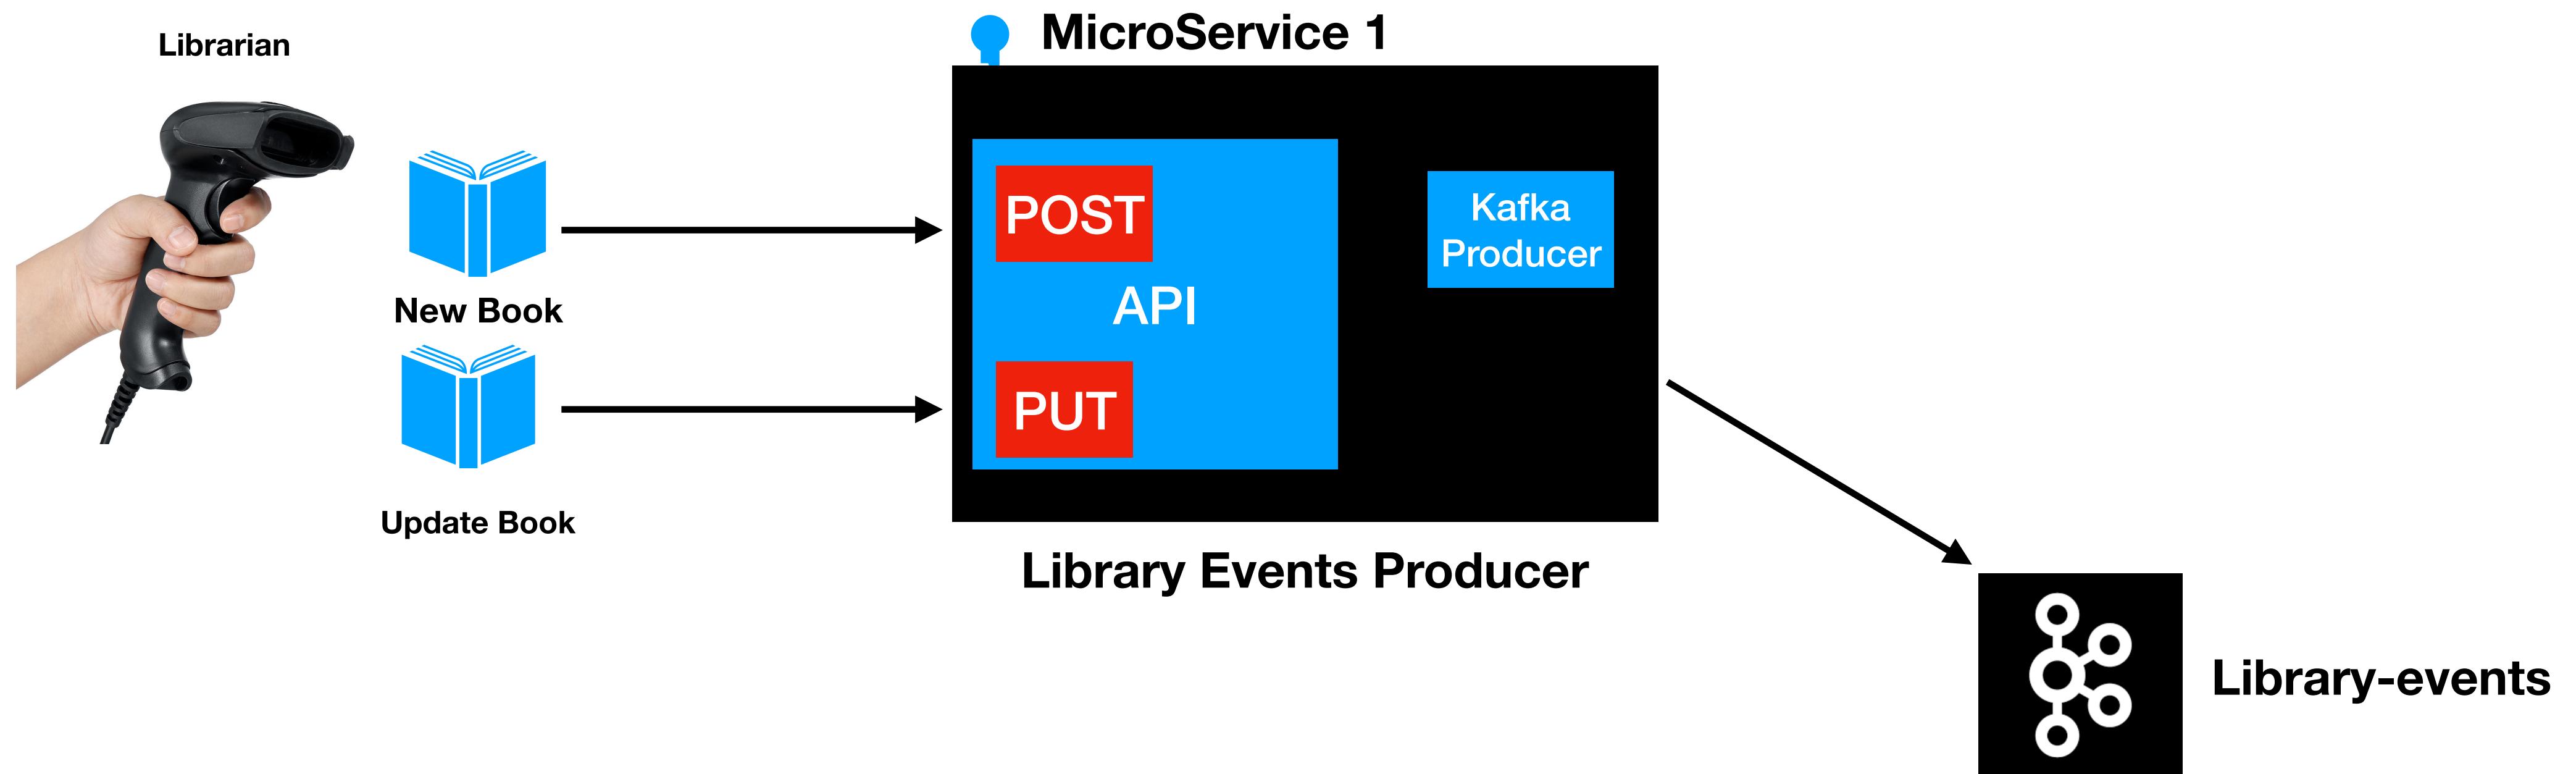
- Test the just focuses on a single unit (method)
- Mocks the external dependencies



# Why Unit Test?

- Unit Tests are handy to mock external dependencies
- Unit Tests are faster compared to Integration tests
- Unit Tests cover scenarios that's not possible with Integration tests

# Library Events Producer API



# PUT - “/v1/libraryevent”

- libraryEventId is a mandatory field

```
{  
  "libraryEventId": 123,  
  "eventStatus": null,  
  "book": {  
    "bookId": 456,  
    "bookName": "Kafka Using Spring Boot",  
    "bookAuthor": "Dilip"  
  }  
}
```

# Kafka Producer Configurations

# Kafka Producer Configurations

- **acks** which is the key configuration for reliable data delivery.
  - acks = 0, 1 and all
    - acks = 1 -> guarantees message is written to a leader ( Default) 
    - acks = all -> guarantees message is written to a leader and to all the replicas 
    - acks=0 -> no guarantee (Not Recommended) 

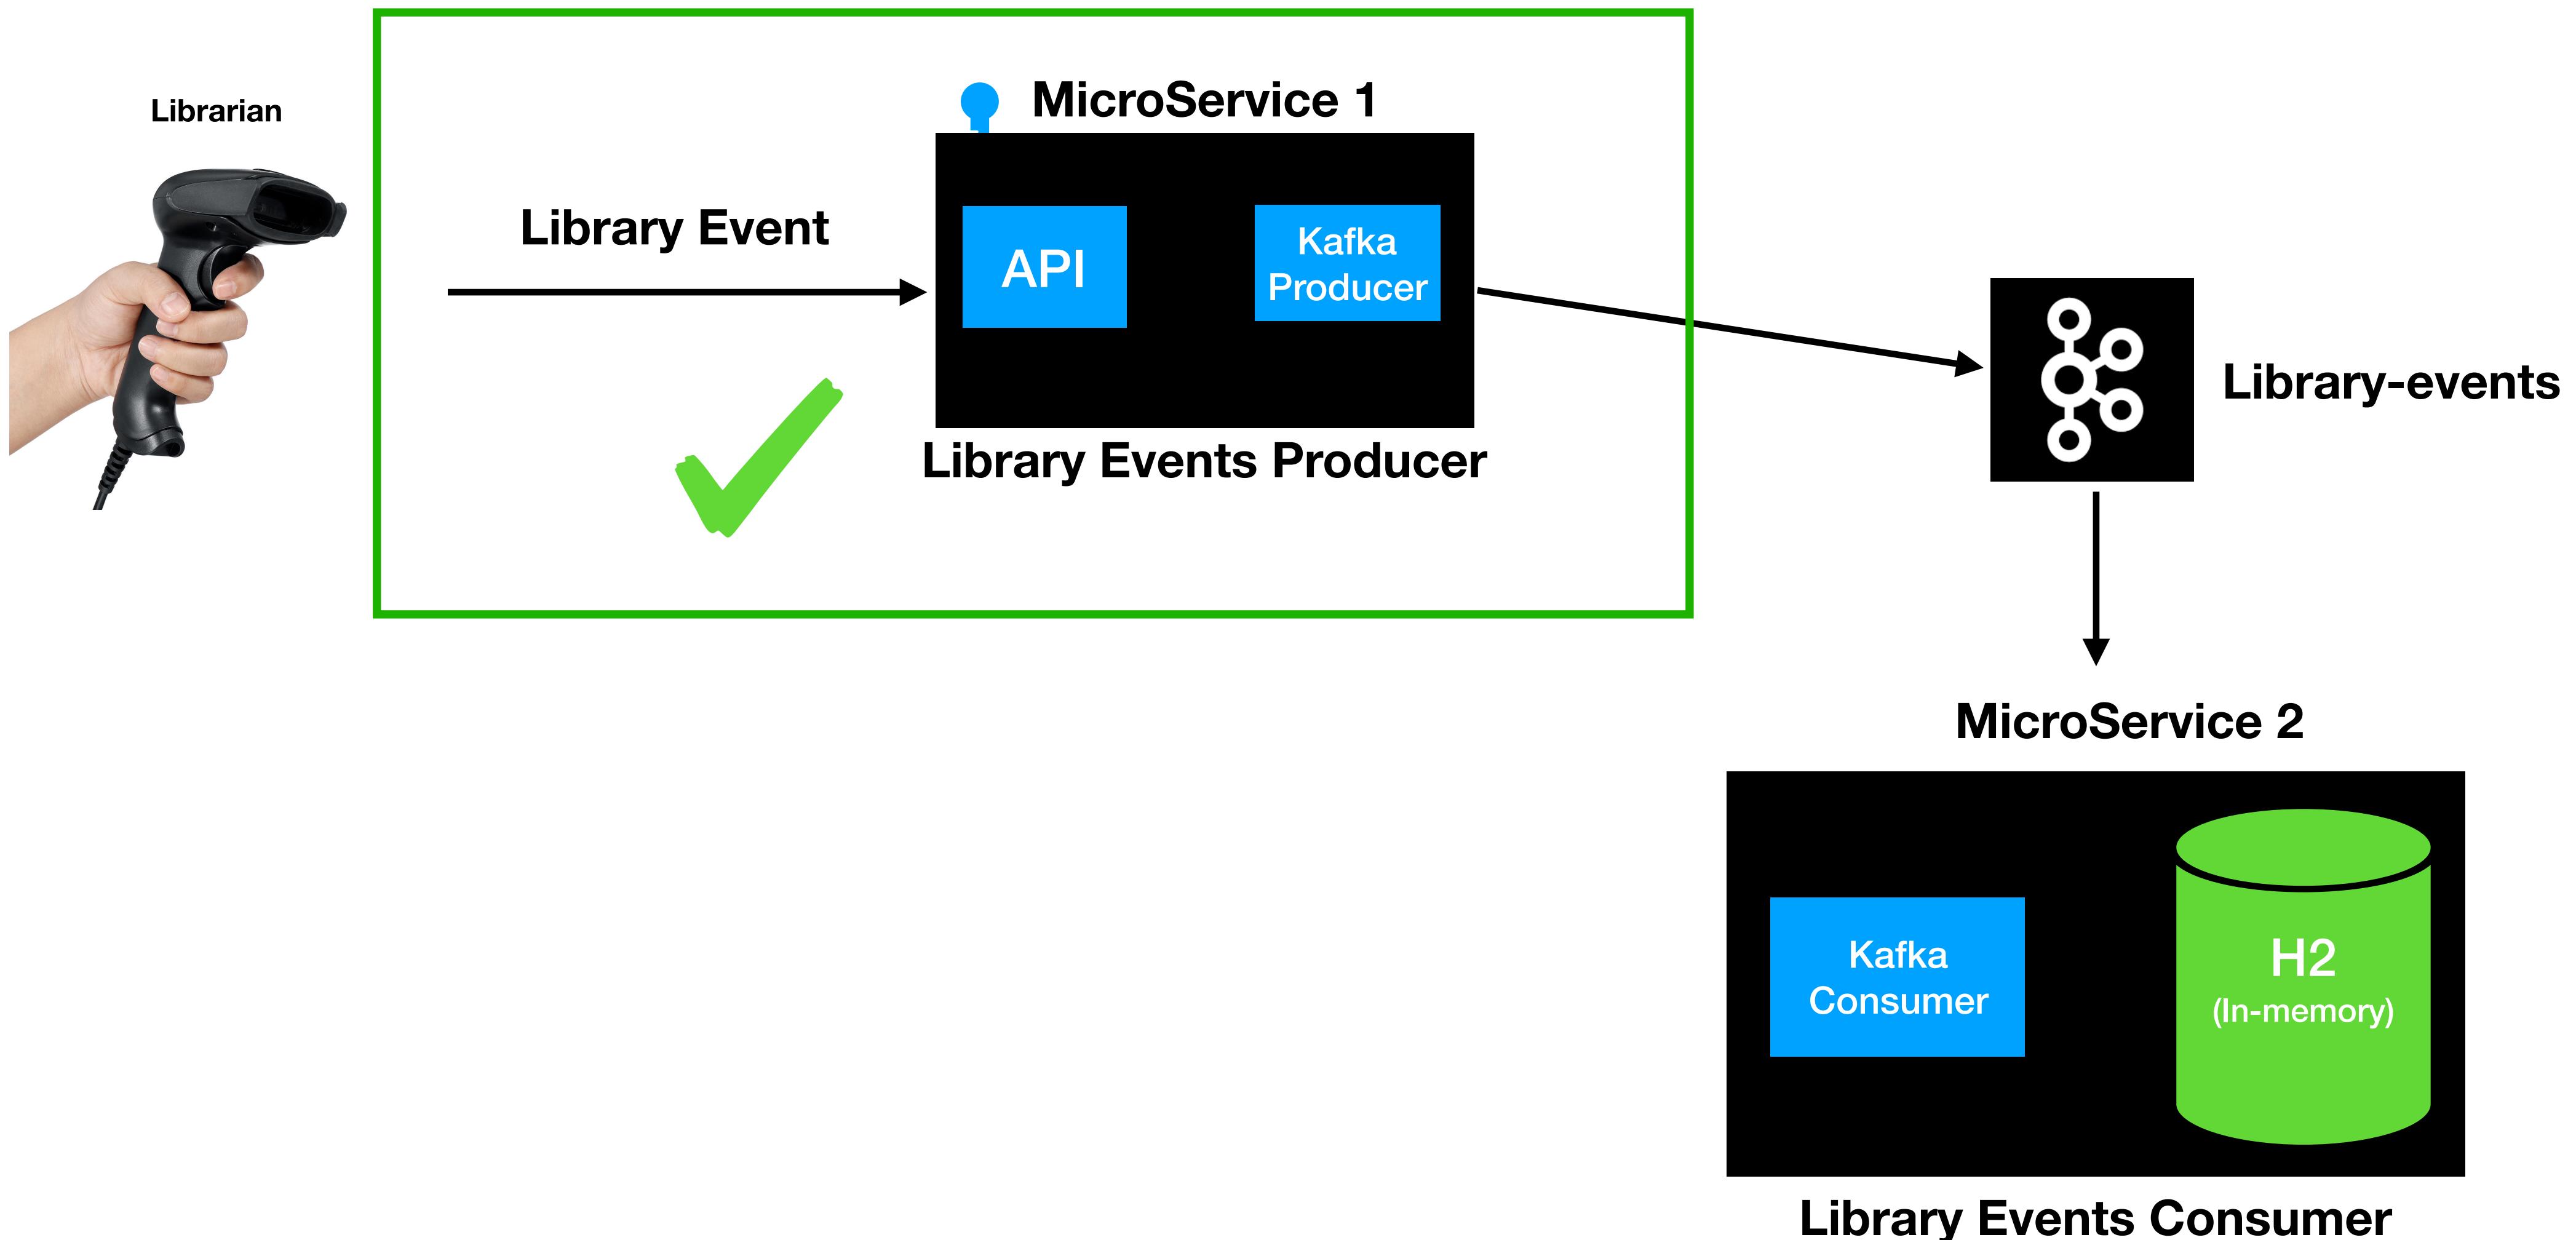
To see these values, execute the integration test and it will display all the producer config related values.

# Kafka Producer Configurations

<https://kafka.apache.org/documentation/#producerconfigs>

- retries
  - Integer value = [0 - 2147483647]
  - In Spring Kafka, the default value is -> **2147483647**
- retry.backoff.ms
  - Integer value represented in milliseconds
  - Default value is 100ms

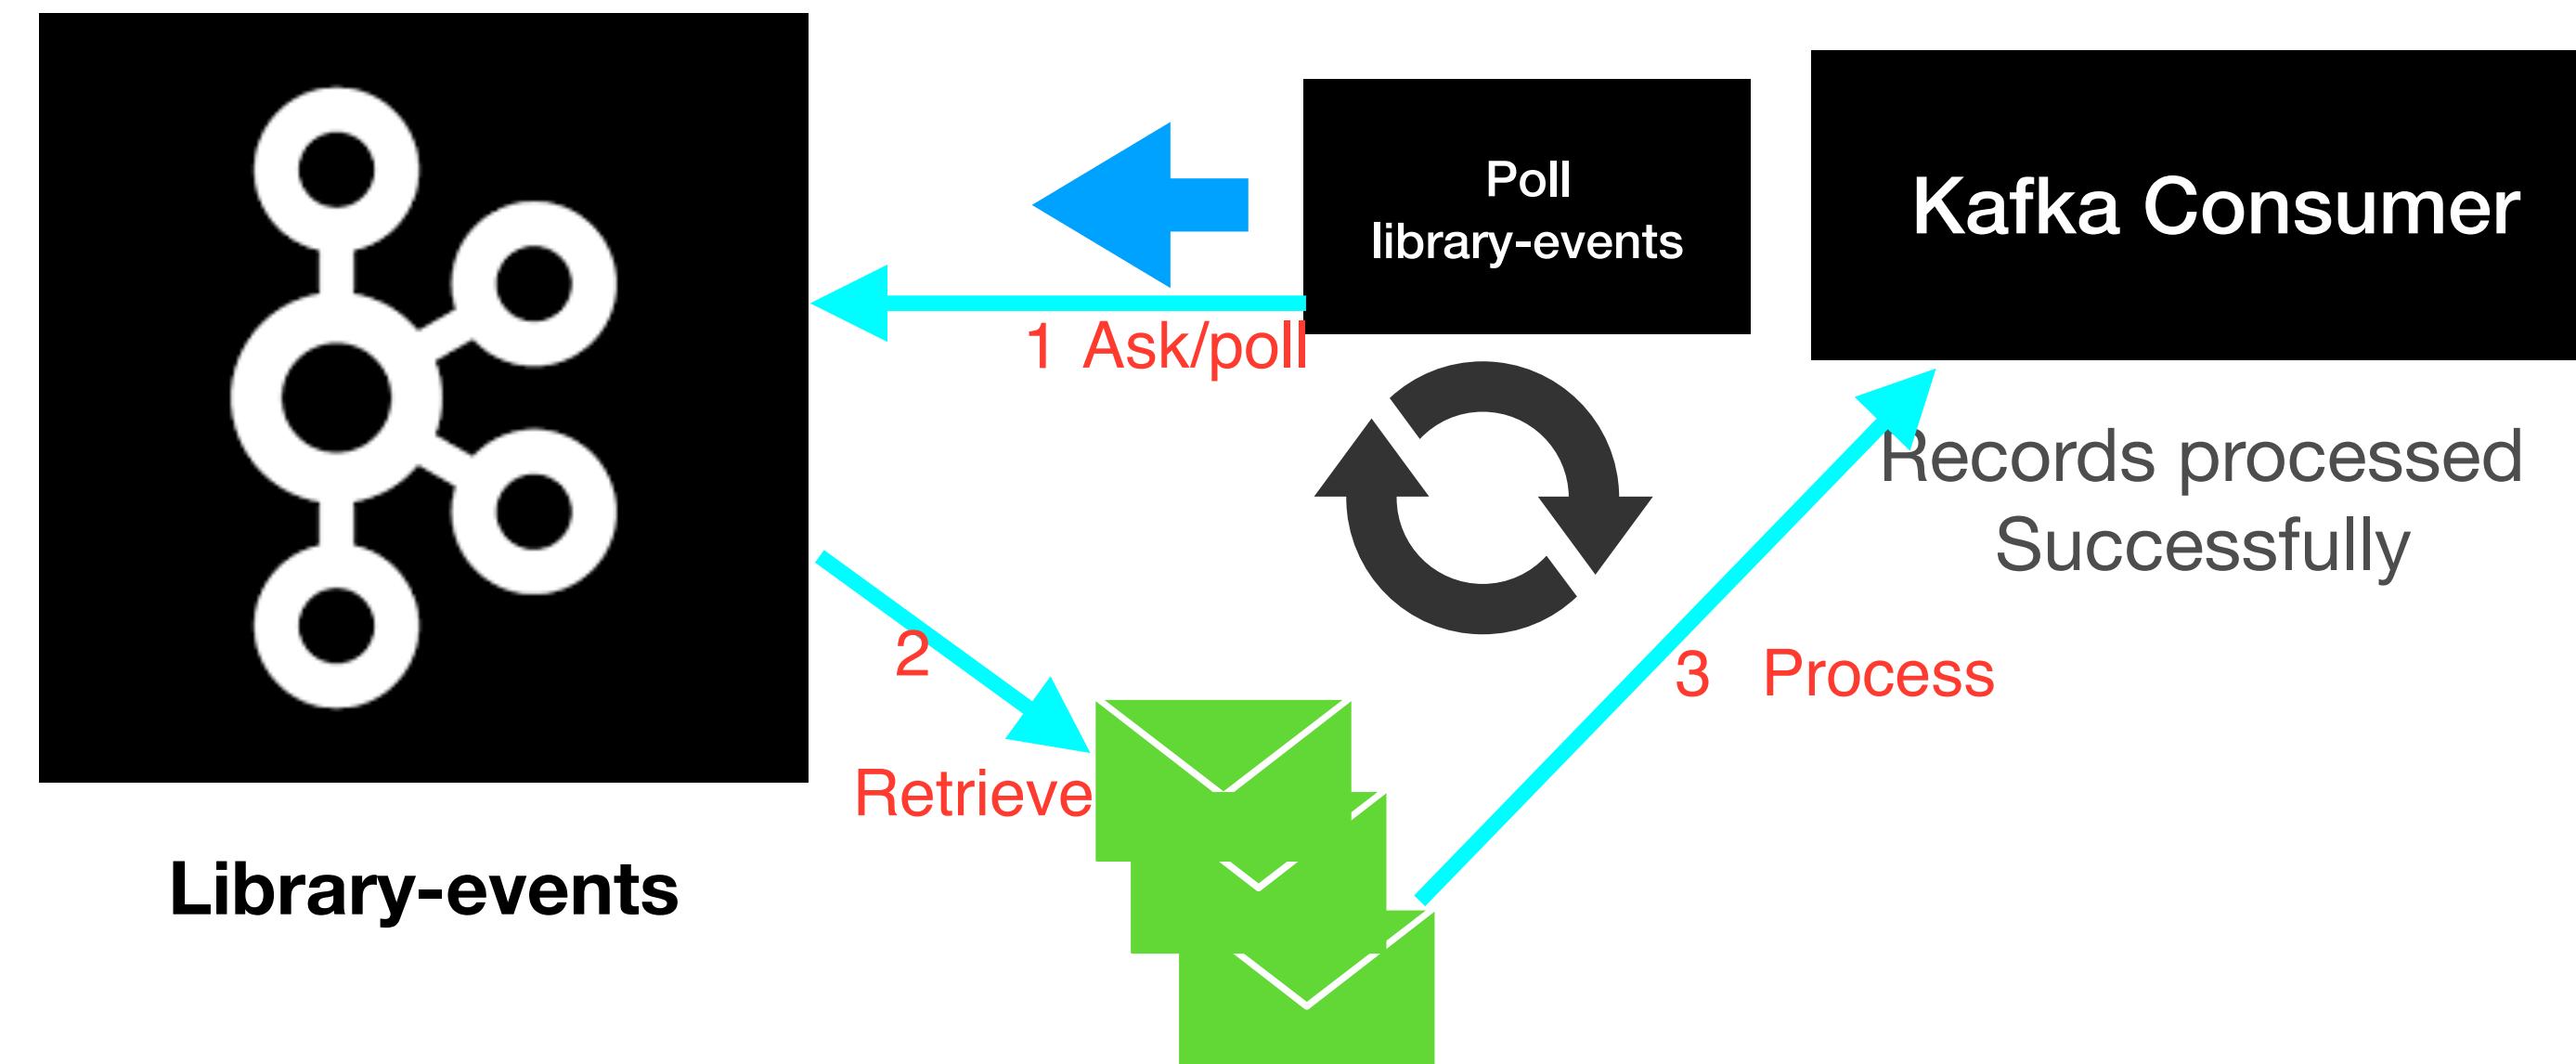
# Library Events Consumer



# **Spring Kafka Consumer**

# Kafka Consumer

TIP: Kafka Consumer, can poll records from multiple topics



# Spring Kafka Consumer

- `MessageListenerContainer` Interface  
Implementation
  - `KafkaMessageListenerContainer`
  - `ConcurrentMessageListenerContainer`
- `@KafkaListener` Annotation
  - Uses `ConcurrentMessageListenerContainer` behind the scenes

# KafkaMessageListenerContainer

- Implementation of MessageListenerContainer
- Polls the records
- Commits the Offsets
- Single Threaded

# **ConcurrentMessageListenerContainer**

Represents multiple **KafkaMessageListenerContainer**

# @KafkaListener

- This is the easiest way to build Kafka Consumer
- KafkaListener Sample Code

```
@KafkaListener(topics = {"${spring.kafka.topic}"})
public void onMessage(ConsumerRecord<Integer, String> consumerRecord) {
    log.info("OnMessage Record : {}", consumerRecord);
}
```

- Configuration Sample Code

```
@Configuration
@EnableKafka
@Slf4j
public class LibraryEventsConsumerConfig {
```

# KafkaConsumer Config

**key-deserializer:** org.apache.kafka.common.serialization.IntegerDeserializer  
**value-deserializer:** org.apache.kafka.common.serialization.StringDeserializer  
**group-id:** library-events-listener-group

# **Consumer Groups & Rebalance**

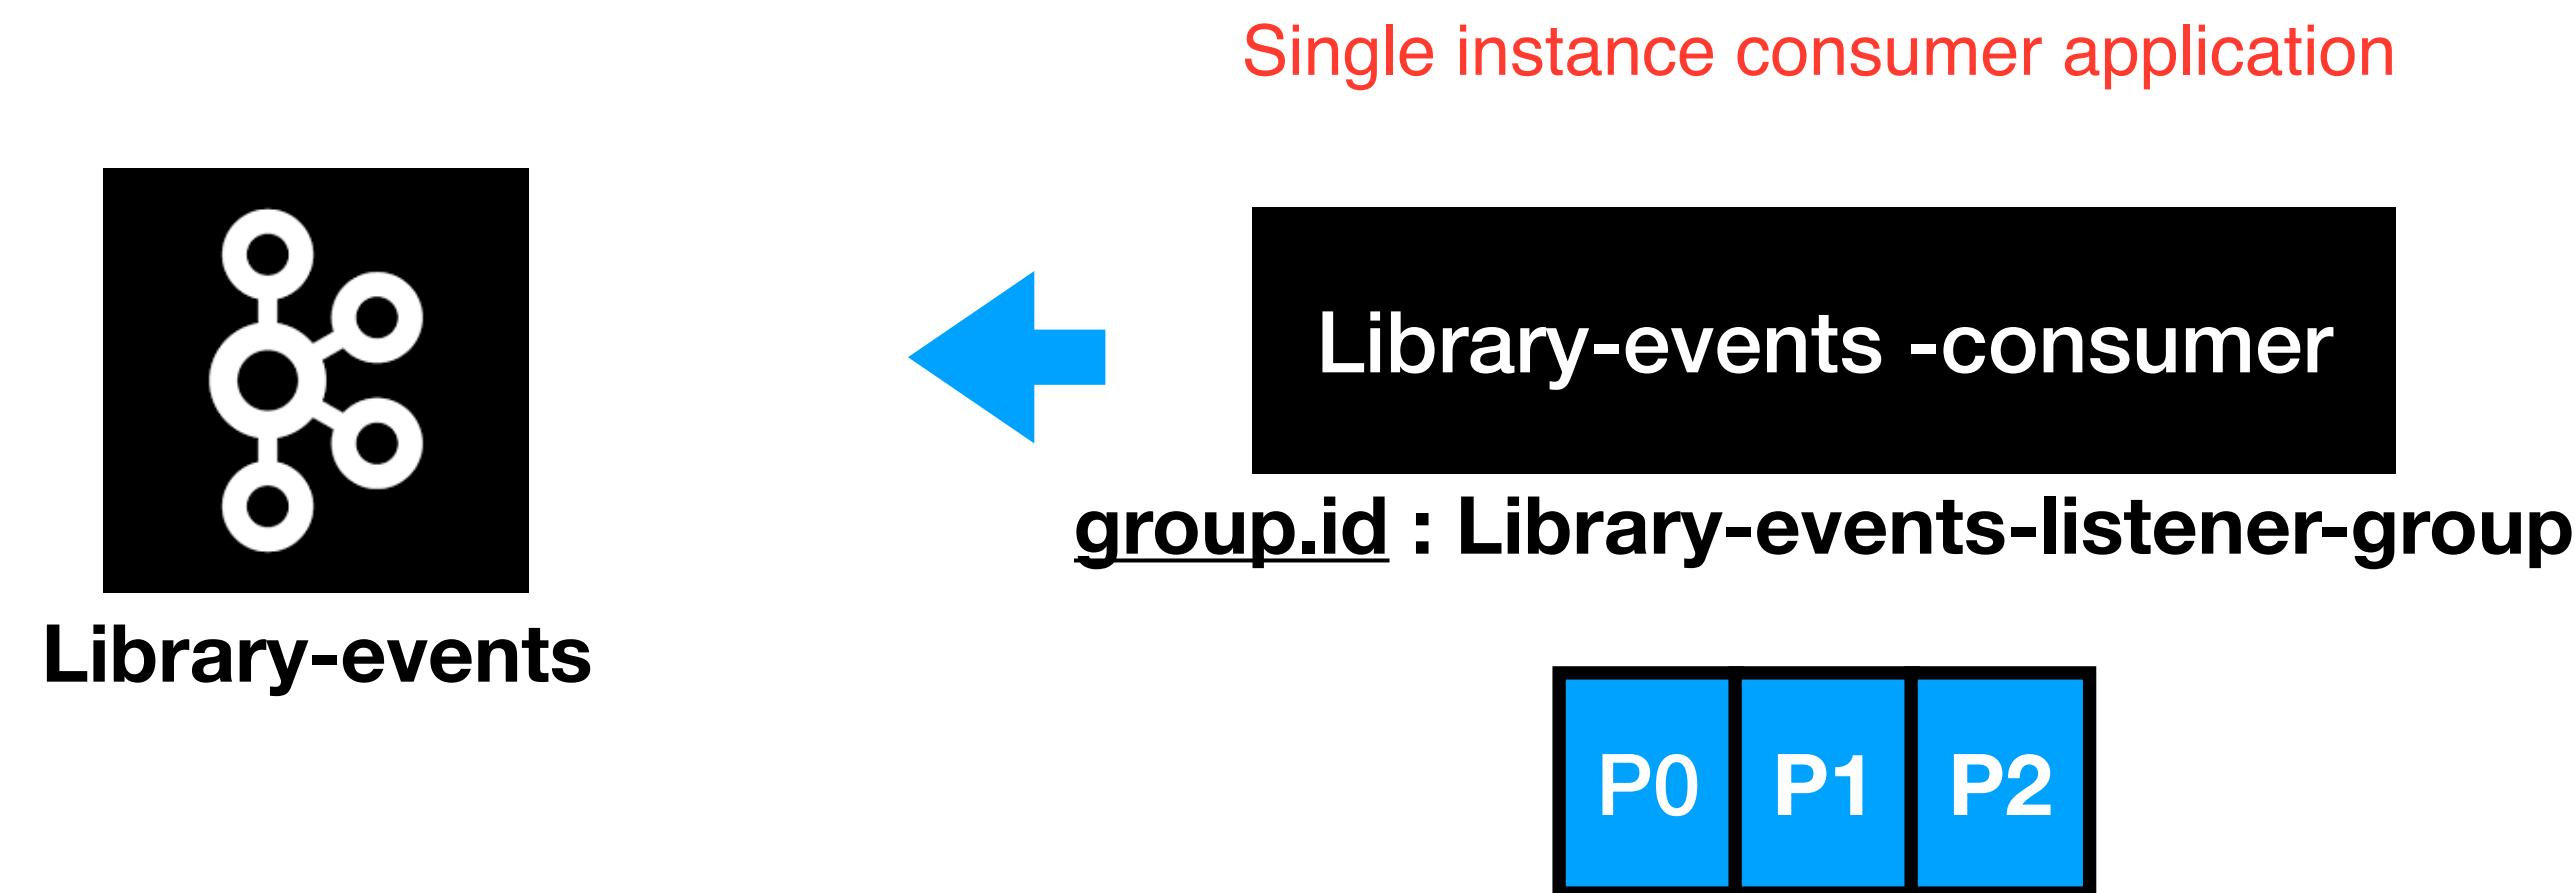
# Consumer Groups

Consumer Groups are the foundation for a scalable message consumption

Multiple instances of the same application with the same group id.

# Rebalance

- Changing the partition ownership from one consumer to another



# Rebalance

- Changing the partition ownership from one consumer to another

Here we have two instances of the same application running on 8081 and 8082

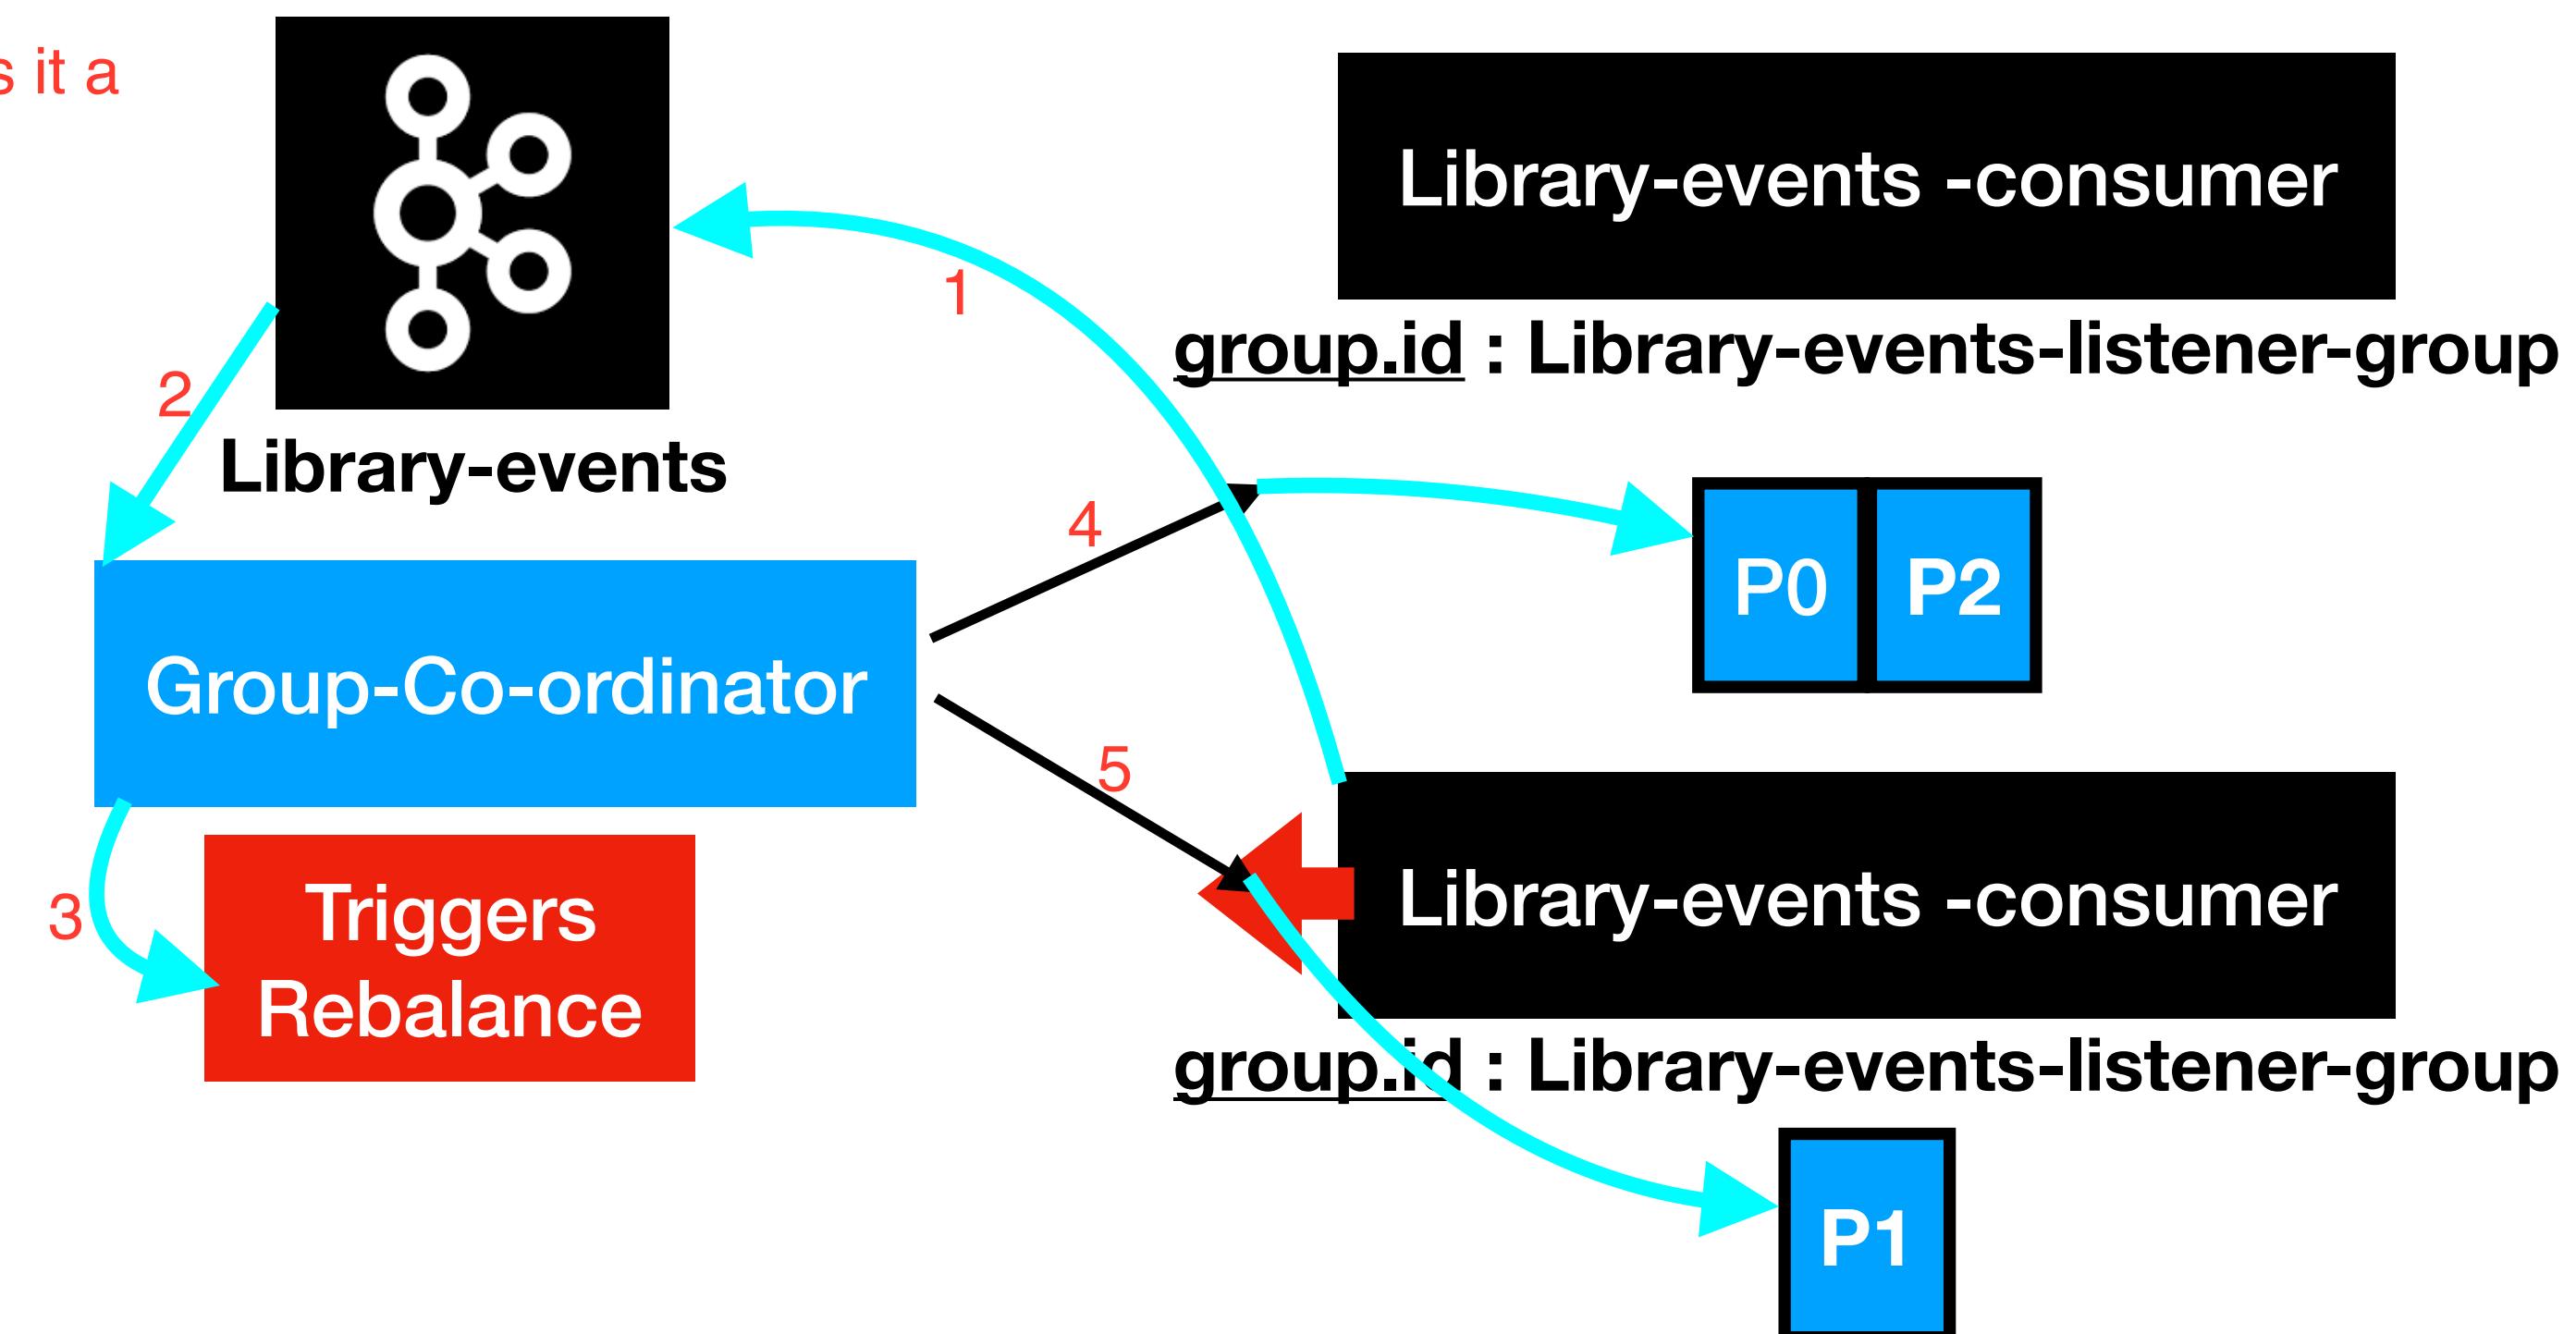
1. Call the Kafka Broker which intern connect to “Group-Co-Ordinator”

2. Group Coordinator checks is it a same group id? if yes

3. It will trigger a “Rebalance”. What it does?

Oh! I have one more instance of the same application in that case, i have to redistribute the partitions

4. Using “Rebalance”, Group-Co-ordinator will assign the odd partitions to 1st consumer (refer 4) and event partitions to 2nd consumer (refer 5)



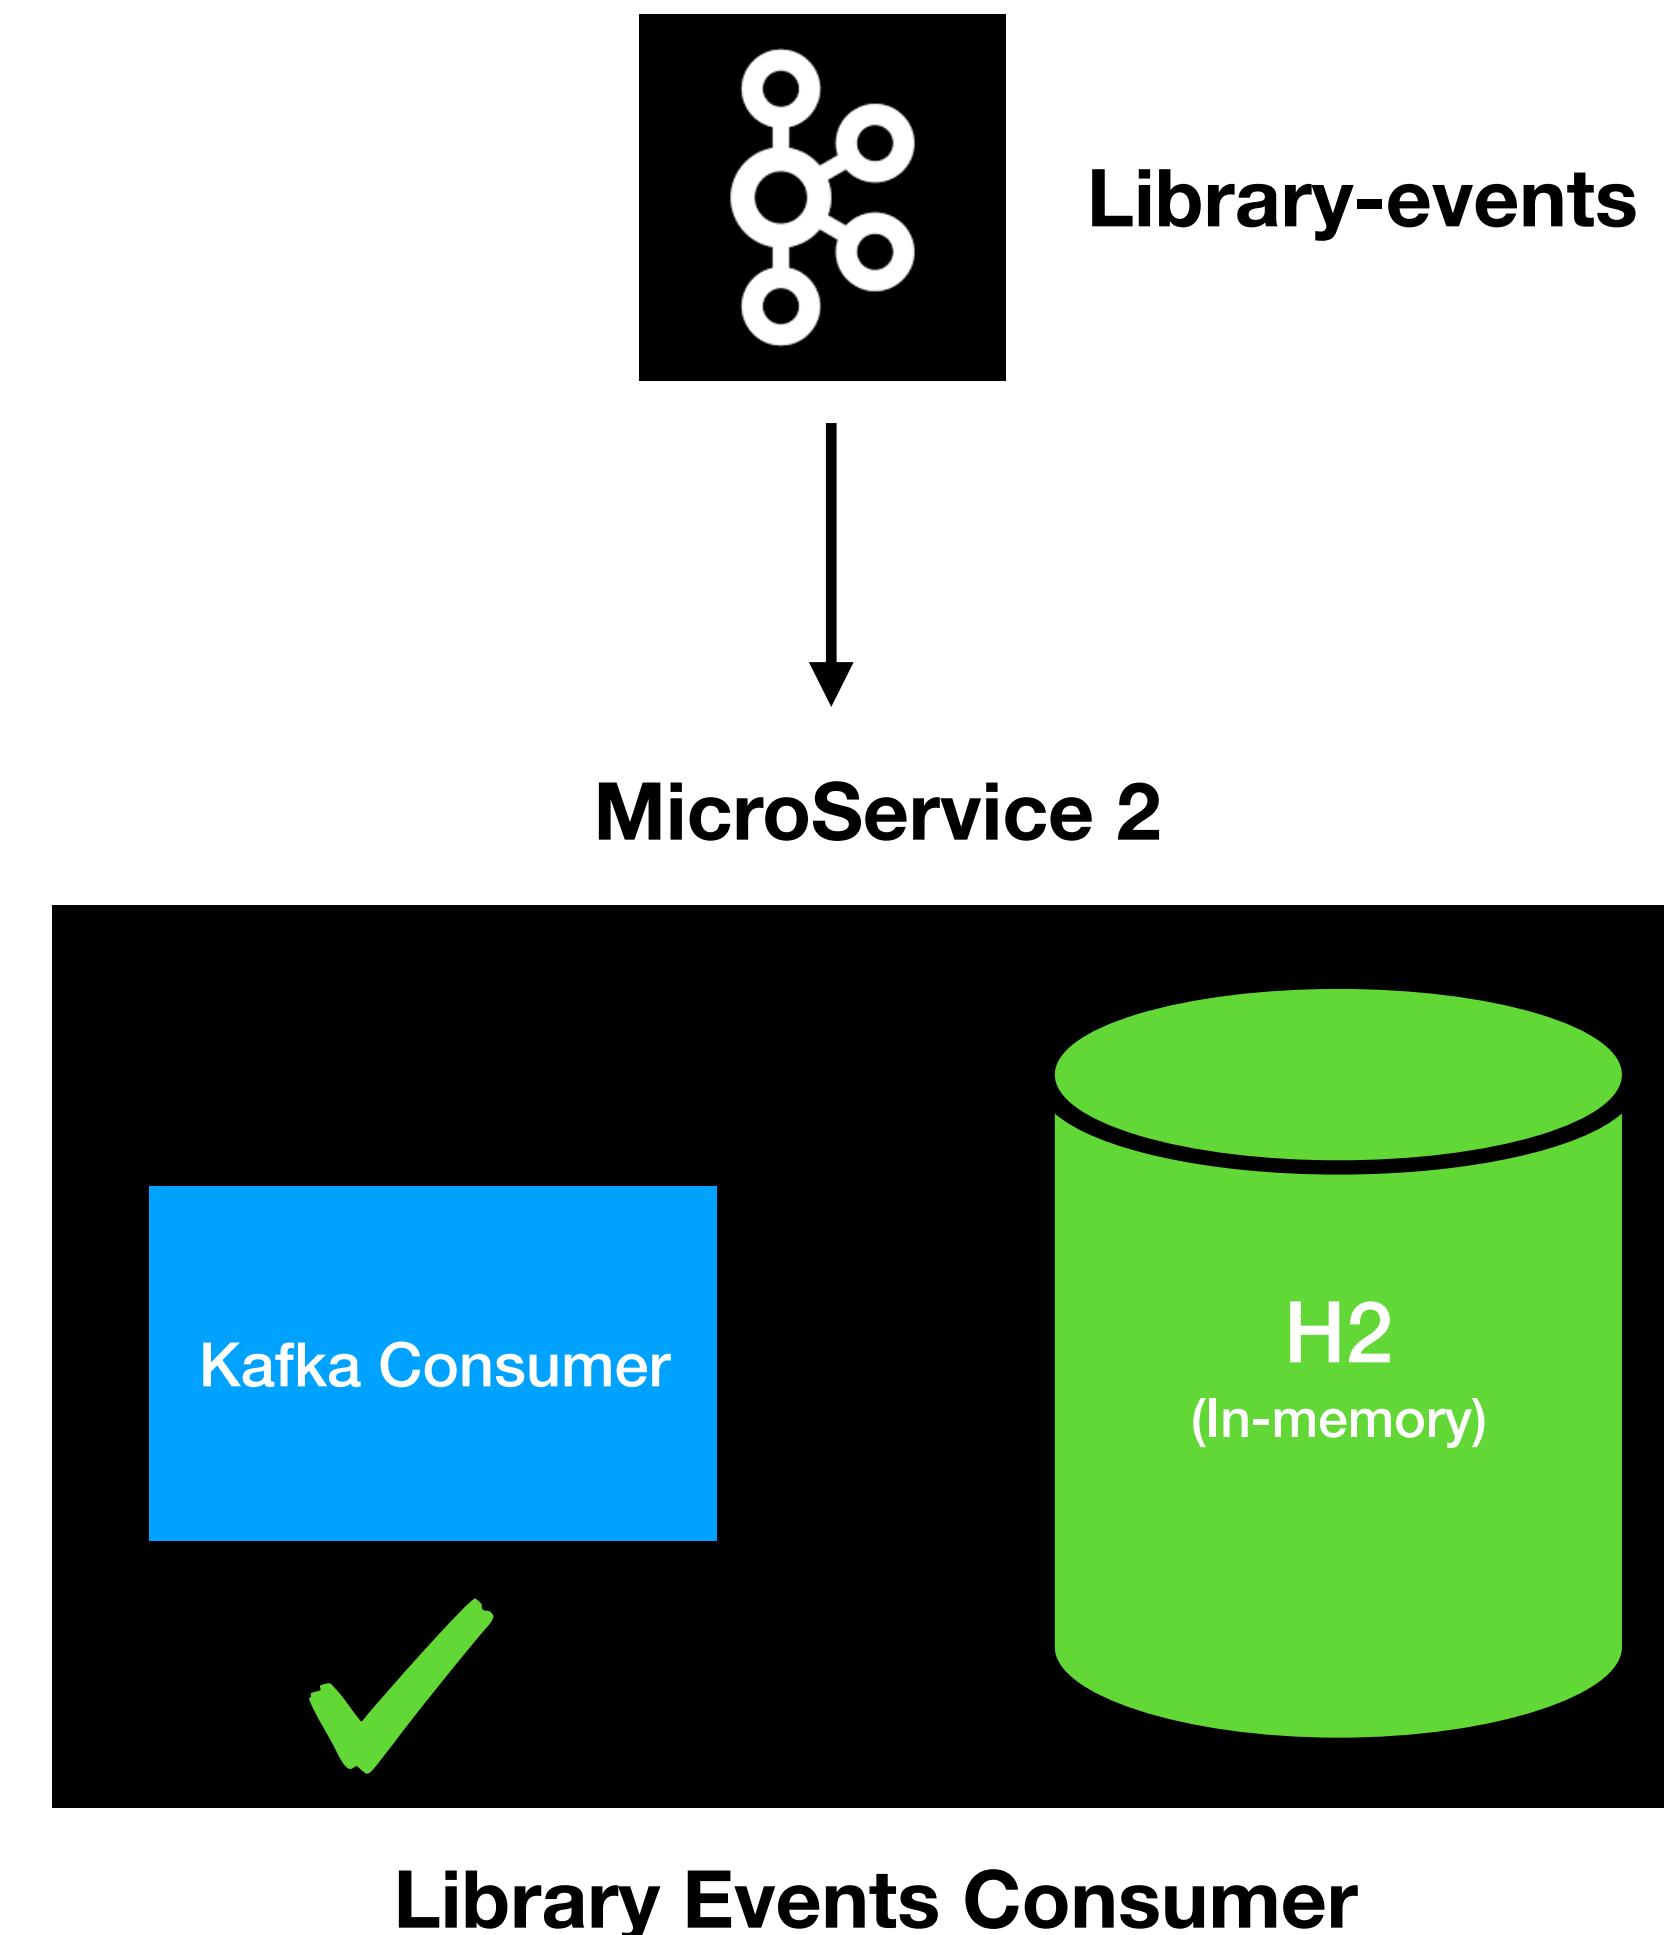
Before look at this slide,

- 1.run your producer application
- 2.run your consumer application
- 3.send messages through REST
- 4.Message reached consumer app
- 5.down your consumer app
- 6.Again start consumer app. Now it won't display the previously sent messages. why?
- 7.Because, once record processed successfully, behind the scene consumer commits the offsets in broker in “`__consumer_offsets`”

# Committing Offsets



# Library Events Consumer



# **Integration Testing For Real DataBases**

# Integration Testing using Real Databases

- Different aspects of writing unit and integration testing
- Integration testing using **TestContainers**



The image is a composite of two parts. On the left, a man with glasses and a blue shirt stands behind a podium, speaking into a microphone. The podium has a purple cloth with the "SpringOne Platform" logo. On the right, a slide from a presentation is shown against a black background. The slide has a teal header with the "SpringOne Platform by Pivotal" logo. The title "Testing Spring Boot Applications" is displayed in large, bold, teal font. Below the title, the name "Andy Wilkinson" and the handle "@ankinson" are listed. At the bottom of the slide, there is another "SpringOne Platform by Pivotal" logo.

SpringOne Platform by Pivotal

## Testing Spring Boot Applications

Andy Wilkinson  
@ankinson

SpringOne Platform by Pivotal

October 7–10 / Austin, TX P

# TestContainers

- What are TestContainers?
  - Testcontainers is a Java library that supports JUnit tests, providing lightweight, throwaway instances of common databases, Selenium web browsers, or anything else that can run in a **Docker** container.
- More Info about TestContainers - <https://www.testcontainers.org/>

# Retry in Kafka Consumer

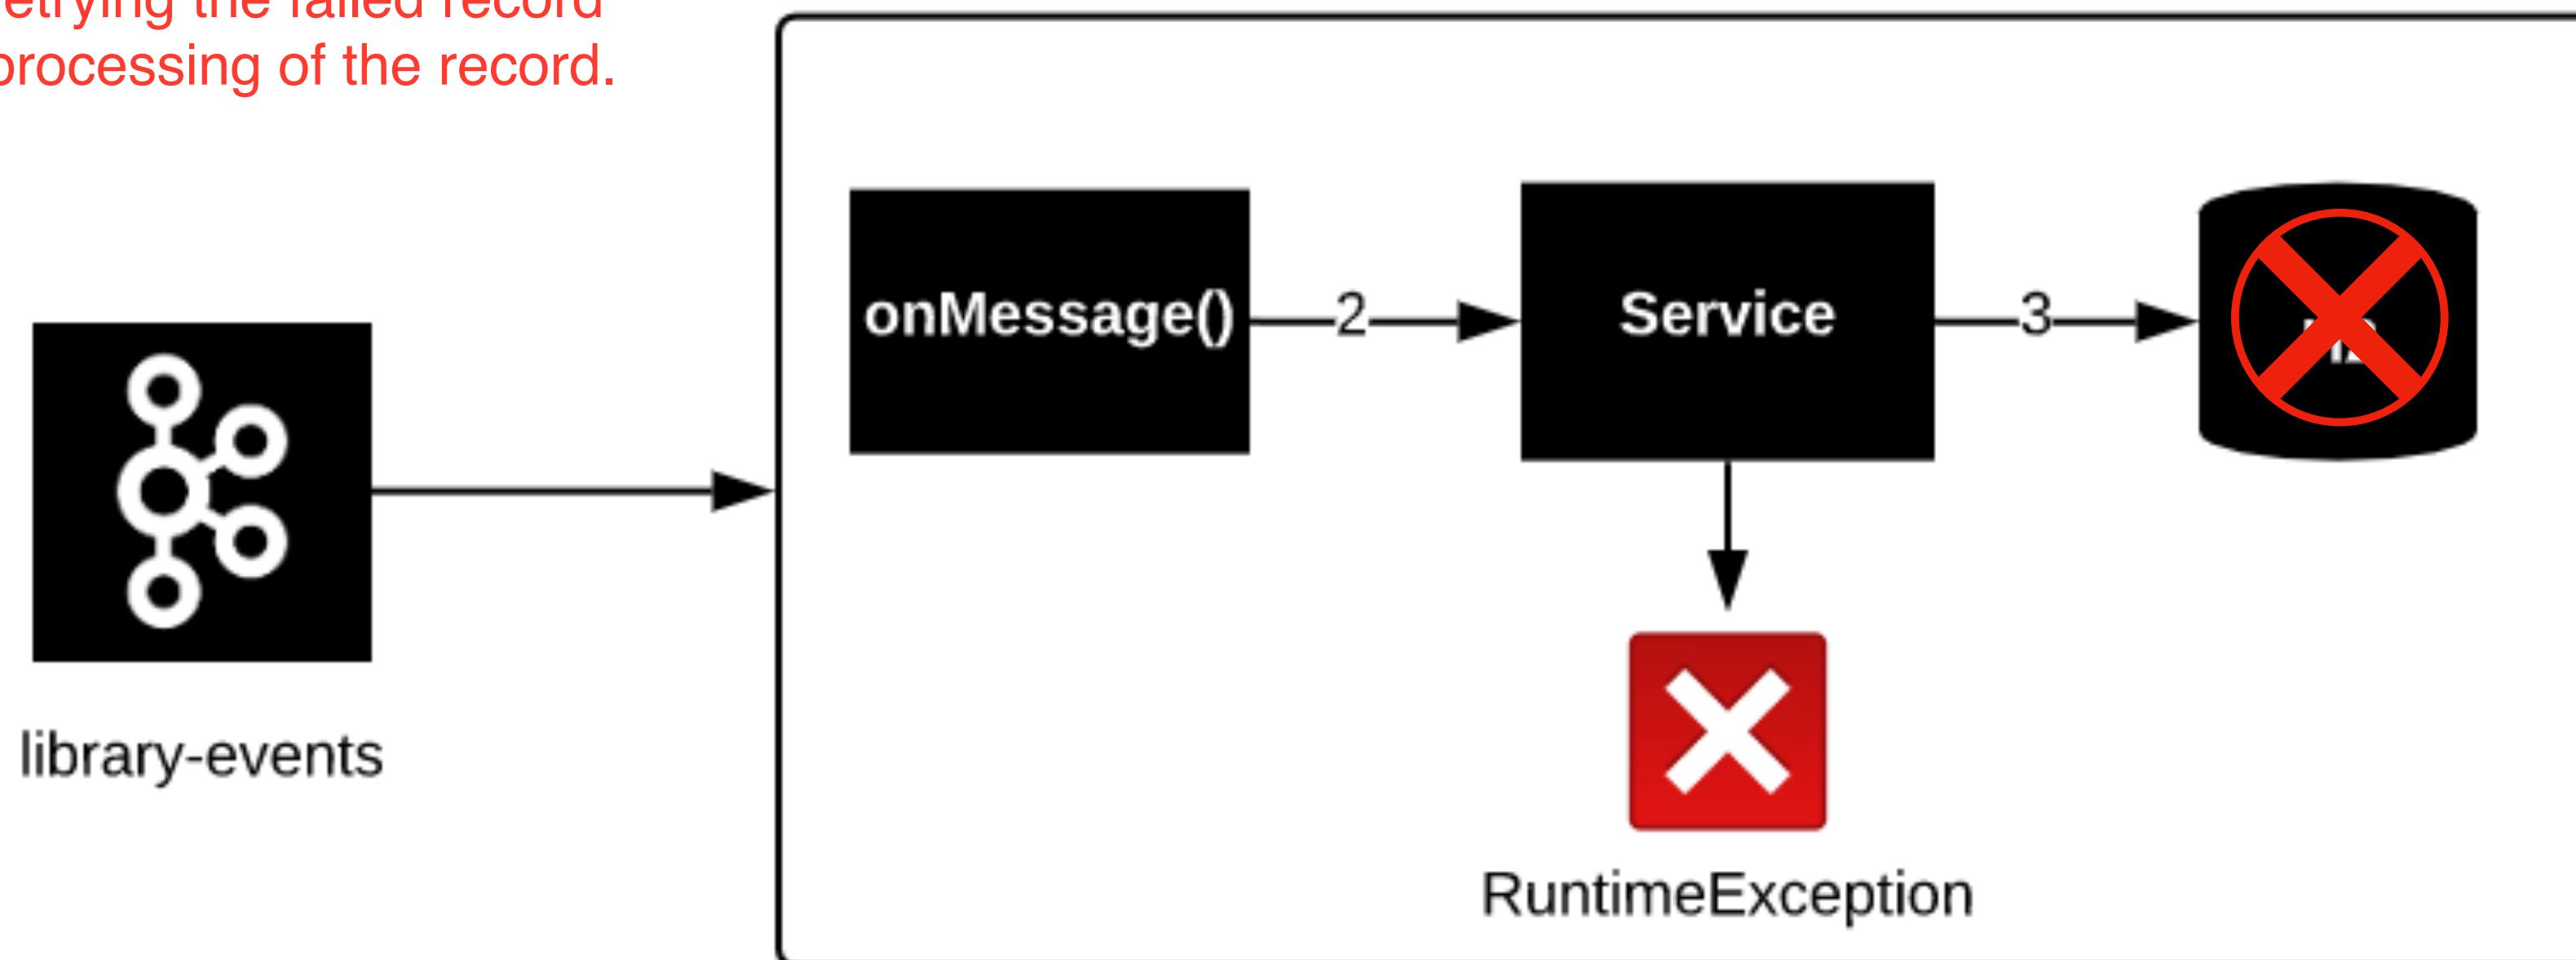
# Error in Kafka Consumer

In this lecture we will learn about the techniques to enable retry for a failed record.

In this example here the consumer reads a message and for some reason the app ran into a temporary network issues with the database which might result in a runtime exception in the service layer here.

In these kind of scenarios retrying the failed record would result in successful processing of the record.

## Library Events Consumer

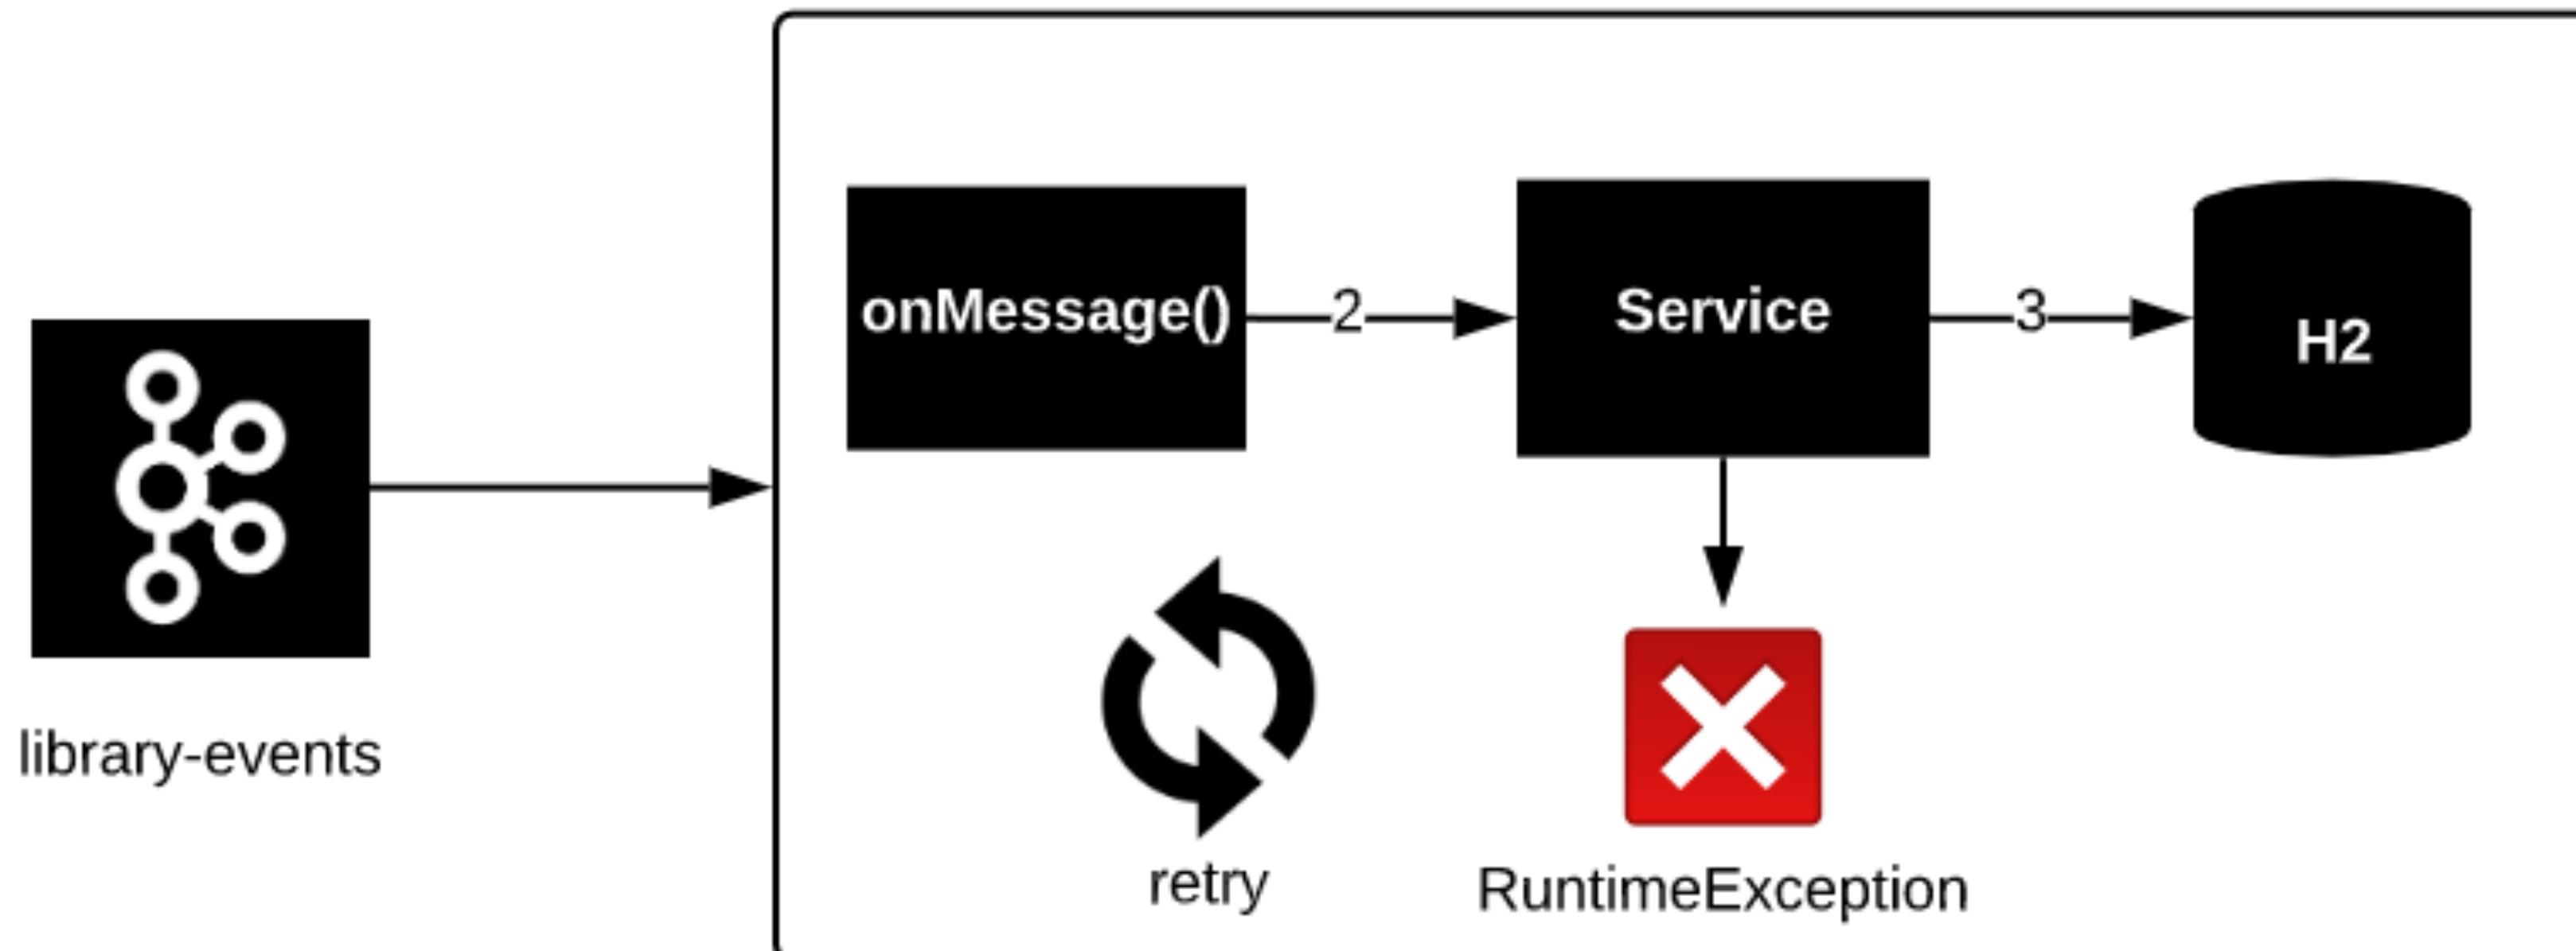


# Retry in Kafka Consumer

The software that we build today is pretty much Microservices the network causes are pretty much everywhere.

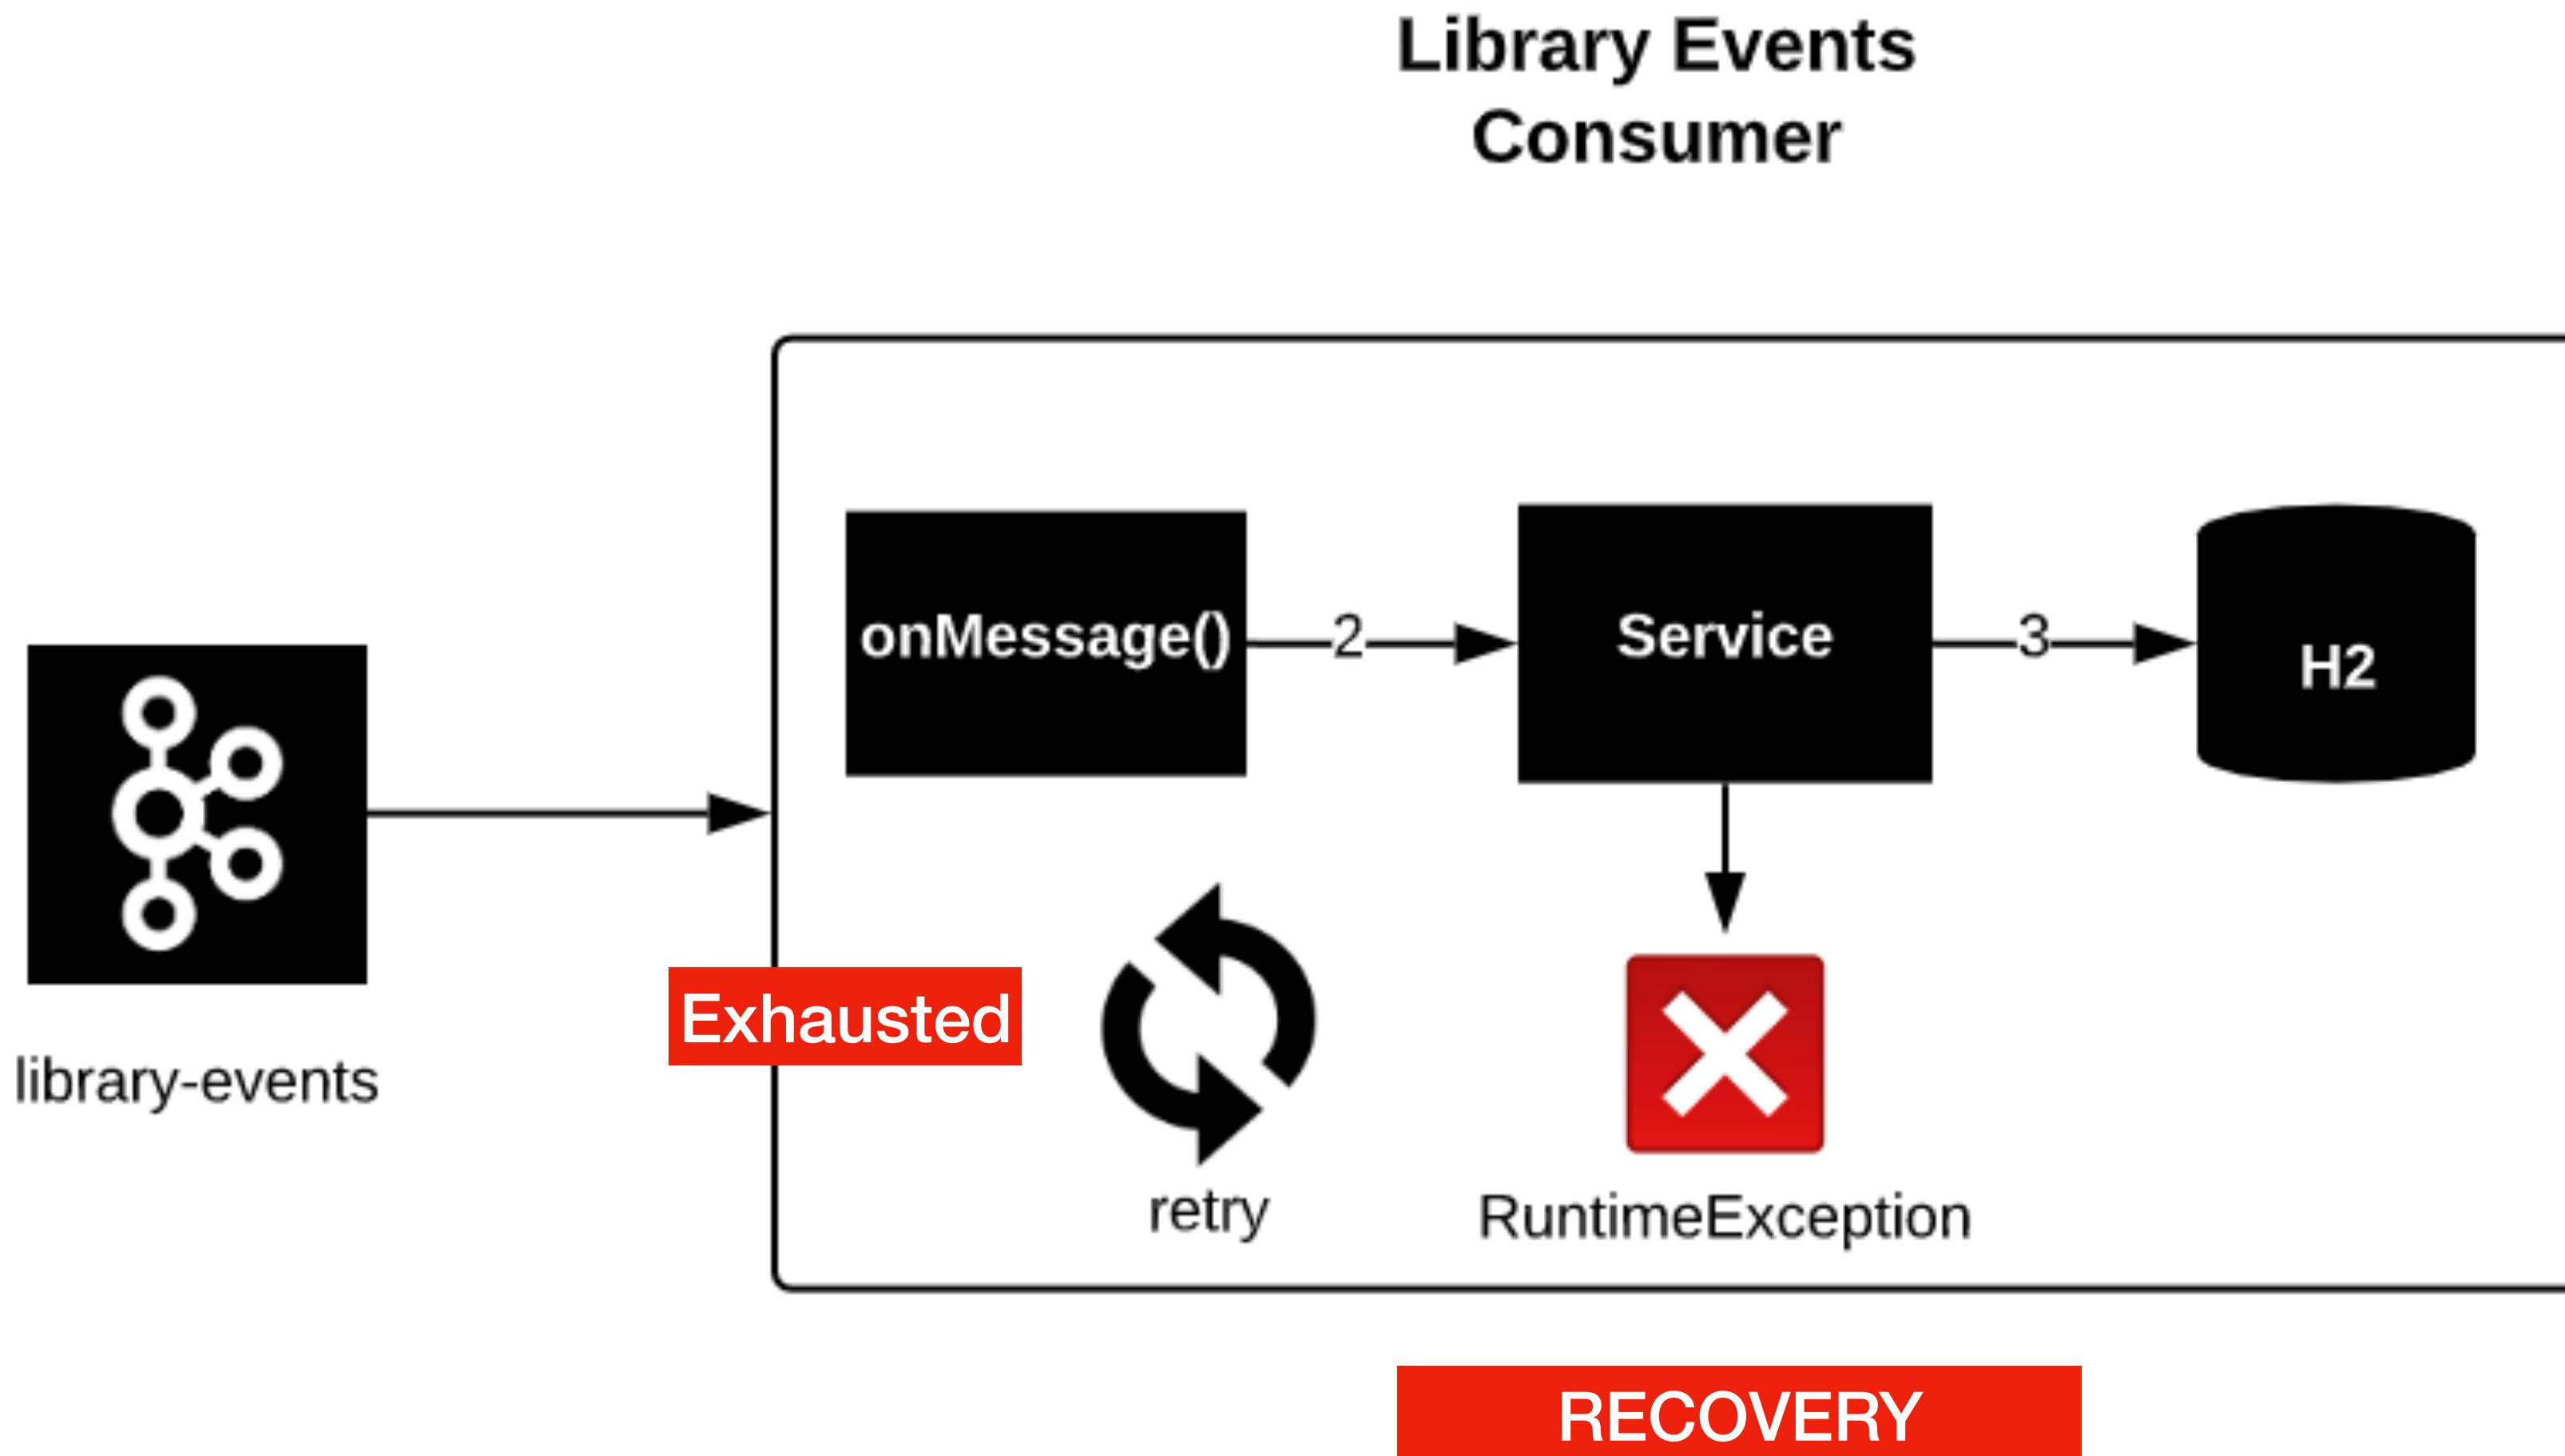
## Library Events Consumer

It's really handy to have retry enabled and the processing logic so that we can avoid a lot of errors.

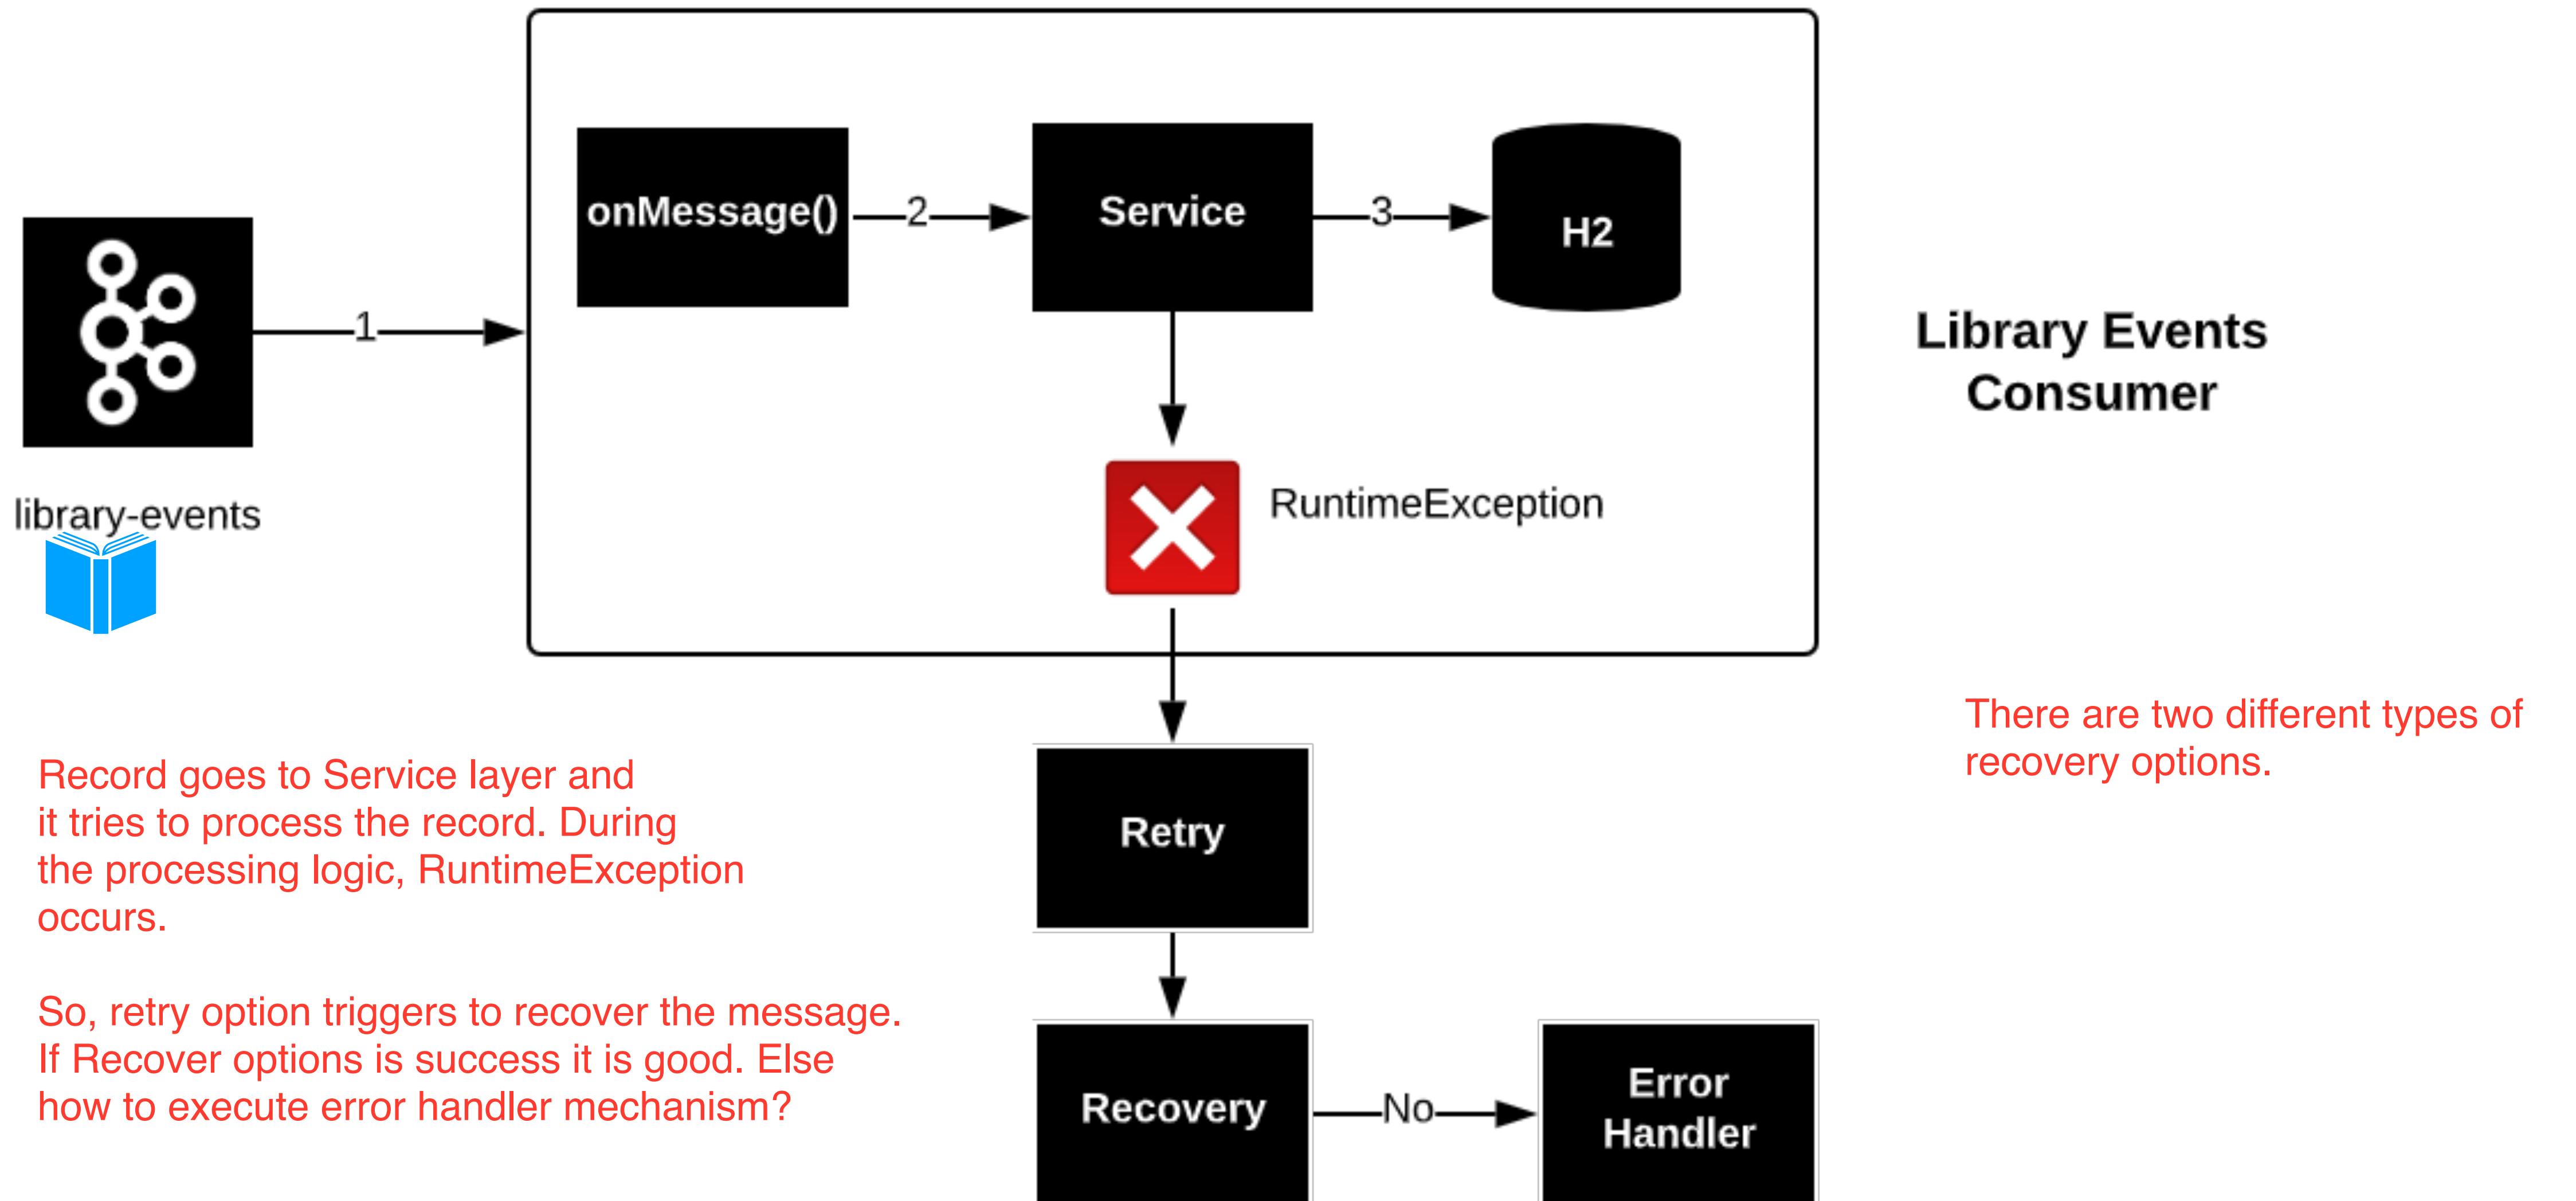


# Recovery in Kafka Consumer

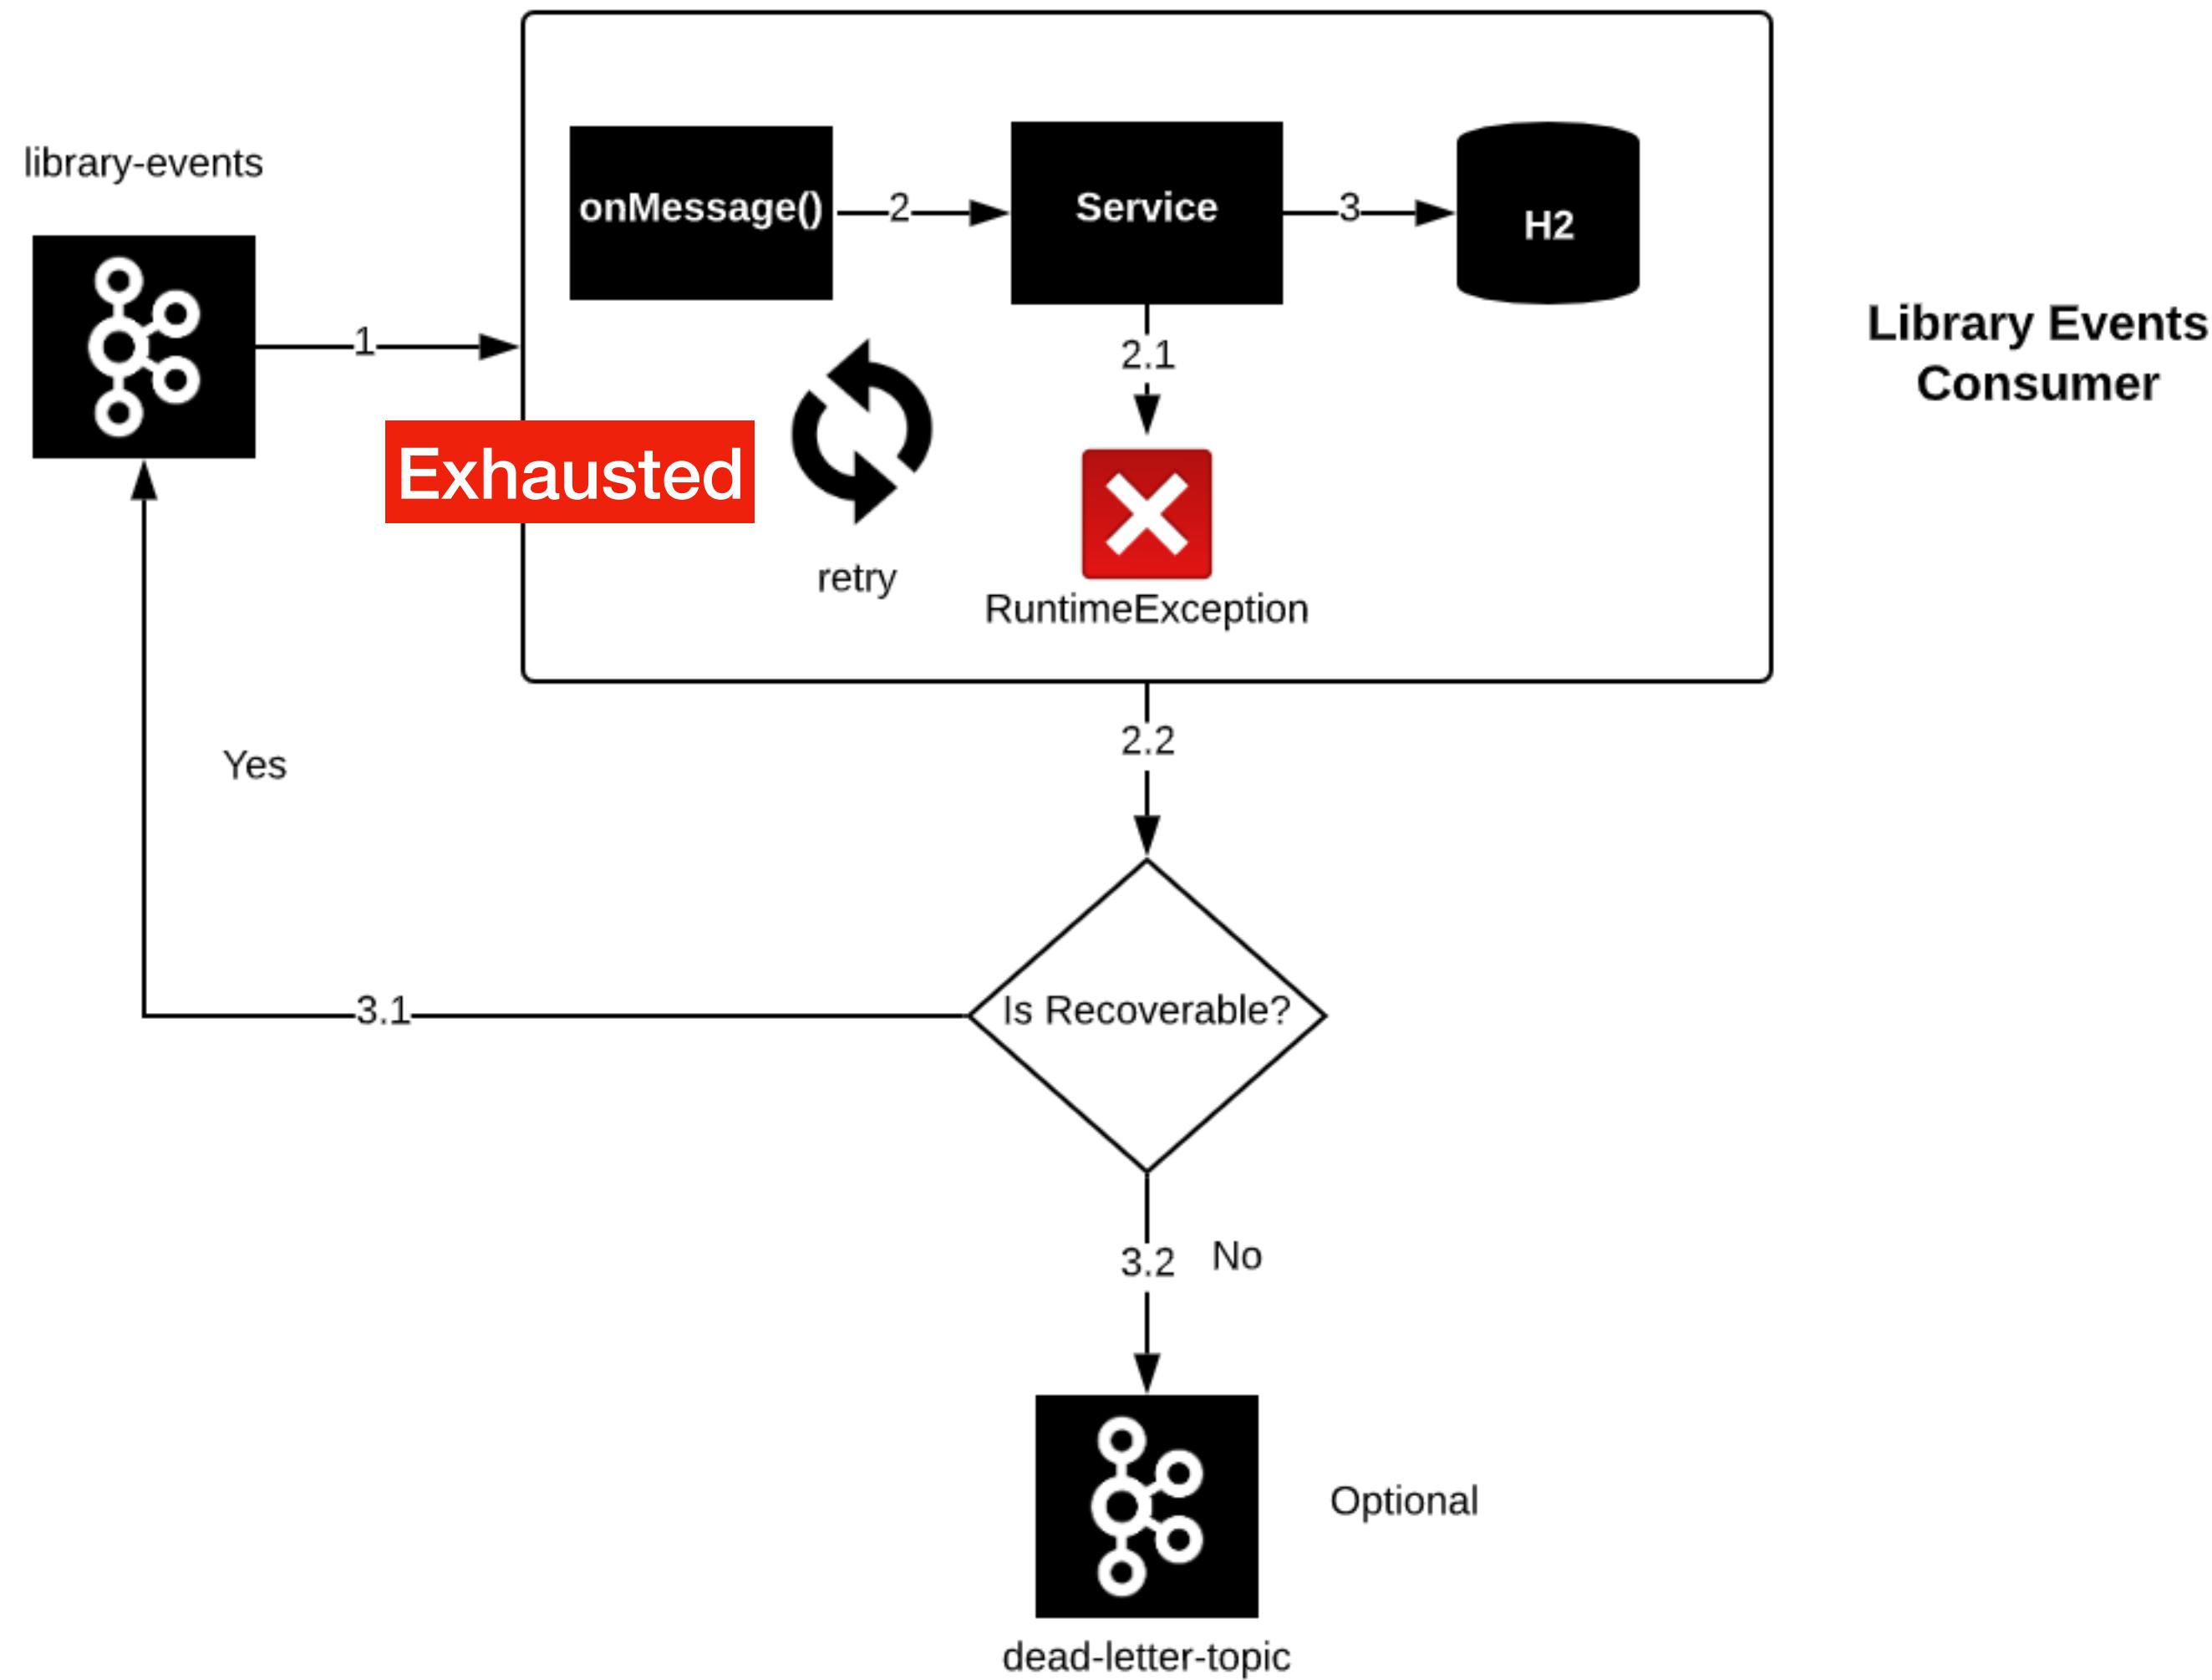
# Recovery in Kafka Consumer



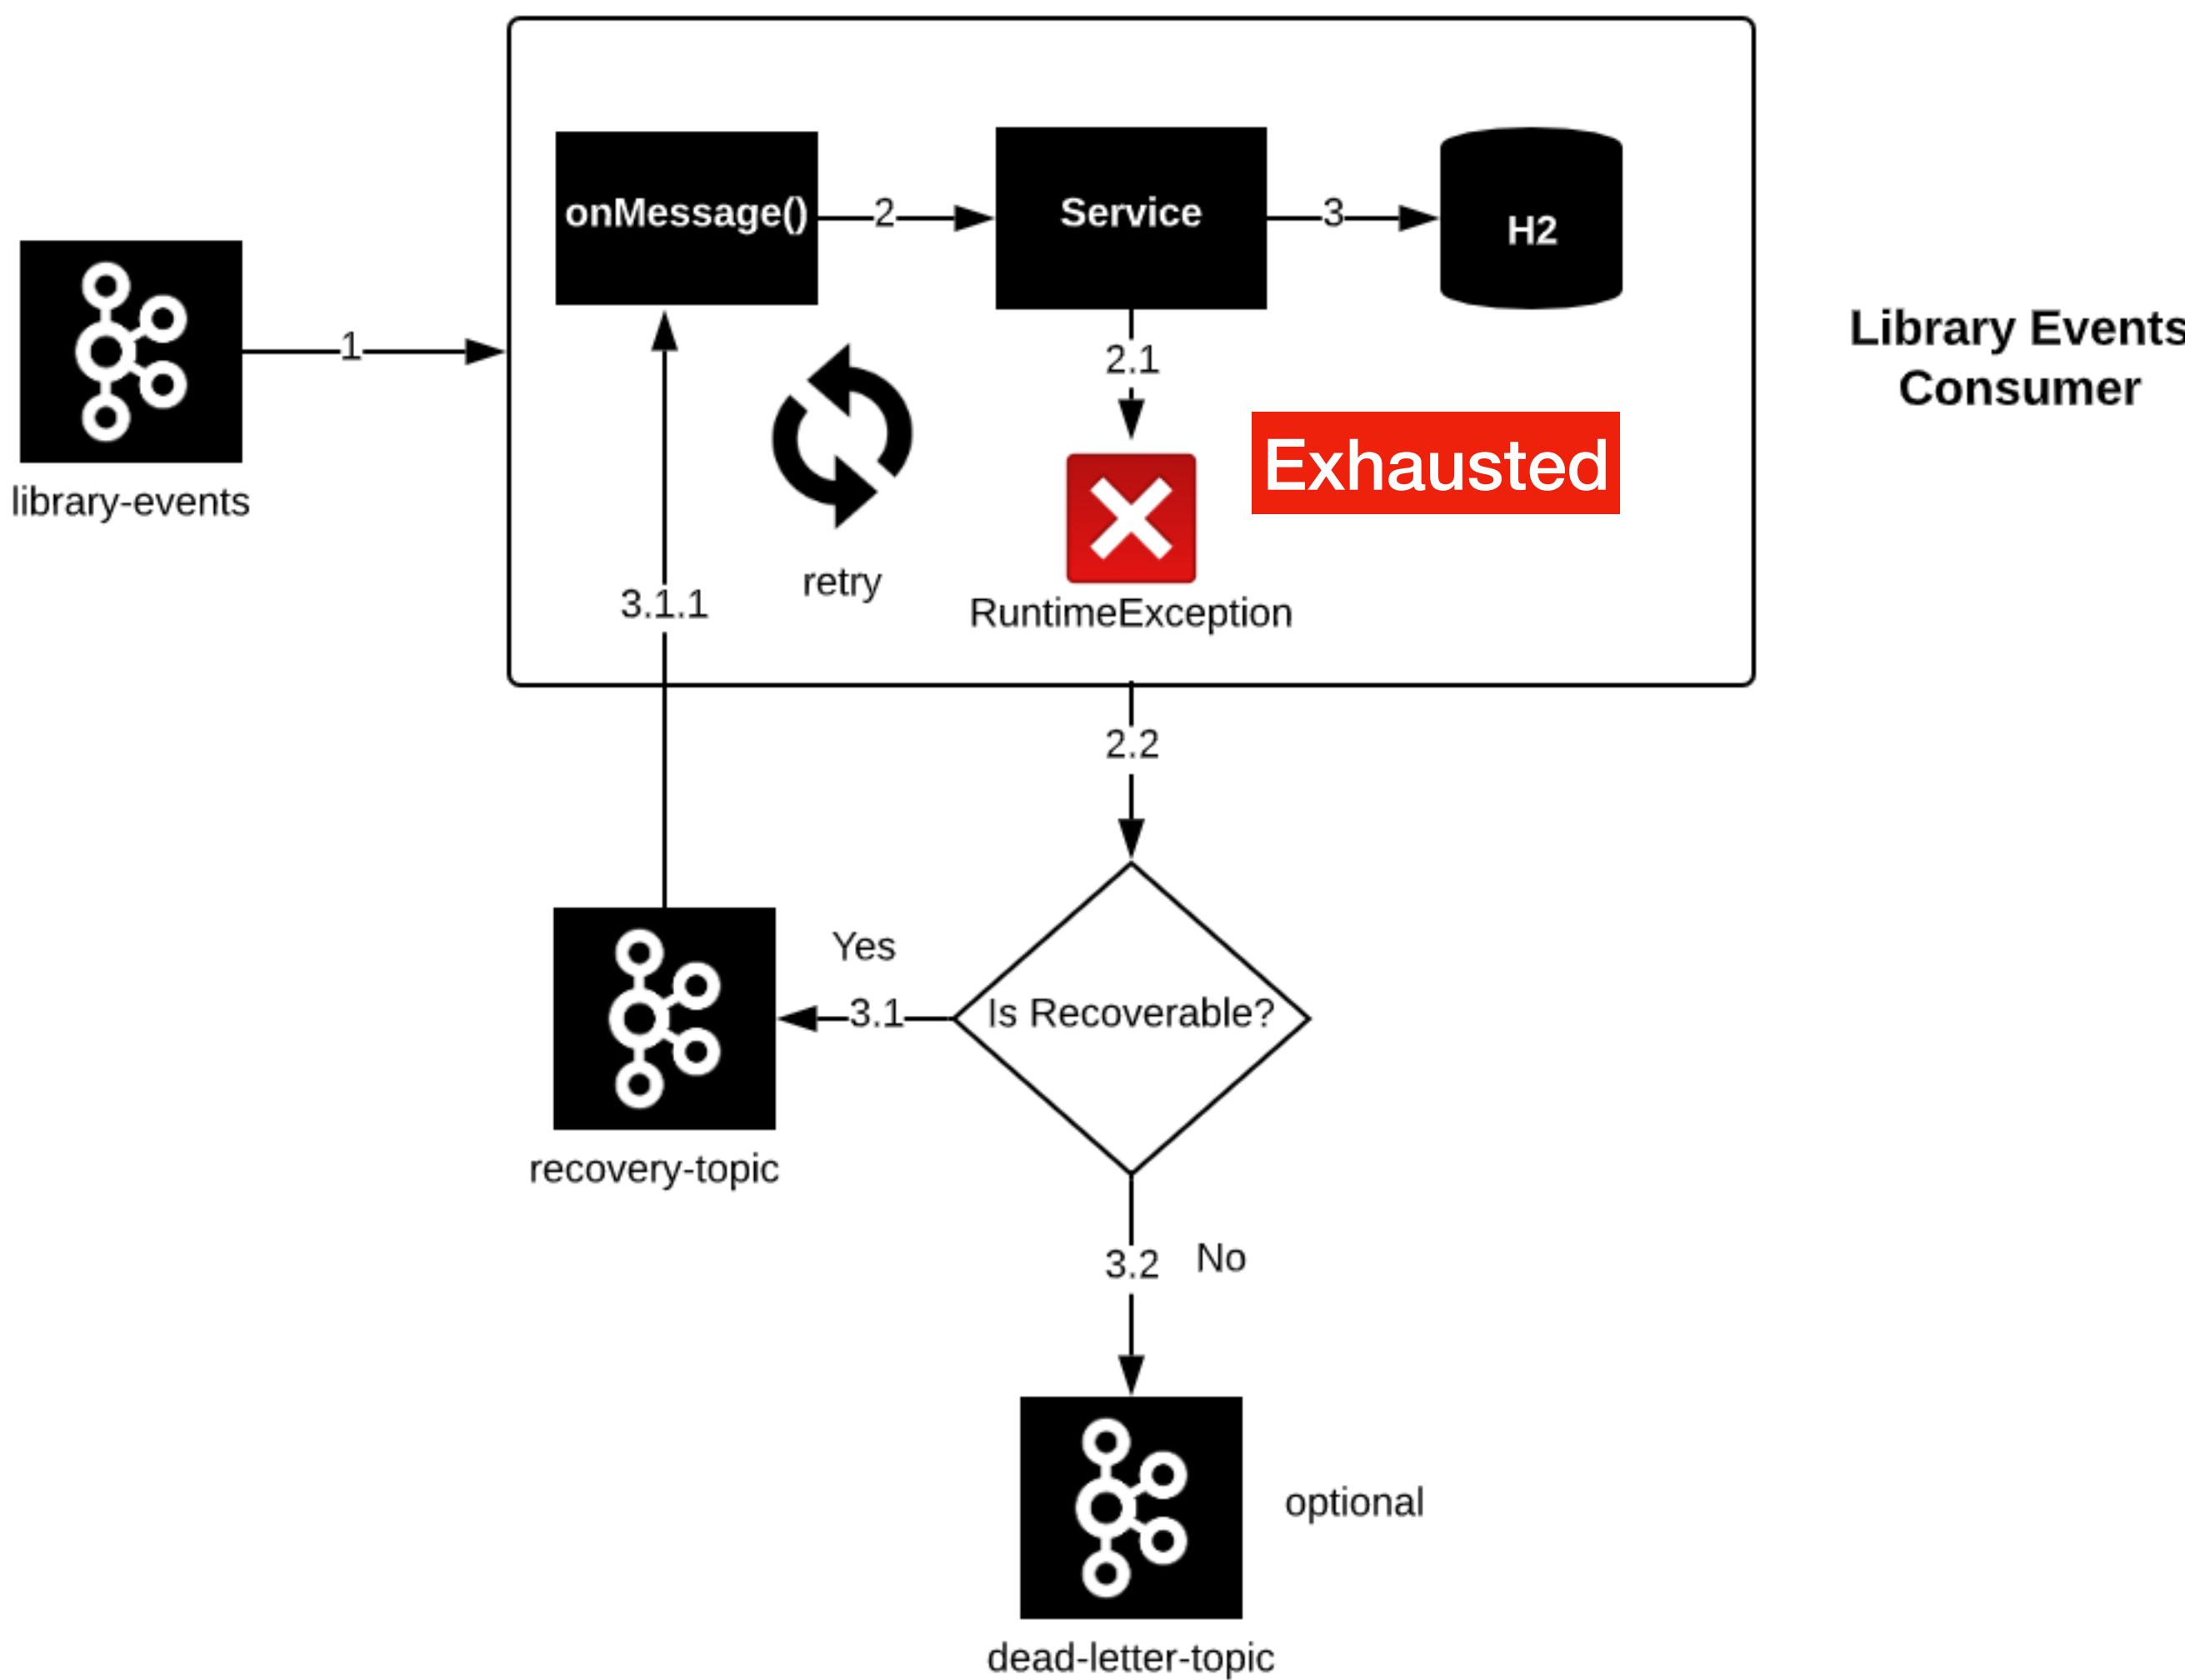
# Retry and Recovery



# Recovery - Type 1



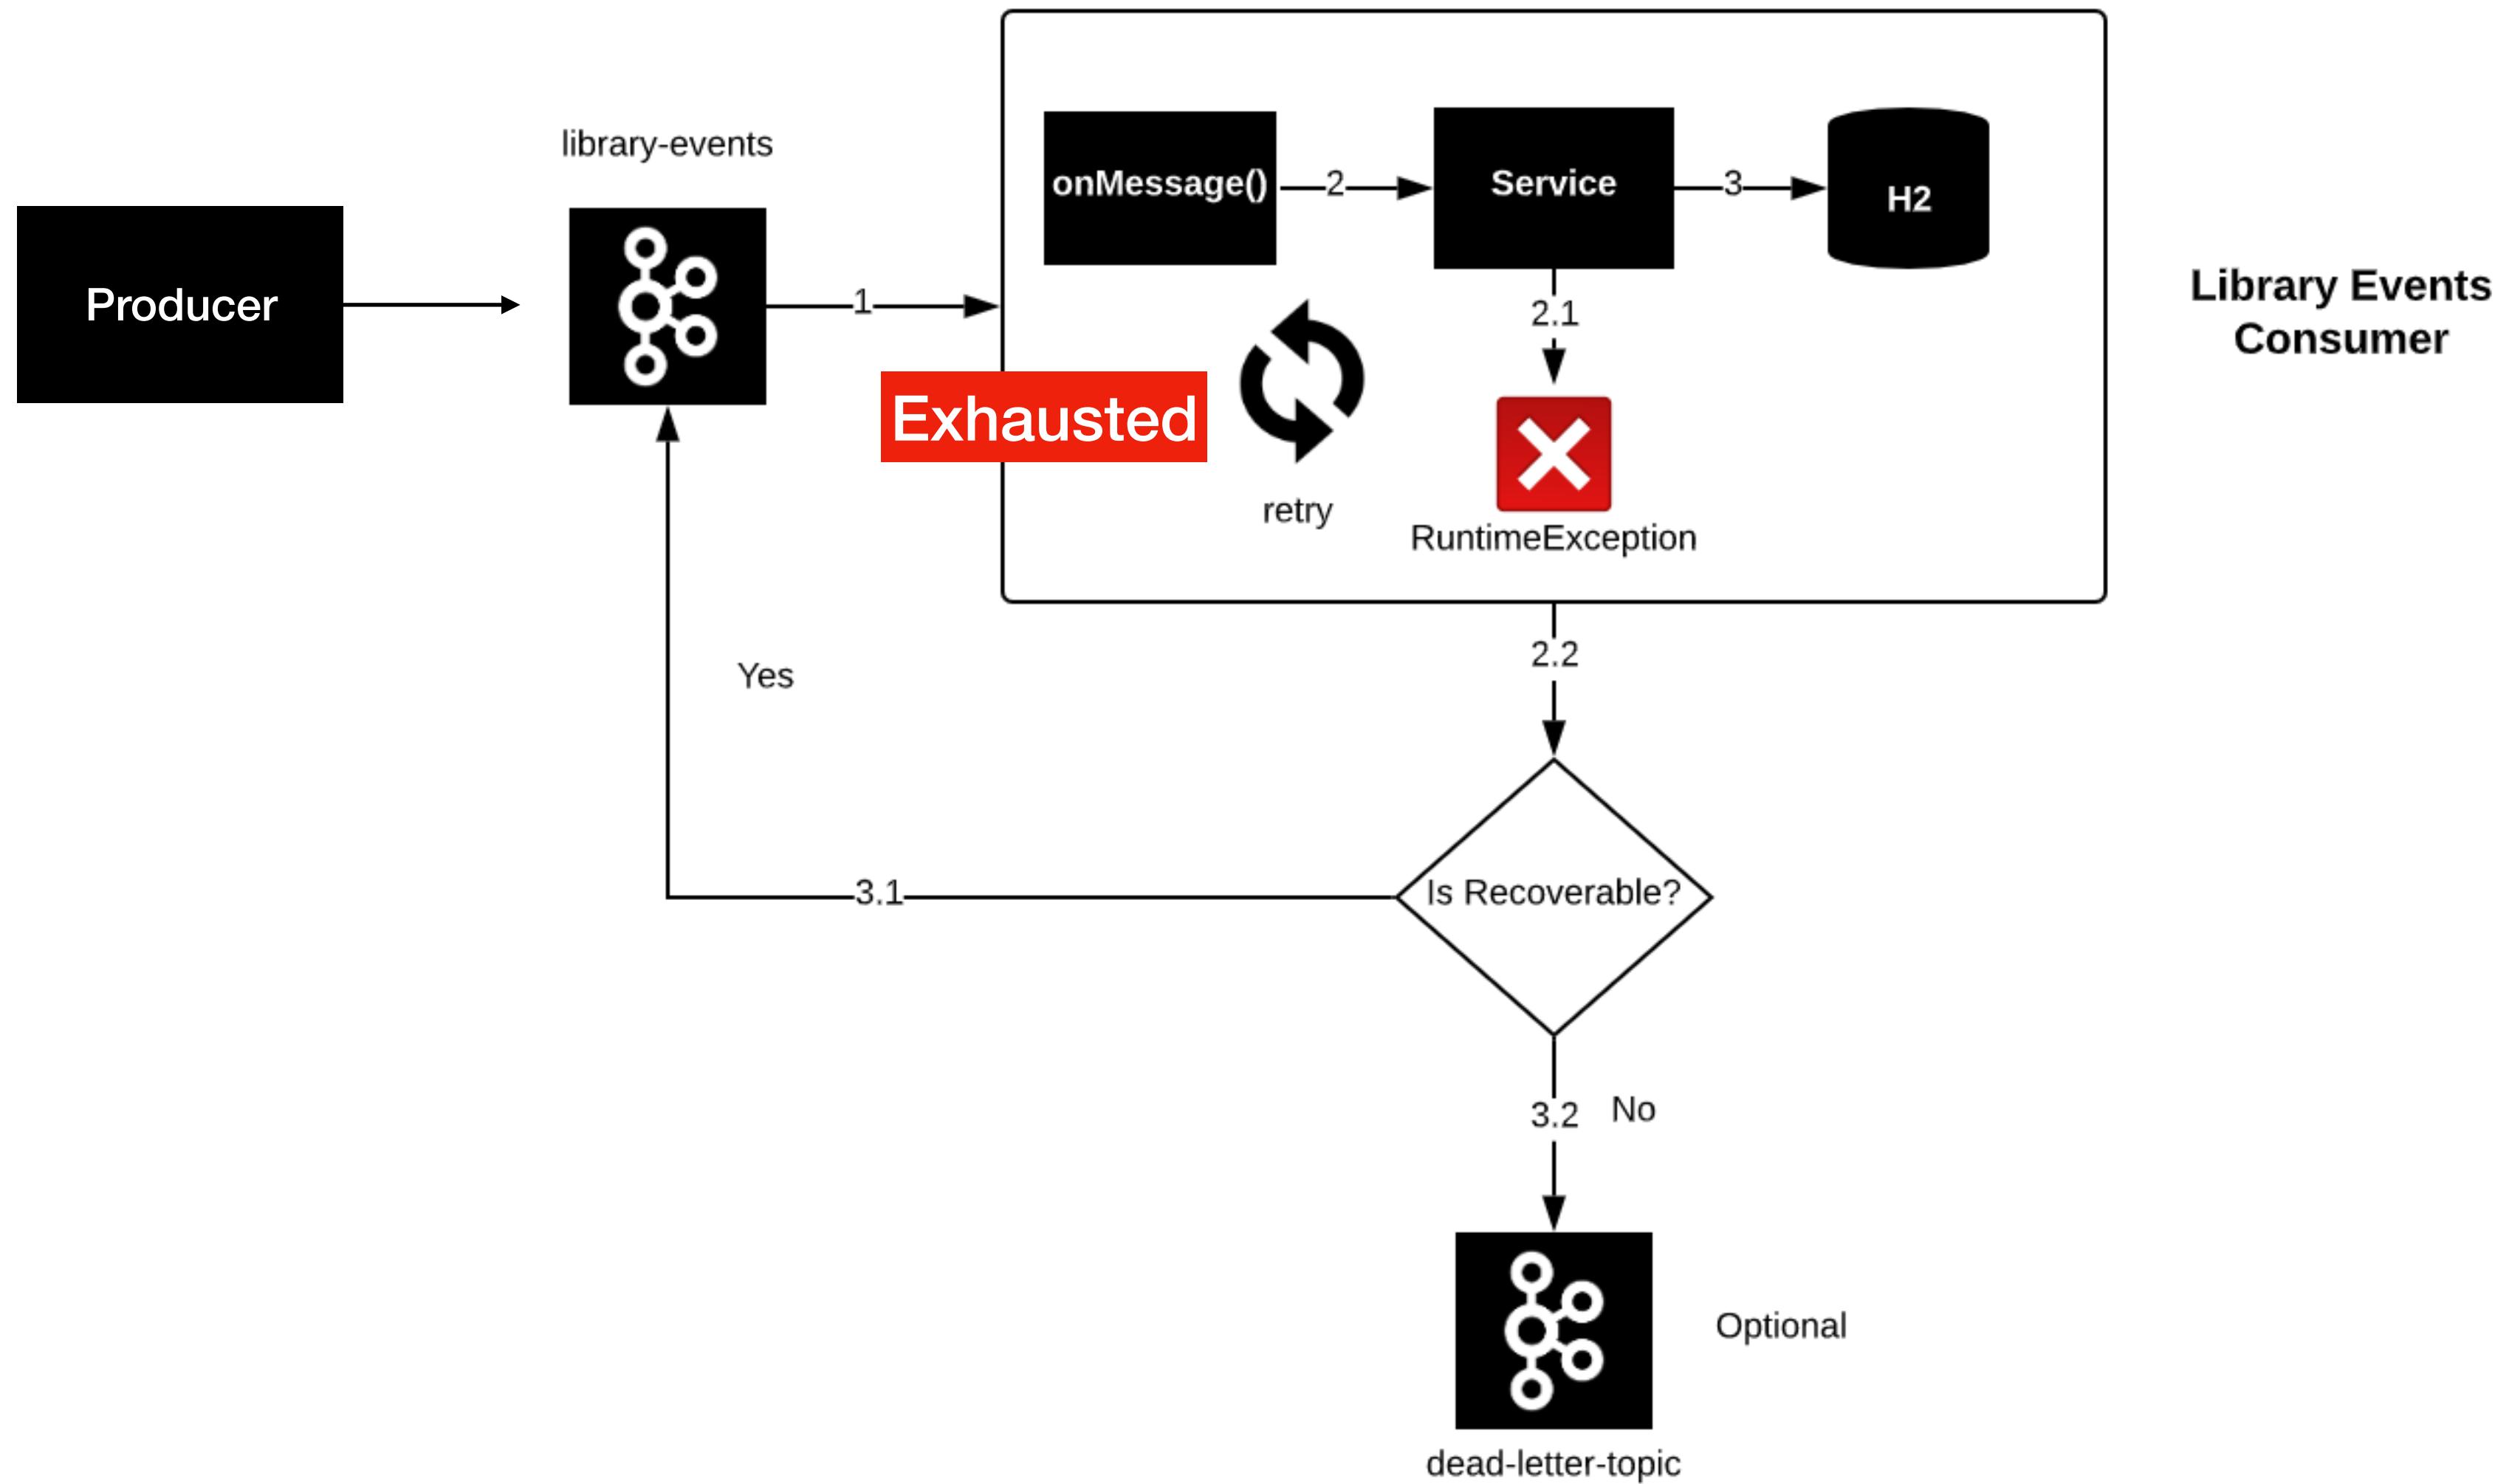
# Recovery - Type 2



# Issues with Recovery ?

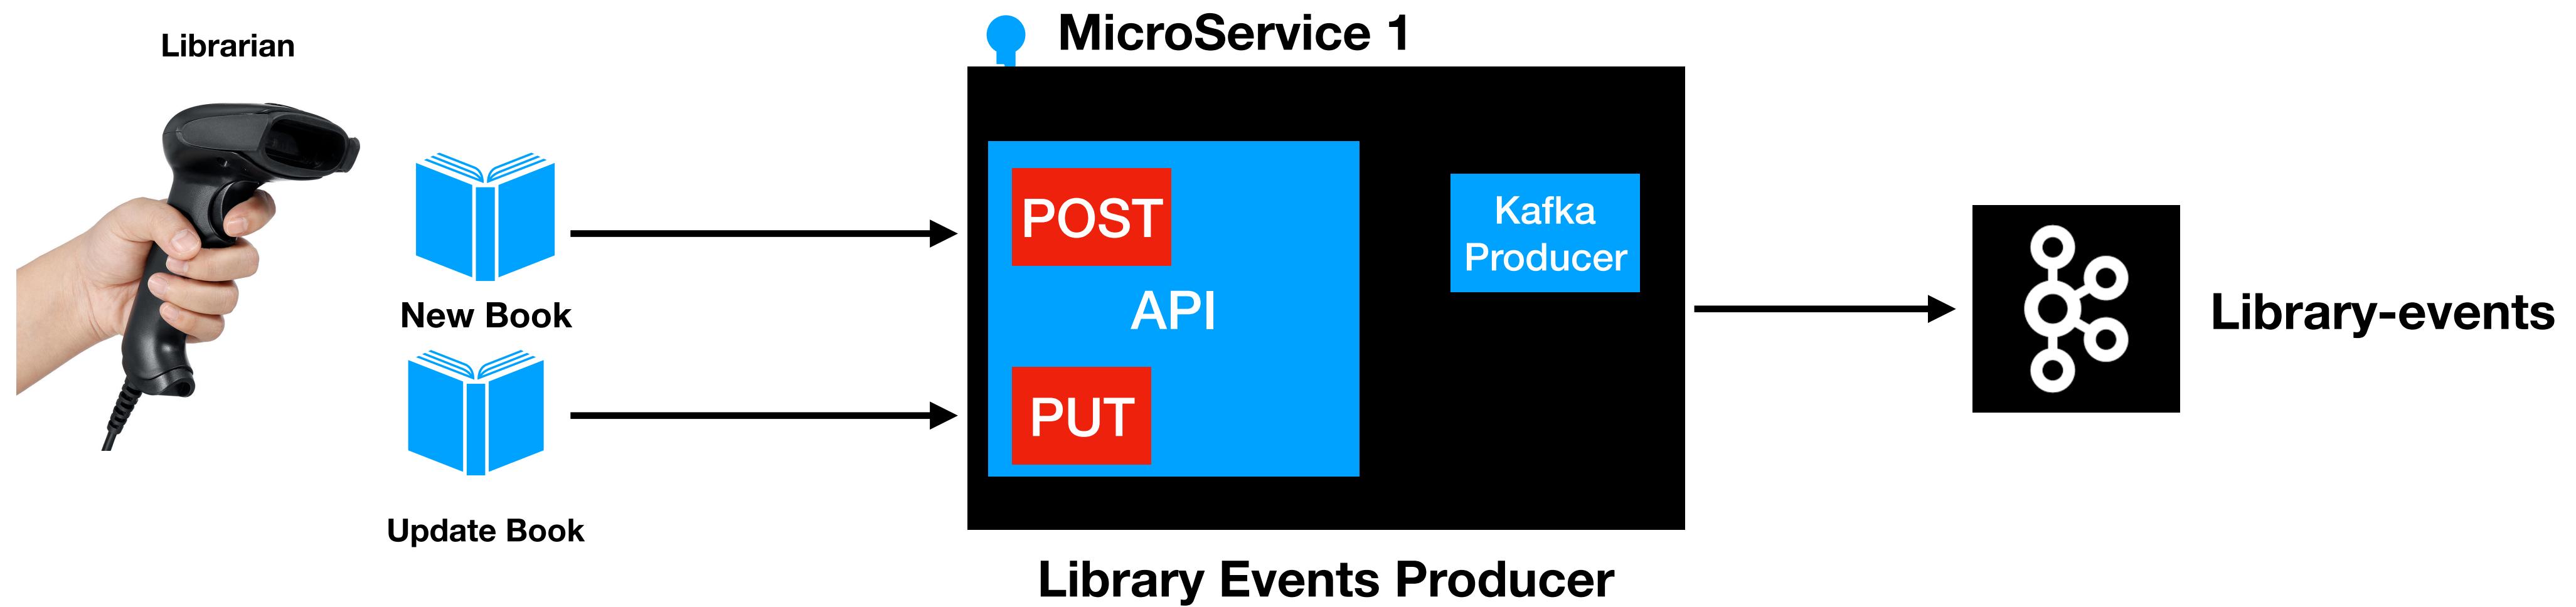
- Recovery can alter the order of events

# Recovery - Type 1



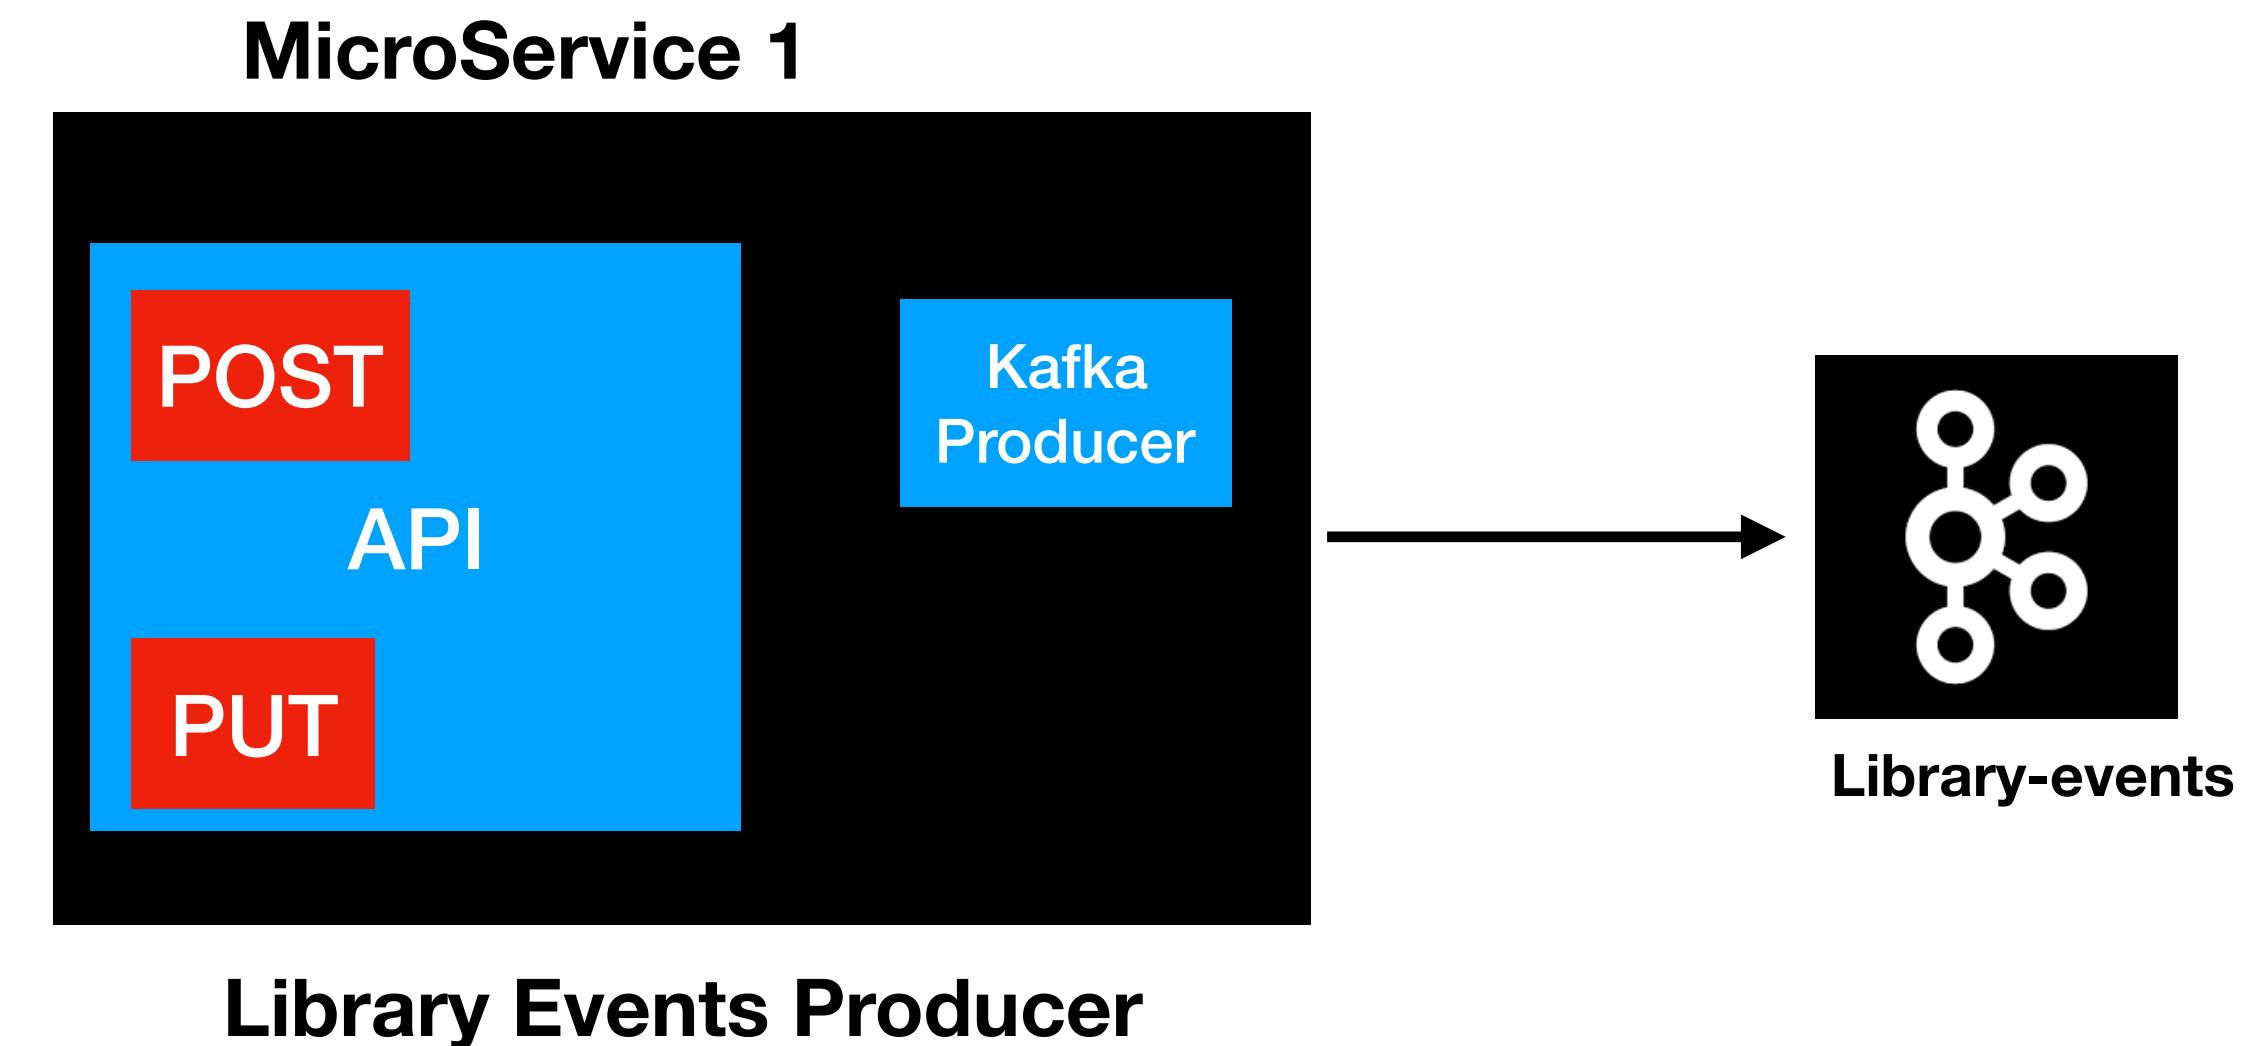
# Error Handling in Kafka Producer

# Library Events Producer API

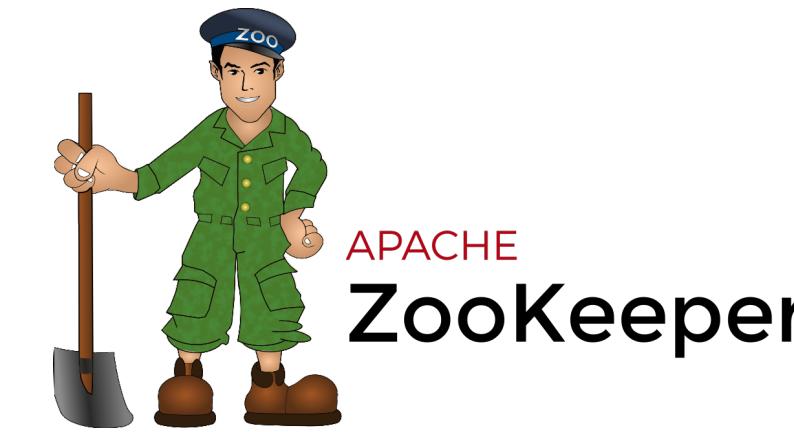


# Kafka Producer Errors

- Kafka Cluster is not available
- If **acks= all** , some brokers are not available
- **min.insync.replicas** config
  - Example : **min.insync.replicas = 2**, But only one broker is available

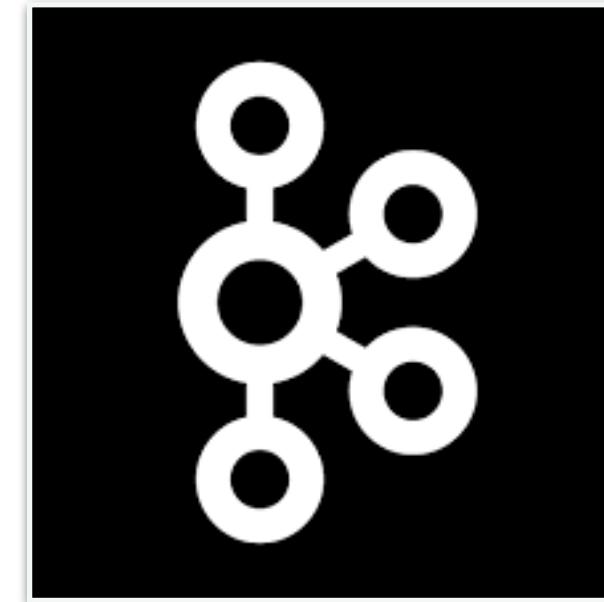


# min.insync.replicas



min.insync.replicas = 2

Kafka Cluster



Broker 1



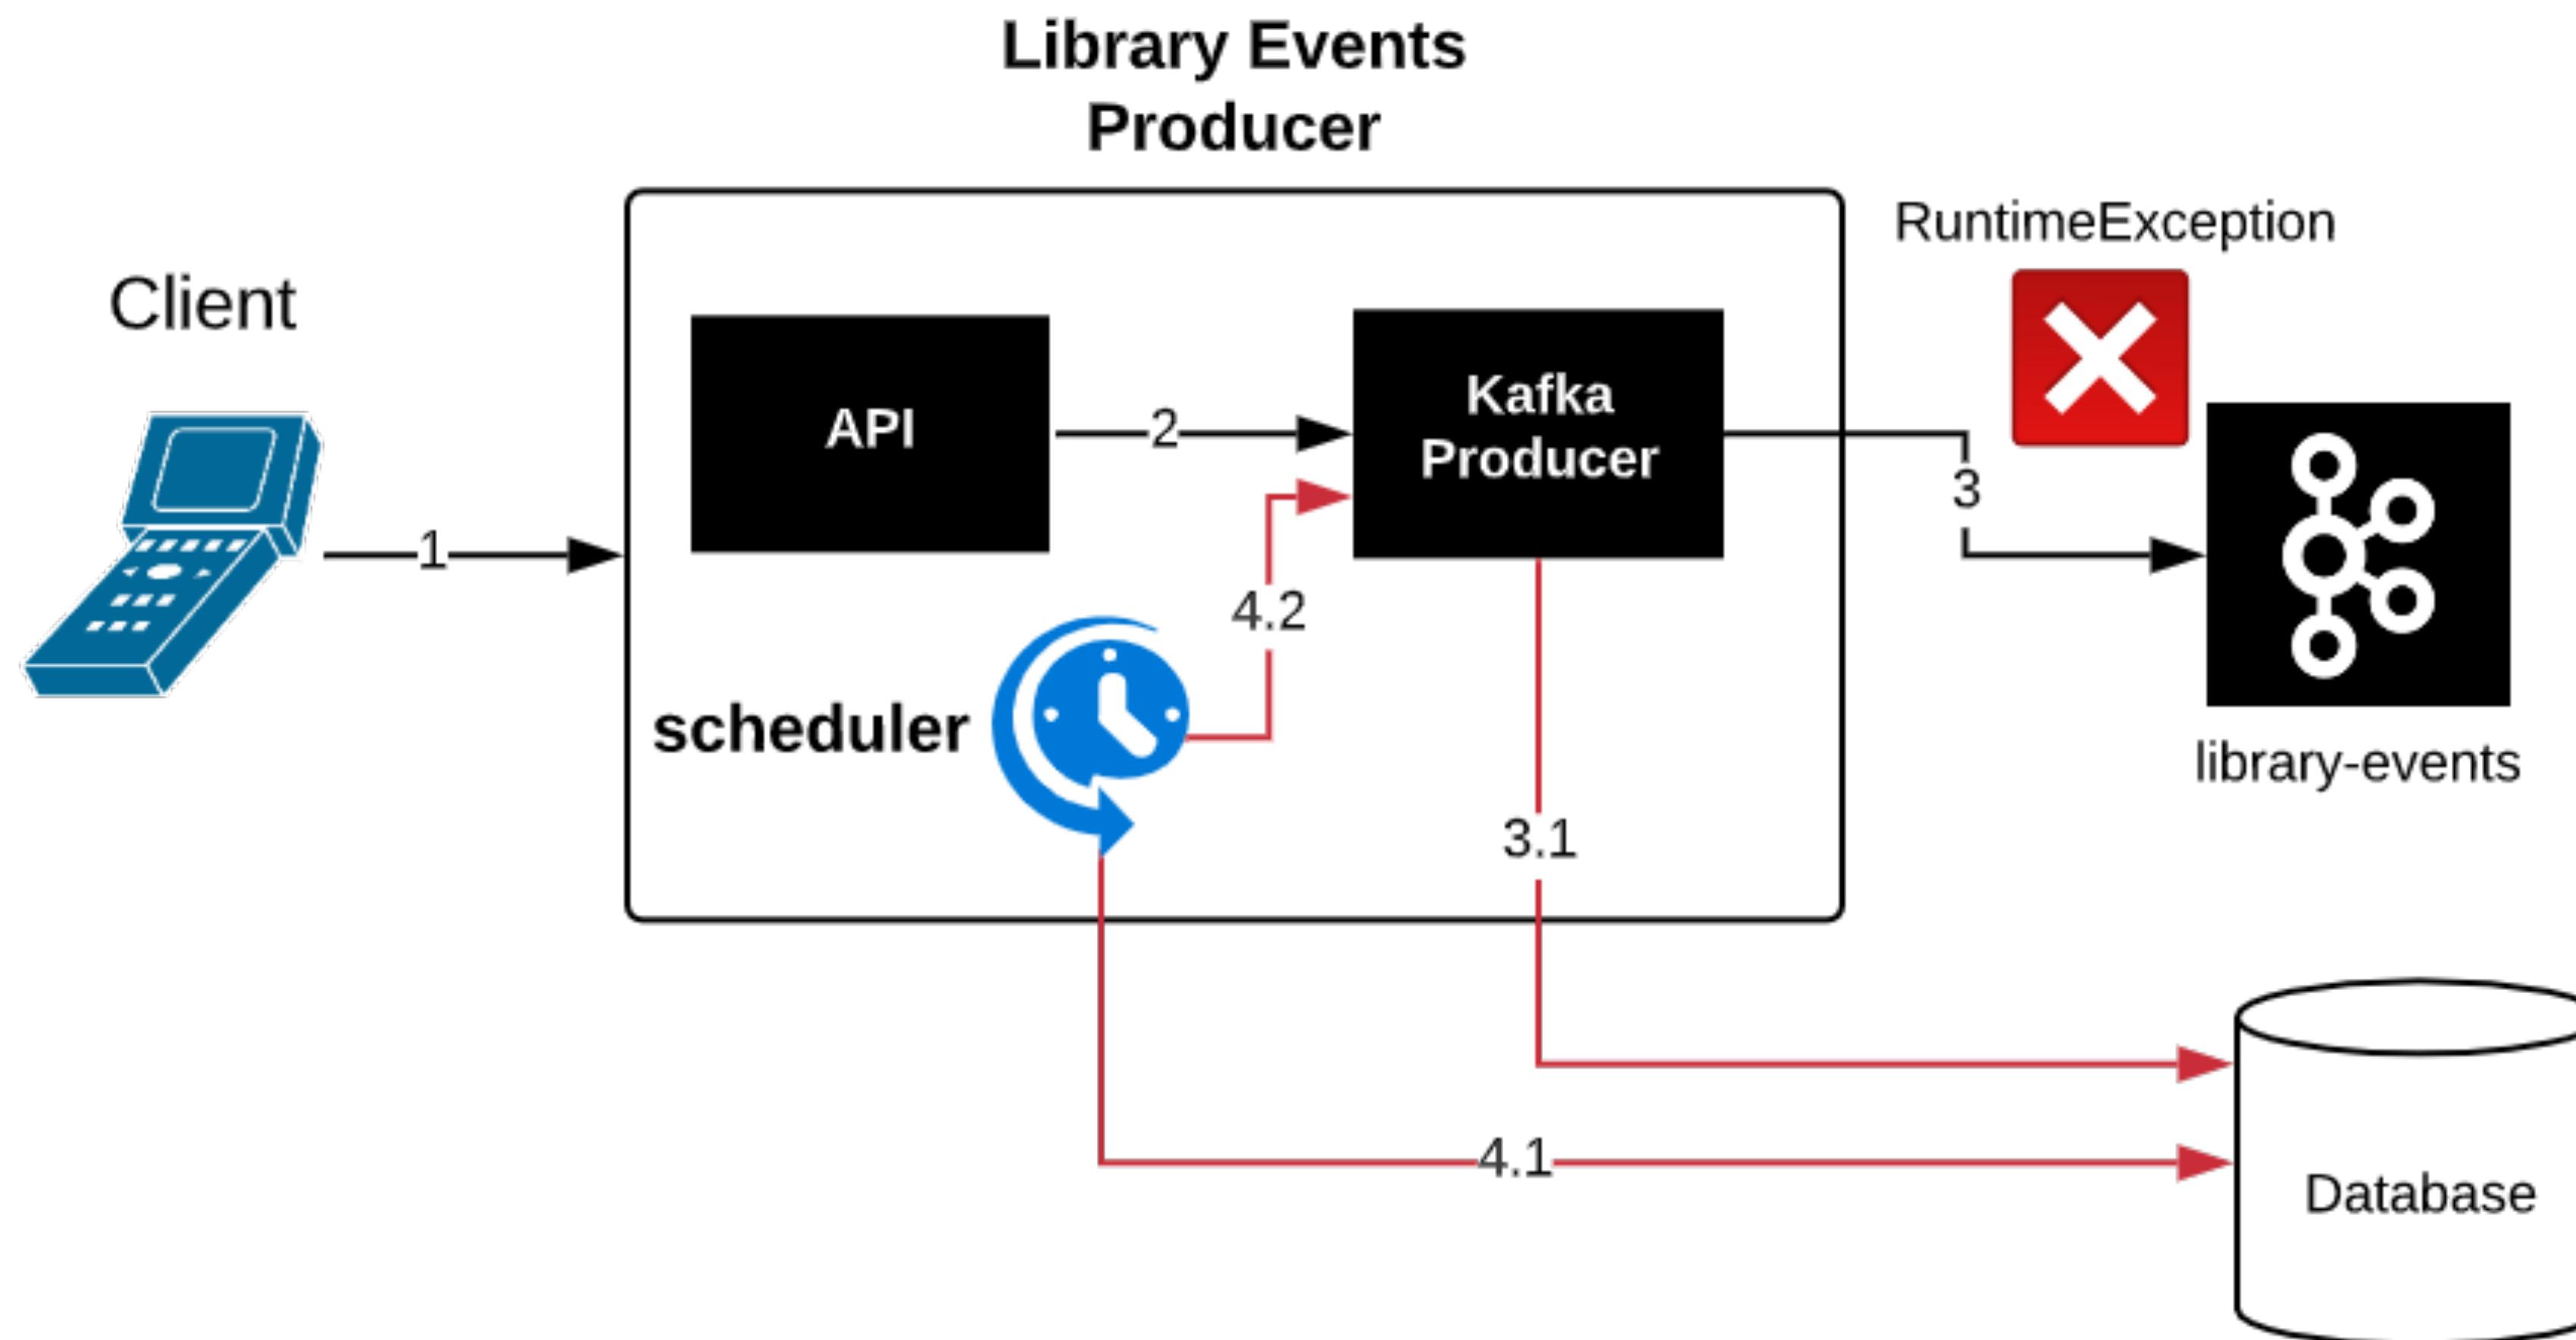
Broker 2



Broker 2

**Retain/Recover  
Failed  
Records**

# Retain/Recover Failed Records



# Retain/Recover Failed Records

## Producer Misconfiguration - Option 2

